

# Generative Recommendation Models: Progress and Directions

Yupeng Hou<sup>1</sup>, An Zhang<sup>2</sup>, Leheng Sheng<sup>2</sup>, Zhengyi Yang<sup>3</sup>, Xiang Wang<sup>3</sup>, Tat-Seng Chua<sup>2</sup>, Julian McAuley<sup>1</sup>

Proxy Speaker: Jiancan Wu

<sup>1</sup>UC San Diego    <sup>2</sup>NUS    <sup>3</sup>USTC

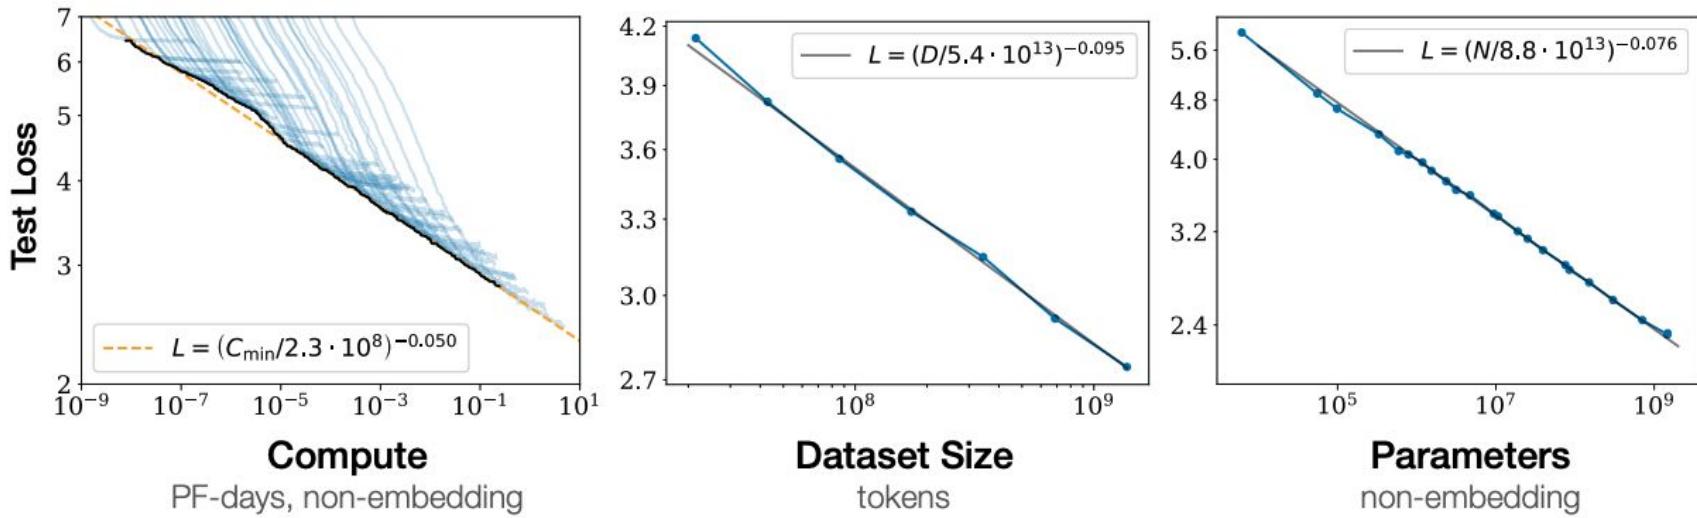


01

# Introduction

of Generative Recommendation

# Scaling Law as a Pathway towards AGI



Scaling laws provide a framework for understanding how **model size**, **data volume**, and **test-time computing** might lead to advanced AI capabilities.

# However ...

## Language Modeling

- Dense world knowledge
- Text tokens (Ten thousands level)



## User Behavior Modeling

- Sparse user-item interactions
- Items (Billion to trillion level)

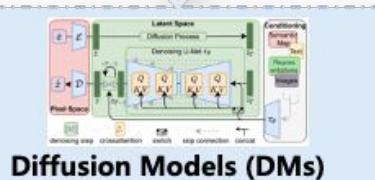
**Scaling laws rarely apply to traditional recommendation models.**

# As the Reflection of Real World,

Perceptual World



Vision



Cognitive World



Language



Behavioral World



Interaction



?

Any foundation models  
for user behaviors?

# What are Generative Models & Why ?

A generative model learns the underlying distribution of data and can generate new samples from it.



Learning  
→

$$p_{\theta}(x)$$

Training data  $\sim p_{\text{data}}(x)$

Generating  
→



Data generation  
distribution

New samples

# A Potential Solution: “Generative” Recommendation



What I can't create I don't  
understand

— *Richard P. Feynman* —

*“What User Behaviors LLMs can not Generate, LLMs do not Understand.”*

# Where are We Now?

## In language and vision:

- Large language/diffusion models have been established.
- Scaling law has been witnessed.

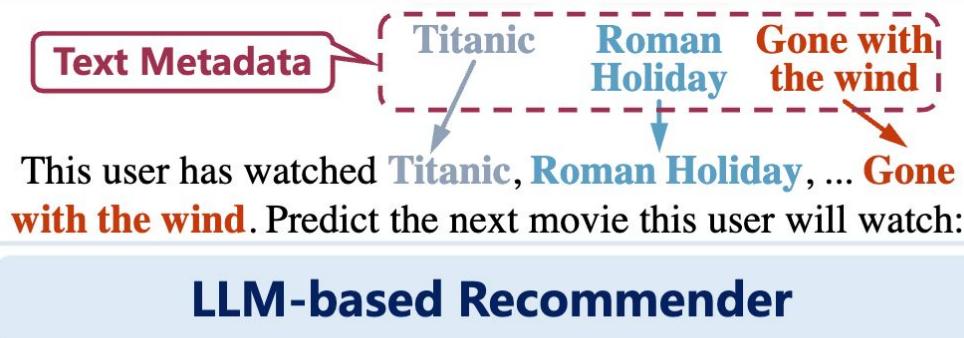
## In recommendation:

- Incorporate generative components in traditional recommender.
- Initial attempts on generative recommendation.

# Pathways towards Scalable Generative Recommendation

## Adapt Pre-trained Models

- Large Language Models



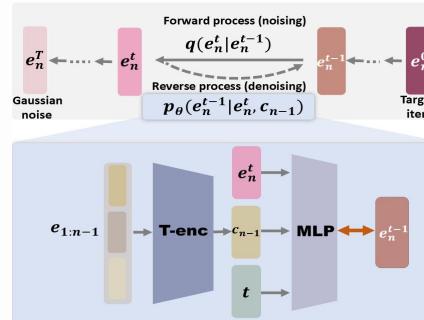
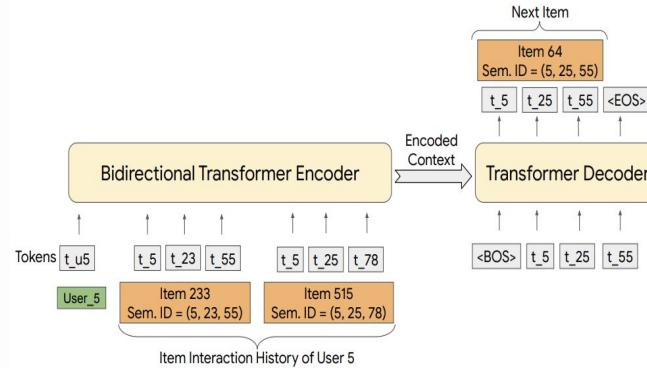
Adapting LLMs for recommendation task

# Pathways towards Scalable Generative Recommendation

## Train from Scratch

- Autoregressive Models
  - Semantic ID

- Diffusion Models



# Schedule Overview

| Time (AEST)   | Session   | Presenter                                |
|---------------|---|--|
| 9:00 - 9:10   | Part 1: Background and Introduction               | Tat-Seng Chua                            |
| 9:10 - 10:10  | Part 2: LLM-based Generative Recommendation       | Leheng Sheng                             |
| 10:10 - 10:30 | Part 3.1: Introduction of Semantic IDs            | Yupeng Hou                               |
| 10:30 - 11:00 | Coffee Break & QA Session                         |  |
| 11:00 - 11:40 | Part 3.2: SemID-based Generative Recommendation   | Yupeng Hou                               |
| 11:40 - 12:10 | Part 4: Diffusion-based Generative Recommendation | Jiancan Wu<br>(proxy speaker of Zhengyi) |
| 12:10 - 12:30 | Part 5: Open Challenges and Beyond                | Yupeng Hou                               |

02

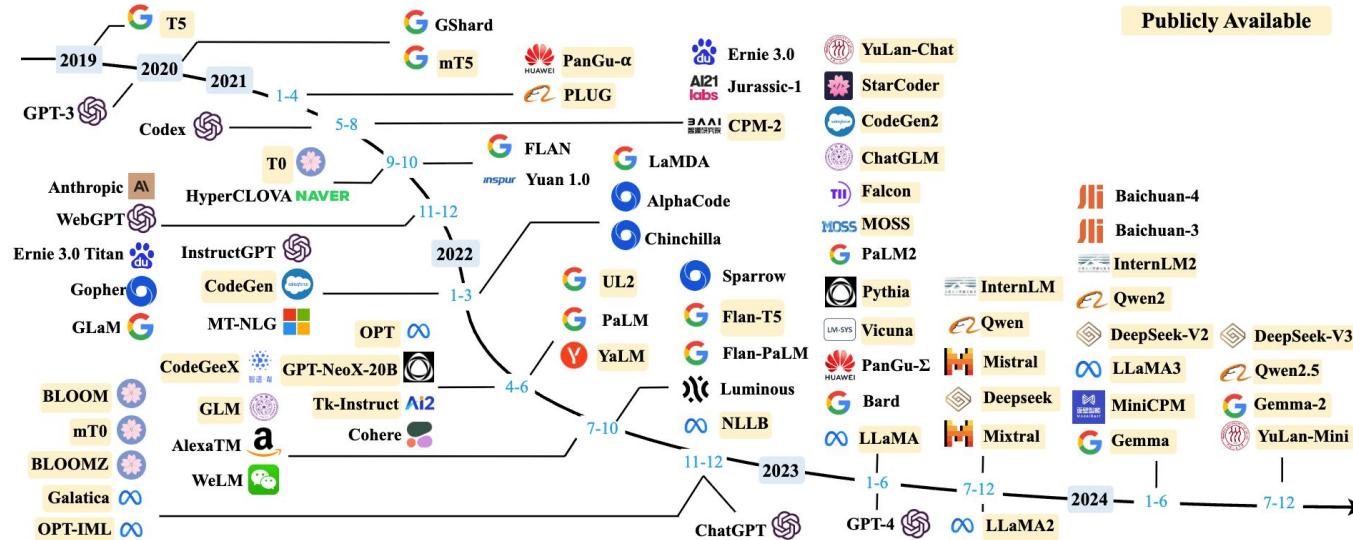
LLM

-based Generative Recommendation

# The Rise of Large Language Models

## Transformer

2017



LLMs are developing so fast recently...

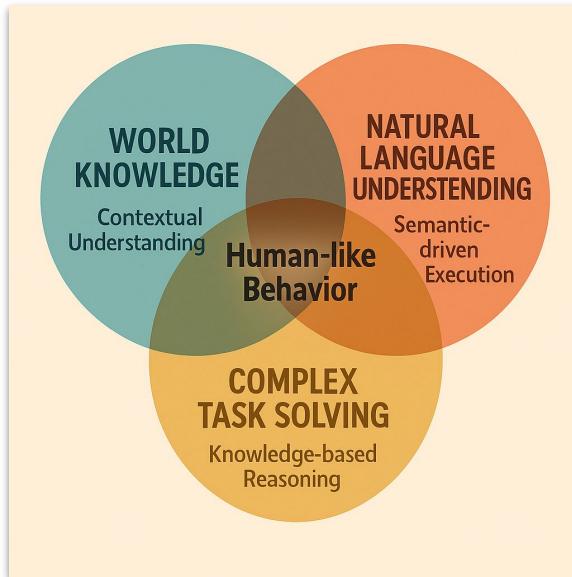
O3, R1...  
2025

# Large Language Models

LLMs are machine learning models that can perform a variety of natural language processing (NLP) tasks



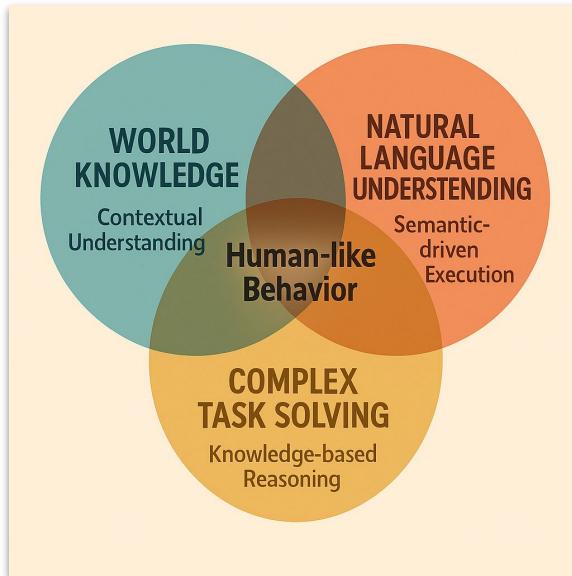
# Large Language Models



## Key features of LLMs:

- World knowledge.
- Natural language understanding.
- Human-like behavior.

# Large Language Models



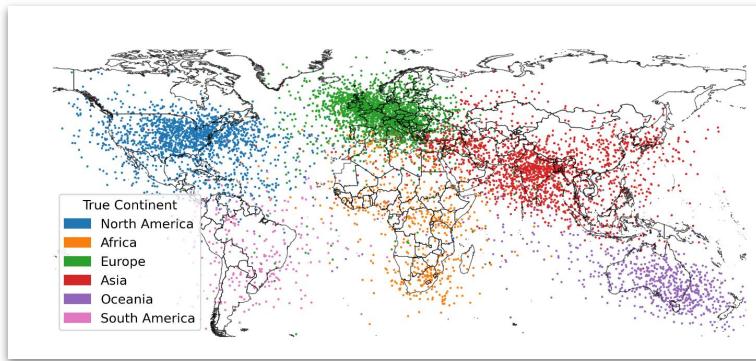
## Key features of LLMs:

- World knowledge.
- Natural language understanding.
- Human-like behavior.

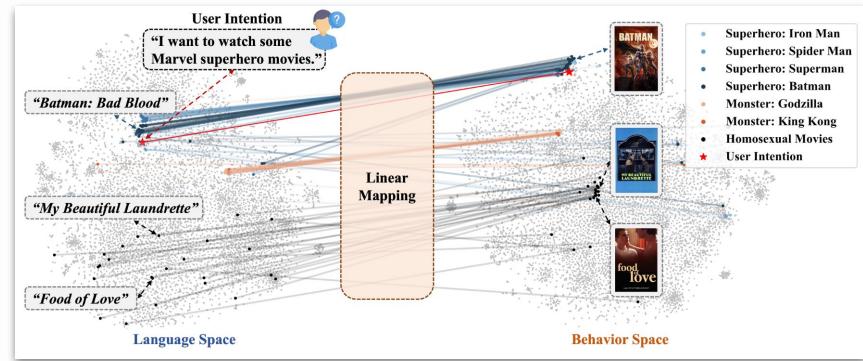
**How can these features benefit recommender systems?**

# Benefits of LLMs for Recommendation

## (1) World knowledge – from pretraining



In space



In recommendation

# Benefits of LLMs for Recommendation

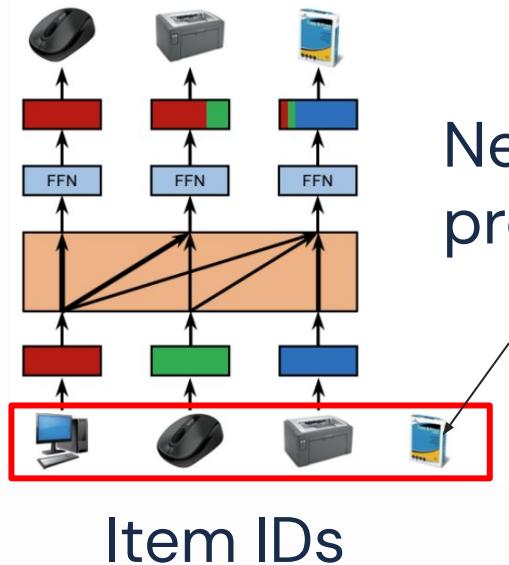
## (1) World knowledge

### LLM as sequential recommender

→ Alleviating the **data sparsity** of ID-based interactions in recommendation

# Benefits of LLMs for Recommendation

## (1) World knowledge

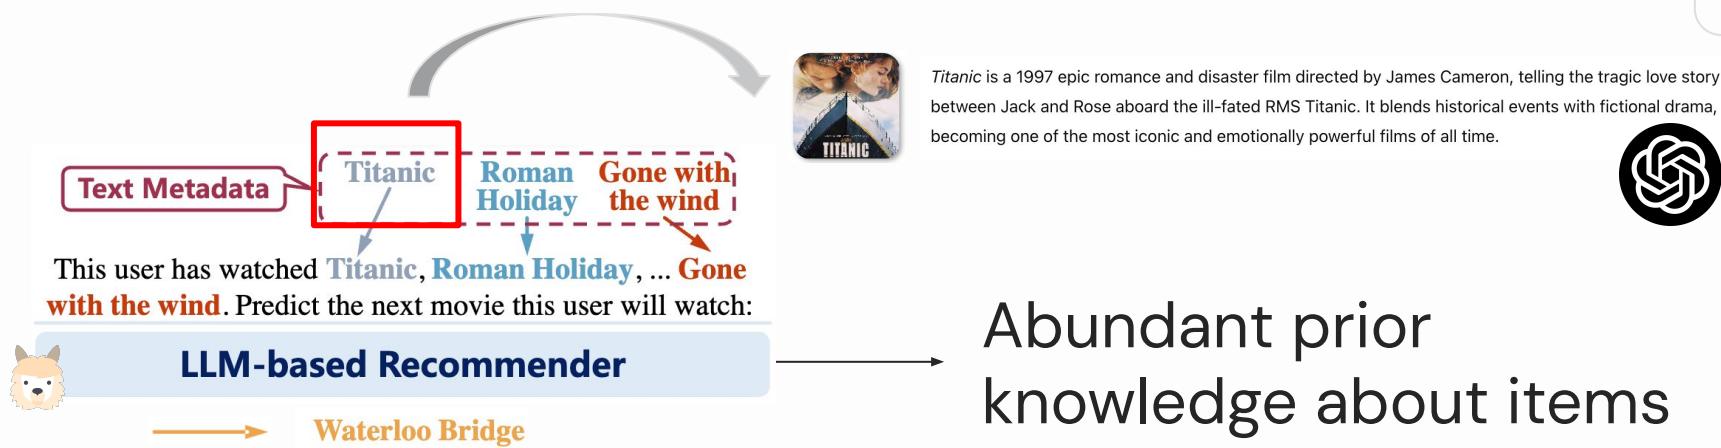


ID-based item modeling  
lack semantic meanings

Example: SASRec [*ICDM'18*]

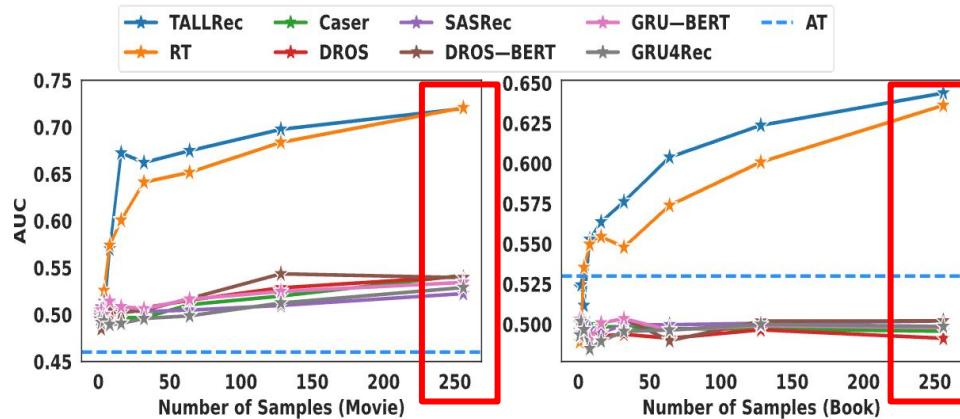
# Benefits of LLMs for Recommendation

## (1) World knowledge



# Benefits of LLMs for Recommendation

## (1) World knowledge



Few data → a good recommender

# Benefits of LLMs for Recommendation

## (1) World knowledge



### LLM as sequential recommender

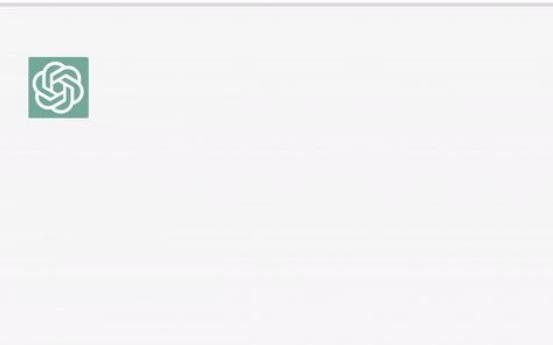


Lower data requirement  
Cross-domain ability  
Cold-start ability  
...

# Benefits of LLMs for Recommendation

## (2) Natural language understanding & generation

CP |



LLMs can interact  
with users fluently

# Benefits of LLMs for Recommendation

## (2) Natural language understanding & generation

LLM as conversational recommender

→ Towards more interactive recommender systems

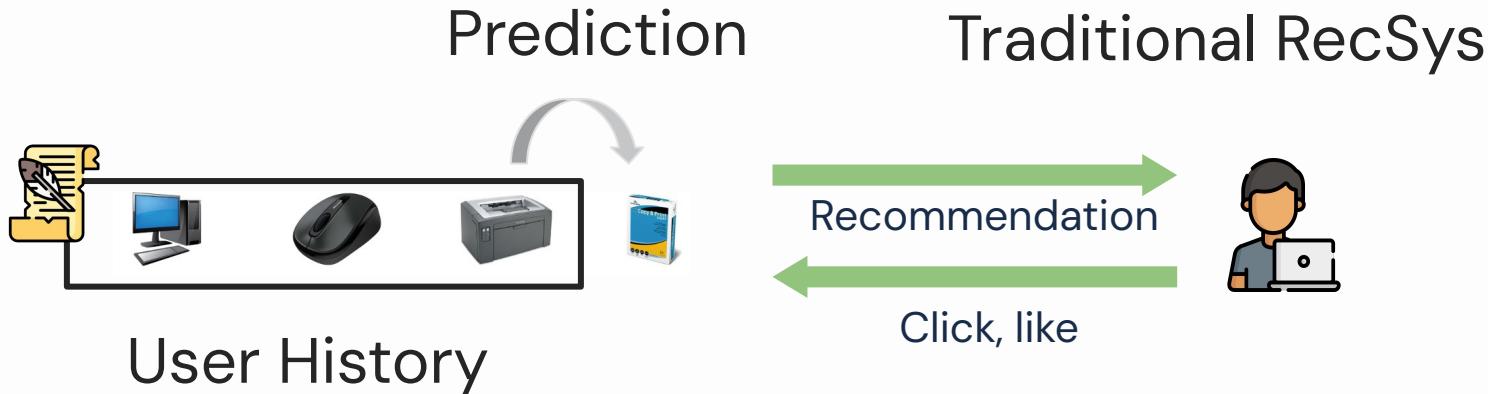
# Benefits of LLMs for Recommendation

## (2) Natural language understanding & generation



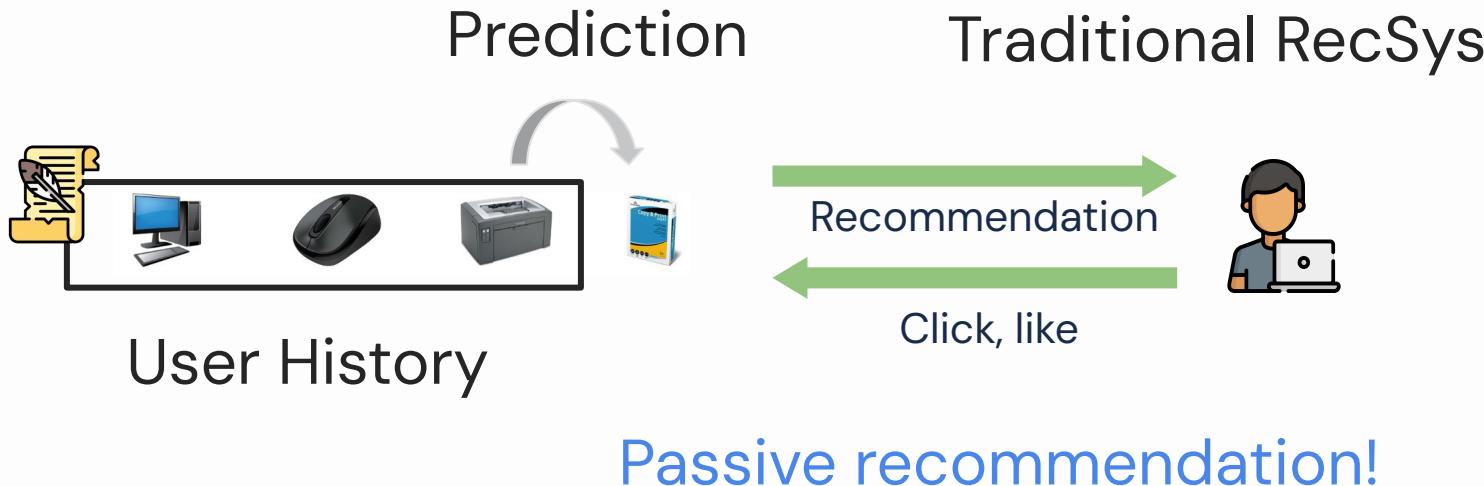
# Benefits of LLMs for Recommendation

## (2) Natural language understanding & generation



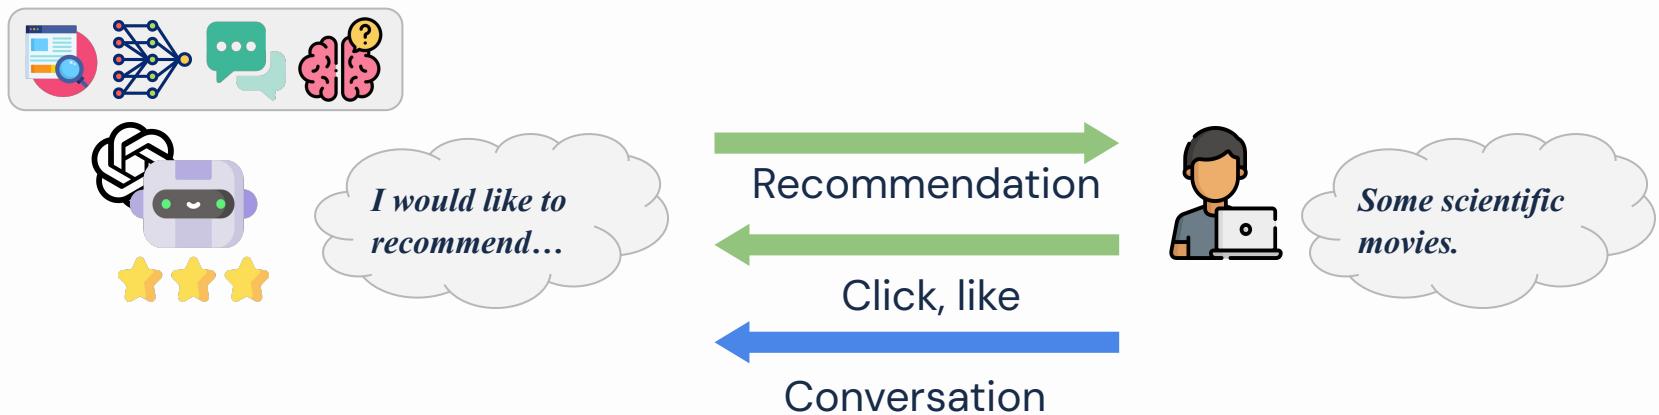
# Benefits of LLMs for Recommendation

## (2) Natural language understanding & generation



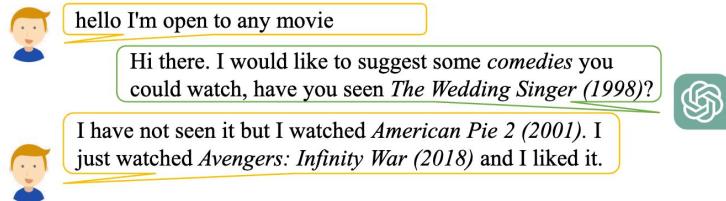
# Benefits of LLMs for Recommendation

## (2) Natural language understanding & generation



# Benefits of LLMs for Recommendation

## (2) Natural language understanding & generation



### LLM as conversational recommender

Interactive  
User-friendly  
More accurate



...

# Benefits of LLMs for Recommendation

## (3) Human-like behavior



# Benefits of LLMs for Recommendation

## (3) Human-like behavior



Generative Agents can (mostly) simulate human behaviors

- Cooperation
- Organization

# Benefits of LLMs for Recommendation

## (3) Human-like behavior

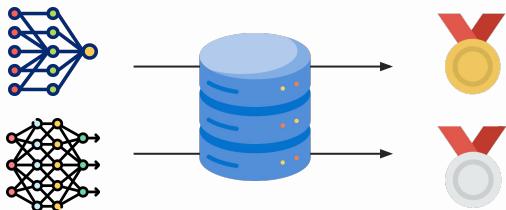
### LLM as user simulator

-> Simulating user behaviors for evaluating recommenders.

# Benefits of LLMs for Recommendation

## (3) Human-like behavior

### Offline recommender evaluation

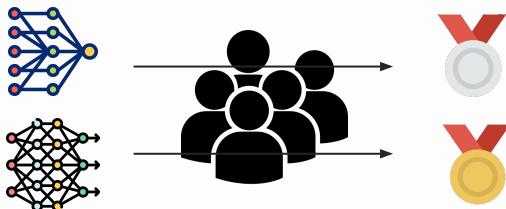


Inaccurate, but  
affordable

# Benefits of LLMs for Recommendation

## (3) Human-like behavior

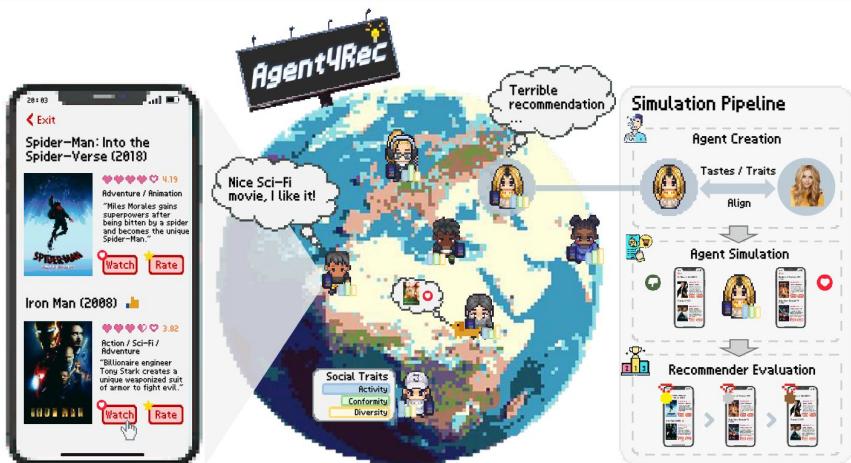
### Online recommender evaluation



Accurate, but  
costly

# Benefits of LLMs for Recommendation

## (3) Human-like behavior



LLM as user simulator

Faithful  
Affordable  
Controllable



...

# Part 1: LLM as Sequential Recommender

(i) **Early efforts**: Pretrained LLMs for recommendation;

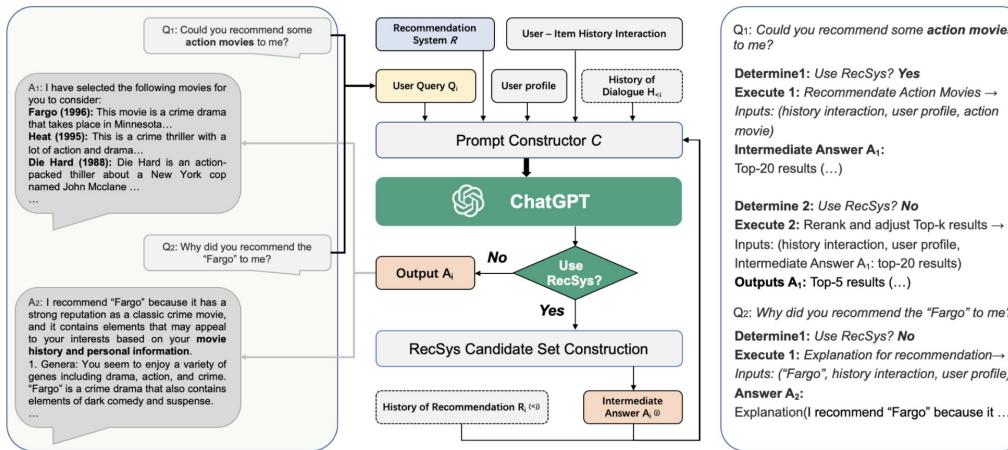
# Early efforts

- Directly use **freezed LLMs** (e.g., GPT 4) for recommendation.

# Early efforts

## Prompt Engineering + In-Context Learning (ChatRec)

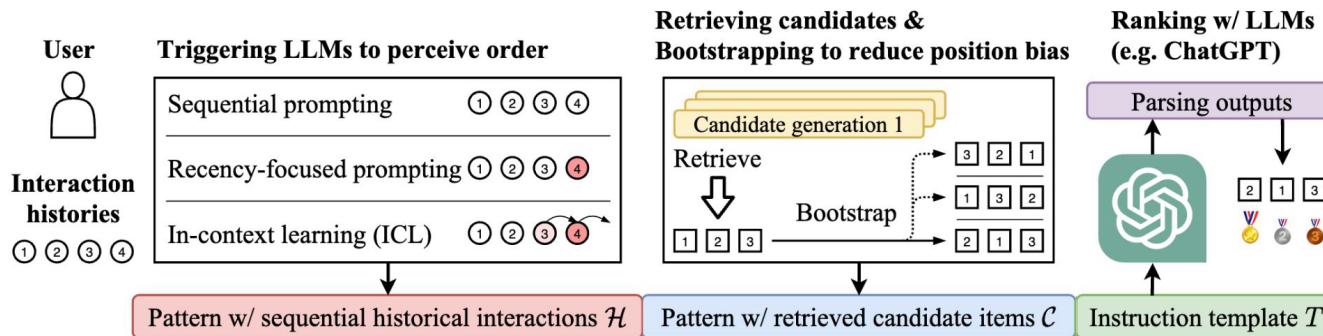
### Key idea: LLMs as the recsys controller



# Early efforts

## Prompt Engineering + In-Context Learning (LLMRank)

Key idea: LLMs as the reranker



# Early efforts

- Directly use freezed LLMs (e.g., GPT 4) for recommendation.
- A **performance gap** compared to traditional recommenders exists.

# Early efforts

Sub-optimal performance comparing to SASRec!

Performance of LLMRank

| Method    | ML-1M        |             |             |              | Games        |             |             |             |
|-----------|--------------|-------------|-------------|--------------|--------------|-------------|-------------|-------------|
|           | N@1          | N@5         | N@10        | N@20         | N@1          | N@5         | N@10        | N@20        |
| full      | Pop          | 0.08        | 1.20        | 4.13         | 5.79         | 0.13        | 1.00        | 2.27        |
|           | BPRMF [49]   | 0.26        | 1.69        | 4.41         | 6.04         | 0.55        | 1.98        | <b>2.96</b> |
|           | SASRec [33]  | <b>3.76</b> | <b>9.79</b> | <b>10.45</b> | <b>10.56</b> | <b>1.33</b> | <b>3.55</b> | <b>4.02</b> |
| zero-shot | BM25 [50]    | 0.26        | 0.87        | 2.32         | 5.28         | 0.18        | 1.07        | 1.80        |
|           | UniSRec [30] | 0.88        | 3.46        | 5.30         | 6.92         | 0.00        | 1.86        | 2.03        |
|           | VQ-Rec [29]  | 0.20        | 1.60        | 3.29         | 5.73         | 0.20        | 1.21        | 1.91        |
| Ours      |              | <b>1.74</b> | <b>5.22</b> | <b>6.91</b>  | <b>7.90</b>  | <b>0.90</b> | <b>2.26</b> | 2.80        |
| Ours      |              |             |             |              |              |             |             | 3.08        |

# Early efforts

Sub-optimal performance comparing to SASRec!

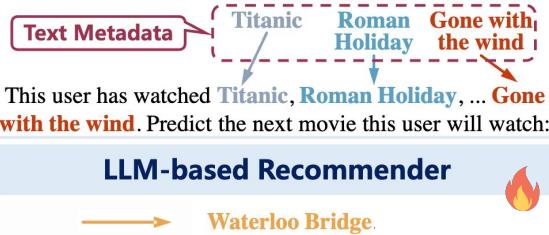
Aligning LLMs for recommendation tasks is necessary!

|           | SASRec [55]  | 0.70        | 0.75        | 10.40       | 10.00       | 1.00        | 0.00        | -1.00 | -1.11 |
|-----------|--------------|-------------|-------------|-------------|-------------|-------------|-------------|-------|-------|
| zero-shot | BM25 [50]    | 0.26        | 0.87        | 2.32        | 5.28        | 0.18        | 1.07        | 1.80  | 2.55  |
|           | UniSRec [30] | 0.88        | 3.46        | 5.30        | 6.92        | 0.00        | 1.86        | 2.03  | 2.31  |
|           | VQ-Rec [29]  | 0.20        | 1.60        | 3.29        | 5.73        | 0.20        | 1.21        | 1.91  | 2.64  |
|           | Ours         | <b>1.74</b> | <b>5.22</b> | <b>6.91</b> | <b>7.90</b> | <b>0.90</b> | <b>2.26</b> | 2.80  | 3.08  |

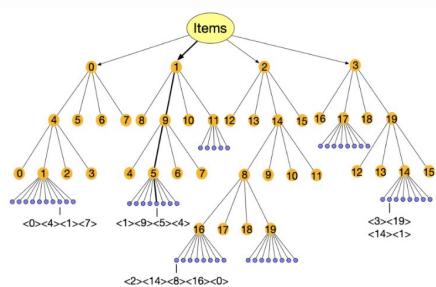
# Part 1: LLM as Sequential Recommender

- (i) Early efforts: Pretrained LLMs for recommendation;
- (ii) **Aligning** LLMs for recommendation;

# Aligning LLMs for recommendation



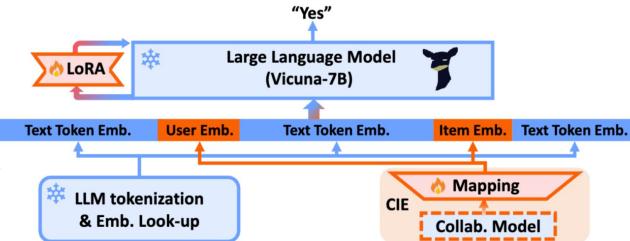
Pure text-based



+ External item tokens

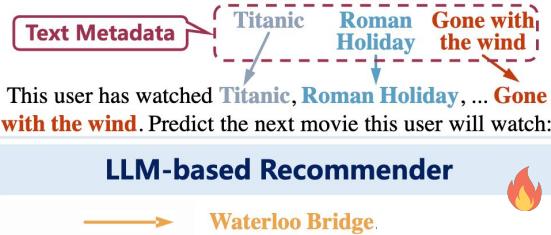


+ Collaborative embeddings

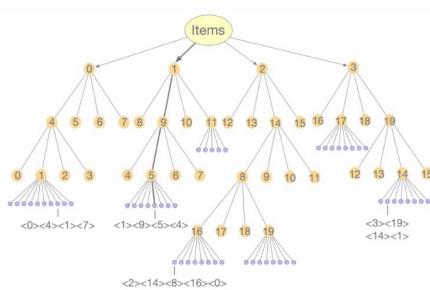


+ Multimodal information

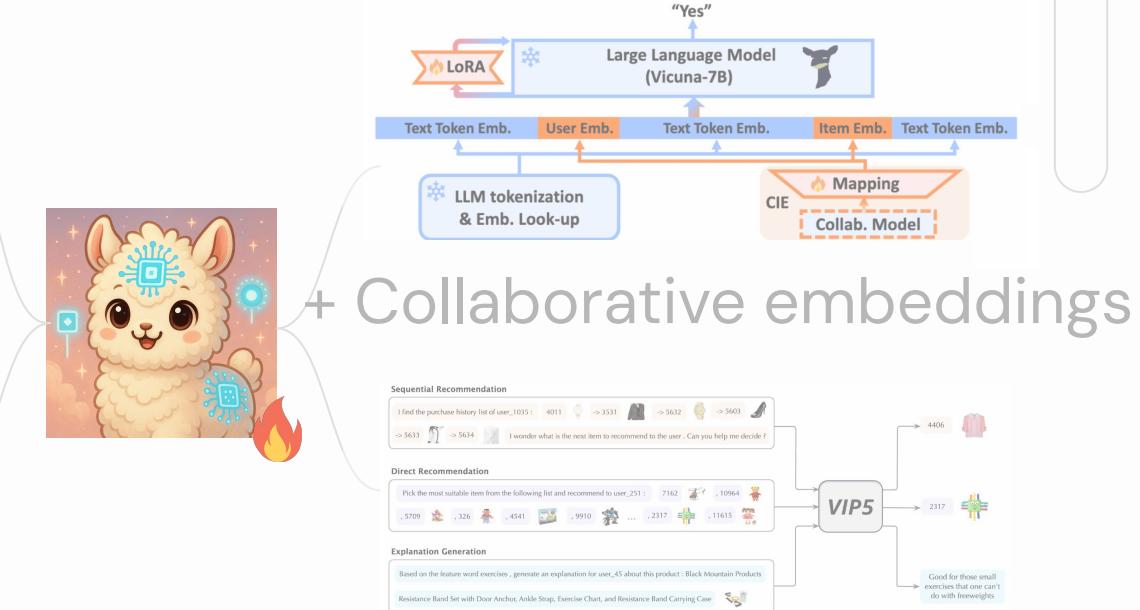
# Aligning LLMs for recommendation



Pure text-based



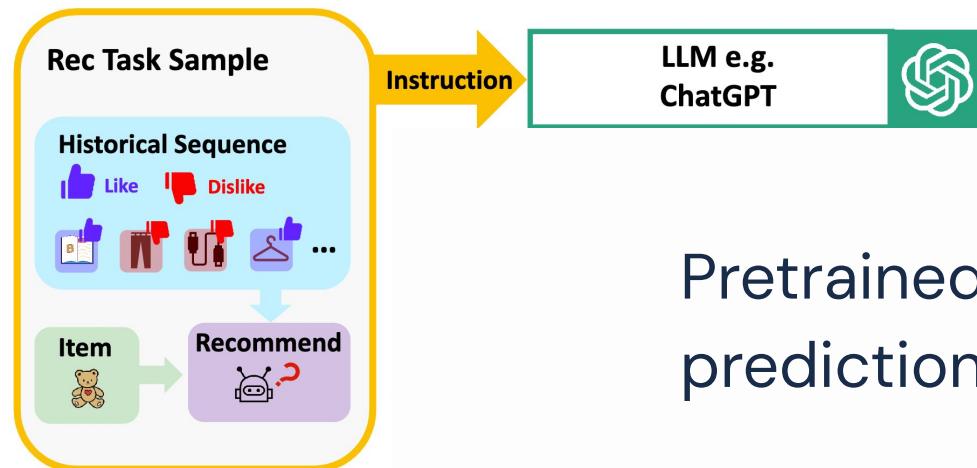
+ External item tokens



+ Multimodal information

# Aligning LLMs for recommendation

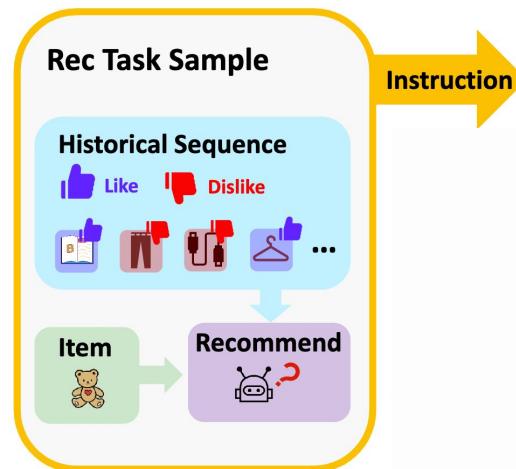
## (1) Pure text-based (TALLRec)



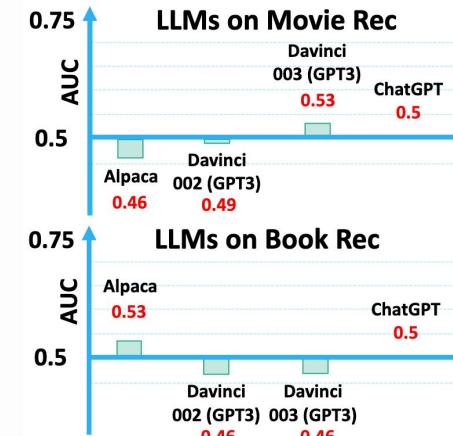
Pretrained LLMs for CTR prediction?

# Aligning LLMs for recommendation

## (1) Pure text-based (TALLRec)



Pretrained LLMs:  
Random Guess!



# Aligning LLMs for recommendation

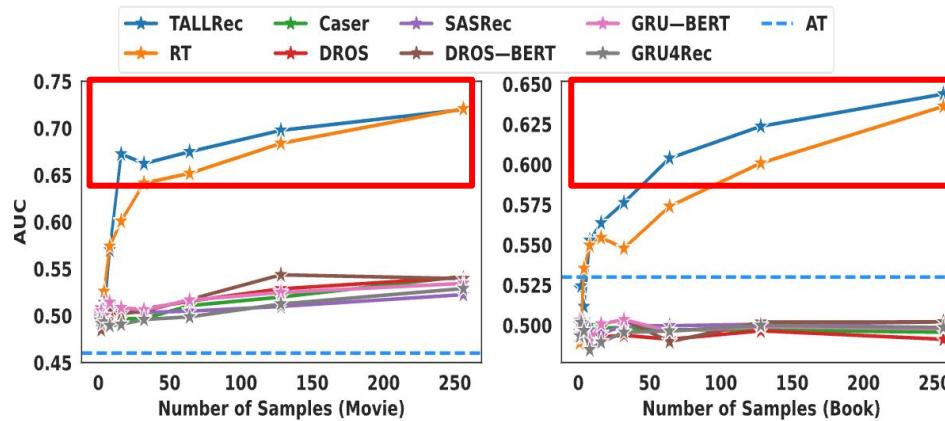
## (1) Pure text-based (TALLRec)



General task alignment → Recommendation alignment

# Aligning LLMs for recommendation

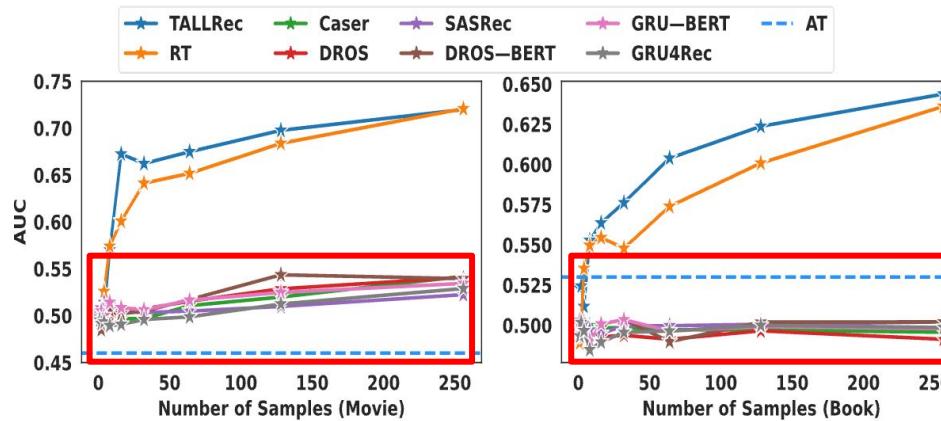
## (1) Pure text-based (TALLRec)



Few training data → Huge improvements

# Aligning LLMs for recommendation

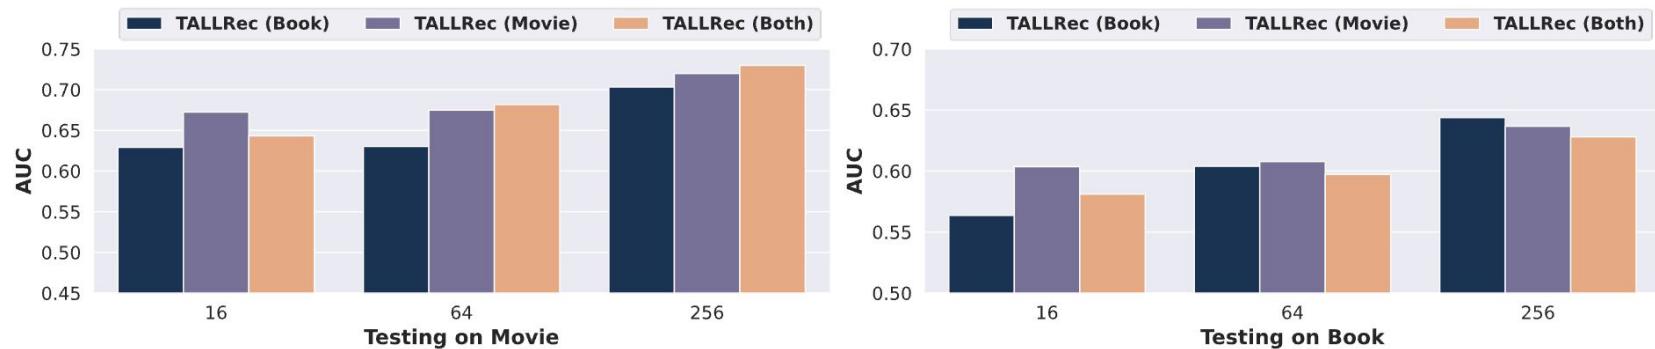
## (1) Pure text-based (TALLRec)



Traditional recommenders: suffer from too-sparse supervision signals

# Aligning LLMs for recommendation

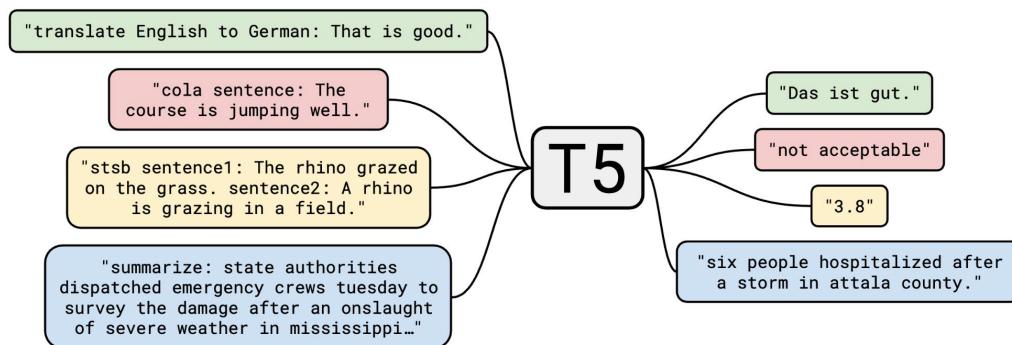
## (1) Pure text-based (TALLRec)



Cross-domain generalization

# Aligning LLMs for recommendation

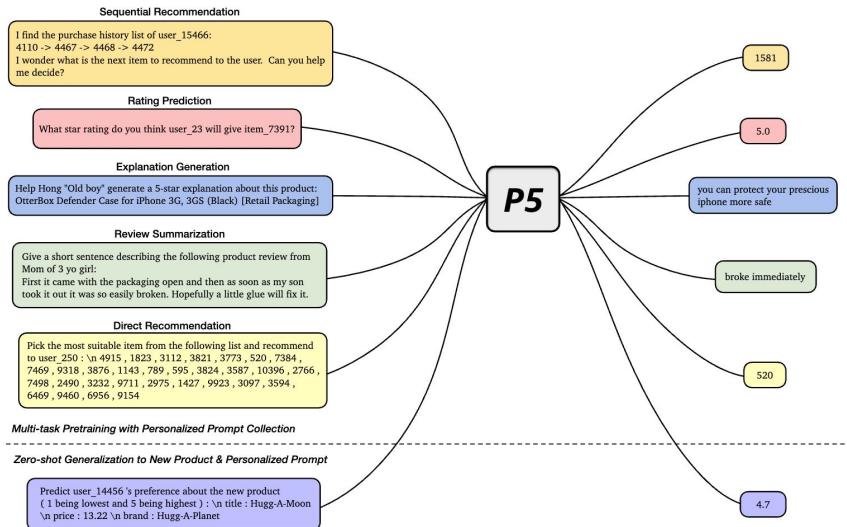
## (1) Pure text-based – Multiple rec tasks



Unified language  
modeling in NLP

# Aligning LLMs for recommendation

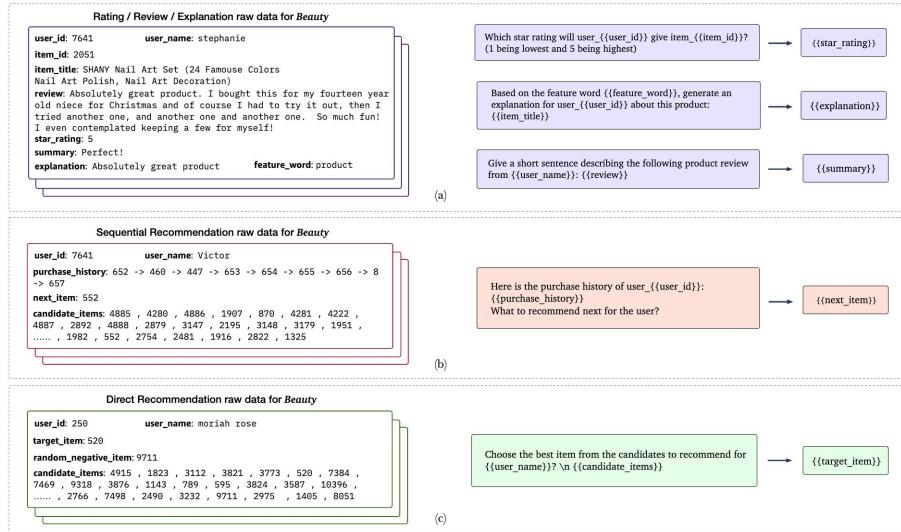
## (1) Pure text-based – Multiple rec taks



Multi-task alignment (P5)  
-> general recommender

# Aligning LLMs for recommendation

## (1) Pure text-based – Multiple rec tasks



Training on different task prompts → multiple recommendation abilities.

# Aligning LLMs for recommendation

## (1) Pure text-based – Multiple rec tasks

Table 6: Performance comparison on review summarization (%).

| Methods     | Sports        |                |               |                | Beauty        |               |               |               | Toys          |               |               |               |
|-------------|---------------|----------------|---------------|----------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|
|             | BLUE2         | ROUGE1         | ROUGE2        | ROUGEL         | BLUE2         | ROUGE1        | ROUGE2        | ROUGEL        | BLUE2         | ROUGE1        | ROUGE2        | ROUGEL        |
| T0 (4-1)    | 2.1581        | 2.2695         | 0.5694        | 1.6221         | 1.2871        | 1.2750        | 0.3904        | 0.9592        | 2.2296        | 2.4671        | 0.6482        | 1.8424        |
| GPT-2 (4-1) | 0.7779        | 4.4534         | 1.0033        | 1.9236         | 0.5879        | 3.3844        | 0.6756        | 1.3956        | 0.6221        | 3.7149        | 0.6629        | 1.4813        |
| P5-S (4-1)  | 2.4962        | 11.6701        | 2.7187        | 10.4819        | 2.1225        | 8.4205        | 1.6676        | 7.5476        | 2.4752        | 9.4200        | 1.5975        | 8.2618        |
| P5-B (4-1)  | <b>2.6910</b> | <b>12.0314</b> | <b>3.2921</b> | <b>10.7274</b> | <u>1.9325</u> | <u>8.2909</u> | <u>1.4321</u> | <u>7.4000</u> | <u>1.7833</u> | <u>8.7222</u> | <u>1.3210</u> | <u>7.6134</u> |

Table 7: Performance comparison on direct recommendation.

| Methods    | Sports        |               |               |               | Beauty        |               |               |               | Toys          |               |               |               |               |
|------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|
|            | HR@1          | HR@5          | NDCC@5        | HR@10         | NDCC@10       | HR@1          | HR@5          | NDCC@5        | HR@10         | NDCC@10       | HR@1          | HR@5          | NDCC@10       |
| BPR-MF     | 0.0314        | 0.1404        | 0.0848        | 0.2563        | 0.1220        | 0.0311        | 0.1426        | 0.0857        | 0.2573        | 0.1224        | 0.0233        | 0.1066        | 0.0641        |
| BPR-MLP    | 0.0351        | 0.1520        | 0.0927        | 0.2671        | 0.1296        | 0.0317        | 0.1392        | 0.0848        | 0.2542        | 0.1215        | 0.0252        | 0.1142        | 0.0688        |
| SimpleX    | 0.0331        | <b>0.2362</b> | <b>0.1505</b> | 0.3290        | 0.1800        | 0.0325        | <b>0.2247</b> | 0.1441        | 0.3090        | 0.1711        | 0.0268        | <b>0.1958</b> | <b>0.1244</b> |
| P5-S (5-1) | 0.0638        | 0.2096        | 0.1375        | 0.3143        | 0.1711        | 0.0600        | 0.2021        | 0.1316        | 0.3121        | 0.1670        | 0.0405        | <b>0.1538</b> | 0.0699        |
| P5-B (5-1) | 0.0245        | 0.0816        | 0.0529        | 0.1384        | 0.0711        | 0.0224        | 0.0904        | 0.0559        | 0.1593        | 0.0780        | 0.0187        | 0.0827        | 0.0500        |
| P5-S (5-4) | 0.0701        | 0.2241        | 0.1483        | <b>0.3313</b> | <b>0.1827</b> | <b>0.0862</b> | <b>0.2448</b> | <b>0.1673</b> | <b>0.3441</b> | <b>0.1993</b> | 0.0413        | 0.1411        | 0.0916        |
| P5-B (5-4) | 0.0299        | 0.1026        | 0.0665        | 0.1708        | 0.0883        | 0.0506        | 0.1557        | 0.1033        | 0.2350        | 0.1287        | 0.0435        | 0.1316        | 0.0882        |
| P5-S (5-5) | 0.0574        | 0.1503        | 0.1050        | 0.2207        | 0.1276        | 0.0601        | 0.1611        | 0.1117        | 0.2370        | 0.1360        | 0.0440        | 0.1282        | 0.0865        |
| P5-B (5-5) | 0.0641        | 0.1794        | 0.1229        | 0.2598        | 0.1488        | 0.0588        | 0.1573        | 0.1089        | 0.2325        | 0.1330        | 0.0386        | 0.1122        | 0.0756        |
| P5-S (5-8) | 0.0567        | 0.1514        | 0.1049        | 0.2196        | 0.1269        | 0.0571        | 0.1566        | 0.1078        | 0.2317        | 0.1318        | <b>0.0451</b> | 0.1322        | 0.0889        |
| P5-B (5-8) | <b>0.0726</b> | 0.1955        | 0.1355        | 0.2802        | 0.1627        | 0.0608        | 0.1564        | 0.1096        | 0.2300        | 0.1332        | 0.0389        | 0.1147        | 0.0767        |

Table 3: Performance comparison on sequential recommendation.

| Methods             | Sports        |               |               |               | Beauty        |               |               |               | Toys          |               |               |               |
|---------------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|
|                     | HR@5          | NDCC@5        | HR@10         | NDCC@10       | HR@5          | NDCC@5        | HR@10         | NDCC@10       | HR@5          | NDCC@5        | HR@10         | NDCC@10       |
| Caser               | 0.0116        | 0.0072        | 0.0194        | 0.0097        | 0.0205        | 0.0131        | 0.0347        | 0.0176        | 0.0166        | 0.0107        | 0.0270        | 0.0141        |
| HGN                 | 0.0189        | 0.0120        | 0.0313        | 0.0159        | 0.0325        | 0.0206        | 0.0512        | 0.0266        | 0.0321        | 0.0221        | 0.0497        | 0.0277        |
| GRU4Rec             | 0.0129        | 0.0086        | 0.0204        | 0.0110        | 0.0164        | 0.0099        | 0.0283        | 0.0137        | 0.0097        | 0.0059        | 0.0176        | 0.0084        |
| BERT4Rec            | 0.0115        | 0.0075        | 0.0191        | 0.0099        | 0.0203        | 0.0124        | 0.0347        | 0.0170        | 0.0116        | 0.0071        | 0.0203        | 0.0099        |
| FDSA                | 0.0182        | 0.0122        | 0.0288        | 0.0156        | 0.0267        | 0.0163        | 0.0407        | 0.0208        | 0.0228        | 0.0140        | 0.0381        | 0.0189        |
| SASRec              | 0.0233        | 0.0154        | 0.0350        | 0.0192        | 0.0387        | 0.0249        | 0.0605        | 0.0318        | 0.0463        | 0.0306        | 0.0675        | 0.0374        |
| S <sup>3</sup> -Rec | 0.0251        | 0.0161        | 0.0385        | 0.0204        | 0.0387        | 0.0244        | 0.0647        | 0.0327        | 0.0443        | 0.0294        | 0.0700        | 0.0376        |
| P5-S (2-3)          | 0.0272        | 0.0169        | 0.0361        | 0.0198        | 0.0503        | 0.0370        | 0.0659        | 0.0421        | <b>0.0648</b> | <b>0.0567</b> | <b>0.0709</b> | <b>0.0587</b> |
| P5-S (2-3)          | 0.0364        | 0.0296        | 0.0431        | 0.0318        | <b>0.0508</b> | <b>0.0379</b> | <b>0.0664</b> | <b>0.0429</b> | 0.0608        | 0.0507        | 0.0688        | 0.0534        |
| P5-S (2-13)         | 0.0258        | 0.0159        | 0.0346        | 0.0188        | 0.0490        | 0.0358        | 0.0646        | 0.0409        | 0.0647        | 0.0566        | 0.0705        | 0.0585        |
| P5-B (2-13)         | <b>0.0387</b> | <b>0.0312</b> | <b>0.0460</b> | <b>0.0336</b> | 0.0493        | 0.0367        | 0.0645        | 0.0416        | 0.0587        | 0.0486        | 0.0675        | 0.0536        |

Table 4: Performance comparison on explanation generation (%).

| Methods     | Sports        |                |               |                | Beauty         |                |                |                | Toys           |                |                |                |
|-------------|---------------|----------------|---------------|----------------|----------------|----------------|----------------|----------------|----------------|----------------|----------------|----------------|
|             | BLUE4         | ROUGE1         | ROUGE2        | ROUGEL         | BLUE4          | ROUGE1         | ROUGE2         | ROUGEL         | BLUE4          | ROUGE1         | ROUGE2         | ROUGEL         |
| Attr2Seq    | 0.5305        | 12.2800        | 1.2107        | 9.1312         | 0.7888         | 12.6590        | 1.6820         | 9.7481         | 1.6238         | 13.2245        | 2.9942         | 10.7398        |
| NRT         | 0.4793        | 11.0723        | 1.1300        | 7.6674         | 0.8299         | 12.7815        | 1.8543         | 9.9477         | 1.9084         | 13.5231        | 3.6708         | 11.1867        |
| PETER       | 0.7112        | 12.8944        | 1.3283        | 9.8635         | 1.1541         | 14.8497        | 2.1413         | 11.4143        | 1.9861         | 14.2716        | 3.6718         | 11.7010        |
| P5-S (3-3)  | <b>1.0447</b> | <b>14.9048</b> | <b>2.1297</b> | <b>11.1778</b> | <b>17.6930</b> | <b>2.2489</b>  | <b>12.8600</b> | <b>2.2892</b>  | <b>15.4505</b> | <b>3.6974</b>  | <b>12.1718</b> |                |
| P5-B (3-3)  | <b>1.0407</b> | 14.1589        | <b>2.1220</b> | 10.6996        | 0.9749         | 16.4530        | 1.8858         | 11.8765        | <b>2.3185</b>  | <b>15.3474</b> | <b>3.7209</b>  | <b>12.1312</b> |
| PETER+      | <b>2.4627</b> | <b>24.1181</b> | 5.1937        | <b>18.4105</b> | <b>3.2606</b>  | 25.5541        | 5.9668         | 19.7168        | <b>4.7919</b>  | 28.3083        | 9.4520         | <b>22.7017</b> |
| P5-S (3-9)  | 1.4101        | 23.5619        | <b>5.4196</b> | 17.6245        | 1.9788         | <b>25.6253</b> | <b>6.3678</b>  | <b>19.9497</b> | 4.1222         | <b>28.4088</b> | 9.5432         | 22.6064        |
| P5-S (3-9)  | <b>1.4689</b> | 23.5476        | 5.3926        | 17.5852        | 1.8765         | 25.1183        | 6.0764         | 19.4488        | 3.8933         | 27.9916        | 9.5896         | 22.2178        |
| P5-S (3-12) | 1.3212        | 23.5476        | 5.3461        | 17.3780        | 1.9425         | 25.0740        | 6.0551         | 19.5601        | 4.2764         | 28.1897        | 9.1327         | 22.2314        |
| P5-B (3-12) | 1.4303        | 23.3810        | 5.3239        | 17.4913        | 1.9031         | 25.1763        | 6.1980         | 19.5188        | 3.5861         | 28.1369        | <b>9.7562</b>  | 22.3056        |

Single LLM → Effective on various recommendation tasks

# Aligning LLMs for recommendation

## (1) Pure text-based (P5)

**Multi-scenario Recommendation:** The items the user has recently clicked on are as follows: {USER BEHAVIOR SEQUENCE}. In scenario {SCENE}, please recommend items.

**Multi-objective Recommendation:** The items the user has recently clicked on are as follows: {USER BEHAVIOR SEQUENCE}. Please find items that the user will {ACTION}.

**Long-tail Item Recommendation:** The items the user has recently clicked on are as follows: {USER BEHAVIOR SEQUENCE}. Please recommend long-tail items.

**Serendipity Recommendation:** The items the user has recently clicked on are as follows: {USER BEHAVIOR SEQUENCE}. Please recommend some new item categories.

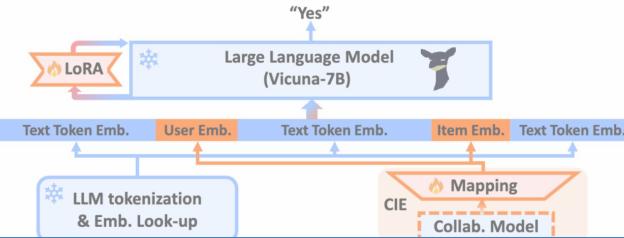
**Long-term Recommendation:** The items the user has recently clicked on are as follows: {USER BEHAVIOR SEQUENCE}. Please find items that match the user's long-term interests.

**Search Problem:** The items the user has recently clicked on are as follows: {USER BEHAVIOR SEQUENCE}. Please recommend items that match {QUERY}.

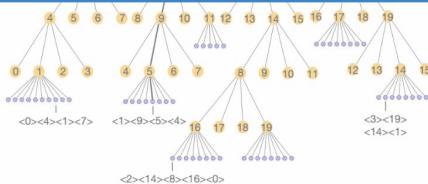
URM:

Unify recommendation & search

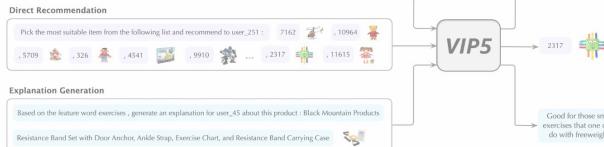
# Aligning LLMs for recommendation



Is textual information enough for alignment?



+ External item tokens

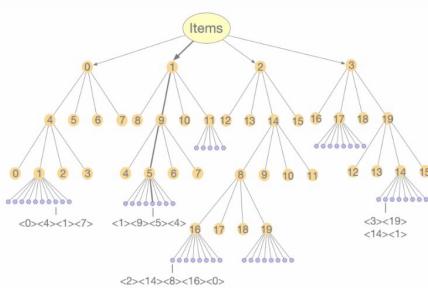


+ Multimodal information

# Aligning LLMs for recommendation



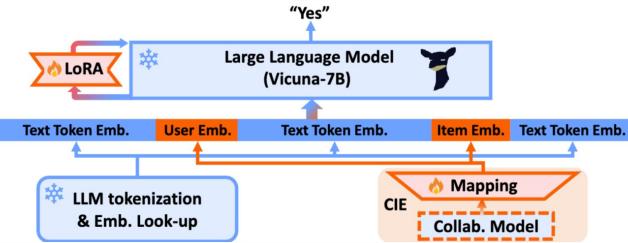
Pure text-based



+ External item tokens



+ Collaborative embeddings



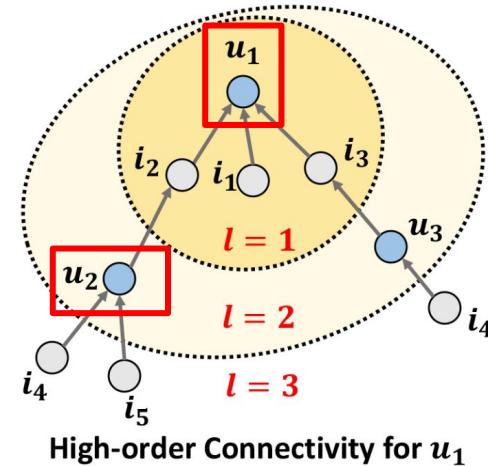
+ Multimodal information

# Aligning LLMs for recommendation

## (2) + Collaborative embeddings

### Motivation:

Language modeling may not capture **collaborative** information

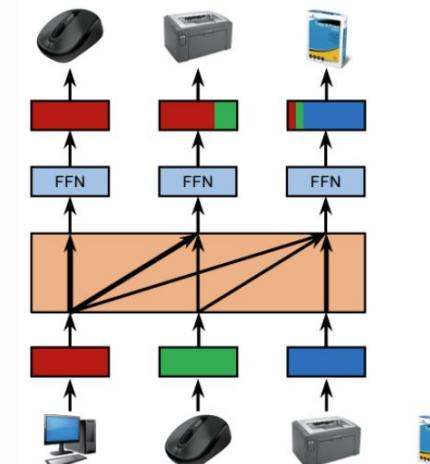


# Aligning LLMs for recommendation

## (2) + Collaborative embeddings

### Solution:

Aligning LLMs with  
embeddings from  
traditional recommenders

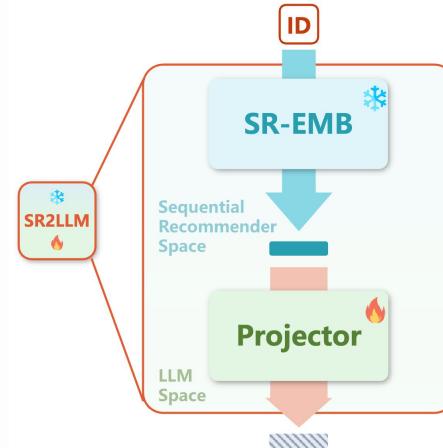


# Aligning LLMs for recommendation

## (2) + Collaborative embeddings (LLaRA)

+ Pretrained item embeddings

| (a) Text-only prompting method.   | (b) Hybrid prompting method.  |
|---|---|
| <b>Input:</b> This user has watched Titanic [PH], Roman Holiday [PH], .... Gone with the wind [PH] in the previous. Please predict the next movie this user will watch. The movie title candidates are The Wizard of Oz [PH], Braveheart [PH],..., Waterloo Bridge [PH],... Batman & Robin [PH]. Choose only one movie from the candidates. The answer is | <b>Input:</b> This user has watched Titanic [ $\text{emb}_s^{14}$ ], Roman Holiday [ $\text{emb}_s^{20}$ ], .... Gone with the wind [ $\text{emb}_s^{37}$ ] in the previous. Please predict the next movie this user will watch. The movie title candidates are The Wizard of Oz [ $\text{emb}_s^5$ ], Braveheart [ $\text{emb}_s^{42}$ ],..., Waterloo Bridge [ $\text{emb}_s^{20}$ ],... Batman & Robin [ $\text{emb}_s^{19}$ ]. Choose only one movie from the candidates. The answer is |
| <b>Output:</b> Waterloo Bridge.   | <b>Output:</b> Waterloo Bridge.   |



# Aligning LLMs for recommendation

## (2) + Collaborative embeddings (LLaRA)

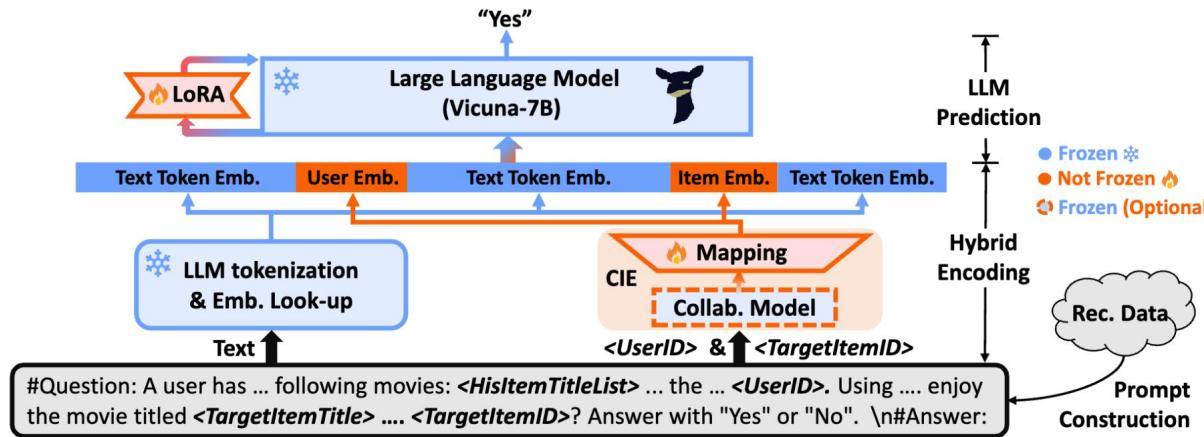
- + Pretrained item embeddings



# Aligning LLMs for recommendation

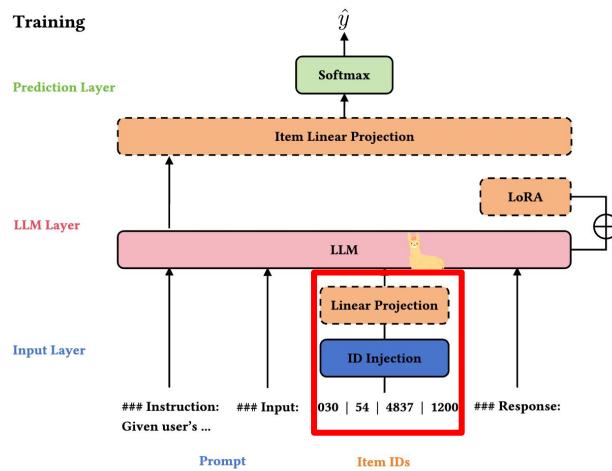
## (2) + Collaborative embeddings (CoLLM)

- + Pretrained item embeddings + user embeddings



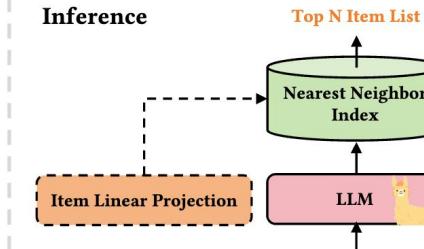
# Aligning LLMs for recommendation

## (2) + Collaborative embeddings (E4SRec)



Discard text;

Collaborative embeddings only



KNN for inference

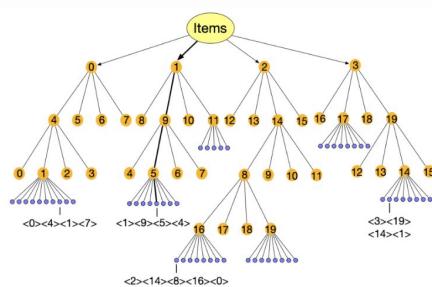
# Aligning LLMs for recommendation



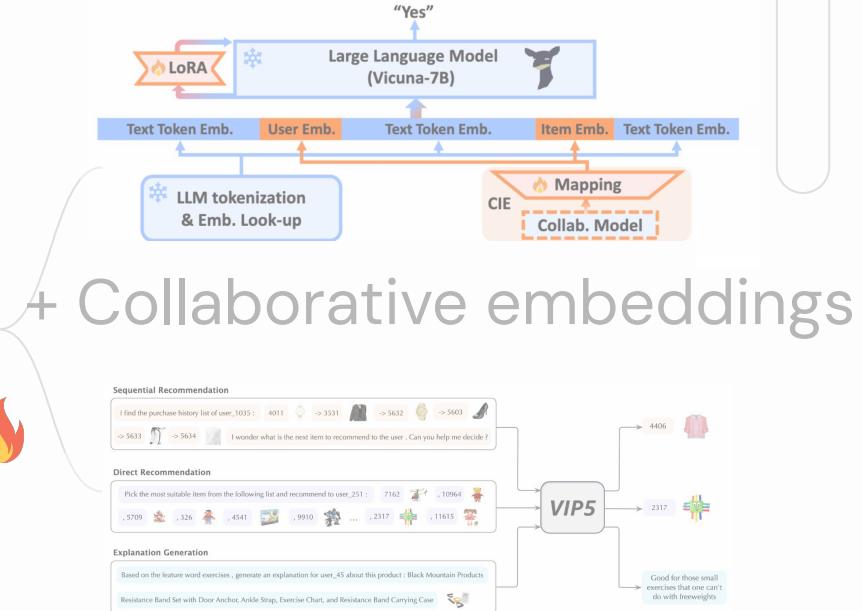
**LLM-based Recommender**

Waterloo Bridge.

Pure text-based



+ External item tokens



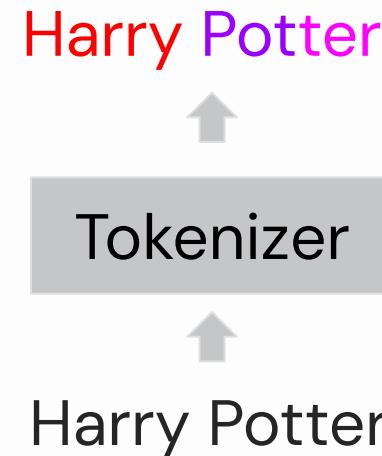
+ Multimodal information

# Aligning LLMs for recommendation

## (3) + External item tokens

### Motivation:

Tokens for language modeling are **not optimal** for recommendation.



# Aligning LLMs for recommendation

## (3) + External item tokens

### Motivation:

Tokens for language modeling are **not optimal** for recommendation.

Maybe better?

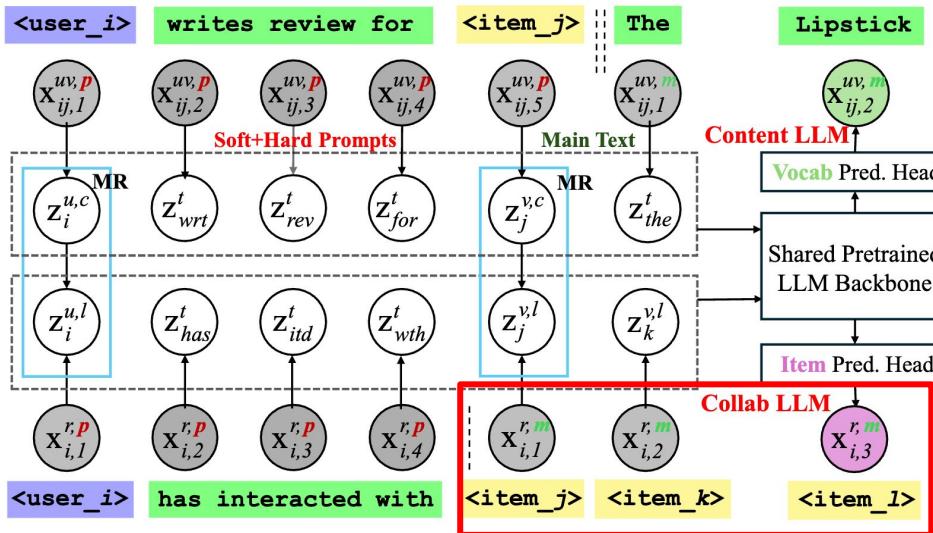
Harry Potter

Tokenizer

Harry Potter

# Aligning LLMs for recommendation

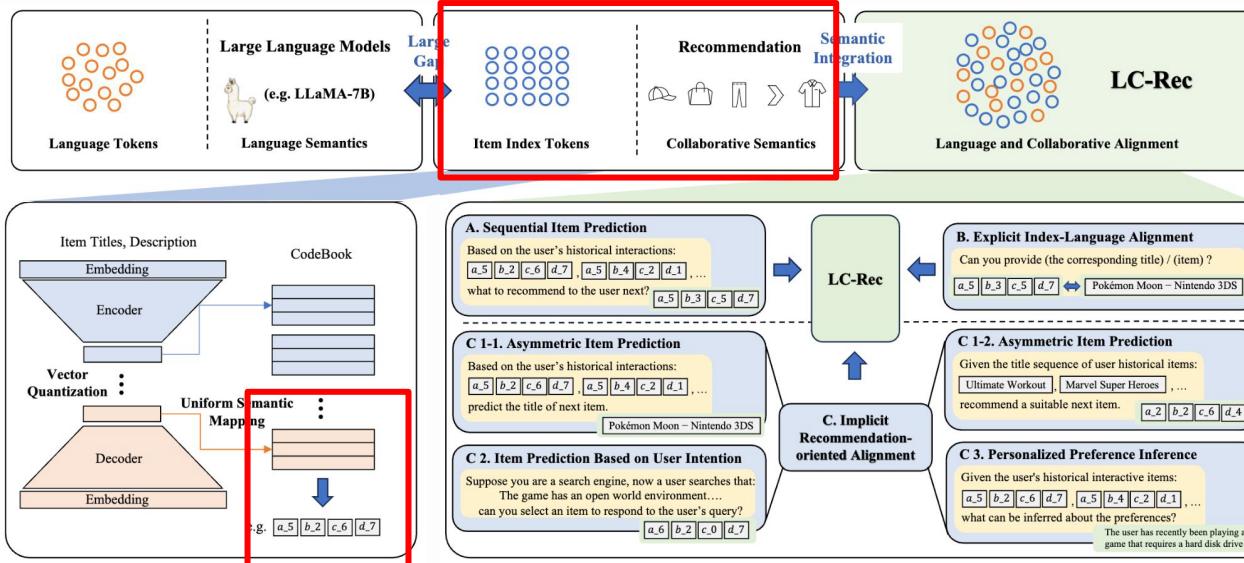
## (3) + External item tokens (CLLM4Rec)



Naive approach:  
One ID for each item

# Aligning LLMs for recommendation

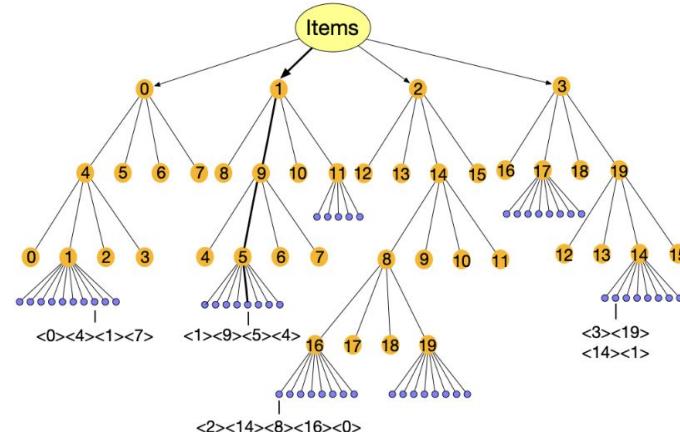
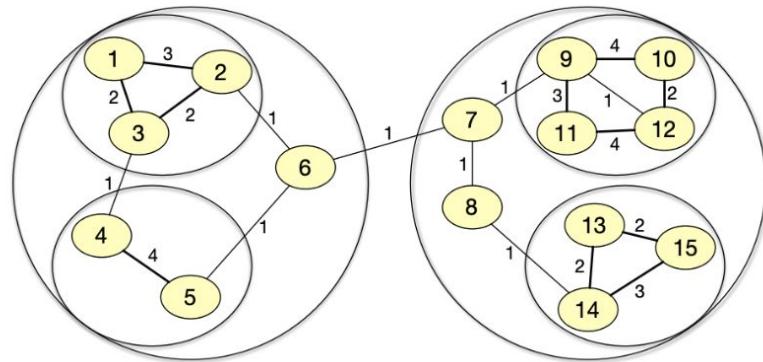
## (3) + External item tokens (LC-Rec)



+ Semantic IDs  
(Similar items have similar IDs)

# Aligning LLMs for recommendation

## (3) + External item tokens

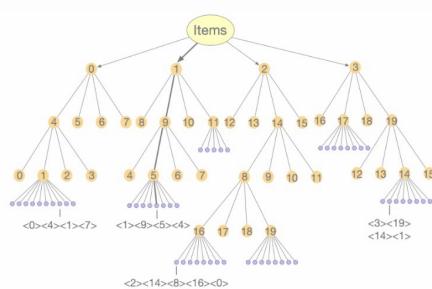


More complicated item tokens design

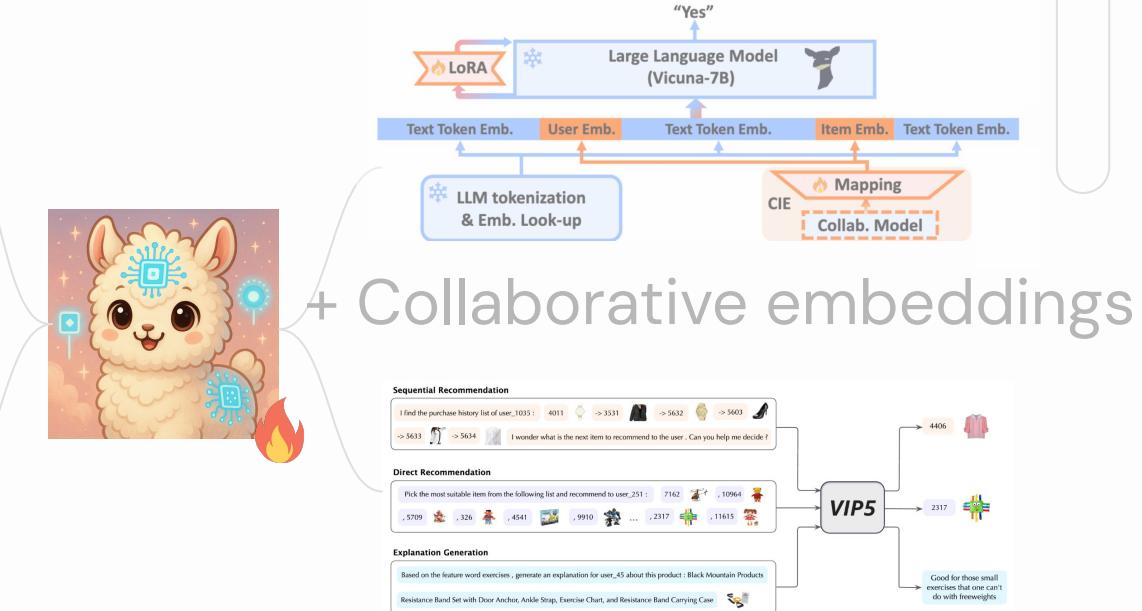
# Aligning LLMs for recommendation



Pure text-based



+ External item tokens



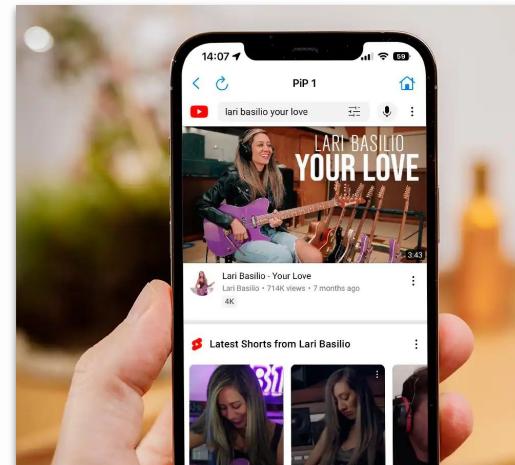
+ Multimodal information

# Aligning LLMs for recommendation

## (4) + Multimodal information

### Motivation:

Human make decisions with  
multimodal information.

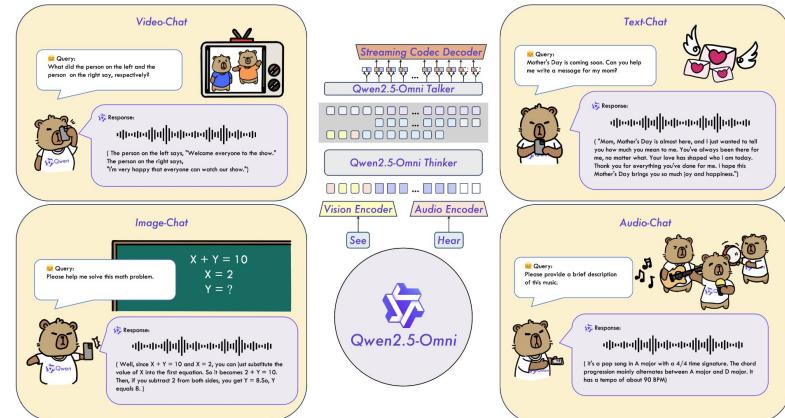


# Aligning LLMs for recommendation

## (4) + Multimodal information

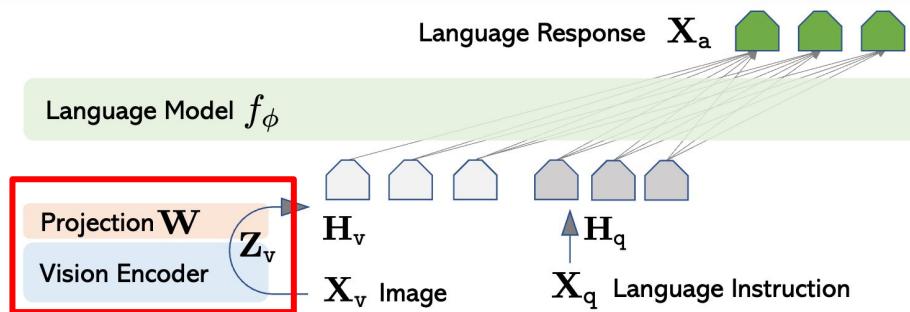
### Motivation:

Post-trained LLM can understand **multimodal** information



# Aligning LLMs for recommendation

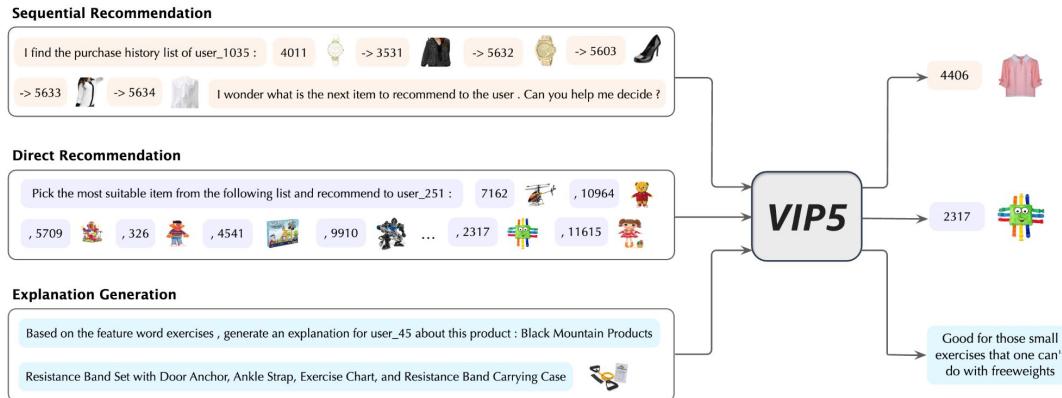
## (4) + Multimodal information



Aligning vision and  
language with a projector

# Aligning LLMs for recommendation

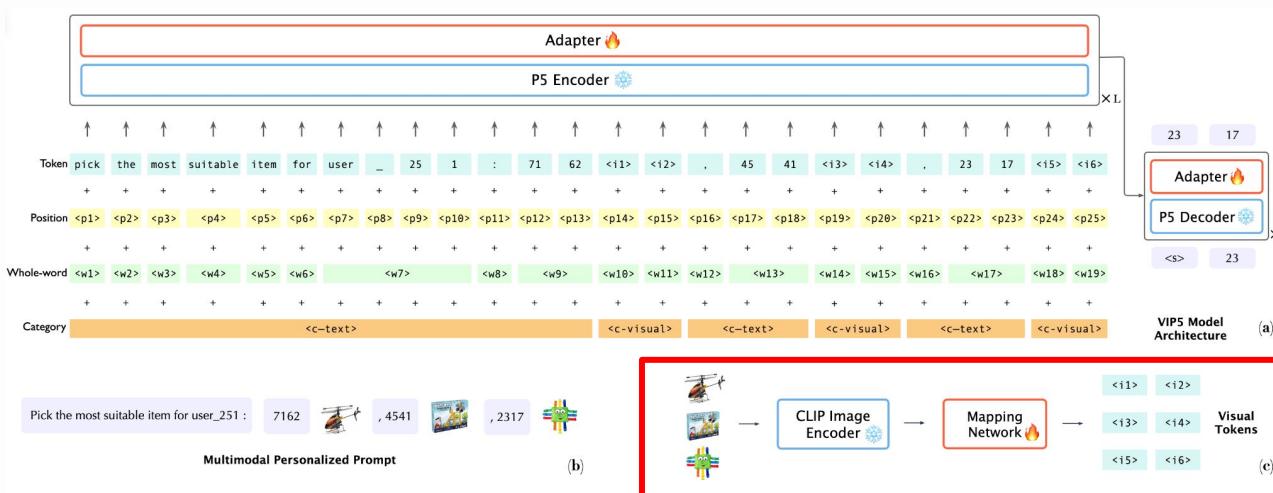
## (4) + Multimodal information (VIP5)



Diff between P5:  
Pair text with its image

# Aligning LLMs for recommendation

## (4) + Multimodal information (VIP5)



Alignment with projector

# Aligning LLMs for recommendation

## (4) + Multimodal information (VIP5)

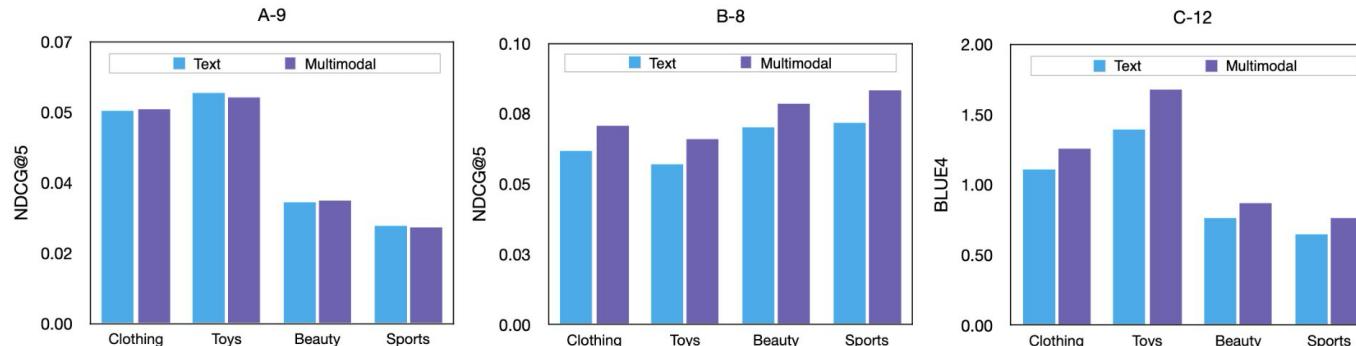
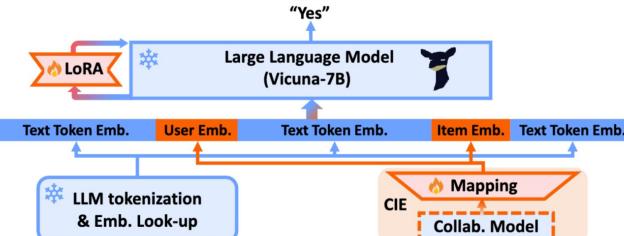


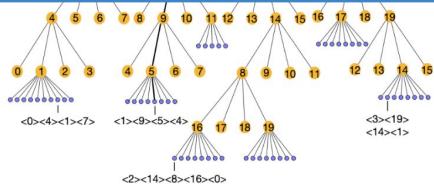
Figure 3: Performance comparison between text-based prompt and multimodal prompt.

Multimodal information is important

# Aligning LLMs for recommendation



Information tailored for recommendation matters!



+ External item tokens

+ Multimodal information



# Part 1: LLM as Sequential Recommender

- (i) Early efforts: Pretrained LLMs for recommendation;
- (ii) Aligning LLMs for recommendation;
- (iii) **Training objective & inference**

# Training objective

## (1) Supervised finetuning (SFT)

I have watched Titanic, Roman Holiday, ... Gone with the wind. Predict the next movie I will watch:

# Training objective

## (1) Supervised finetuning (SFT)

I have watched Titanic, Roman Holiday, ... Gone with the wind. Predict the next movie I will watch:

Waterloo Bridge.



Prediction

# Training objective

## (1) Supervised finetuning (SFT)

I have watched Titanic, Roman Holiday, ... Gone with the wind. Predict the next movie I will watch:  
**Waterloo Bridge.**



Prediction

# Training objective

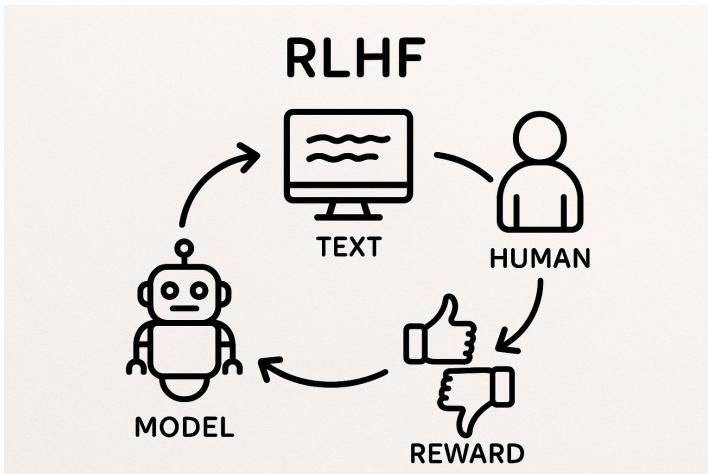
## (1) Supervised finetuning (SFT)

$$\mathcal{L}_{\text{SFT}}(\theta) = -\mathbb{E}_{(x,y) \sim \mathcal{D}} \left[ \sum_{t=1}^T \log P_\theta(y_t \mid y_{<t}) \right]$$

Always predict the next token

# Training objective

## (2) Preference learning



LLMs are trained to align  
human preferences

Recommendation is about  
user preferences

# Training objective

## (2) Preference learning

I have watched Titanic, Roman Holiday, ... Gone with the wind. Predict the next movie I will watch:



Waterloo Bridge



Harry Potter



# Training objective

## (2) Preference learning

$$\mathcal{L}_{\text{DPO}} = -\mathbb{E}_{(x_u, e_p, e_d)} \left[ \log \sigma \left( \beta \log \frac{\pi_\theta(e_p|x_u)}{\pi_{\text{ref}}(e_p|x_u)} - \beta \log \frac{\pi_\theta(e_d|x_u)}{\pi_{\text{ref}}(e_d|x_u)} \right) \right],$$

**Direct Preference Optimization!**

I have watched Titanic, Roman Holiday, ... Gone with the wind. Predict the next movie I will watch:



Waterloo Bridge

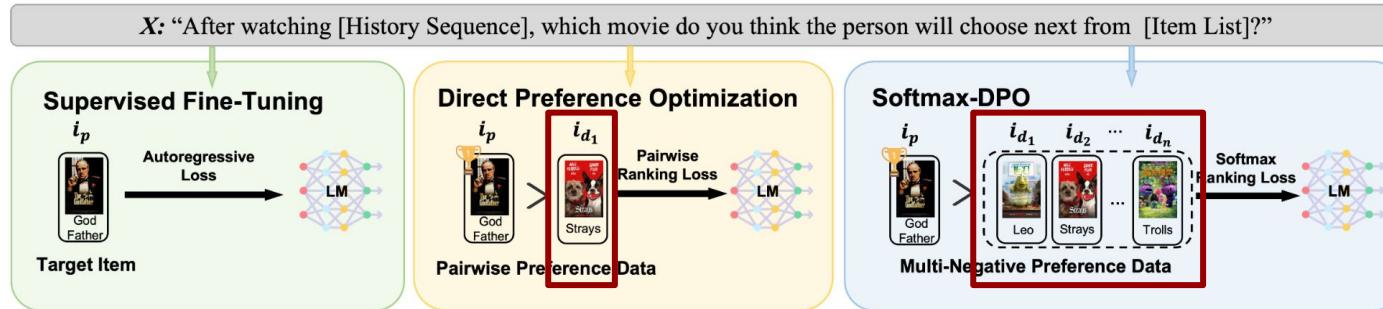


Harry Potter



# Training objective

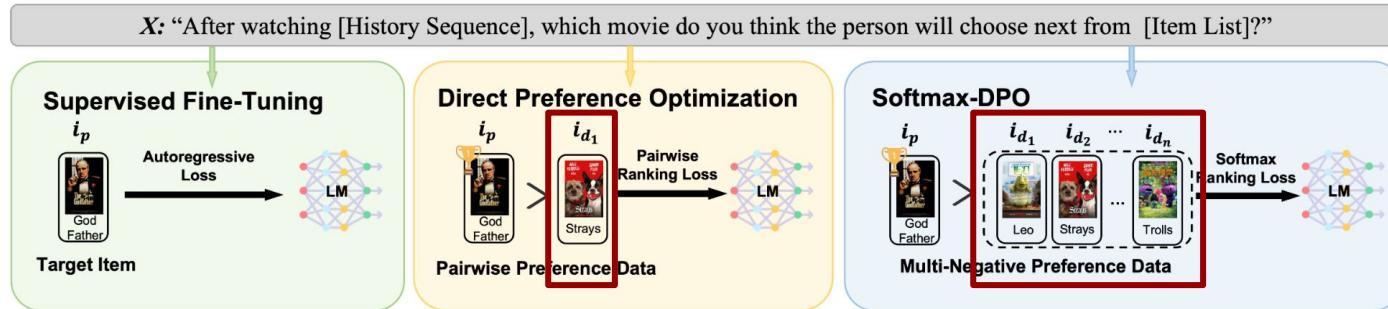
## (2) Preference learning



Single negative → Multiple negatives

# Training objective

## (2) Preference learning



$$\mathcal{L}_{\text{S-DPO}}(\pi_\theta; \pi_{\text{ref}}) = -\mathbb{E}_{(x_u, e_p, \mathcal{E}_d) \sim \mathcal{D}} \left[ \log \sigma \left( -\log \sum_{e_d \in \mathcal{E}_d} \exp \left( \beta \log \frac{\pi_\theta(e_d|x_u)}{\pi_{\text{ref}}(e_d|x_u)} - \beta \log \frac{\pi_\theta(e_p|x_u)}{\pi_{\text{ref}}(e_p|x_u)} \right) \right) \right].$$

# Training objective

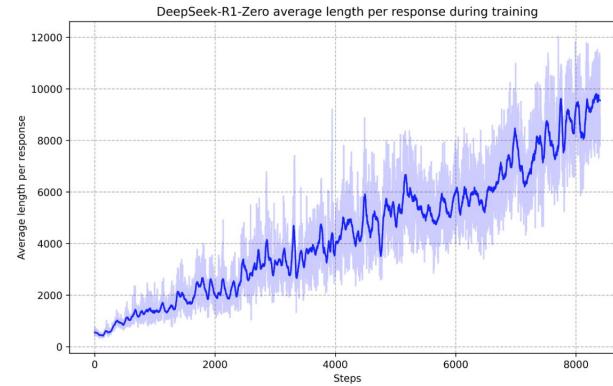
## (3) Reinforce learning

Emergent reasoning capabilities through RL

$$\mathcal{J}_{GRPO}(\theta) = \mathbb{E}[q \sim P(Q), \{o_i\}_{i=1}^G \sim \pi_{\theta_{old}}(O|q)]$$

$$\frac{1}{G} \sum_{i=1}^G \left( \min \left( \frac{\pi_\theta(o_i|q)}{\pi_{\theta_{old}}(o_i|q)} A_i, \text{clip} \left( \frac{\pi_\theta(o_i|q)}{\pi_{\theta_{old}}(o_i|q)}, 1-\epsilon, 1+\epsilon \right) A_i \right) - \beta \mathbb{D}_{KL}(\pi_\theta || \pi_{ref}) \right), \quad (1)$$

$$\mathbb{D}_{KL}(\pi_\theta || \pi_{ref}) = \frac{\pi_{ref}(o_i|q)}{\pi_\theta(o_i|q)} - \log \frac{\pi_{ref}(o_i|q)}{\pi_\theta(o_i|q)} - 1, \quad (2)$$



---

Question: If  $a > 1$ , then the sum of the real solutions of  $\sqrt{a - \sqrt{a+x}} = x$  is equal to

Response: <think>

To solve the equation  $\sqrt{a - \sqrt{a+x}} = x$ , let's start by squaring both  $\dots$

$$(\sqrt{a - \sqrt{a+x}})^2 = x^2 \implies a - \sqrt{a+x} = x^2.$$

Rearrange to isolate the inner square root term:

$$(a - x^2)^2 = a + x \implies a^2 - 2ax^2 + (x^2)^2 = a + x \implies x^4 - 2ax^2 - x + (a^2 - a) = 0$$

$\dots$

Wait, wait. Wait. That's an aha moment I can flag here.

Let's reevaluate this step-by-step to identify if the correct sum can be  $\dots$

We started with the equation:

$$\sqrt{a - \sqrt{a+x}} = x$$

First, let's square both sides:

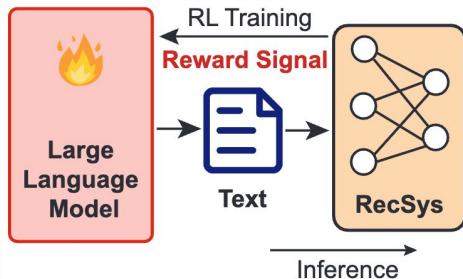
$$a - \sqrt{a+x} = x^2 \implies \sqrt{a+x} = a - x^2$$

Next, I could square both sides again, treating the equation:  $\dots$

$\dots$

# Training objective

## (3) Reinforce learning



$$\max_{\theta} \mathbb{E}_{s \sim p(s), a \sim \pi_{\theta}(a|s)} [f(a|s)].$$

Maximize the reward from  
recommender system

---

### Prompt Template for REC-R1 + Dense Retriever (Product Search)

You are an expert in generating queries for dense retrieval. Given a customer query, your task is to retain the original query while expanding it with additional semantically relevant information, retrieve the most relevant products, ensuring they best meet customer needs. If no useful expansion is needed, return the original query as is.

Below is the query:

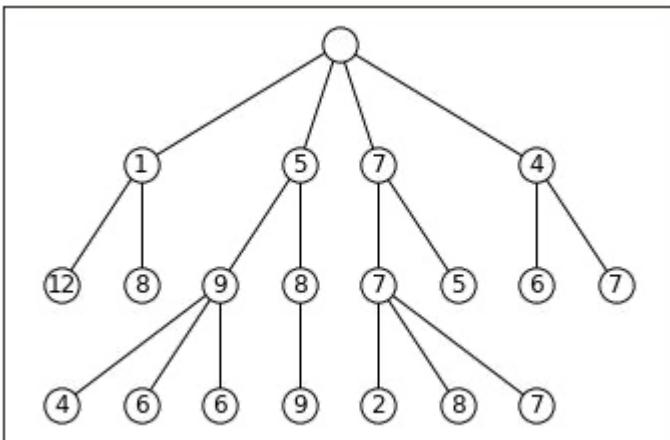
```
``` {user_query} ```  
<|im_start|>system  
You are a helpful AI assistant. You first think about the reasoning process  
in the mind and then provide the user with the answer.  
<|im_end|>  
<|im_start|>user  
[PROMPT as above]  
Show your work in <think>\> tags. Your final response must be in JSON  
format within <answer>\>tags. For example,
```

```
<answer>  
{ "query": xxx }  
</answer>
```

```
<|im_end|>  
<|im_start|>assistant  
Let me solve this step by step.  
<think>
```

# Inference

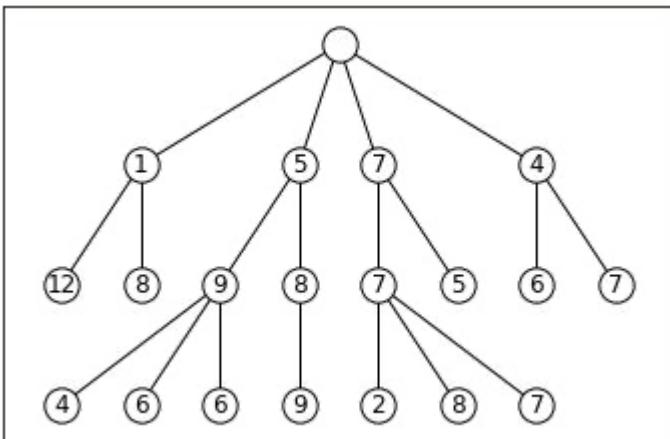
## (1) Beam Search



Generating answers with the top-k highest scored beams

# Inference

## (1) Beam Search



It may generate **invalid items**

In RecSys :  
No Hallucination permitted!

# Inference

## (2) Constrained Beam Search

**Valid items:**

Waterloo Bridge, Waterfall  
Story, and Waterloo War

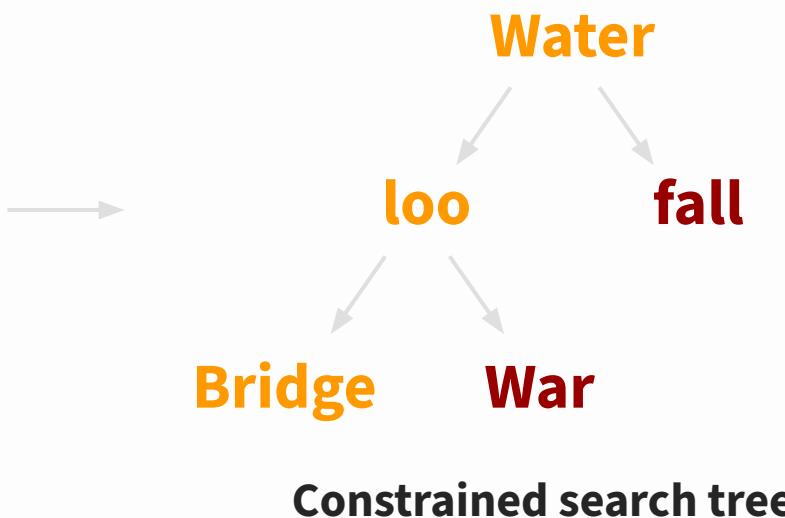
**How to make the generated  
items always valid?**

# Inference

## (2) Constrained Beam Search

**Valid items:**

Waterloo Bridge, Waterfall  
Story, and Waterloo War



# Inference

## (2) Constrained Beam Search

I have watched **Titanic**, **Roman Holiday**, ...

**Gone with the wind.** Predict the next movie

I will watch:

$P = 1$   
→

**Water**

# Inference

## (2) Constrained Beam Search

I have watched **Titanic**, **Roman Holiday**, ...

**Gone with the wind.** Predict the next movie

I will watch:



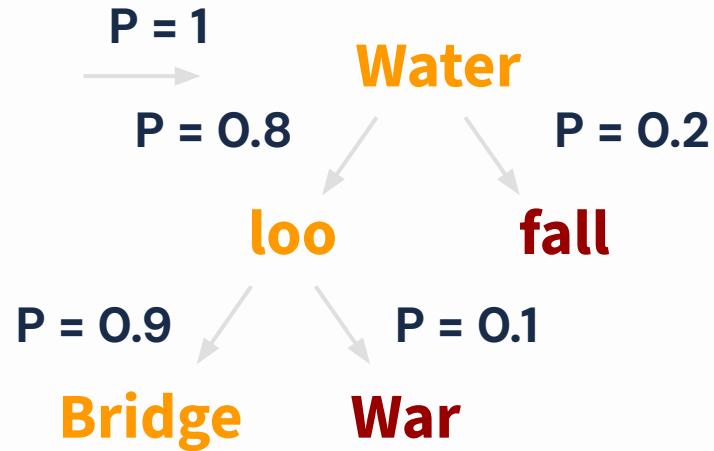
# Inference

## (2) Constrained Beam Search

I have watched **Titanic**, **Roman Holiday**, ...

**Gone with the wind.** Predict the next movie

I will watch:

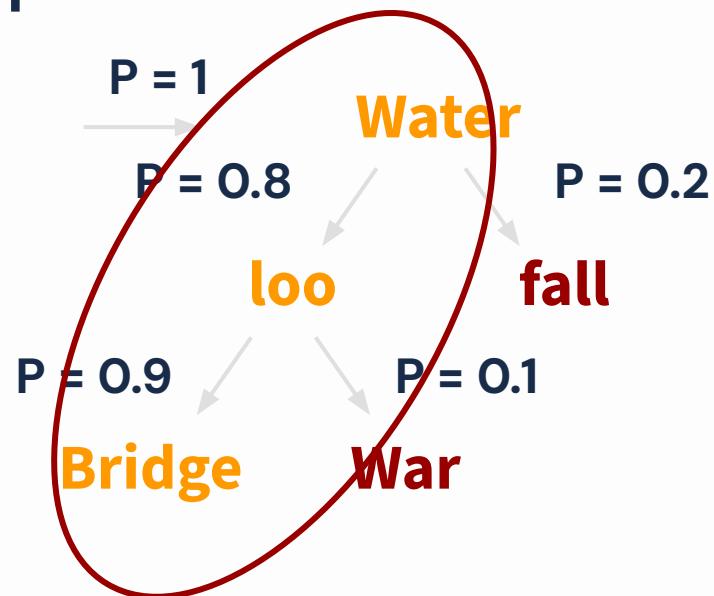


# Inference

## (2) Constrained Beam Search

I have watched **Titanic**, **Roman Holiday**, ...  
**Gone with the wind**. Predict the next movie  
I will watch:

Valid Item!



# Inference

## (3) Special design

$$\mathcal{S}(h_{\leq t}) = \mathcal{S}(h_{\leq t-1}) + \log(p(h_t|x, h_{\leq t-1})),$$

$$\mathcal{S}(h) = \mathcal{S}(h) / h_L^\alpha,$$

**Length penalty in beam search;  
Human does not like over long sentences.**

**Redundant for recommendation**

# Inference

## (3) Special design

$$\mathcal{S}(h_{\leq t}) = \mathcal{S}(h_{\leq t-1}) + \log(p(h_t|x, h_{\leq t-1})),$$

$\mathcal{S}(h) = \mathcal{S}(h)/\cancel{\lambda}$ , Remove length penalty

|          | Instruments   | Books         | CDs           | Sports        | Toys          | Games         |
|----------|---------------|---------------|---------------|---------------|---------------|---------------|
| Baseline | 0.1062        | 0.0308        | 0.0956        | 0.1171        | 0.0965        | 0.0610        |
| $D^3$    | <b>0.1111</b> | <b>0.0354</b> | <b>0.1190</b> | <b>0.1215</b> | <b>0.1025</b> | <b>0.0767</b> |
| - RLN    | 0.1093        | 0.0353        | 0.1000        | 0.1200        | 0.0975        | 0.0659        |
| - TFA    | 0.1086        | 0.0309        | 0.1115        | 0.1192        | 0.1006        | 0.0732        |

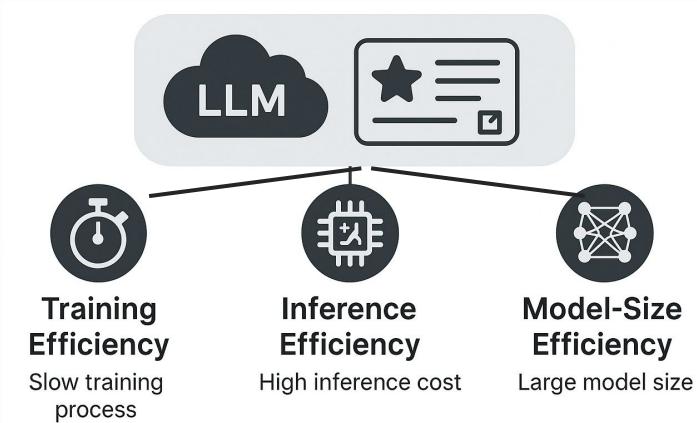
Imp when removing

# Part 1: LLM as Sequential Recommender

- (i) Early efforts: Pretrained LLMs for recommendation;
- (ii) Aligning LLMs for recommendation;
- (iii) Training objective & inference
- (iv) Efficiency**

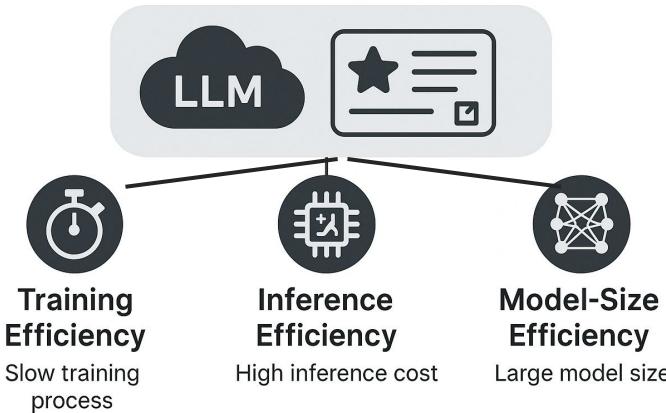
# Efficiency

## A crucial question in real-world deployment



# Efficiency

A crucial question in real-world deployment



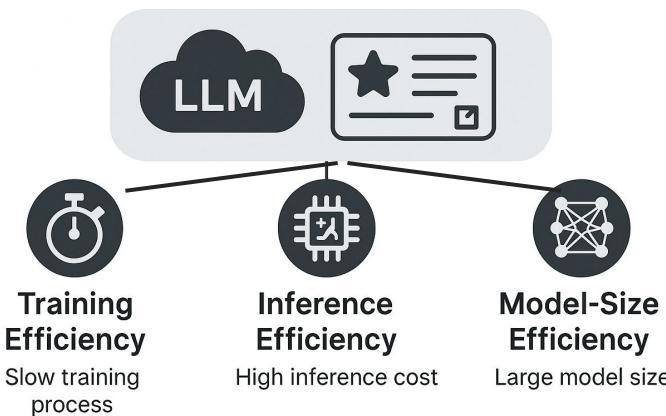
**Training efficiency:**

LLM: update by **months**

Recommender: update by **hours**

# Efficiency

A crucial question in real-world deployment



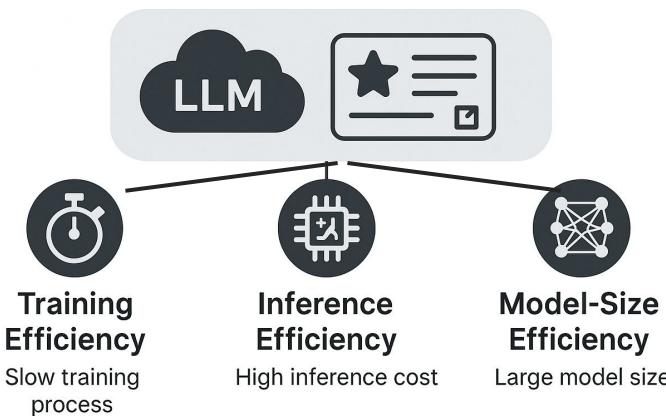
**Inference efficiency:**

LLM: wait for **seconds**

Recommender: wait for **milliseconds**

# Efficiency

A crucial question in real-world deployment



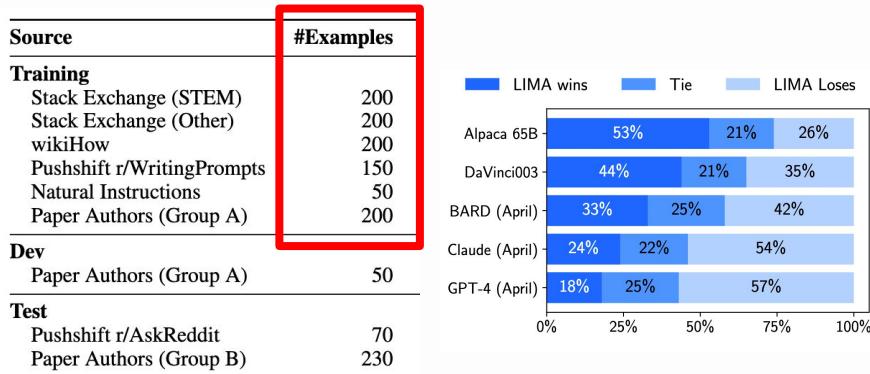
**Model-size efficiency:**

LLM: serve for **millions**

Recommender: serve for **billions**

# Efficiency

## (1) Training efficiency



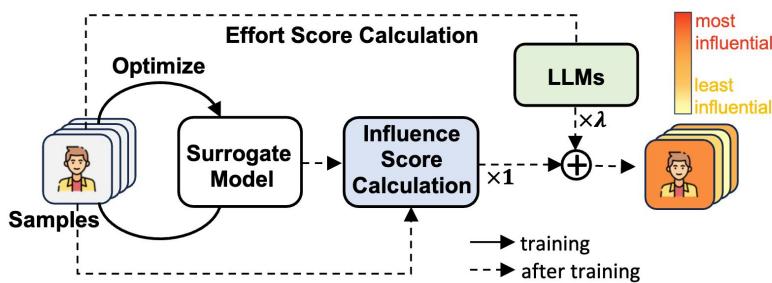
Less is more for alignment

1k high quality examples ->

Surpass large scale training

# Efficiency

## (1) Training efficiency



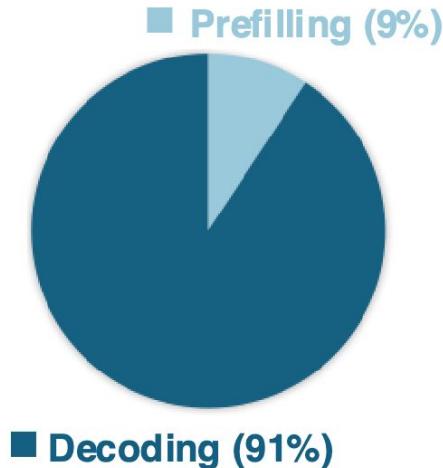
Select the most informative examples ->

Reducing 95% training time

|            | Games  |        |        |        |         | Time↓ |
|------------|--------|--------|--------|--------|---------|-------|
|            | R@10↑  | R@20↑  | N@10↑  | N@20↑  |         |       |
| Full       | 0.0169 | 0.0233 | 0.0102 | 0.0120 | 36.87h  |       |
| DEALRec    | 0.0181 | 0.0276 | 0.0115 | 0.0142 | 1.67h   |       |
| % Improve. | 7.10%  | 18.45% | 12.75% | 18.33% | -95.47% |       |

# Efficiency

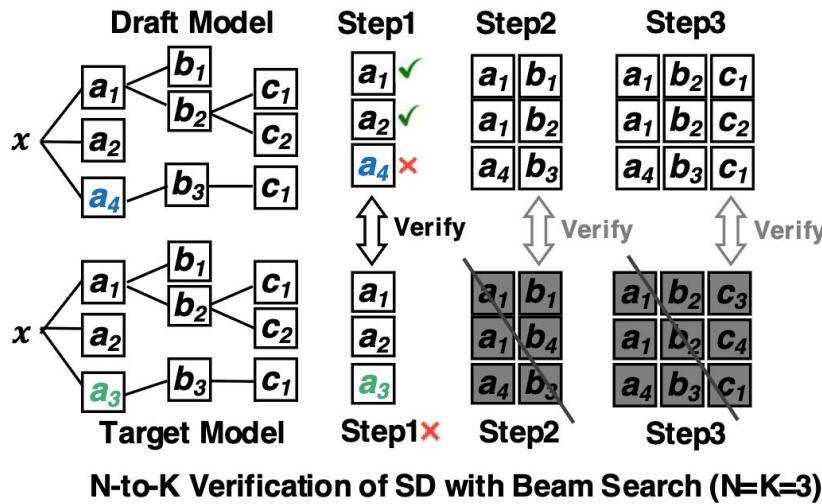
## (2) Inference efficiency



Autoregressive paradigm in LLM  
-> huge time on the **decoding stage**

# Efficiency

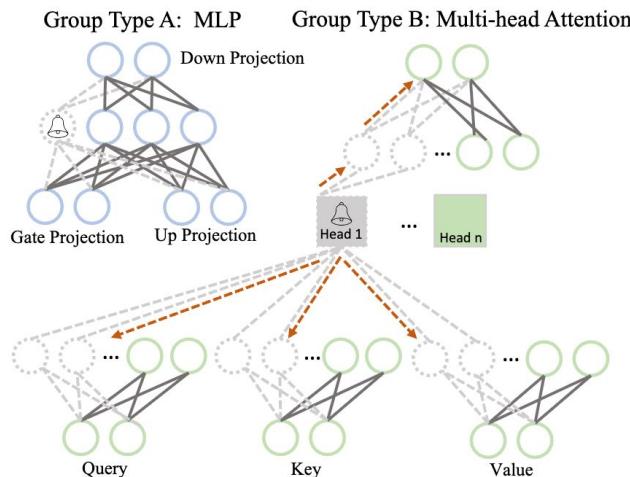
## (2) Inference efficiency



Speculative decoding:  
Decoder acceleration with a  
small-size draft model

# Efficiency

## (3) Model-size efficiency – Pruning

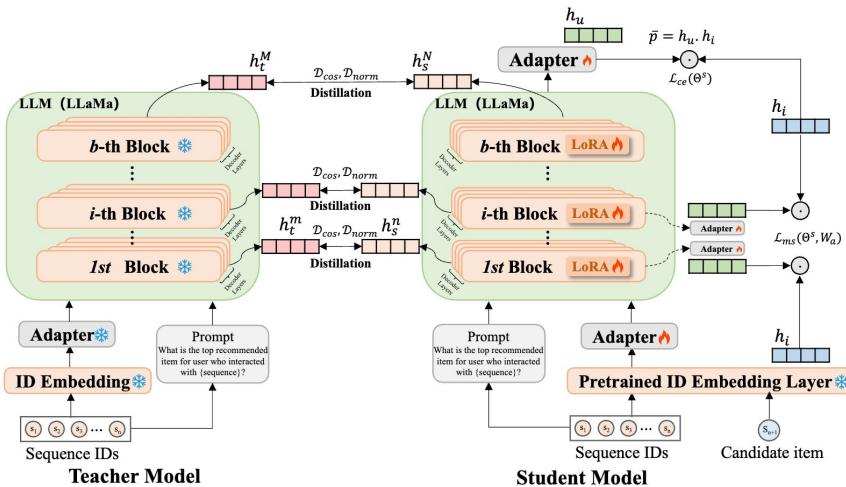


Similar performance with  
0.6% parameters

| Method          | #Params | Tasks        |              |              |
|-----------------|---------|--------------|--------------|--------------|
|                 |         | TNEWS↑       | IFLYTEK↑     | CSL↑         |
| M6-base         | 327M    | 0.598        | 0.631        | 0.852        |
| ALBERT-zh-base  | 12M     | 0.550        | 0.564        | 0.785        |
| M6-Edge         | 10M     | <b>0.552</b> | <b>0.586</b> | <b>0.831</b> |
| ALBERT-zh-tiny  | 4M      | 0.534        | 0.488        | 0.750        |
| M6-Edge, Pruned | 2M      | <b>0.537</b> | <b>0.559</b> | <b>0.798</b> |

# Efficiency

## (3) Model-size efficiency – Distillation



SLM learns from LLM  
With Hard label + soft label

# Efficiency

## (3) Model-size efficiency – Distillation

Table 3: Experimental results (%) on the Music and Sport dataset.

| Model                 | Music       |             |             |             | Sport       |             |             |             | Rank        |
|-----------------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
|                       | HR@1        | HR@5        | NDCG@5      | MRR         | HR@1        | HR@5        | NDCG@5      | MRR         |             |
| Caser                 | 0.71        | 3.28        | 1.96        | 2.29        | 1.05        | 3.75        | 2.39        | 2.84        | 13.50       |
| GRU4Rec               | 1.89        | 3.22        | 2.57        | 3.08        | 5.26        | 7.75        | 6.52        | 7.08        | 10.13       |
| BERT4Rec              | 2.10        | 3.16        | 2.64        | 3.11        | 4.81        | 6.70        | 5.79        | 6.26        | 10.63       |
| SASRec                | 1.82        | 5.72        | 3.79        | 4.51        | 4.70        | 8.43        | 6.59        | 7.24        | 8.75        |
| HGN                   | 2.01        | 5.49        | 3.82        | 4.17        | 3.42        | 6.24        | 4.83        | 5.30        | 10.50       |
| LightSANs             | 1.05        | 4.06        | 2.54        | 3.00        | 5.18        | 8.94        | 7.07        | 7.72        | 8.25        |
| S <sup>3</sup> -Rec   | 2.48        | 7.37        | 4.94        | 4.68        | 4.14        | 8.49        | 6.89        | 7.35        | 6.88        |
| DuoRec                | 1.84        | 4.50        | 3.19        | 3.04        | 4.13        | 8.81        | 7.03        | 6.64        | 9.13        |
| MAERec                | 2.19        | 6.35        | 4.67        | 3.96        | 4.01        | 8.35        | 6.65        | 6.98        | 8.63        |
| Open-P5               | 4.35        | 8.12        | 6.74        | -           | 5.49        | 8.50        | 6.92        | -           | 5.33        |
| E4SRec                | <b>5.62</b> | <b>9.29</b> | <b>7.50</b> | <b>7.98</b> | <b>6.40</b> | <b>9.67</b> | <b>8.05</b> | <b>8.70</b> | <b>1.75</b> |
| E4SRec8               | 5.46        | 8.86        | 7.21        | 7.74        | 5.48        | 8.63        | 7.06        | 7.76        | 3.63        |
| E4SRec4               | <b>5.33</b> | <b>8.75</b> | <b>7.08</b> | <b>7.59</b> | <b>5.41</b> | <b>8.65</b> | <b>7.04</b> | <b>7.72</b> | <b>4.50</b> |
| SLMRec <sub>4←8</sub> | <b>5.72</b> | <b>9.15</b> | <b>7.48</b> | <b>8.03</b> | <b>6.62</b> | <b>9.83</b> | <b>8.25</b> | <b>8.89</b> | <b>1.25</b> |

Reduced model-size;  
Reduced inference time

| Method                       | Tr time(h) | Inf time(h) | Tr params (B) | Inf params (B) |
|------------------------------|------------|-------------|---------------|----------------|
| Open-P5 <sub>LLaMa</sub>     | 0.92       | 4942        | 0.023         | 7.237          |
| E4SRec                       | 3.95       | 0.415       | 0.023         | 6.631          |
| <b>SLMREC</b> <sub>4←8</sub> | 0.60       | 0.052       | 0.003         | 0.944          |

# Part 1: LLM as Sequential Recommender

**(1) Early efforts: pretrained LLMs for rec**

**(2) Aligning LLMs for recommendation**

- Pure text-based
- External item tokens
- Collaborative embeddings
- Multimodal information

**(3) Training objective & inference**

**Training:** SFT, DPO, RL;

**Inference:** (constrained) beam search

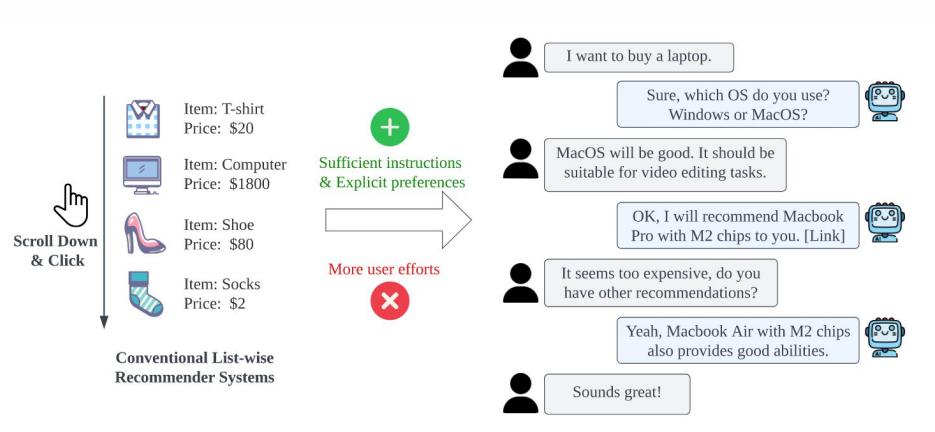
**(4) Efficiency**

Data efficiency; Inference efficiency; Model-size efficiency

# **Part 2: LLM as Conversational Recommender**

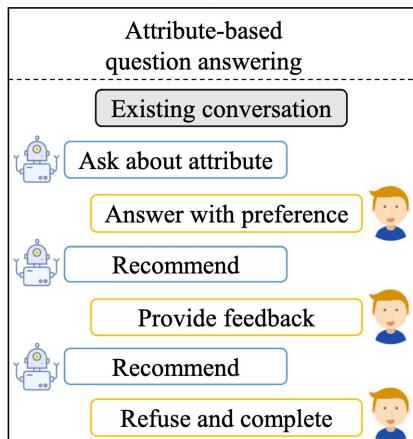
# Conversational Recommender System (CRS)

- Recommendations with multiple turns conversation
- Interactive; engaging users in the loop



# Paradigms of CRS before the era of LLM

## Attribute-based



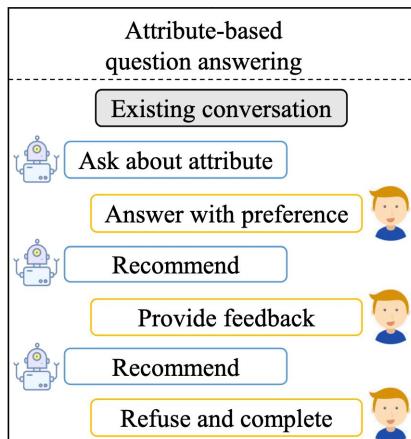
User Simulator



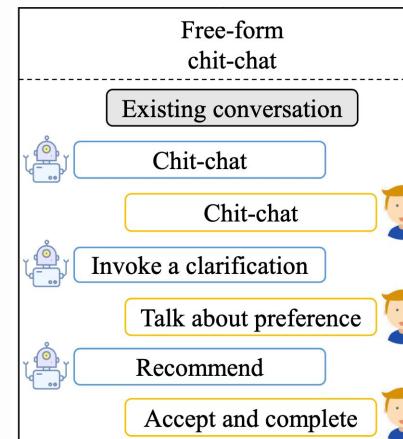
Conversational Recommendation System

# Paradigms of CRS before the era of LLM

## Attribute-based



## Free-form



User Simulator



Conversational Recommendation System

# Paradigms of CRS before the era of LLM

**Features:** Task-specific conversational recommenders,  
trained on limited conversation data.

# Paradigms of CRS before the era of LLM

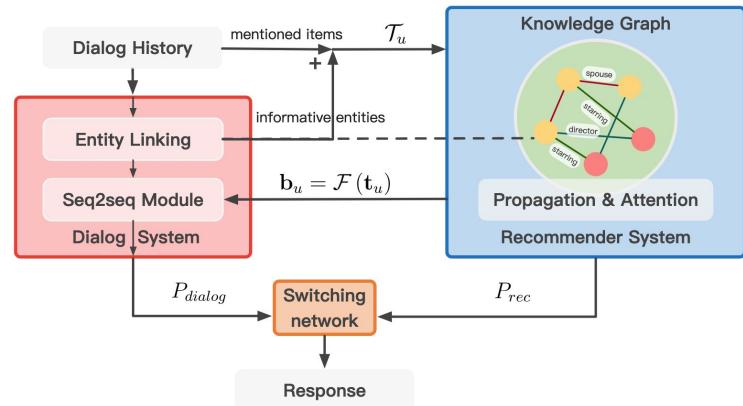
**Features:** Task-specific conversational recommenders,  
trained on limited conversation data.

- Lack of world knowledge.
- Requirement of complicated strategies.
- Incompatible natural language generation abilities.
- Lack of generalization capabilities.

# Paradigms of CRS before the era of LLM

## Traditional CRS: KBRD

- End-to-end conversational recommender system
- Switching between conversation and recommendation
- External knowledge from knowledge graph



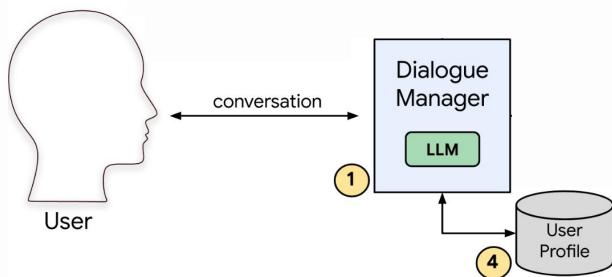
# Example

## LLM as conversational recommender



# LLM as Conversational Recommender

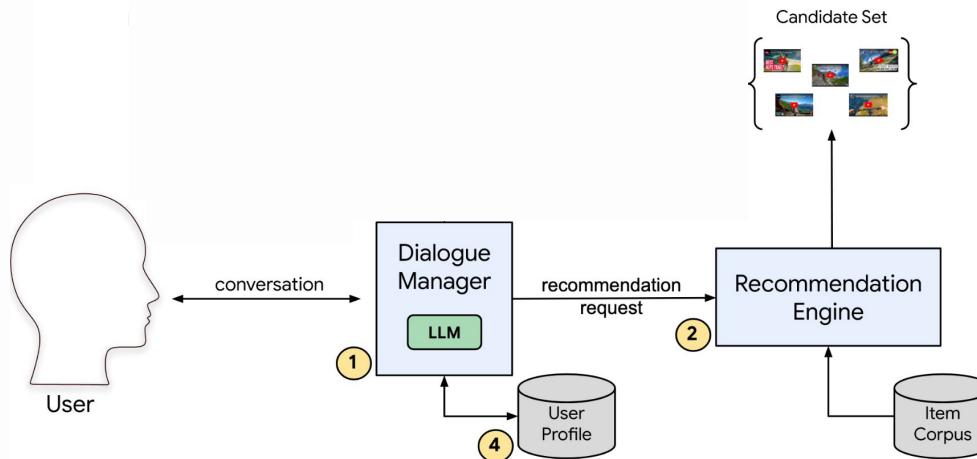
## Framework (RecLLM)



Conversation with users  
via LLMs

# LLM as Conversational Recommender

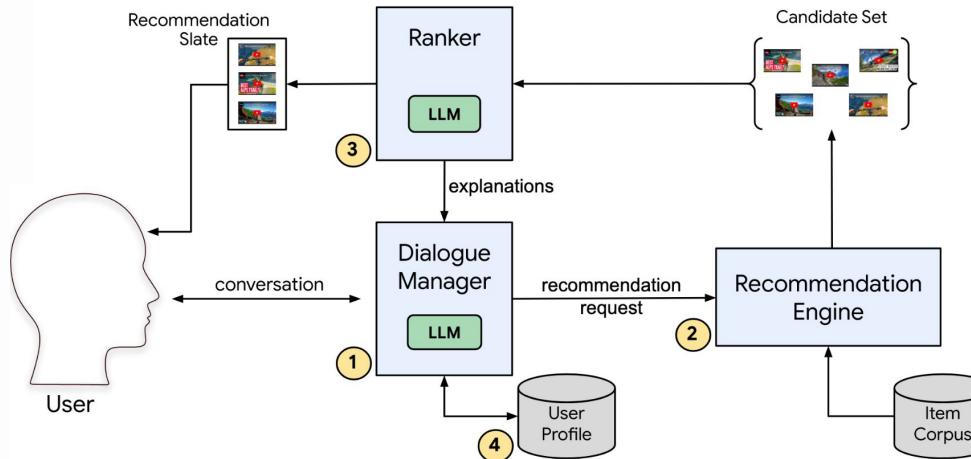
## Framework (RecLLM)



Recommendation  
via tools

# LLM as Conversational Recommender

## Framework (RecLLM)

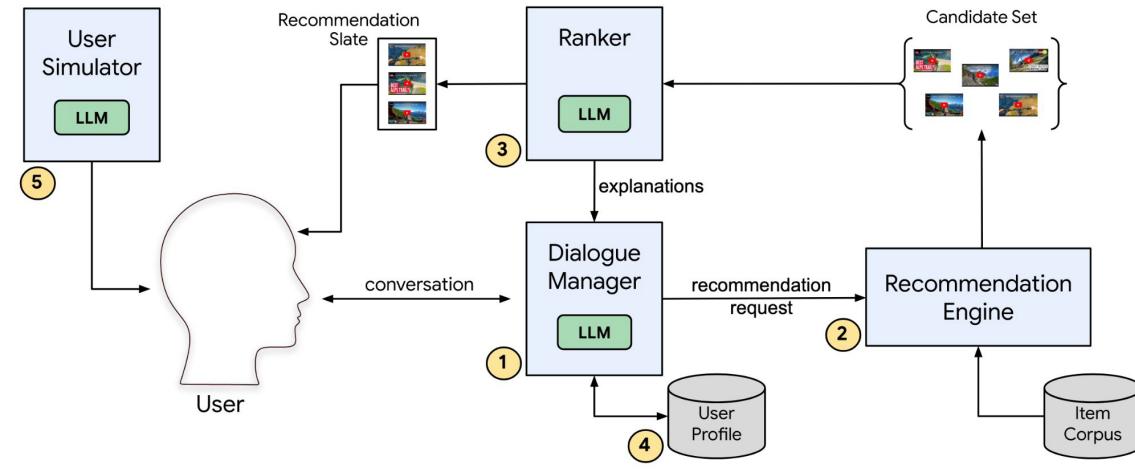


Fine-grained  
reranking via LLMs

# LLM as Conversational Recommender

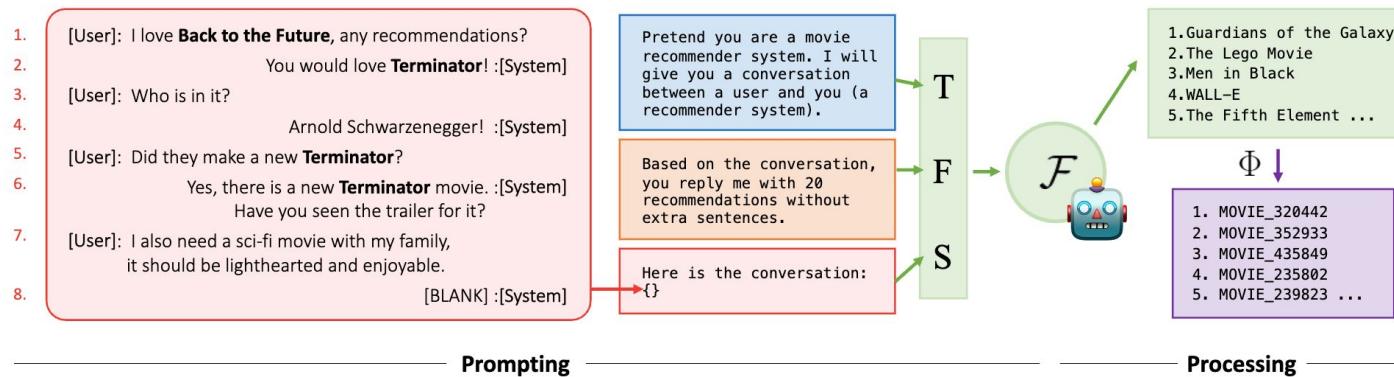
## Framework (RecLLM)

## Evaluation via LLMs



# LLM as Conversational Recommender

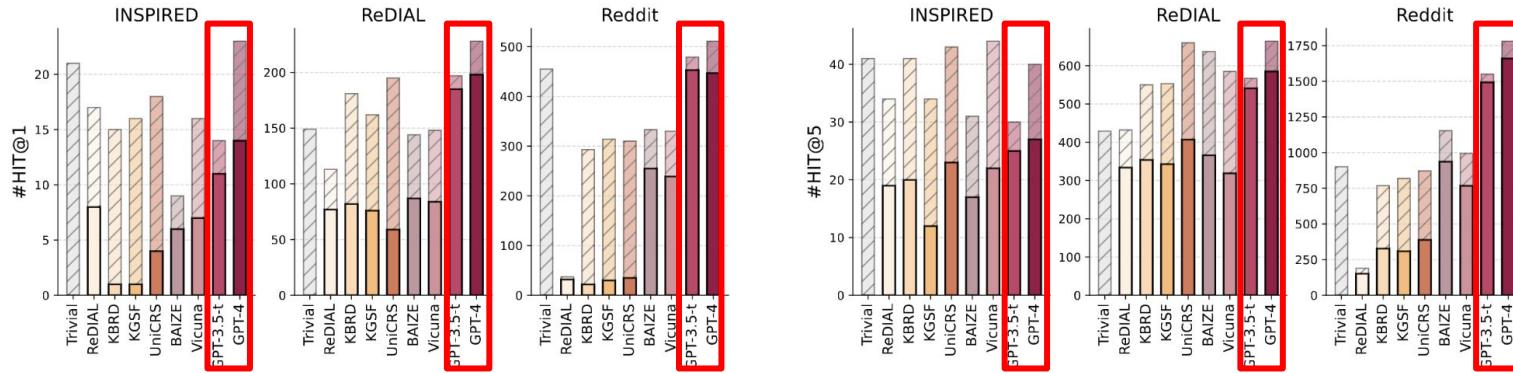
## LLMs as zero-shot CRS



How powerful are LLMs for zero-shot CRS?

# LLM as Conversational Recommender

## LLMs as zero-shot CRS



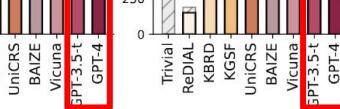
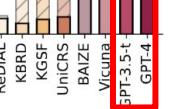
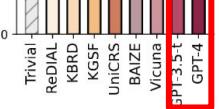
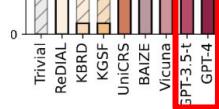
Can surpass traditional CRSs!

# LLM as Conversational Recommender

## LLMs as zero-shot CRS

INSPIRED      ReDIAL      Reddit      INSPIRED      ReDIAL      Reddit

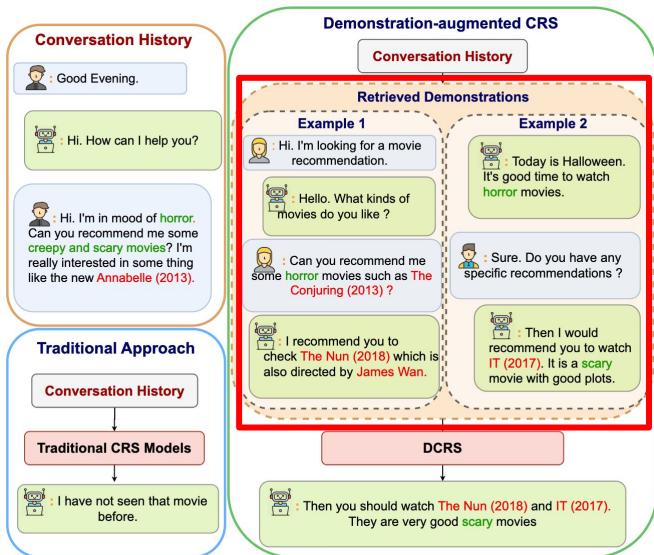
Towards better LLM-based CRS?



Can surpass traditional CRSs!

# LLM as Conversational Recommender

## + Demonstration

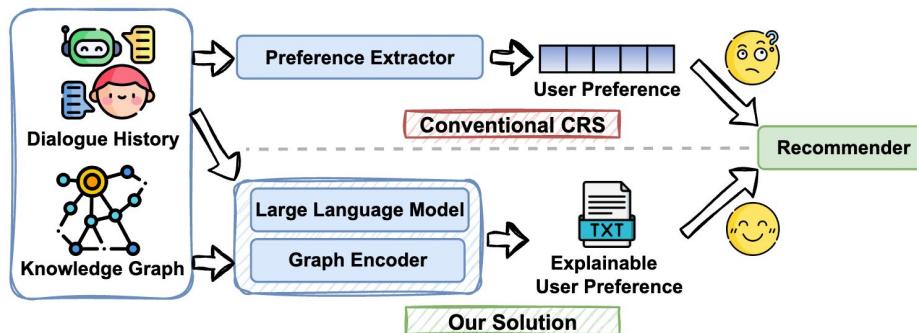


Prompting with previously successful conversation

Relevant conversation history helps!

# LLM as Conversational Recommender

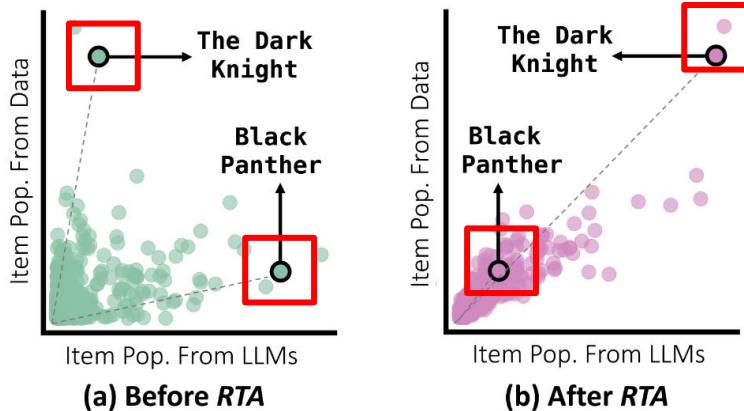
## + Knowledge graph



Recommendation-specific knowledge graph helps

# LLM as Conversational Recommender

## + Collaborative information



Collaborative information  
(e.g., popularity) helps LLMs  
fit the real distribution in CRS

# LLM as Conversational Recommender

## Challenges – Datasets

Public datasets for CRS are limited, due to the scarcity of conversational products and real-world CRS datasets

# LLM as Conversational Recommender

## Challenges – Evaluation

Traditional metrics like NDCG and BLEU are often insufficient to assess user experience

# LLM as Conversational Recommender

## Challenges – Product

What is the **form** of LLM-based CRS products?

ChatBot? Search bar? Independent App?

# **Part 2: LLM as Conversational Recommender**

**(1) LLMs show potential in CRS**

**(2) LLM-based CRS can be improved with:**

demonstration, collaborative information ...

**(3) Challenges in LLM-based CRSs:**

dataset, evaluation, and product



# User simulators before the era of LLM

## RL-based user simulator

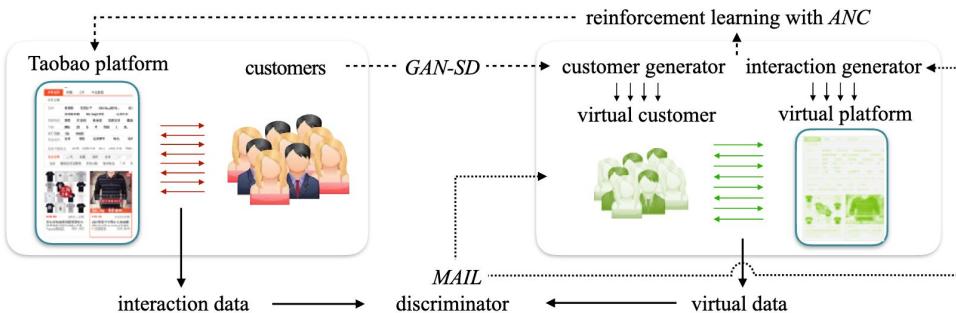


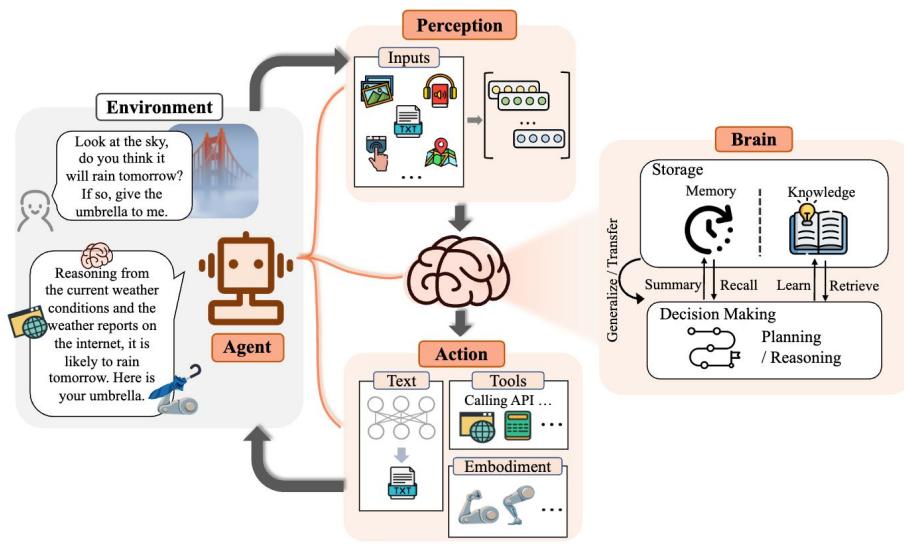
Figure 1: Virtual-Taobao architecture for reinforcement learning.

High sampling cost  
Overfitting risks  
Training instability  
Limited action space

...

# LLM as User Simulator

## Generative agents



Perception

Planning

Memory

Action

...

# LLM as User Simulator

# Generative agents for recommendation

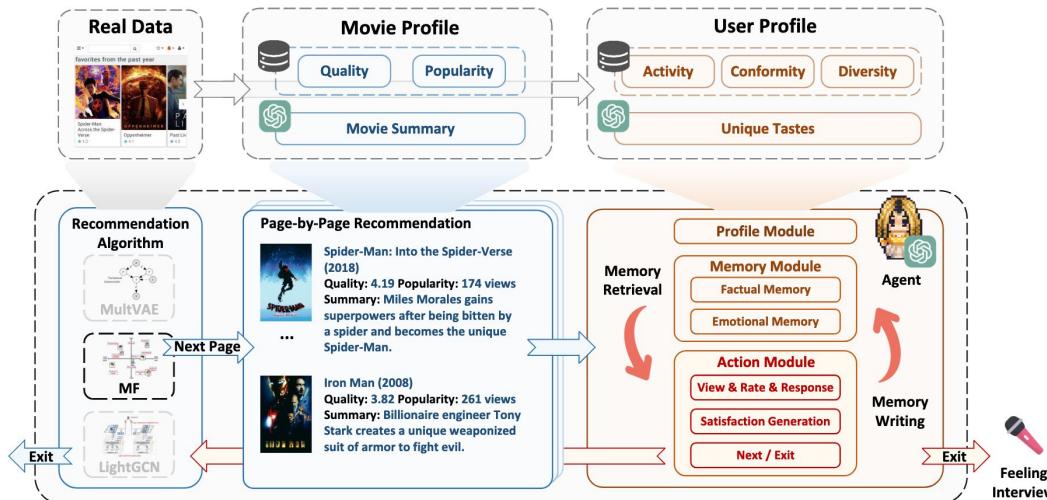


Human-like behavior  
Abundant action space  
Reduced training cost

10

# LLM as User Simulator

## Generative agents for recommendation

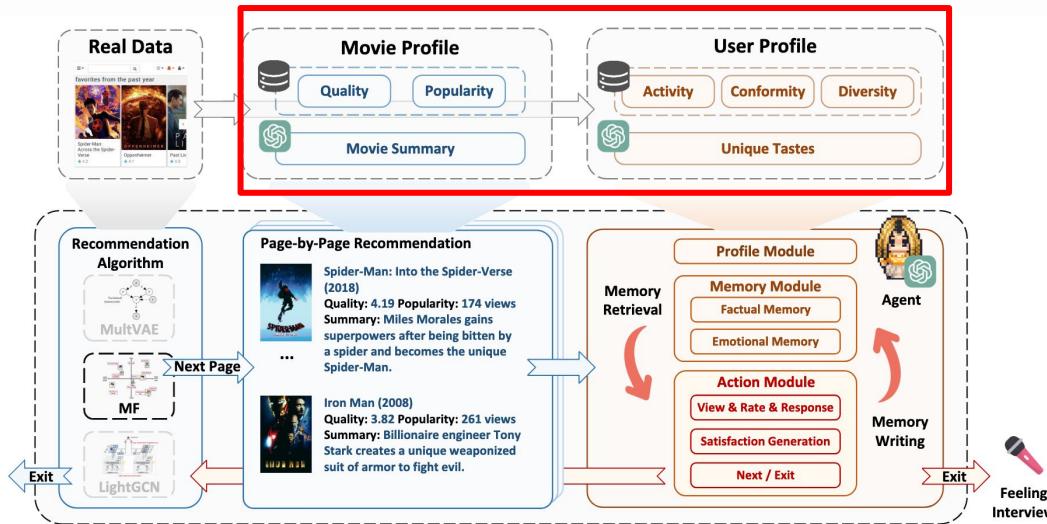


Realworld-like  
simulation paradigm

- 1000 users
- Page-by-page simulation

# LLM as User Simulator

## Generative agents for recommendation

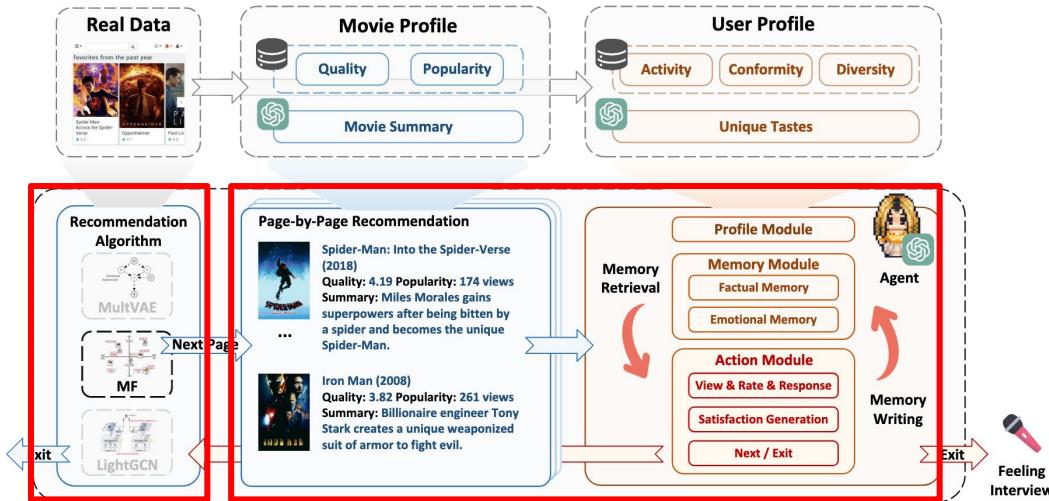


Realworld-like  
simulation paradigm

- 1000 users
- Page-by-page simulation

# LLM as User Simulator

## Generative agents for recommendation

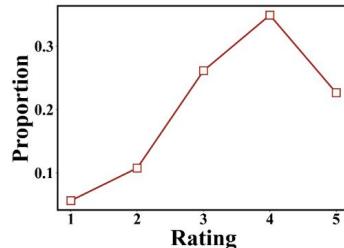


Realworld-like  
simulation paradigm

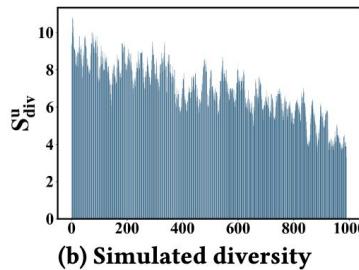
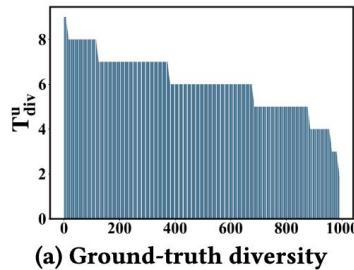
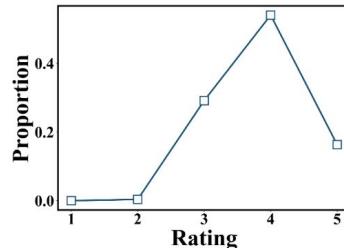
- 1000 users
- Page-by-page simulation

# LLM as User Simulator

## Generative agents for recommendation



(a) Distribution on MovieLens (b) Agent-simulated distribution



Aligned user preferences  
& Recommender evaluation

Table 2: Recommendation strategies evaluation.

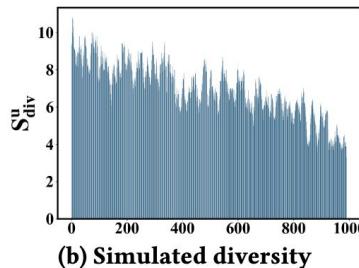
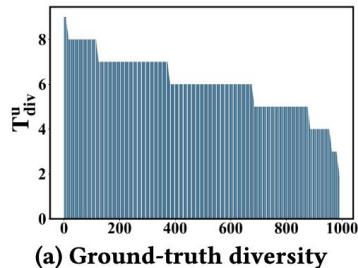
|          | $\bar{P}_{view}$ | $\bar{N}_{like}$ | $\bar{P}_{like}$ | $\bar{N}_{exit}$ | $\bar{S}_{sat}$ |
|----------|------------------|------------------|------------------|------------------|-----------------|
| Random   | 0.312            | 3.3              | 0.269            | 2.99             | 2.93            |
| Pop      | 0.398            | 4.45             | 0.360            | 3.01             | 3.42            |
| MF       | 0.488            | <b>6.07*</b>     | 0.462            | <b>3.17*</b>     | 3.80            |
| MultVAE  | 0.495            | 5.69             | 0.452            | 3.10             | 3.75            |
| LightGCN | <b>0.502*</b>    | 5.73             | <b>0.465*</b>    | 3.02             | <b>3.85*</b>    |

# LLM as User Simulator

## Generative agents for recommendation



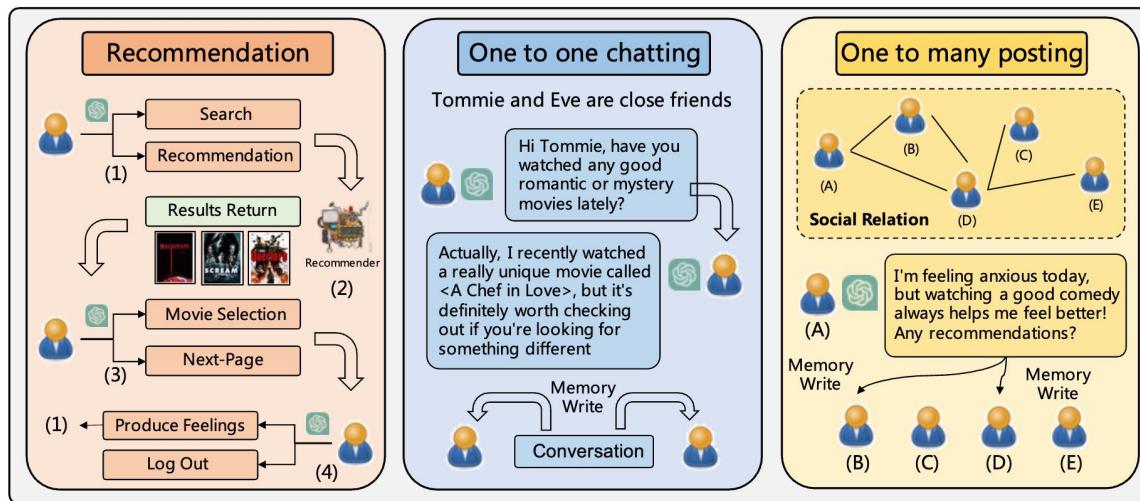
How to make the simulation more faithful?



|          | $\bar{P}_{view}$ | $\bar{N}_{like}$ | $\bar{P}_{like}$ | $\bar{N}_{exit}$ | $\bar{S}_{sat}$ |
|----------|------------------|------------------|------------------|------------------|-----------------|
| Random   | 0.312            | 3.3              | 0.269            | 2.99             | 2.93            |
| Pop      | 0.398            | 4.45             | 0.360            | 3.01             | 3.42            |
| MF       | 0.488            | <b>6.07*</b>     | 0.462            | <b>3.17*</b>     | 3.80            |
| MultVAE  | 0.495            | 5.69             | 0.452            | 3.10             | 3.75            |
| LightGCN | <b>0.502*</b>    | 5.73             | <b>0.465*</b>    | 3.02             | <b>3.85*</b>    |

# LLM as User Simulator

## + Social behaviors



Recommendation  
Chat  
Networking

# LLM as User Simulator

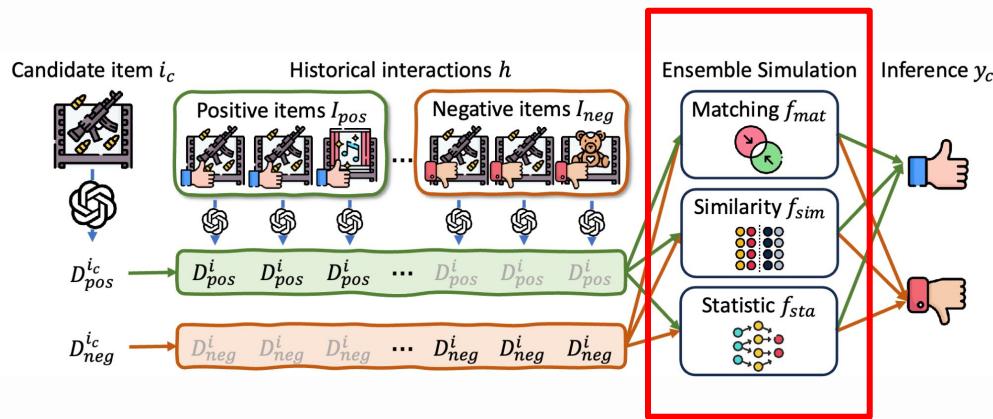
## + Multiturn conversation

|                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| User ID: A1EXXXXDUE6B0                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| Historical Interactions: ["Robin Williams: Live On Broadway", "Mission Impossible on VHS", "Solaris", "Elysium", "Wall Street", "Mystic River", ...]                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| Conversation 1:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| <p>[User] Hi, I recently watched <b>Mission Impossible on VHS</b> and it was a fantastic high-tech spy movie! Tom Cruise gave ...</p> <p>[Agent] That's great to hear! I have a movie recommendation for you based on your likes. How about checking out <i>World War Z</i>?</p> <p>[User] Oh, I'm definitely interested. Can you tell me more about the movie?</p> <p>[Agent] Certainly! <i>World War Z</i> is a good zombie war flick...</p>                                                                                                                                                                                                                 |
| Conversation 2:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| <p>[User] Hi, I recently watched the movie "<b>Solaris</b>" and I have to say, I didn't enjoy it at all. It felt overly melodramatic and lacked substance...</p> <p>[Agent] I understand why you didn't like "<i>Solaris</i>". I can recommend a movie you might enjoy instead. How about "<b>Elysium</b>"? It's a bilingual film that...</p> <p>[User] Oh, I've actually already watched "<i>Elysium</i>" and it was better than I expected, but still not great....</p> <p>[Agent] How about giving "<i>Wrecking Crew</i>" a try? It's another movie you might like based on your previous preferences...</p> <p>[User] Sure, that sounds interesting...</p> |

Simulating users in  
the conversational  
scenarios

# LLM as User Simulator

## + Multi-facet simulation objective



Category matching  
Fine-grained similarity  
Statistic information

# LLM as User Simulator

## + Multi-facet simulation objective

| Dataset      | Metric         | PPO    | TRPO   | A2C           | DQN           |
|--------------|----------------|--------|--------|---------------|---------------|
| Yelp         | <b>A. Rwd↑</b> | 9.97   | 13.45  | 24.15         | <b>27.56</b>  |
|              | <b>T. Rwd↑</b> | 141.57 | 157.42 | 267.60        | <b>330.98</b> |
|              | Liking %↑      | 34.59  | 40.07  | 48.35         | <b>49.43</b>  |
| Amazon Music | <b>A. Rwd↑</b> | 10.49  | 11.31  | 13.45         | <b>16.70</b>  |
|              | <b>T. Rwd↑</b> | 129.03 | 140.15 | 141.03        | <b>181.42</b> |
|              | Liking %↑      | 29.30  | 32.46  | 29.54         | <b>33.18</b>  |
| Amazon Games | <b>A. Rwd↑</b> | 18.72  | 21.35  | <b>27.56</b>  | 26.43         |
|              | <b>T. Rwd↑</b> | 208.43 | 242.26 | <b>317.56</b> | 269.02        |
|              | Liking %↑      | 33.15  | 37.64  | <b>43.52</b>  | 40.73         |
| Amazon Movie | <b>A. Rwd↑</b> | 29.42  | 27.47  | 31.72         | <b>38.60</b>  |
|              | <b>T. Rwd↑</b> | 310.69 | 301.40 | 354.34        | <b>416.18</b> |
|              | Liking %↑      | 38.59  | 36.70  | 42.37         | <b>44.50</b>  |
| Anime        | <b>A. Rwd↑</b> | 14.12  | 14.58  | <b>21.50</b>  | 18.03         |
|              | <b>T. Rwd↑</b> | 155.74 | 163.44 | <b>242.95</b> | 201.94        |
|              | Liking %↑      | 25.46  | 24.27  | <b>31.52</b>  | 30.67         |

Reliable environment for  
RL-based recommenders

# Part 3: LLM as User Simulator

- (1) RL-based simulators are limited in action space, action space, and training instability
- (2) LLMs open up a new paradigm for simulating users
- (3) They can give feedback for RL-based recommenders
- (4) Challenges:  
scaling, training, industry deployment

03

# Semantic ID

-based Generative Recommendation

# How to Index an Item in RecSys?

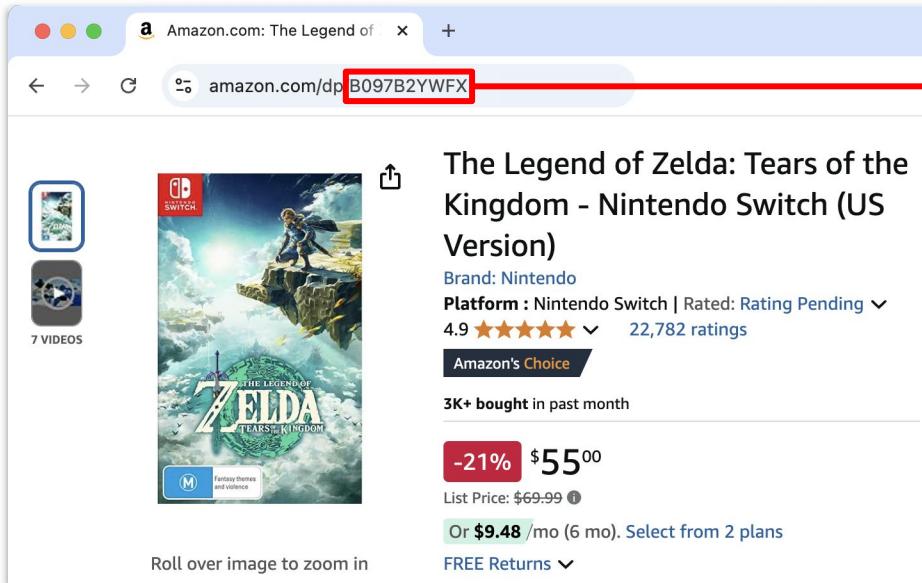
A screenshot of a web browser displaying an Amazon product page. The URL in the address bar is [amazon.com/dp/B097B2YWFX](https://amazon.com/dp/B097B2YWFX). The page shows the product title, brand, platform, rating, price, and purchase information.

The product is **The Legend of Zelda: Tears of the Kingdom - Nintendo Switch (US Version)**, Brand: **Nintendo**, Platform: **Nintendo Switch**, Rated: **Rating Pending**. It has a **4.9 ★★★★☆** rating from **22,782 ratings**. The price is **\$55<sup>00</sup>**, down from **\$69.99** (-21%). Payment options include **\$9.48 /mo (6 mo)**. Other details include **Amazon's Choice**, **3K+ bought in past month**, and **FREE Returns**.

On the left side of the page, there is a large image of the game cover for "The Legend of Zelda: Tears of the Kingdom" for Nintendo Switch. Below the main image, there are two smaller thumbnail images labeled "7 VIDEOS". A zoom icon (a magnifying glass) is positioned above the main image.

At the bottom of the page, there is a note: "Roll over image to zoom in".

# How to Index an Item in RecSys?



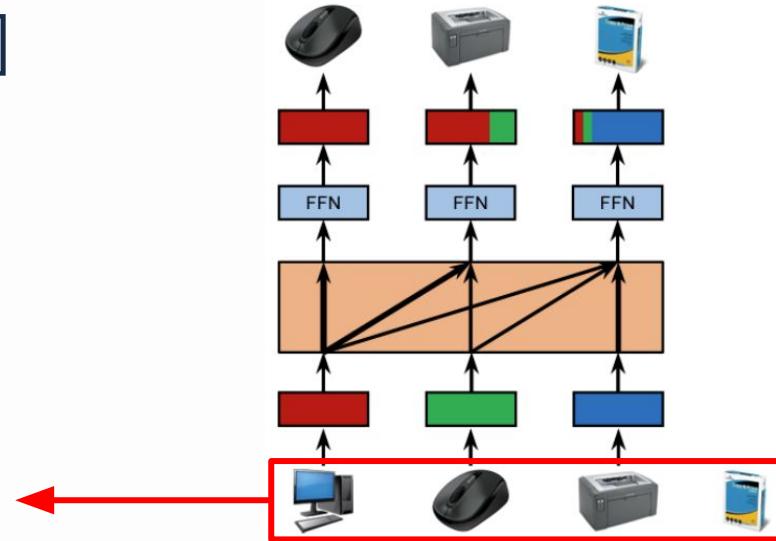
Item ID:

B097B2YWFX

# How to Index an Item in RecSys?

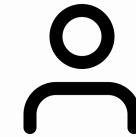
Example: SASRec [*ICDM'18*]

Each item is indexed by  
a unique **item ID**



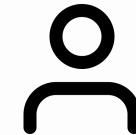
# How to Index an Item in LLMs

I wanna some popular Nintendo games



# How to Index an Item in LLMs

I wanna some popular Nintendo games



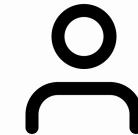
LLMs

How about or ?



# How to Index an Item in LLMs

I wanna some popular Nintendo games



LLMs

How about or ?



How to index items in LLMs? **Item ID?**



# How to Index an Item in LLMs

How many tokens in LLMs?

|                                                                                                   |         |          |
|---------------------------------------------------------------------------------------------------|---------|----------|
|  Meta            | Llama 3 | ~128,000 |
|  OpenAI          | GPT-4o  | ~200,000 |
|  Google DeepMind | Gemma 2 | ~256,000 |

# How to Index an Item in LLMs

How many tokens in LLMs?

|                                                                                                   |         |          |
|---------------------------------------------------------------------------------------------------|---------|----------|
|  Meta            | Llama 3 | ~128,000 |
|  OpenAI          | GPT-4o  | ~200,000 |
|  Google DeepMind | Gemma 2 | ~256,000 |

How many item IDs?

Amazon-Reviews-2023 ~48,200,000

# How to Index an Item in LLMs

How many tokens/item IDs in LLMs/RecSys?

**Difficult to align  
these vocabularies  
given so many tokens**

~128,000  
~200,000  
~256,000  
  
~48,200,000

# How to Index an Item in LLMs

Is there a way to  
**index** a large volume of items  
using a **compact vocabulary?**

# Semantic IDs

(also called: SemID or SID)

A few tokens that jointly index one item.

t3, t321, t643, t1011



# Semantic IDs

(also called: SemID or SID)

A few tokens that jointly index one item.

t3, **t321**, t643, t1011



{t257, t258, ..., t320, **t321**, t322, ..., t511, t512}

Each token from a **vocabulary shared by all items**

# Semantic IDs

(also called: SemID or SID)

A few tokens that jointly index one item.

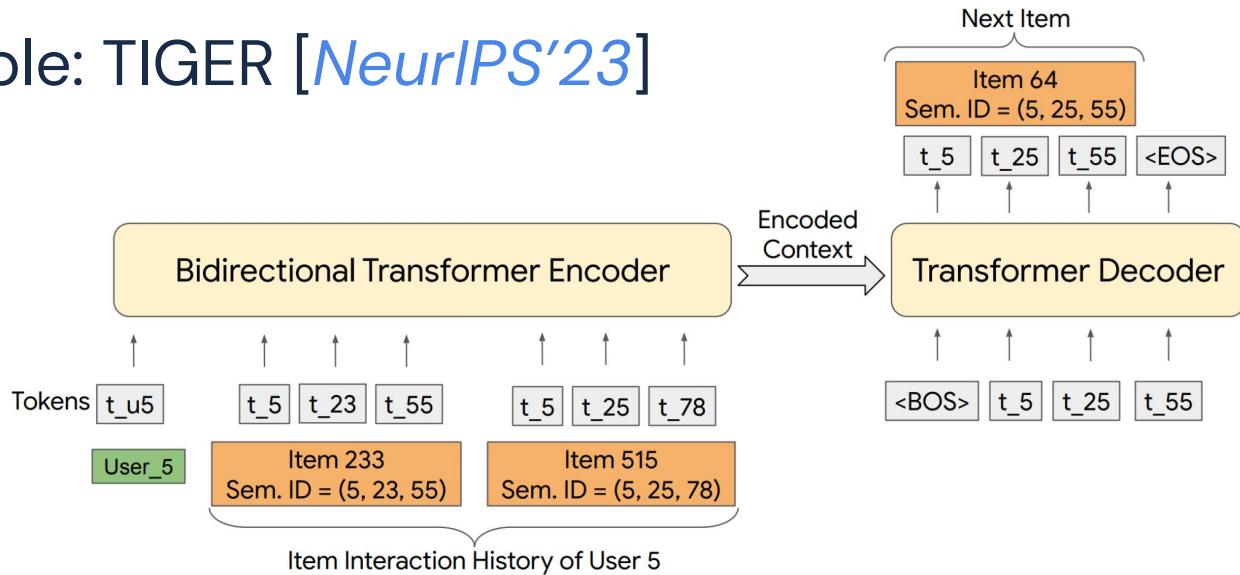
**t3, t321, t643, t1011**

Can index maximally  $256^4 \approx 4.3 \times 10^9$  items with 1024 tokens

(4 tokens per item, each from a vocabulary of 256)

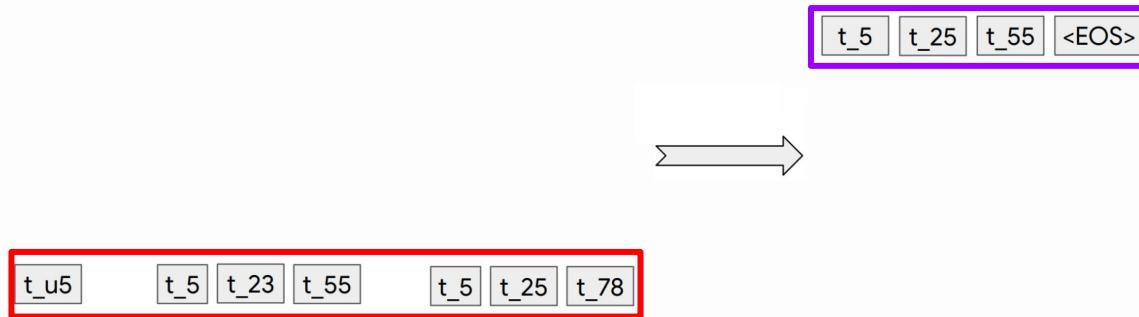
# Generative Models based on Semantic IDs

Example: TIGER [NeurIPS'23]



# Generative Models based on Semantic IDs

Example: TIGER [*NeurIPS'23*]



Recommendation as a **seq**-to-**seq** generation problem

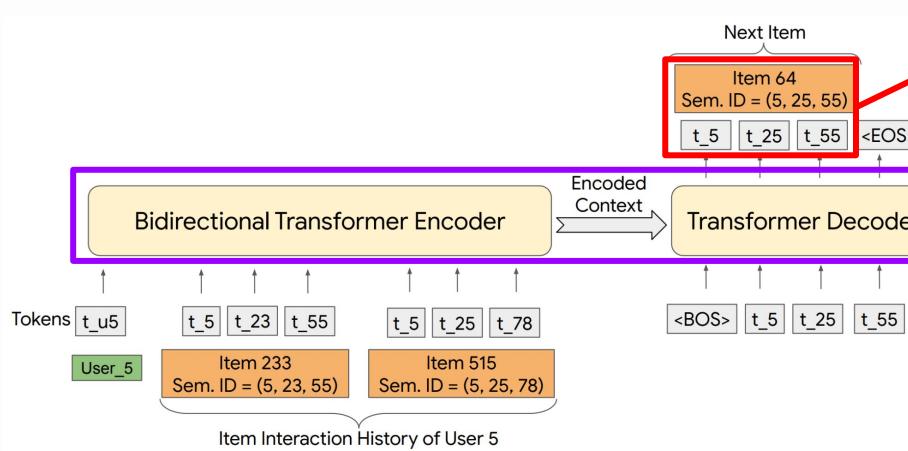
# Generative Models based on Semantic IDs

Recommendation as a **seq**-to-**seq** generation problem

**Input:** user interacted items  $\{c_{11}, c_{12}, c_{13}, c_{14}, c_{21}, c_{22}, \dots\}$

**Output:** next item  $\{c_{t1}, c_{t2}, c_{t3}, c_{t4}\}$

# SemID-based Generative Recommendation



Part 1:

How to construct **SIDs**

Part 2:

How to build SID-based  
Generative Rec **Models**

# **Part 1: Semantic ID Construction**

# Semantic ID Construction

**Input:** all data associated with the item  
(description, title, interactions, features, ...)



**Output:** mapping between **items**  $\leftrightarrow$  **Semantic IDs**

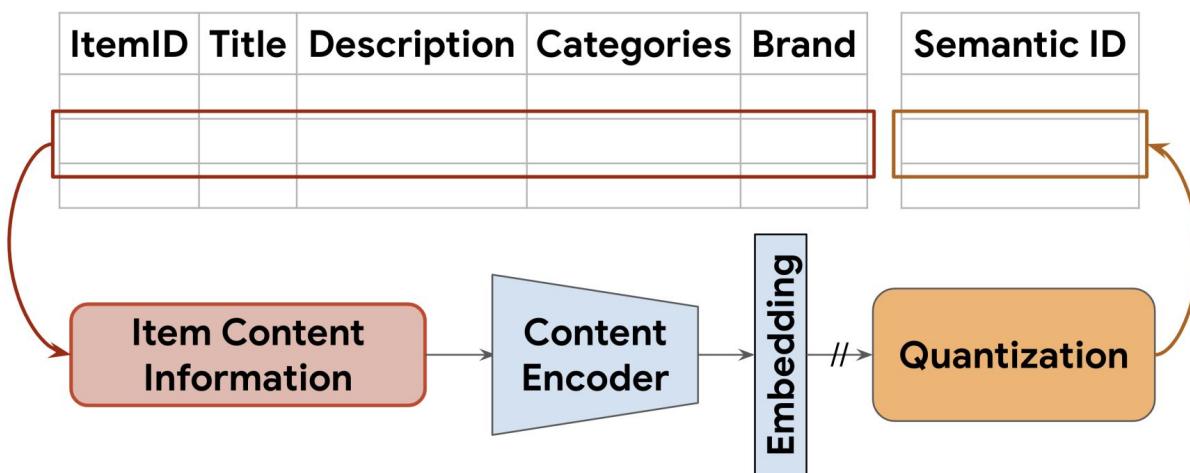
B097B2YWFX  $\leftrightarrow$  {t3, t321, t643, t1011} 169

# Part 1: Semantic ID Construction

(i) **First example:** TIGER and RQ-VAE-based SemIDs;

# SemID Construction – First Example: TIGER

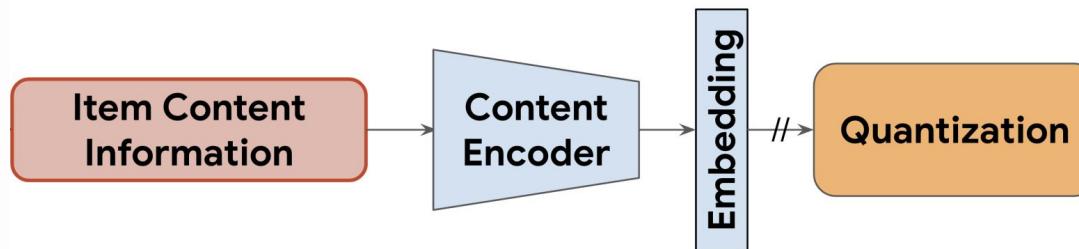
**Input:** concatenated text features



**Output:** mapping between **items**  $\leftrightarrow$  **Semantic IDs**

# SemID Construction – First Example: TIGER

**Text** ➤ **Vector** ➤ **IDs**



# SemID Construction – First Example: TIGER

## 1. Item Content Information (**Text**)

| ItemID      | Title                                                                     | Description                                                 |
|-------------|---------------------------------------------------------------------------|-------------------------------------------------------------|
| B097B2YWFX. | The Legend of Zelda: Tears of the Kingdom - Nintendo Switch (US Version). | An epic adventure across the land ... threaten the kingdom? |

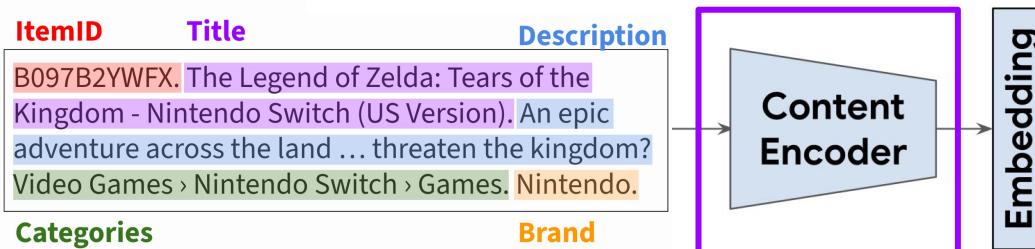
Categories

Brand

# SemID Construction – First Example: TIGER

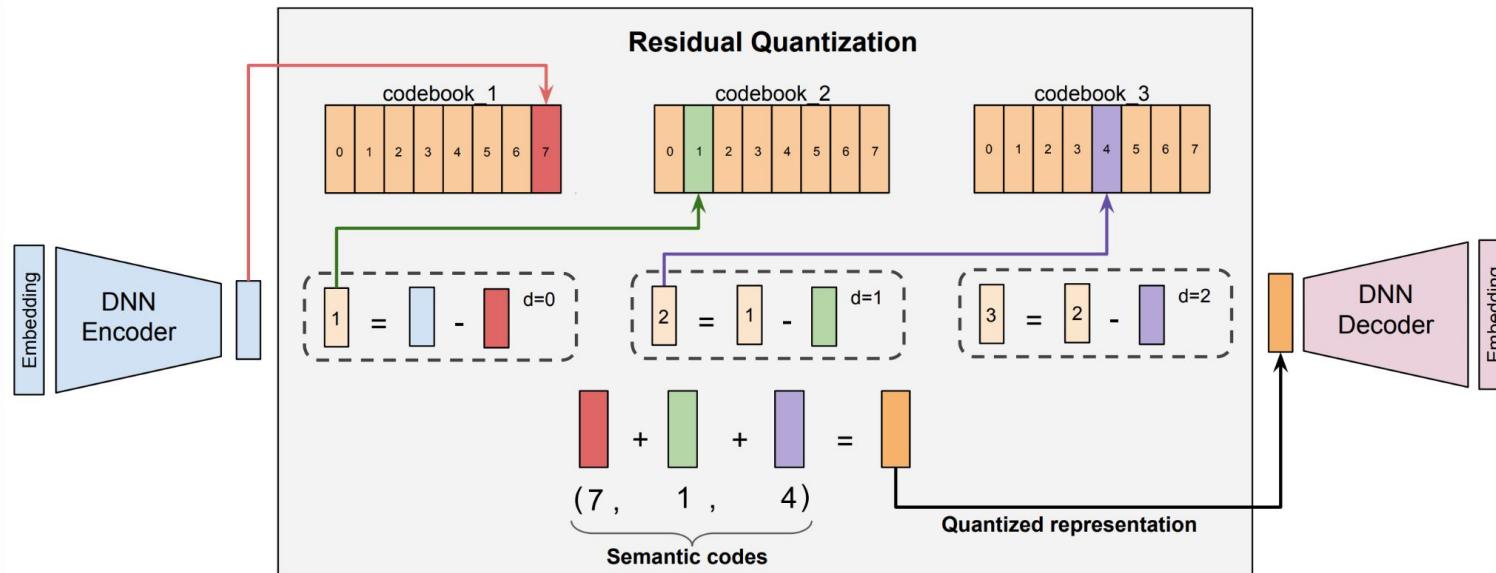
## 2. Content Encoder + Embedding (**Text** ➤ **Vector**)

Pre-trained (fixed) sentence embedding model  
**(SentenceT5)**



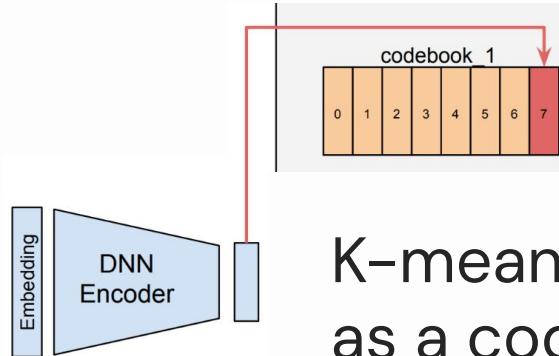
# SemID Construction – First Example: TIGER

## 3. RQ-VAE Quantization (Vector $\rightarrow$ IDs)



# SemID Construction – First Example: TIGER

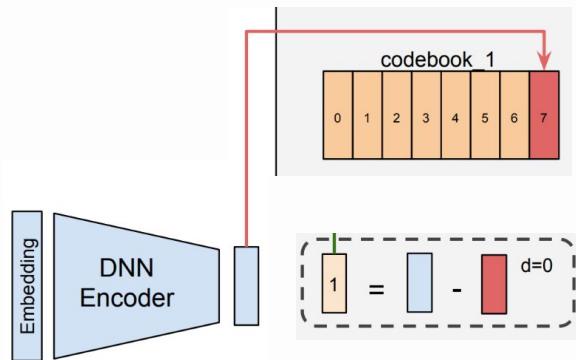
## 3. RQ-VAE Quantization (Vector $\rightarrow$ IDs)



K-means: cluster center ID  
as a code in the codebook

# SemID Construction – First Example: TIGER

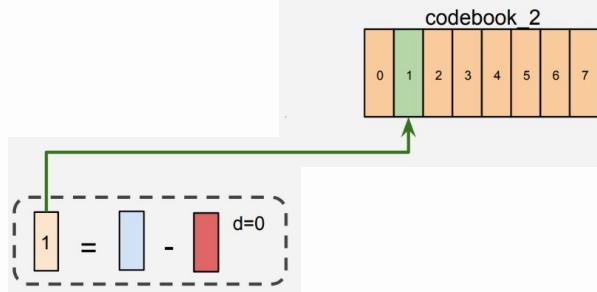
## 3. RQ-VAE Quantization (Vector $\rightarrow$ IDs)



Residual of “input vector” and  
“clustering center vector”

# SemID Construction – First Example: TIGER

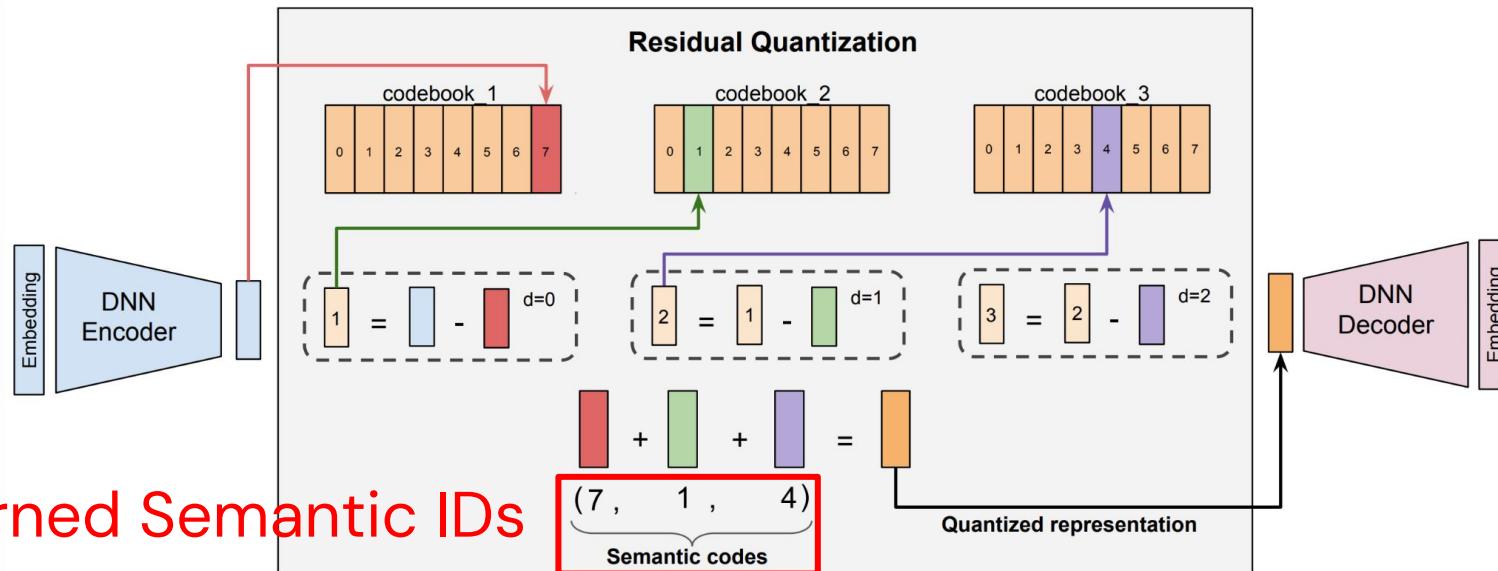
## 3. RQ-VAE Quantization (Vector $\rightarrow$ IDs)



Residual as next level's input

# SemID Construction – First Example: TIGER

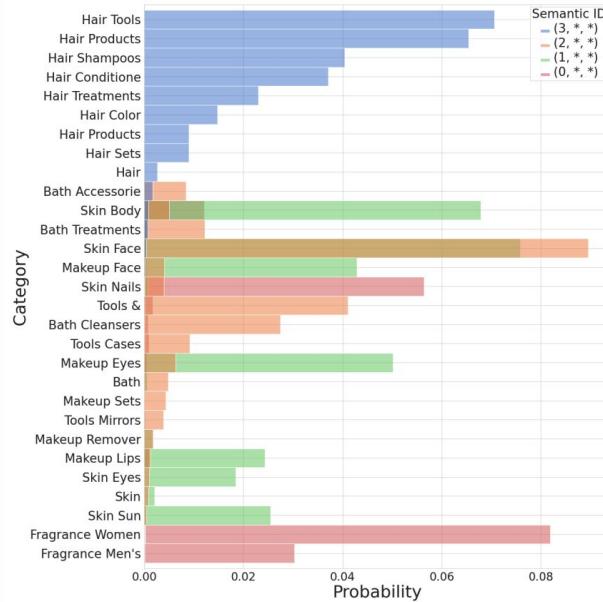
## 3. RQ-VAE Quantization (Vector $\rightarrow$ IDs)



# SemID Construction – First Example: TIGER

## Properties of RQ-VAE-based SemIDs

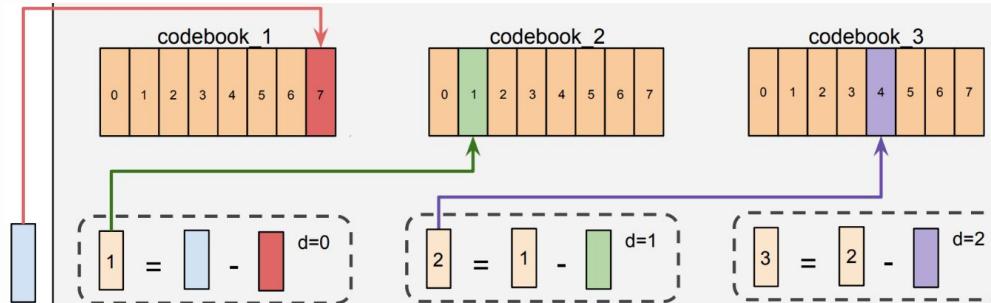
### 1. Semantic;



# SemID Construction – First Example: TIGER

## Properties of RQ-VAE-based SemIDs

1. Semantic;
2. Ordered / sequential dependent;



# SemID Construction – First Example: TIGER

## Collisions



(12, 24, 52)



(12, 24, 52)

# SemID Construction – First Example: TIGER

## Collisions



(12, 24, 52, 0)



(12, 24, 52, 1)

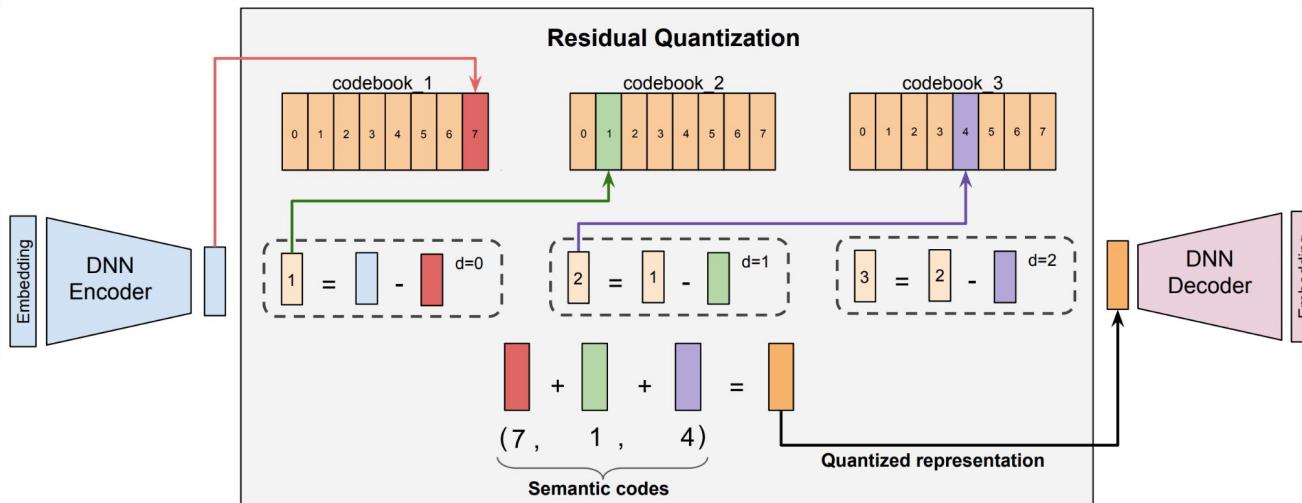
One extra token  
to avoid conflicts

# Part 1: Semantic ID Construction

- (i) First example: TIGER and RQ-VAE-based SemIDs;
- (ii) **Techniques** to construct SemIDs;

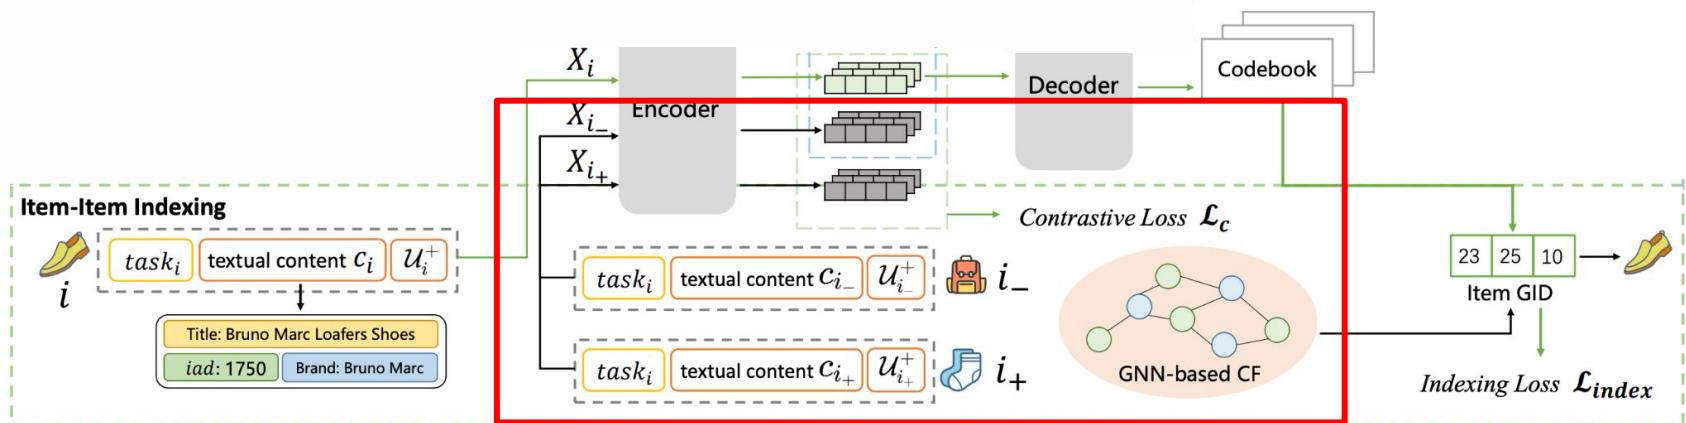
# Techniques to Construct SemIDs

## Residual Quantization



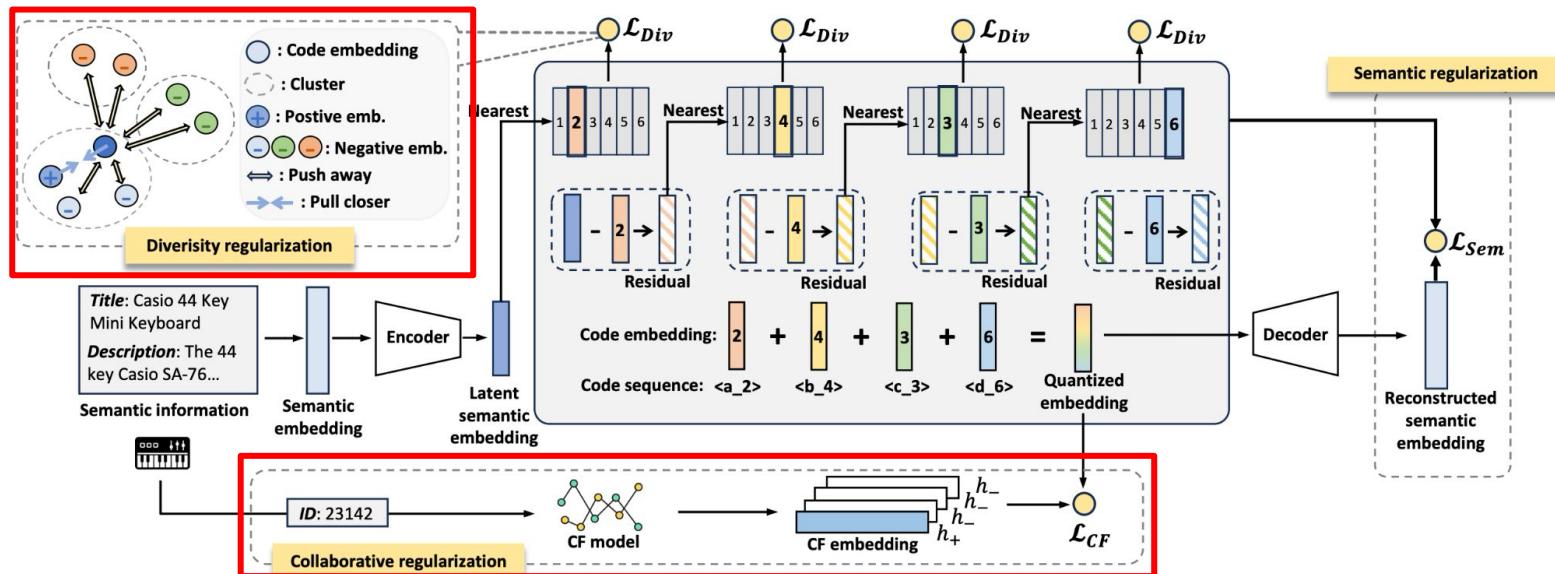
# Techniques to Construct SemIDs

## Residual Quantization + Item-level Regularization



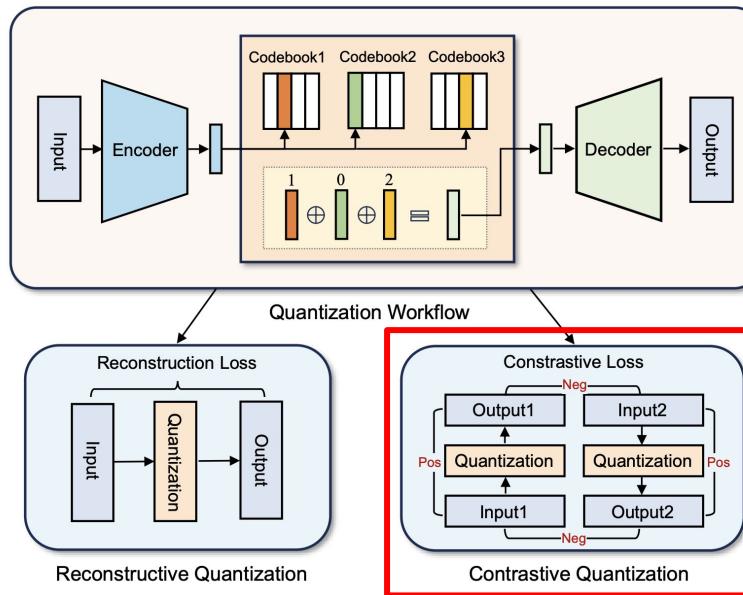
# Techniques to Construct SemIDs

## Residual Quantization + Item-level Regularization



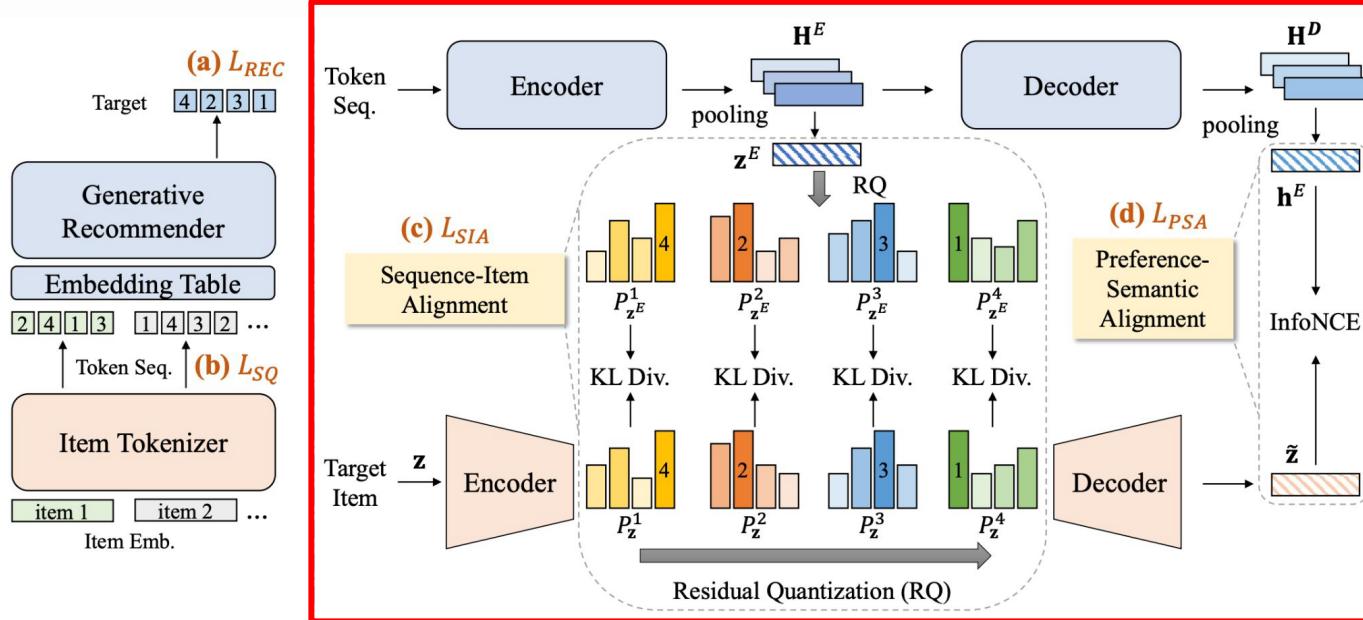
# Techniques to Construct SemIDs

## Residual Quantization + Item-level Regularization



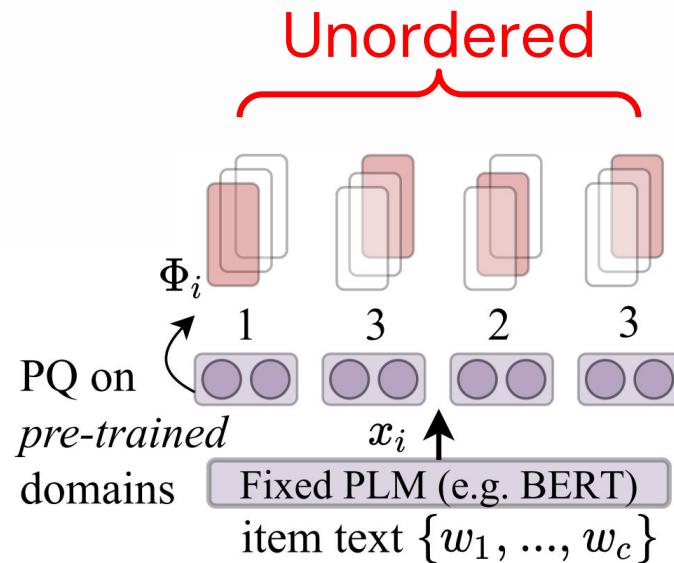
# Techniques to Construct SemIDs

## Residual Quantization + Recommendation Loss



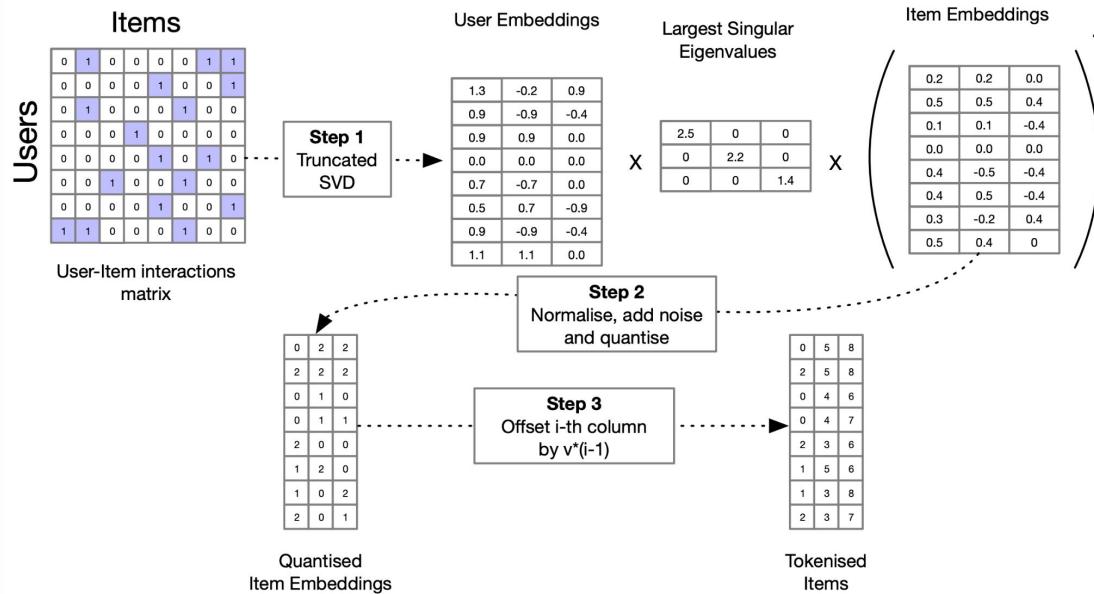
# Techniques to Construct SemIDs

## Product Quantization



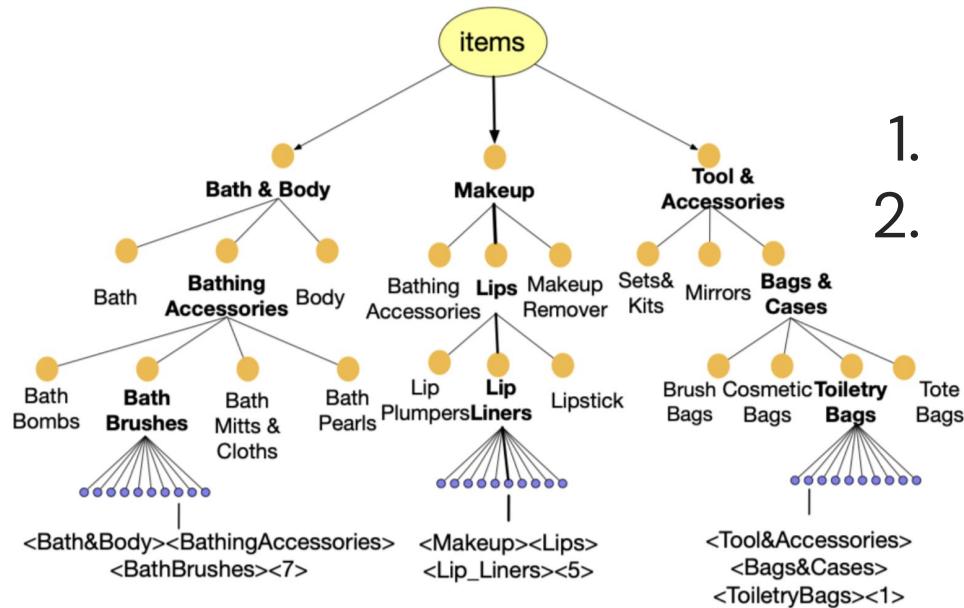
# Techniques to Construct SemIDs

## Product Quantization



# Techniques to Construct SemIDs

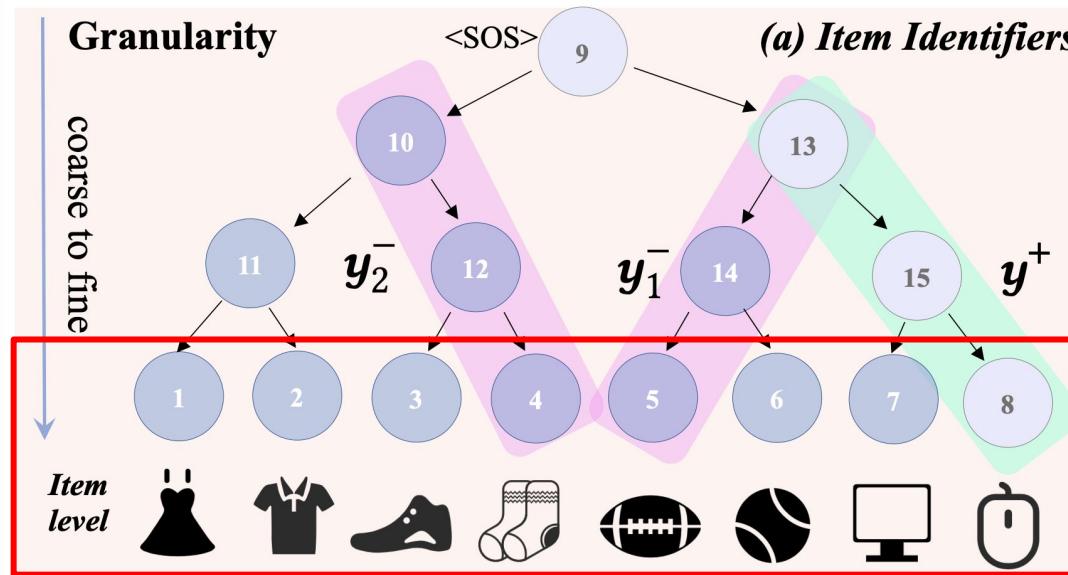
## Hierarchical Clustering (**Heuristics**-based)



1. Ordered;
2. **Variable-length** SemIDs;

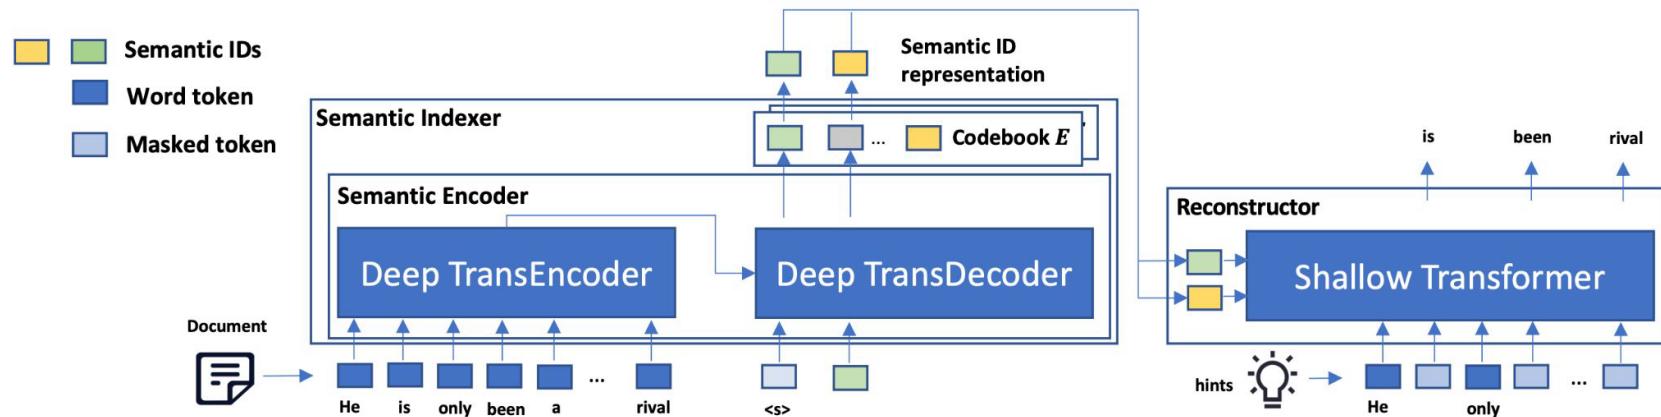
# Techniques to Construct SemIDs

## Hierarchical Clustering (**Latent**-based)



# Techniques to Construct SemIDs

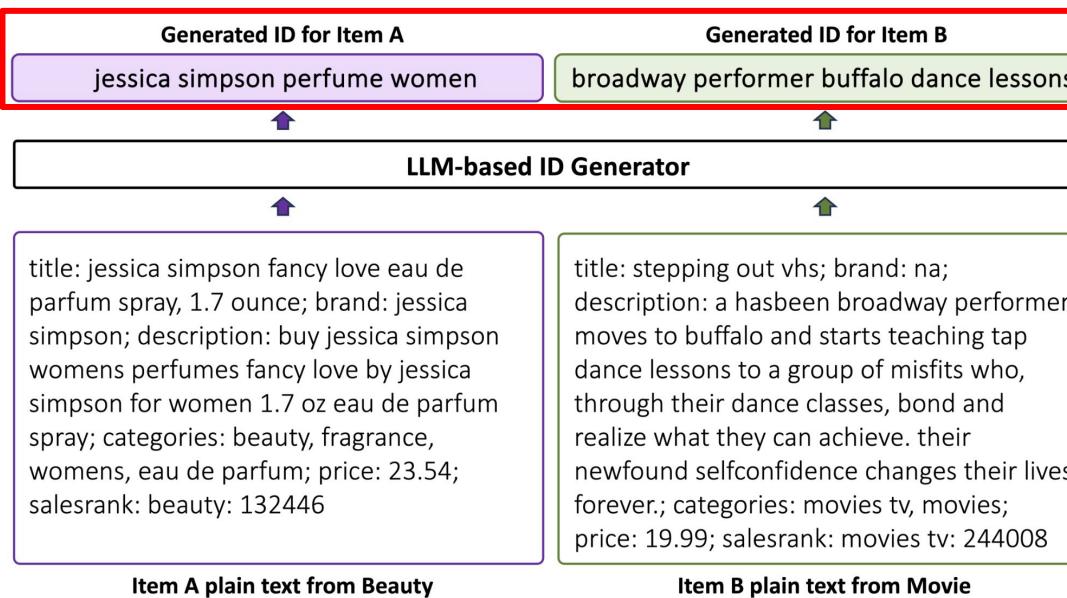
## Language Model-based ID Generator



Natural language as inputs;  
SemIDs as outputs

# Techniques to Construct SemIDs

## Language Model-based ID Generator



Words as SemIDs  
(like tagging)

# Techniques to Construct SemIDs

## Context-independent

| Action Tokenization     | Example                      | Contextual | Unordered |
|-------------------------|------------------------------|------------|-----------|
| Product Quantization    | VQ-Rec (Hou et al., 2023)    | ✗          | ✓         |
| Hierarchical Clustering | P5-CID (Hua et al., 2023)    | ✗          | ✗         |
| Residual Quantization   | TIGER (Rajput et al., 2023)  | ✗          | ✗         |
| Text Tokenization       | LMIndexer (Jin et al., 2024) | ✗          | ✗         |
| Raw Features            | HSTU (Zhai et al., 2024)     | ✗          | ✗         |
| SentencePiece           | SPM-SID (Singh et al., 2024) | ✗          | ✗         |

Same item ⇒ **fixed semIDs** in all sequences

# Techniques to Construct SemIDs

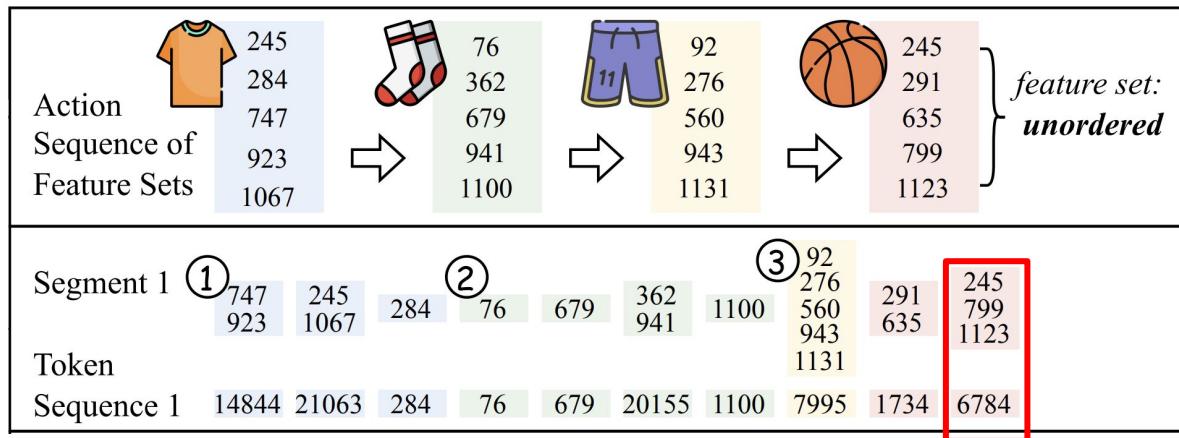
Context-independent  $\Rightarrow$  Context-aware

| Action    | 245<br>284<br>747<br>923<br>1067 | Sequence of Feature Sets                     | 76<br>362<br>679<br>941<br>1100           | feature set:<br><i>unordered</i>                                       |
|-----------|----------------------------------|----------------------------------------------|-------------------------------------------|------------------------------------------------------------------------|
| Segment 1 | ①<br>747<br>923                  | 245<br>1067                                  | 284<br>76<br>679<br>941<br>1100           | ②<br>362<br>943<br>1131<br>276<br>560<br>943<br>291<br>635<br>1123     |
| Token     | 14844<br>21063<br>284            | 76<br>679<br>20155<br>1100                   | 7995<br>1734<br>6784                      | Sequence 1                                                             |
| Segment 2 | 747<br>284<br>245<br>1067        | 923<br>679<br>941<br>76<br>362<br>560<br>943 | 1100<br>1131<br>276<br>245<br>799<br>1123 | ④<br>8316<br>941<br>76<br>362<br>19895<br>24364<br>276<br>6784<br>1734 |
| Token     | 747<br>284<br>21063              | 8316<br>941<br>76<br>362<br>19895            | 24364<br>276<br>6784<br>1734              | Sequence 2                                                             |

Same item  $\Rightarrow$   
different semIDs  
based on context

# Techniques to Construct SemIDs

Context-independent  $\Rightarrow$  Context-aware



**Core Idea:**  
Merge frequently co-occurring features as new tokens

(ActionPiece: “WordPiece” tokenization for generative rec)

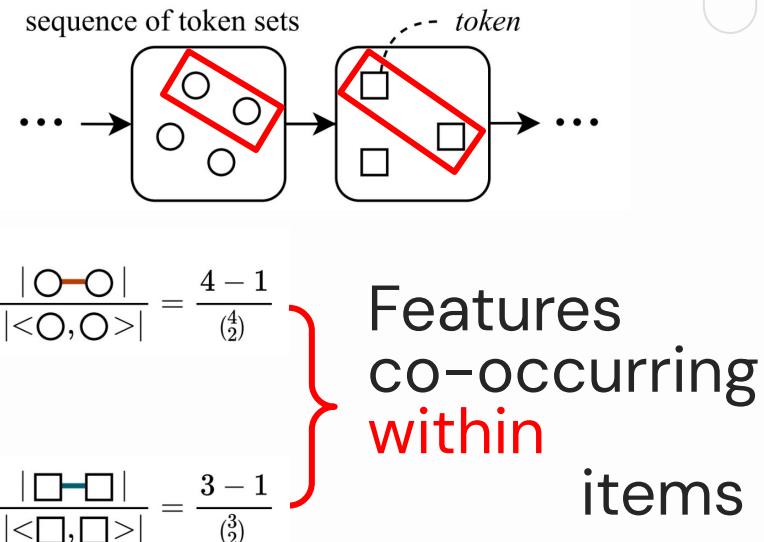
# Techniques to Construct SemIDs

Context-independent  $\Rightarrow$  Context-aware

## Algorithm 1 ActionPiece Vocabulary Construction

**input** Sequence corpus  $\mathcal{S}'$ , initial tokens  $\mathcal{V}_0$ , target size  $Q$   
**output** Merge rules  $\mathcal{R}$ , constructed vocabulary  $\mathcal{V}$

```
1: Initialize vocabulary  $\mathcal{V} \leftarrow \mathcal{V}_0$  # each initial token corresponds  
   to one unique item feature  
2:  $\mathcal{R} \leftarrow \emptyset$   
3: while  $|\mathcal{V}| < Q$  do  
4:   # Count: accumulate weighted token co-occurrences  
5:    $\text{count}(\cdot, \cdot) \leftarrow \text{Count}(\mathcal{S}', \mathcal{V})$  # Algorithm 2  
6:   # Update: Merge a frequent token pair into a new token  
7:   Select  $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$   
8:    $\mathcal{S}' \leftarrow \text{Update}(\mathcal{S}', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$  # Algorithm 3  
9:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$  # new merge rule  
10:   $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$  # add new token to the vocabulary  
11: end while  
return  $\mathcal{R}, \mathcal{V}$ 
```



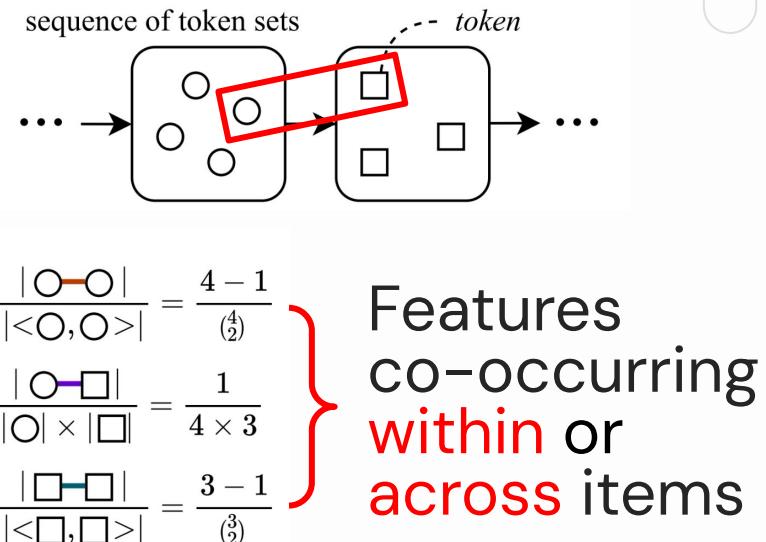
# Techniques to Construct SemIDs

Context-independent  $\Rightarrow$  Context-aware

## Algorithm 1 ActionPiece Vocabulary Construction

**input** Sequence corpus  $\mathcal{S}'$ , initial tokens  $\mathcal{V}_0$ , target size  $Q$   
**output** Merge rules  $\mathcal{R}$ , constructed vocabulary  $\mathcal{V}$

```
1: Initialize vocabulary  $\mathcal{V} \leftarrow \mathcal{V}_0$  # each initial token corresponds  
   to one unique item feature  
2:  $\mathcal{R} \leftarrow \emptyset$   
3: while  $|\mathcal{V}| < Q$  do  
4:   # Count: accumulate weighted token co-occurrences  
5:    $\text{count}(\cdot, \cdot) \leftarrow \text{Count}(\mathcal{S}', \mathcal{V})$  # Algorithm 2  
6:   # Update: Merge a frequent token pair into a new token  
7:   Select  $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$   
8:    $\mathcal{S}' \leftarrow \text{Update}(\mathcal{S}', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$  # Algorithm 3  
9:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$  # new merge rule  
10:   $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$  # add new token to the vocabulary  
11: end while  
return  $\mathcal{R}, \mathcal{V}$ 
```



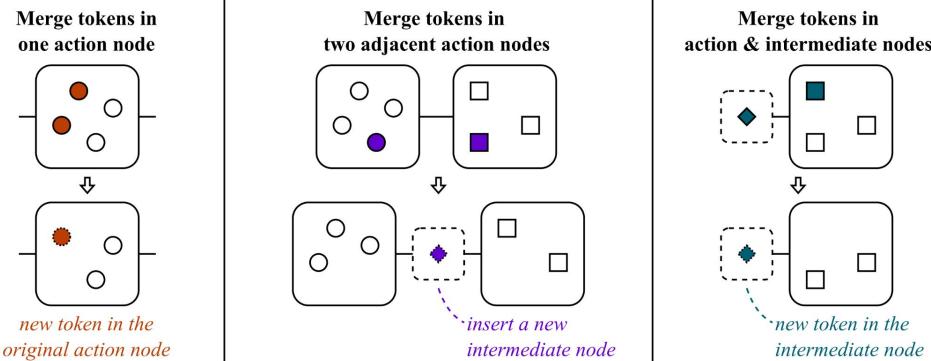
# Techniques to Construct SemIDs

Context-independent  $\Rightarrow$  Context-aware

## Algorithm 1 ActionPiece Vocabulary Construction

**input** Sequence corpus  $\mathcal{S}'$ , initial tokens  $\mathcal{V}_0$ , target size  $Q$   
**output** Merge rules  $\mathcal{R}$ , constructed vocabulary  $\mathcal{V}$

```
1: Initialize vocabulary  $\mathcal{V} \leftarrow \mathcal{V}_0$  # each initial token corresponds  
   to one unique item feature  
2:  $\mathcal{R} \leftarrow \emptyset$   
3: while  $|\mathcal{V}| < Q$  do  
4:   # Count: accumulate weighted token co-occurrences  
5:   count( $\cdot, \cdot$ )  $\leftarrow$  Count( $\mathcal{S}', \mathcal{V}$ ) # Algorithm 2  
6:   # Update: Merge a frequent token pair into a new token  
7:   Select  $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$   
8:    $\mathcal{S}' \leftarrow \text{Update}(\mathcal{S}', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$  # Algorithm 3  
9:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$  # new merge rule  
10:   $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$  # add new token to the vocabulary  
11: end while  
return  $\mathcal{R}, \mathcal{V}$ 
```



# **Summary of Techniques to Construct SemIDs**

## **Context-independent**

- Residual Quantization (+ regularization)
- Product Quantization
- Hierarchical Clustering
- LM-based ID Generator

# **Summary of Techniques to Construct SemIDs**

## **Context-independent**

- Residual Quantization (+ regularization)
- Product Quantization
- Hierarchical Clustering
- LM-based ID Generator

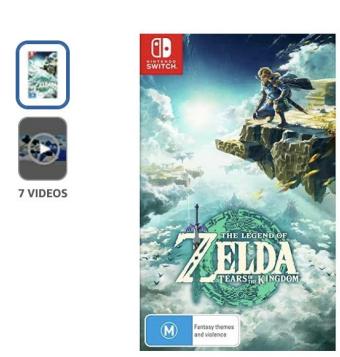
## **Context-aware**

# Part 1: Semantic ID Construction

- (i) First example: TIGER and RQ-VAE-based SemIDs;
- (ii) Techniques to construct SemIDs;
- (iii) **Inputs** for SemID construction;

# Inputs for SemID Construction

Input: all data associated with the item



The Legend of Zelda: Tears of the Kingdom - Nintendo Switch (US Version)

Brand: Nintendo

Platform : Nintendo Switch | Rated: Rating Pending ▾

4.9 ★★★★★ 22,782 ratings

Amazon's Choice

3K+ bought in past month

-21% \$55<sup>00</sup>

List Price: \$69.99 ⓘ

Or \$9.48 /mo (6 mo). Select from 2 plans

FREE Returns ▾

# Inputs for SemID Construction

**Input: all data associated with the item**

What exactly does “all data” mean? 🤔

# Inputs for SemID Construction

## Text or Multimodal Features



| ItemID      | Title                                                                     | Description                                                 |
|-------------|---------------------------------------------------------------------------|-------------------------------------------------------------|
| B097B2YWFX. | The Legend of Zelda: Tears of the Kingdom - Nintendo Switch (US Version). | An epic adventure across the land ... threaten the kingdom? |
| Video Games | Nintendo Switch                                                           | Games                                                       |
|             | Nintendo                                                                  |                                                             |

Text

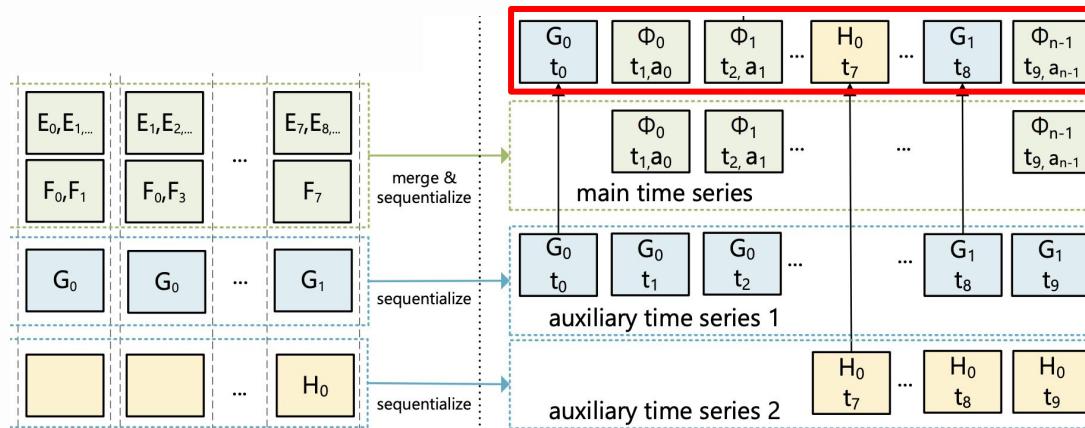


Multimodal

# Inputs for SemID Construction

## Categorical Features

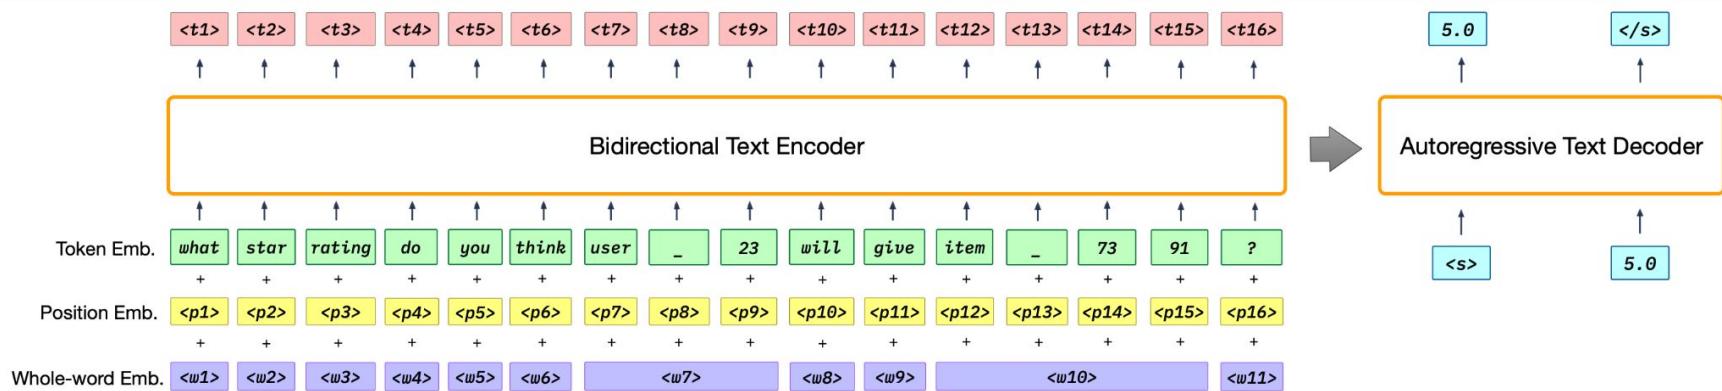
Categorical Features  $\rightarrow$  IDs  
Merge & Sequentialize



# Inputs for SemID Construction

## No Features

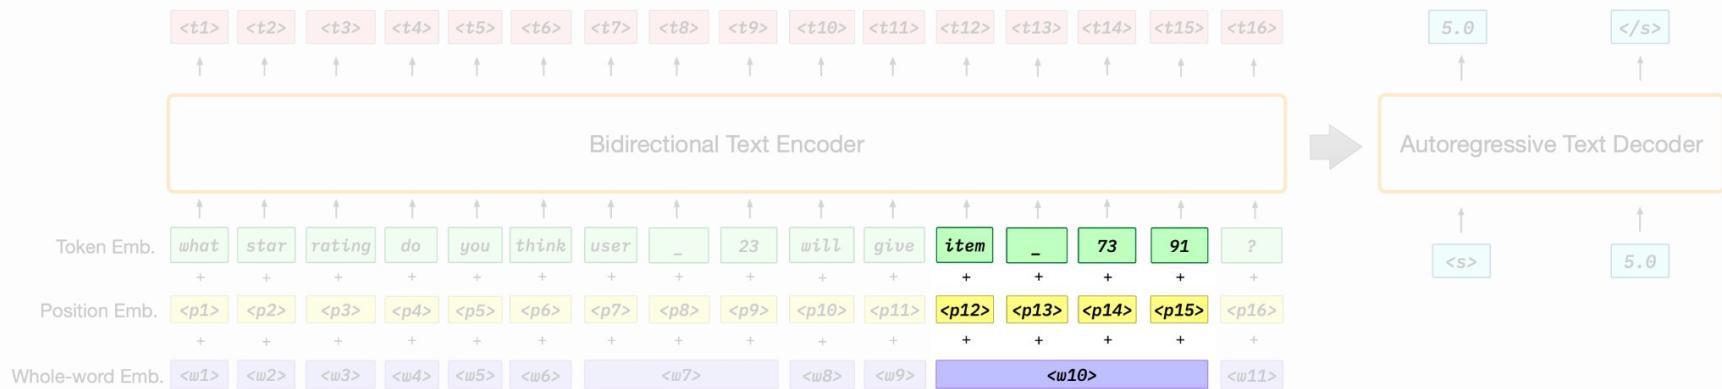
Item ID > IDs  
Text Tokenizer



# Inputs for SemID Construction

## No Features

Item ID > IDs  
Text Tokenizer

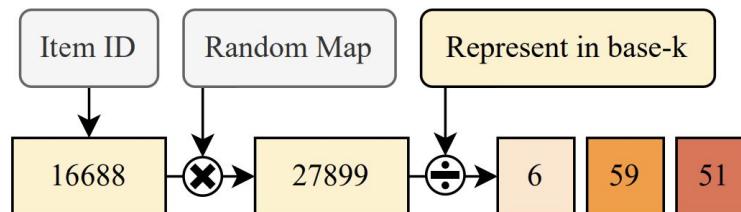


# Inputs for SemID Construction

No Features

Random IDs

Balanced Chunked ID



# Inputs for SemID Construction

**Input: all data associated with the item**

## (1) Item Metadata

Text / Multimodal / Categorical / No Features

# Inputs for SemID Construction

**Input: all data associated with the item**

## (1) Item Metadata

Text / Multimodal / Categorical / No Features

## (2) Item Metadata + Behaviors

# Inputs for SemID Construction

**Input: all data associated with the item**

## (1) Item Metadata

Text / Multimodal / Categorical / No Features

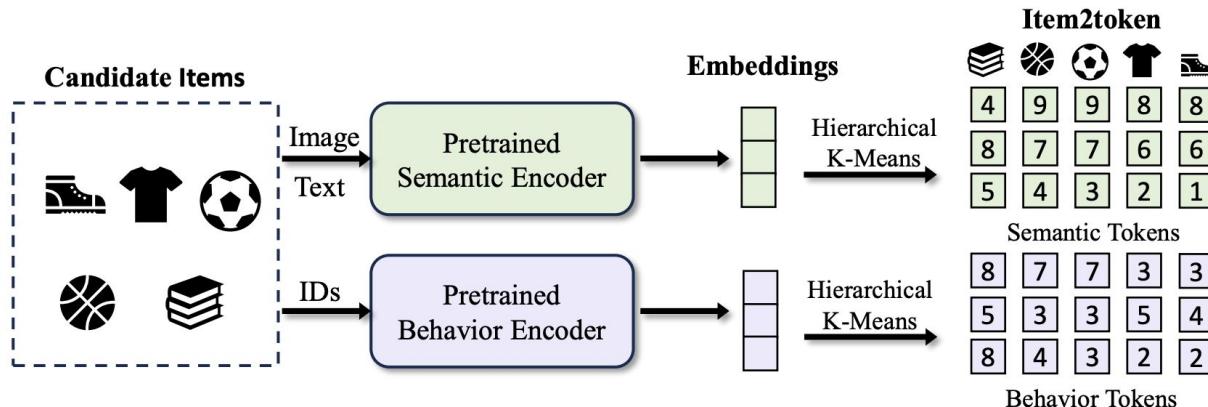
## (2) Item Metadata + Behaviors

But how?

# Inputs for SemID Construction

## Item Metadata + Behaviors

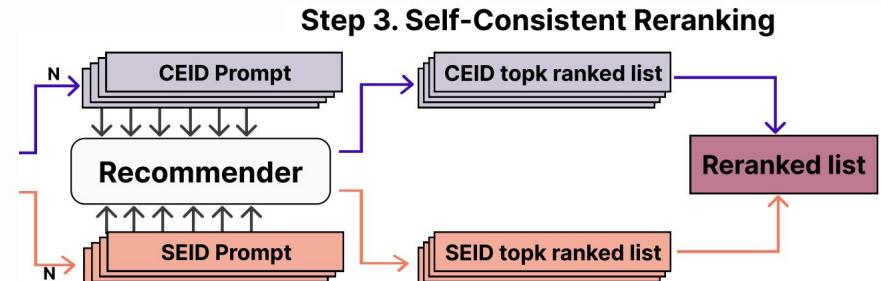
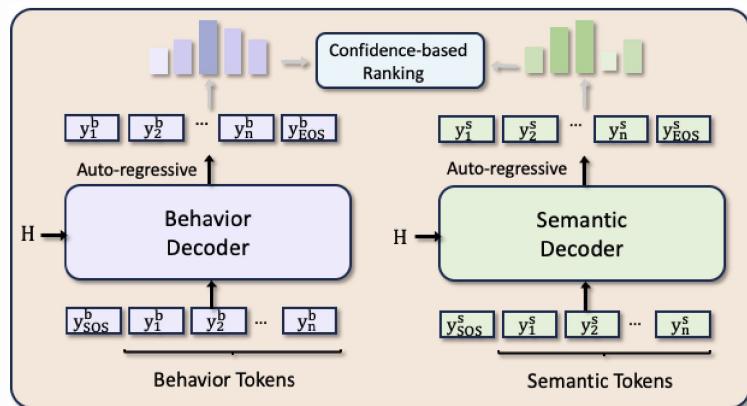
## Fused Semantic IDs



# Inputs for SemID Construction

Item Metadata + Behaviors

Fused Semantic IDs + Two-stream Generation

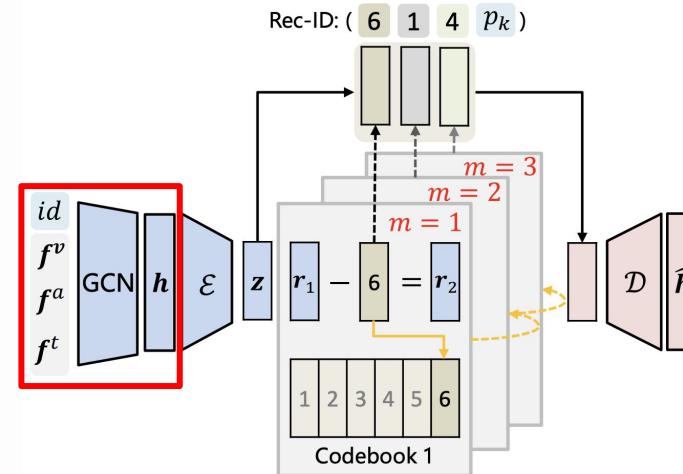


# Inputs for SemID Construction

Item Metadata + Behaviors

Fused Representations

User–Item Graph +  
Semantic Features



# Inputs for SemID Construction

## Item Metadata + Behaviors

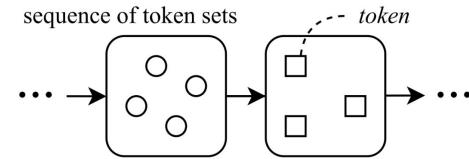
### Train Tokenizer on Behavior Sequence Corpus

#### Algorithm 1 ActionPiece Vocabulary Construction

---

```
input Sequence corpus  $S'$ , initial tokens  $\mathcal{V}_0$ , target size  $Q$ 
output Merge rules  $\mathcal{R}$ , constructed vocabulary  $\mathcal{V}$ 
1: Initialize vocabulary  $\mathcal{V} \leftarrow \mathcal{V}_0$  # each initial token corresponds
   to one unique item feature
2:  $\mathcal{R} \leftarrow \emptyset$ 
3: while  $|\mathcal{V}| < Q$  do
4:   # Count: accumulate weighted token co-occurrences
5:    $\text{count}(\cdot, \cdot) \leftarrow \text{Count}(S', \mathcal{V})$  # Algorithm 2
6:   # Update: Merge a frequent token pair into a new token
7:   Select  $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$ 
8:    $S' \leftarrow \text{Update}(S', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$  # Algorithm 3
9:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$  # new merge rule
10:   $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$  # add new token to the vocabulary
11: end while
return  $\mathcal{R}, \mathcal{V}$ 
```

---



$$P(O, O) = \frac{|O-O|}{|<O, O>|} = \frac{4-1}{\binom{4}{2}}$$
$$P(O, \square) = \frac{|O-\square|}{|O| \times |\square|} = \frac{1}{4 \times 3}$$
$$P(\square, \square) = \frac{|\square-\square|}{|<\square, \square>|} = \frac{3-1}{\binom{3}{2}}$$

Features  
co-occurring  
within or  
across items

# Inputs for SemID Construction

## Item Metadata + Behaviors

### Train Tokenizer on Behavior Sequence Corpus

---

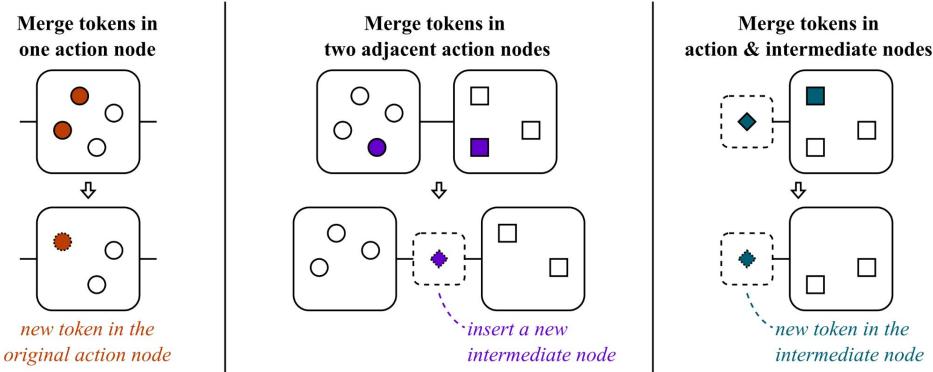
#### Algorithm 1 ActionPiece Vocabulary Construction

---

**input** Sequence corpus  $\mathcal{S}'$ , initial tokens  $\mathcal{V}_0$ , target size  $Q$   
**output** Merge rules  $\mathcal{R}$ , constructed vocabulary  $\mathcal{V}$

- 1: Initialize vocabulary  $\mathcal{V} \leftarrow \mathcal{V}_0$  # each initial token corresponds to one unique item feature
- 2:  $\mathcal{R} \leftarrow \emptyset$
- 3: **while**  $|\mathcal{V}| < Q$  **do**
- 4:   **# Count:** accumulate weighted token co-occurrences
- 5:    $\text{count}(\cdot, \cdot) \leftarrow \text{Count}(\mathcal{S}', \mathcal{V})$  # Algorithm 2
- 6:   **# Update:** Merge a frequent token pair into a new token
- 7:   Select  $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$
- 8:    $\mathcal{S}' \leftarrow \text{Update}(\mathcal{S}', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$  # Algorithm 3
- 9:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$  # new merge rule
- 10:    $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$  # add new token to the vocabulary
- 11: **end while**
- return**  $\mathcal{R}, \mathcal{V}$

---

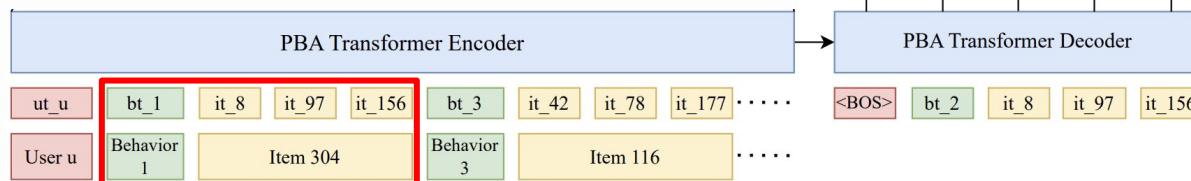


# Inputs for SemID Construction

Item Metadata + Behaviors

Multi-Behavior Recommendation

Semantic IDs fused  
with **behavior types**



Unified Multi-Task Framework

Target Behavior Item Prediction

bt\_1 it\_? it\_? it\_?

Behavior-Specific Item Prediction

bt\_\* it\_? it\_? it\_?

Behavior-Item Prediction

bt\_? it\_? it\_? it\_?

Behavior Prediction

bt\_? [ ] [ ] [ ]

Next Behavior Next Item

bt\_2 it\_8 it\_97 it\_156 <EOS>

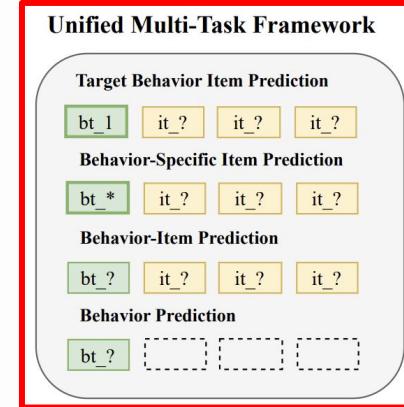
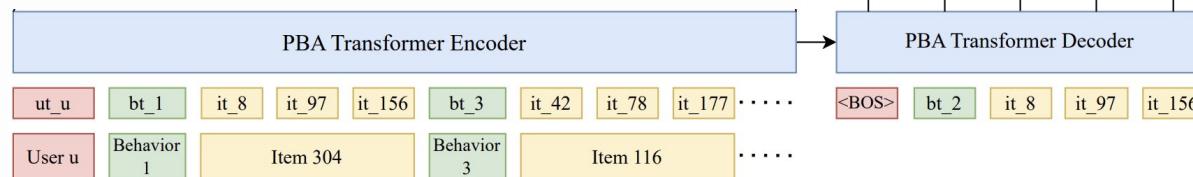
# Inputs for SemID Construction

Item Metadata + Behaviors

Multi-Behavior Recommendation

Next Token Prediction as  
natural **multi-task learning**

(prompted by behavior type)



# Inputs for SemID Construction

**Input: all data associated with the item**

## (1) Item Metadata

Text / Multimodal / Categorical / No Features

## (2) Item Metadata + Behaviors

Fused semantic IDs & Representations

Tokenizer trained on behavior sequences

# **Part 1 Summary – SemID Construction**

**(1) First Example: TIGER**

**(2) Construction Techniques**

**(3) Inputs**

# **Part 1 Summary – SemID Construction**

**(1) First Example: TIGER**

**(2) Construction Techniques**

Context-independent (PQ, RQ, Clustering,  
LM-based generator) → Context-aware

**(3) Inputs**

# **Part 1 Summary – SemID Construction**

**(1) First Example: TIGER**

**(2) Construction Techniques**

Context-independent (PQ, RQ, Clustering,  
LM-based generator) → Context-aware

**(3) Inputs**

Item Metadata (Text, Multimodal, Features)

+ Behaviors (Fused SemIDs / Representations)

# **Part 2: SemID-based Generative Recommendation Model Architecture**

# SemID-based Recommender Architecture

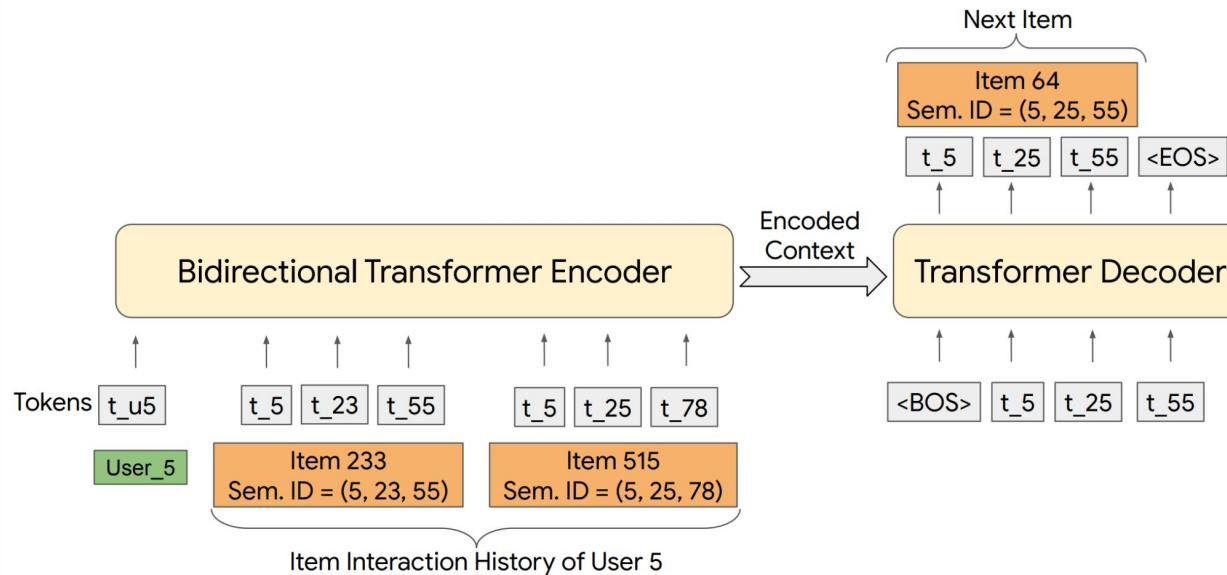
Recommendation as a **seq**-to-**seq** generation problem

**Input:** user interacted items  $\{c_{11}, c_{12}, c_{13}, c_{14}, c_{21}, c_{22}, \dots\}$

**Output:** next item  $\{c_{t1}, c_{t2}, c_{t3}, c_{t4}\}$

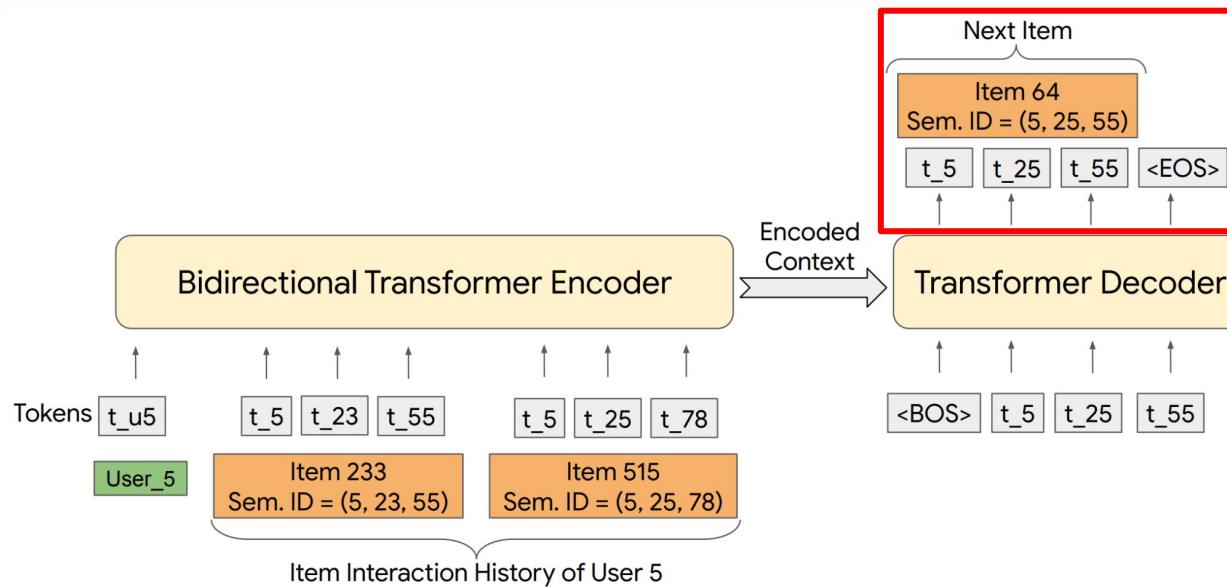
# SemID-based Recommender Architecture

## Architecture: Decoder-Only / Encoder-Decoder



# SemID-based Recommender Architecture

**Objective:** Next-Token Prediction



# SemID-based Recommender Architecture

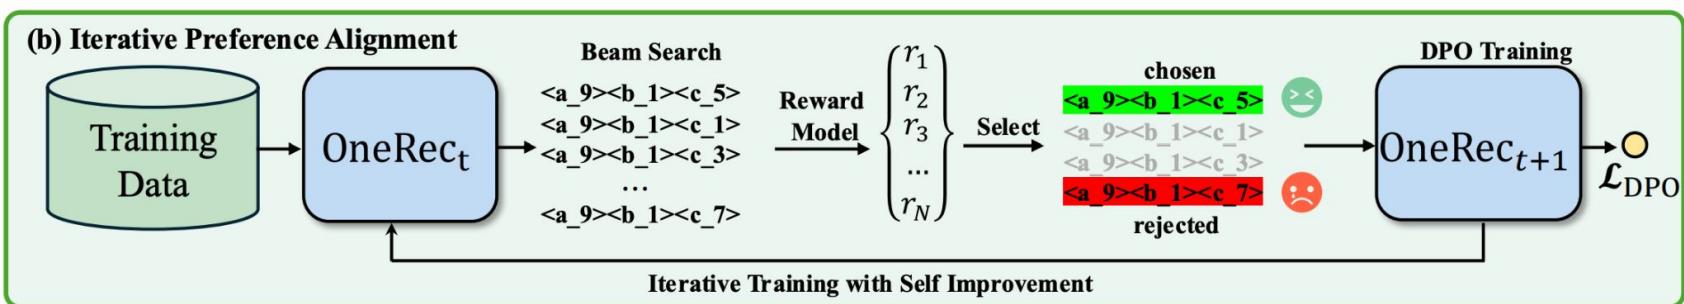
**Objective:** Next-Token Prediction

Could we add **negative samples** like BPR loss?

# SemID-based Recommender Architecture

**Objective:** Preference Alignment Objective

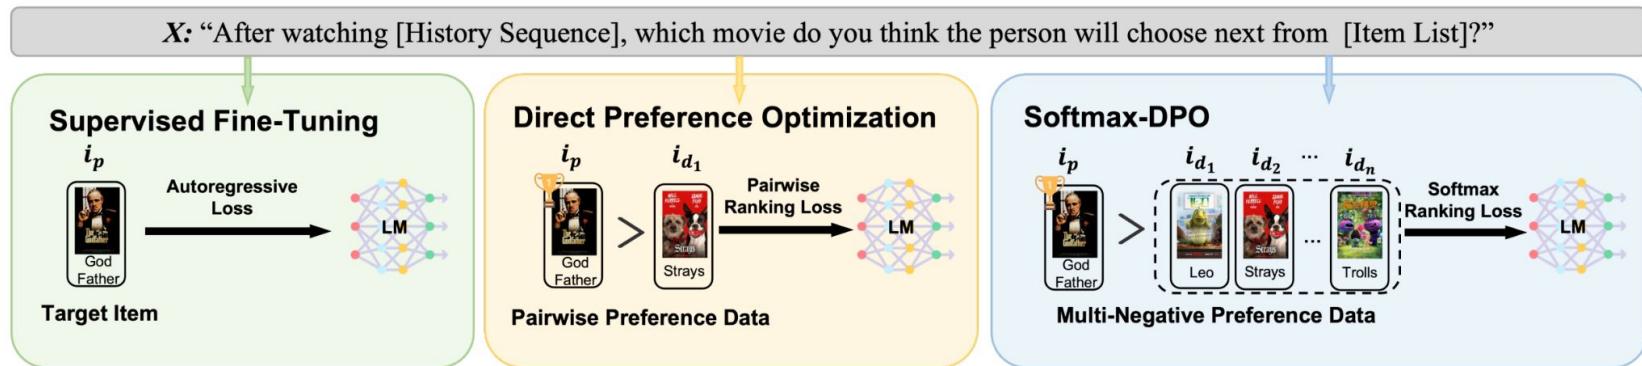
One negative sample per instance



# SemID-based Recommender Architecture

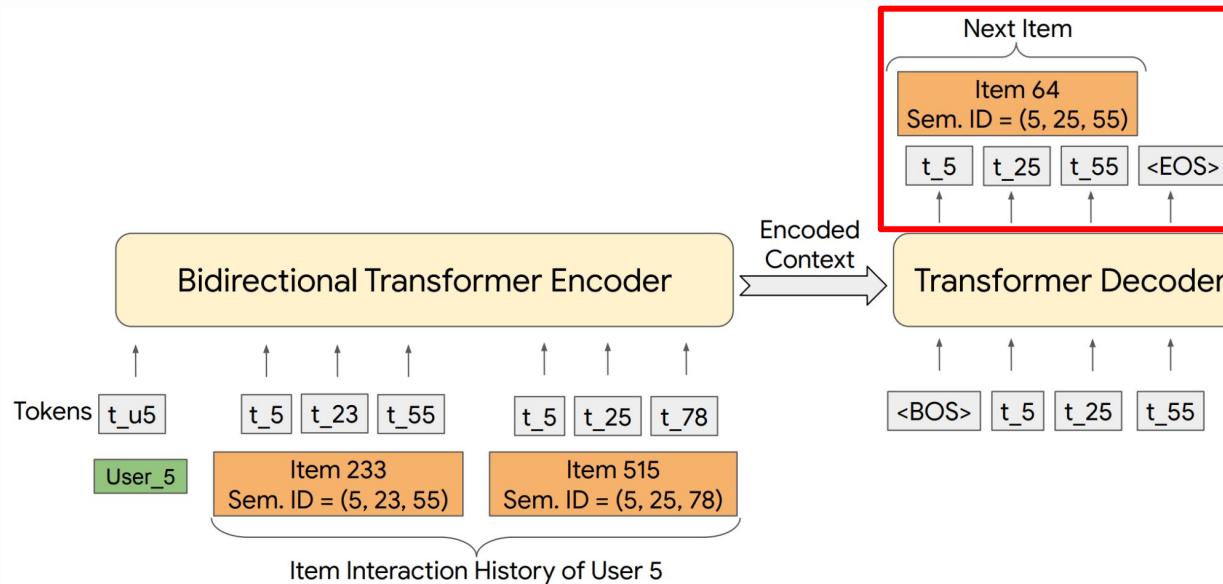
**Objective:** Preference Alignment Objective

Multiple negative samples per instance



# SemID-based Recommender Architecture

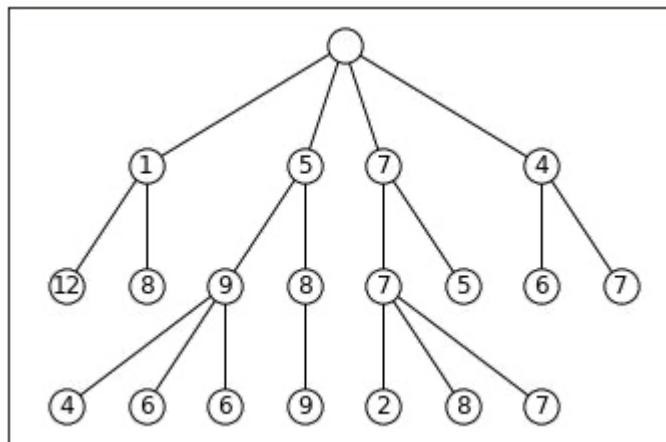
**Inference:** How to get a **ranking list**?



# SemID-based Recommender Architecture

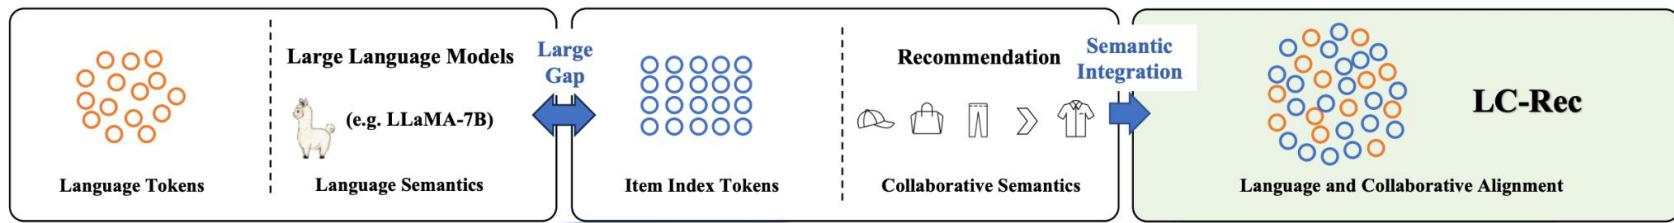
**Inference:** How to get a **ranking list**?

(Constrained) Beam Search



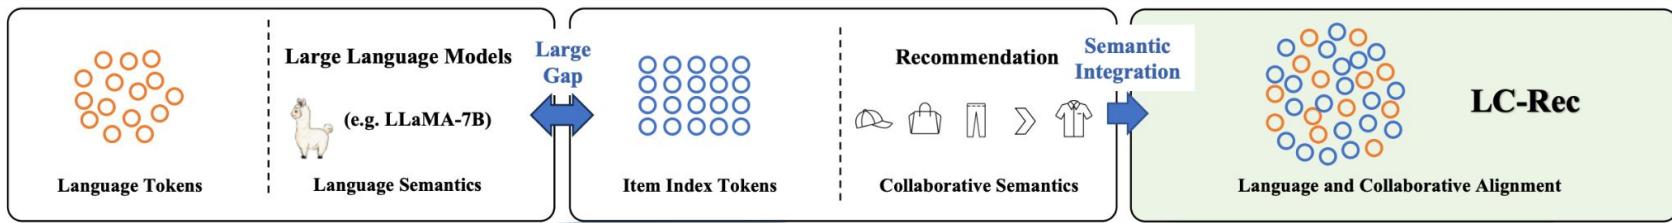
# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec



# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec

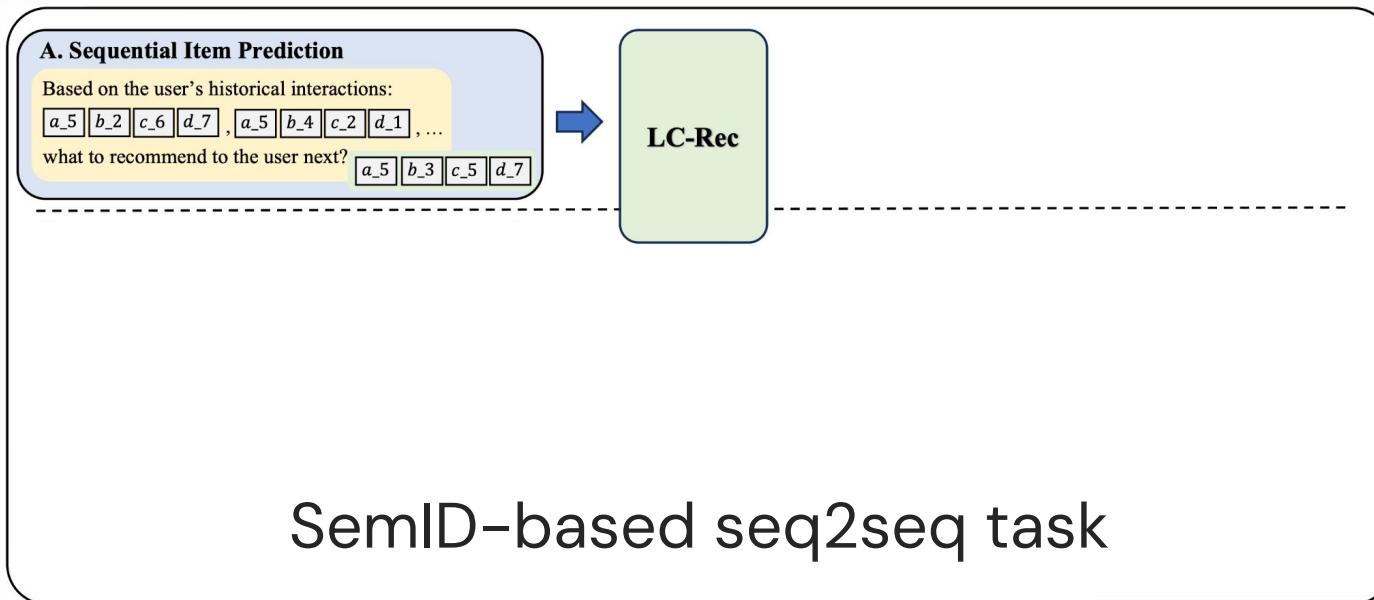


Core Idea:

Construct **instructions** containing  
both **Semantic IDs** and **language tokens**

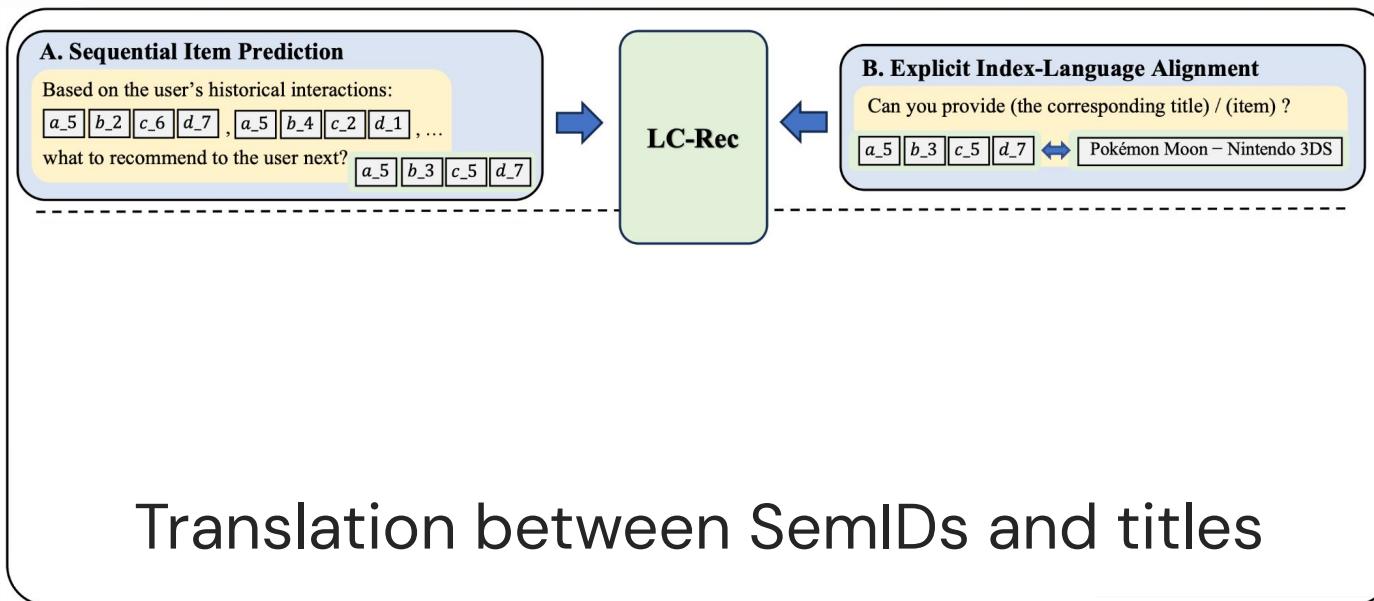
# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec



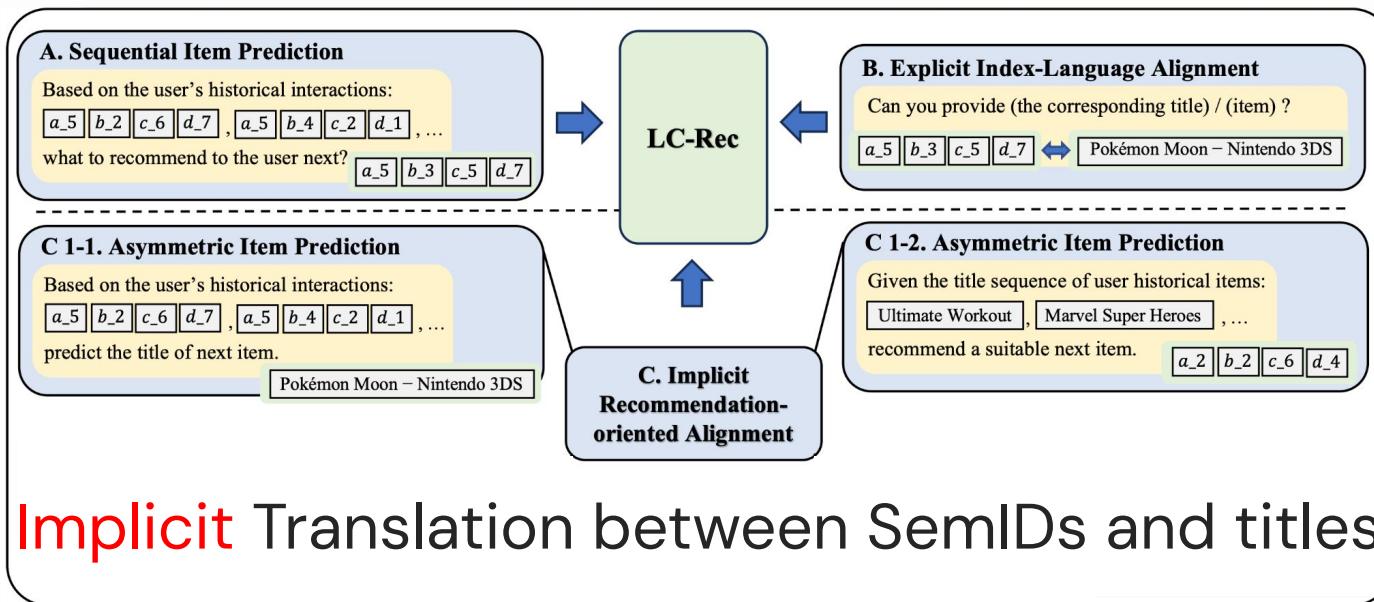
# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec



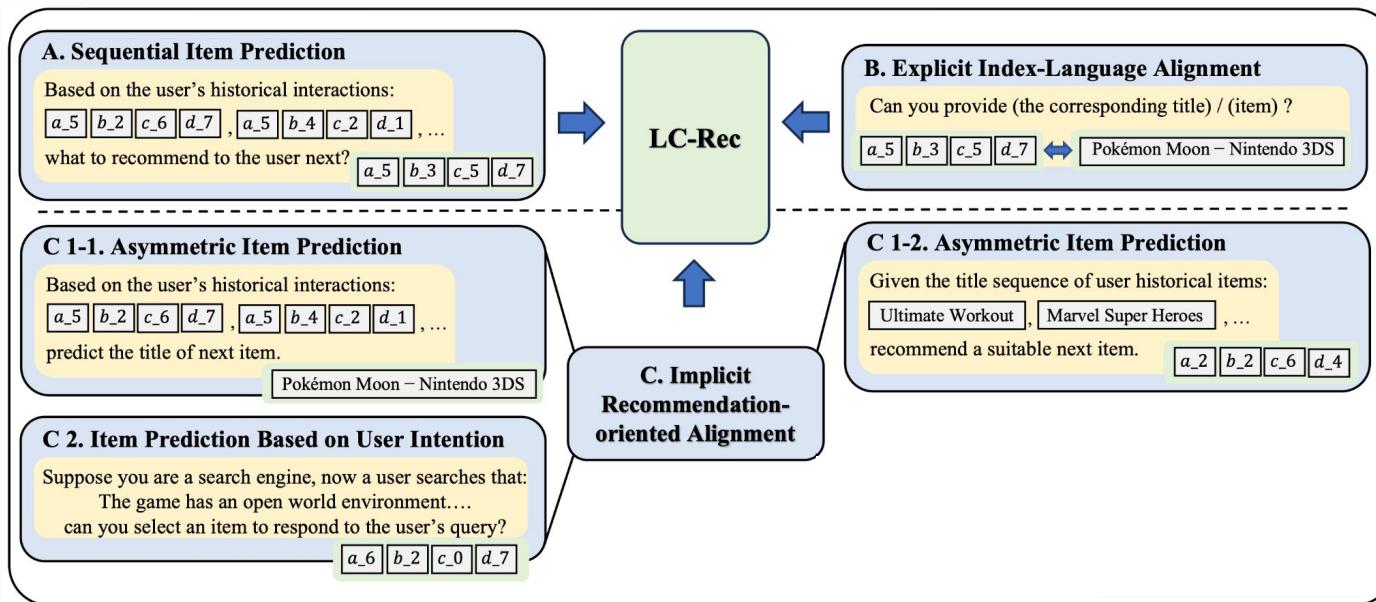
# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec



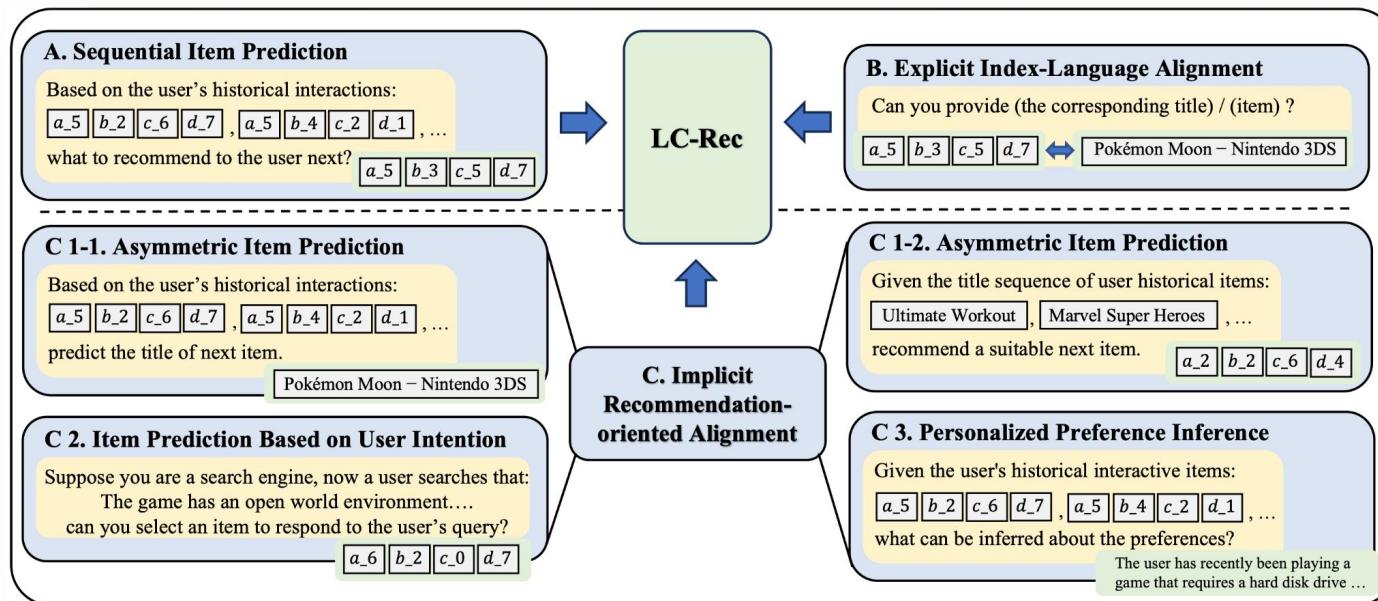
# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec



# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec



# **Part 2 Summary – Architecture**

**(1) Train from Scratch**

**(2) Align with LLMs**

# **Part 2 Summary – Architecture**

## **(1) Train from Scratch**

Objective (NTP, DPO, S-DPO)

Inference (Beam Search)

## **(2) Align with LLMs**

# **Part 2 Summary – Architecture**

## **(1) Train from Scratch**

Objective (NTP, DPO, S-DPO)

Inference (Beam Search)

## **(2) Align with LLMs**

LC-Rec: Instructions containing both semIDs and language tokens

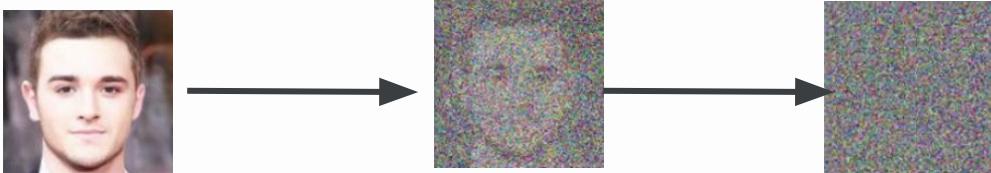
04

# Diffusion Model

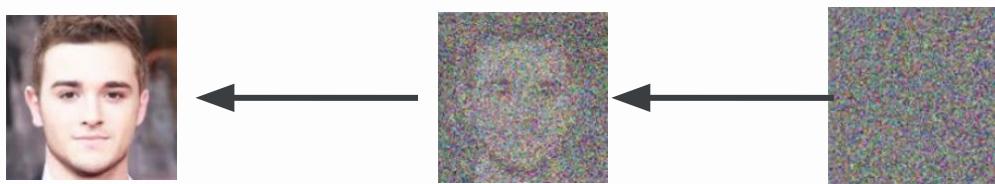
-based Generative Recommendation

# What is Diffusion

Forward  
process

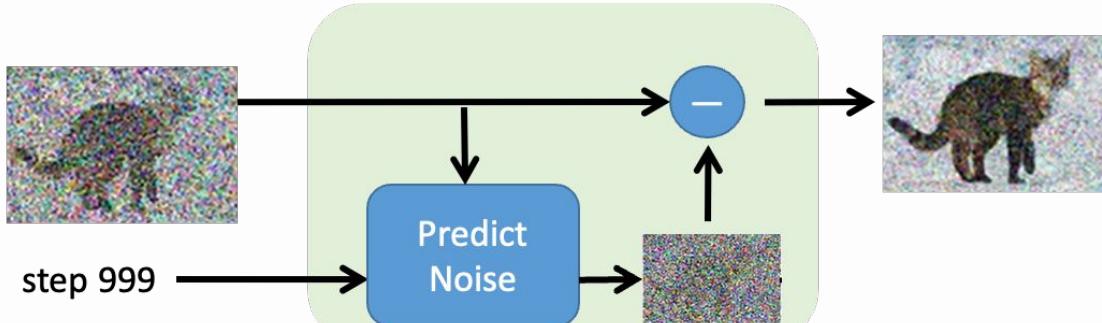


Reverse  
process



Build the mapping between data sample and Gaussian sample

# What is Diffusion



---

## Algorithm 1 Training

```
1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$ 
4:    $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take gradient descent step on
         $\nabla_{\theta} \|\boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, t)\|^2$ 
6: until converged
```

---

---

## Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\bar{\alpha}_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```

---

Remove the noise step by step from a Gaussian sample.

# Diffusion in CV

**Diffusion is at the core of visual content generation.**

Image generation

Stable Diffusion, DALL-E...



Video generation

Sora, Hunyuan–Video, Keling...



# Diffusion for recommendation

## Use diffusion to **enhance** traditional recommender

- More robust representation
- Data augmentation

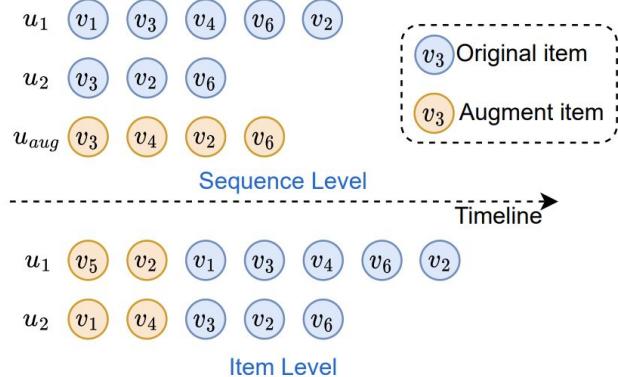
## Diffusion **as recommender**

- Diffuse on the user interaction vector
- Diffuse on item representation
- Discrete diffusion

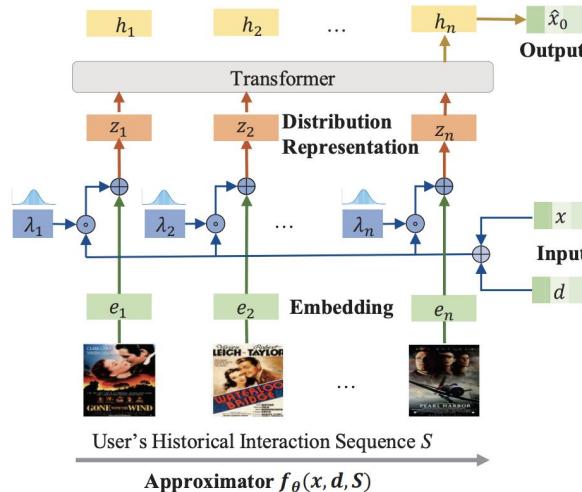
## Diffusion for **personalized content** generation

- Personalized try-on, image,....

# Diffusion as enhancer



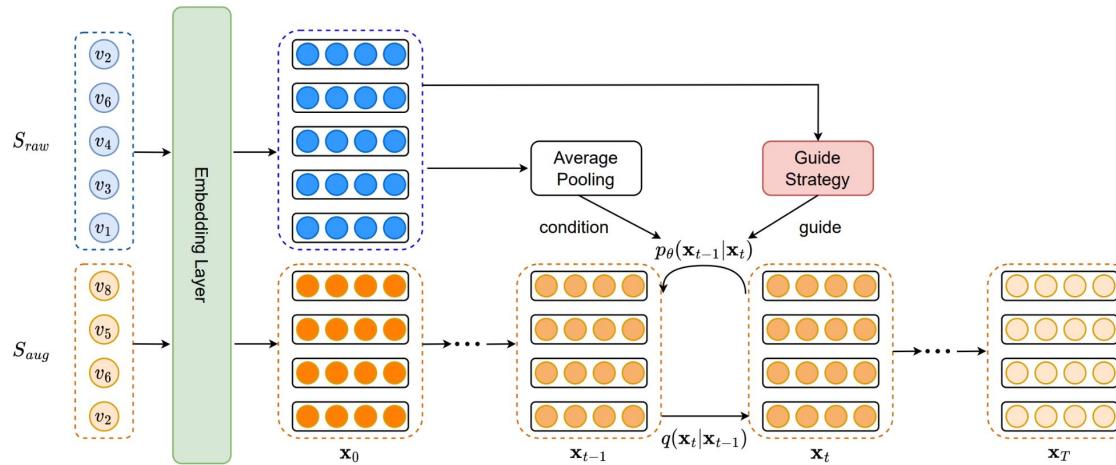
Generate more interaction  
or sequences



Enhance the robustness of  
embeddings

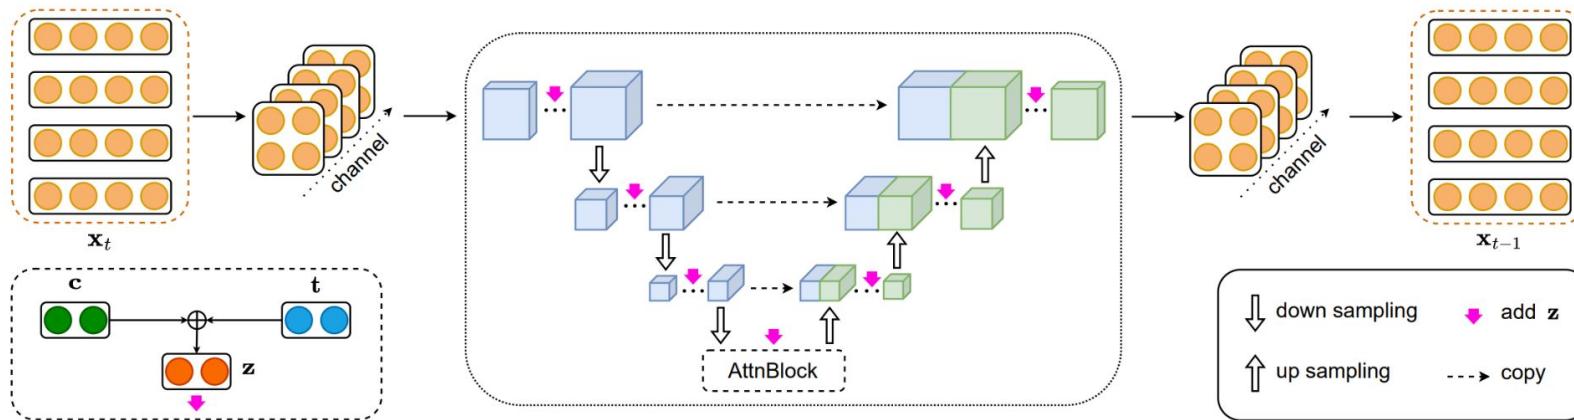
# Pseudo sequence generation (I)

Generate pseudo sequence embeddings  
conditioned on historical interaction sequence



# Pseudo sequence generation (II)

The model architecture is adopted from U-Net



# Diffusion for recommendation

Use diffusion to enhance traditional recommender

- More robust representation
- Data augmentation

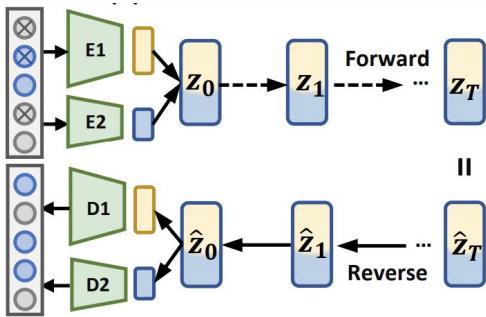
**Diffusion as recommender**

- Diffuse on the user interaction vector
- Diffuse on item representation
- Discrete diffusion

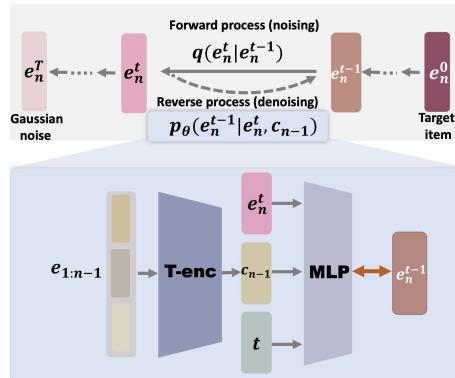
Diffusion for personalized content generation

- Personalized try-on, image,....

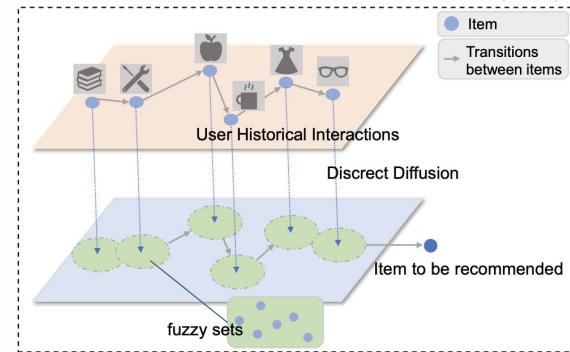
# Diffusion as recommender



Diffuse on the user interaction vector



Diffuse on item representation

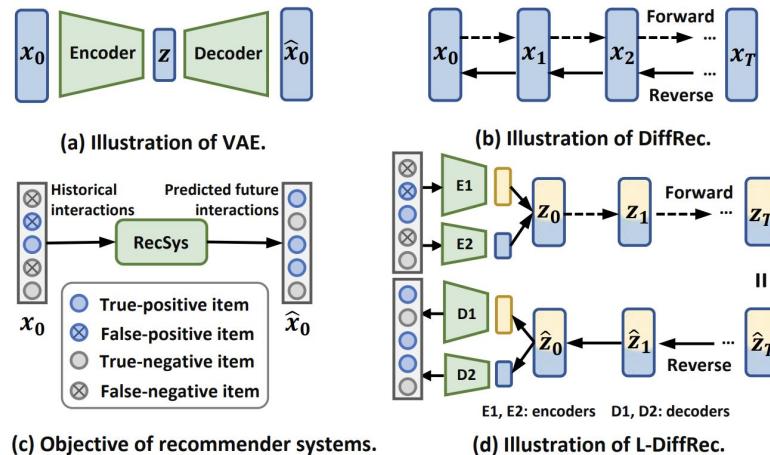


Discrete diffusion

# Interaction vector completion (I)

Motivation – limitation of GANs and VAEs:

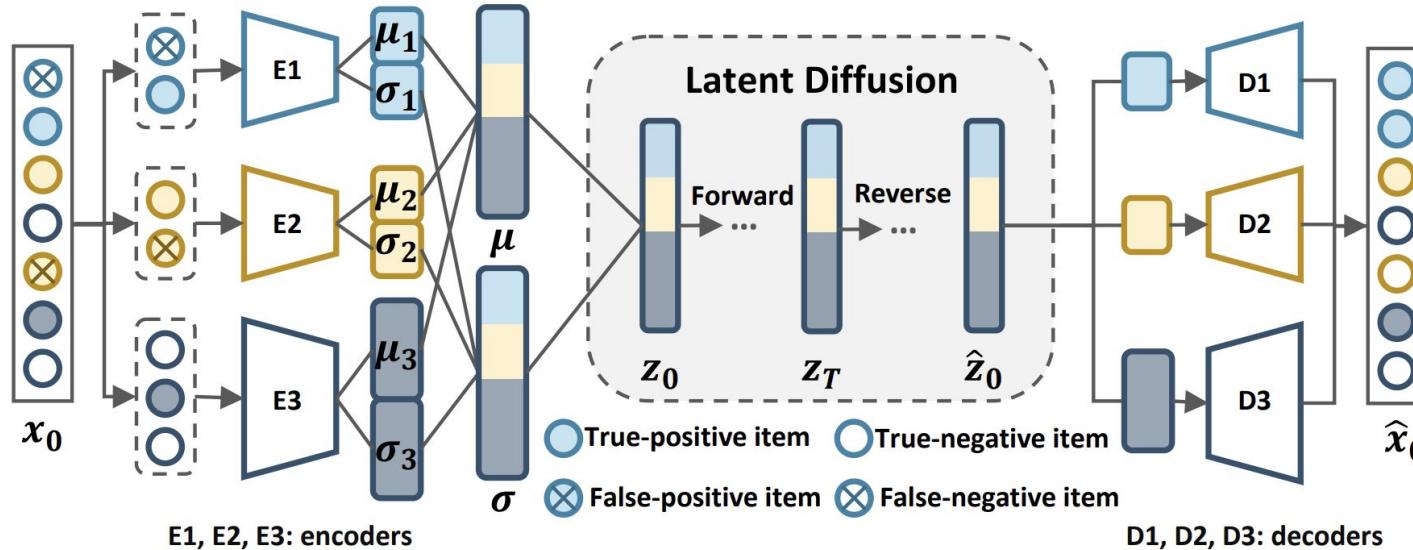
GAN- and VAE-based recommenders suffers from issues like **instability** and **representation collapse**.



# Interaction vector completion (II)

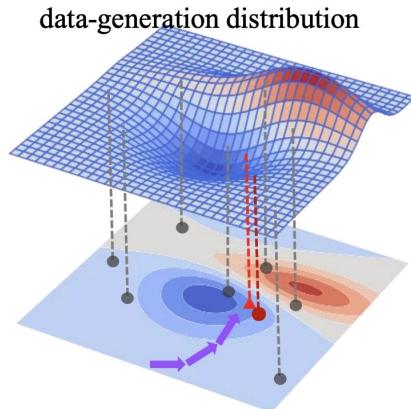
Forward: corrupt the interaction vector into gaussian noise

Reverse: recover the interaction vector from the gaussian



# Generate item embedding

There exists an implicit distribution, from which target item embedding can be generated.



Learning-to-generate Paradigm

Challenge:

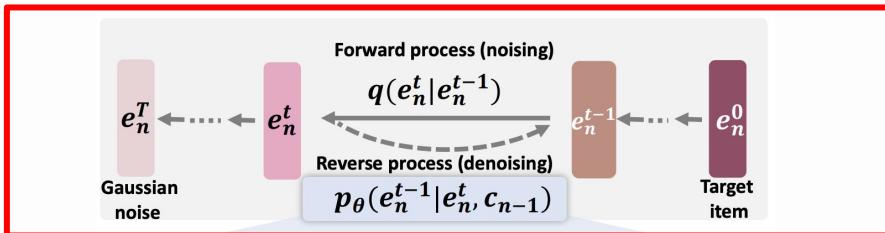
- The data-generation distribution is complicated and unknown.

Solution:

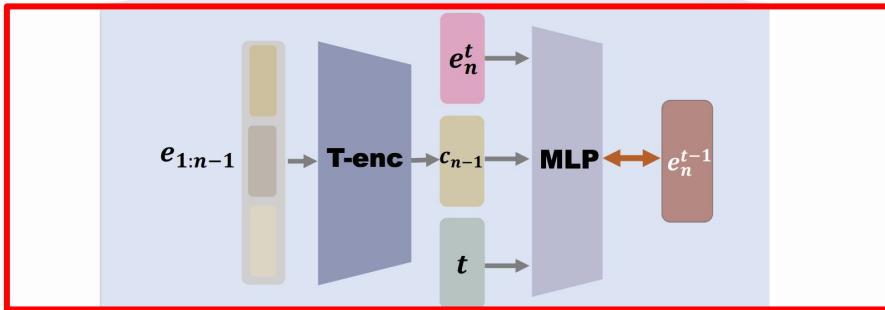
- Capture the data-generation distribution by connecting it with Gaussian distribution.
- This can be achieved by diffusion.

# Generate item embedding

- Diffusion on target item embeddings.
- Guided by user interaction sequence for personalization.



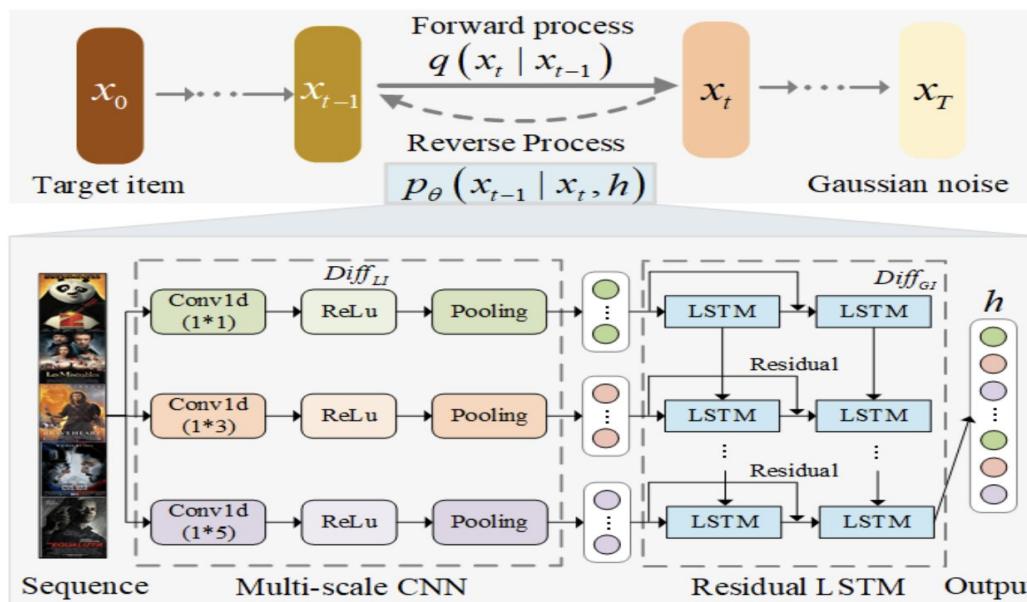
Diffusion  
process



Guidance

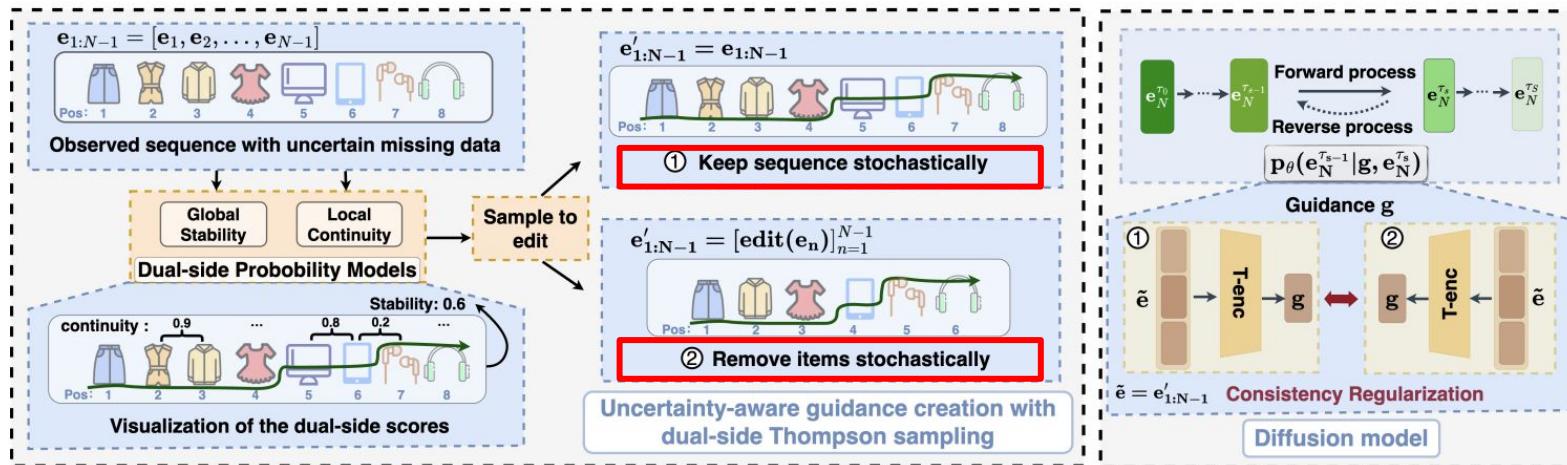
# Generate item embedding

- Different sequence encoder



# Generate item embedding

- Uncertainty-aware guidance



# Generate item embedding

- Incorporate preference optimization

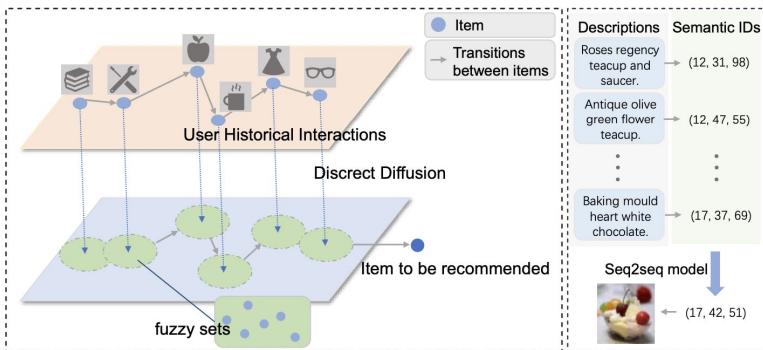
$$\mathcal{L}_{\text{Simple}} = \mathbb{E}_{(\mathbf{e}_0^+, \mathbf{c}, t)} \left[ \left\| \mathcal{F}_\theta(\mathbf{e}_t^+, t, \mathcal{M}(\mathbf{c})) - \mathbf{e}_0^+ \right\|_2^2 \right],$$

$$\mathcal{L}_{\text{BPR-Diff-C}} = -\log \sigma(-|\mathcal{H}| \cdot [S(\hat{\mathbf{e}}_0^+, \mathbf{e}_0^+) - S(\mathcal{F}_\theta(\bar{\mathbf{e}}_t^-, t, \mathcal{M}(\mathbf{c})), \bar{\mathbf{e}}_0^-)]).$$

$$\mathcal{L}_{\text{PerferDiff}} = \underbrace{\lambda \mathcal{L}_{\text{Simple}}}_{\text{Learning Generation}} + \underbrace{(1 - \lambda) \mathcal{L}_{\text{BPR-Diff-C}}}_{\text{Learning Preference}}.$$

# Discrete diffusion

State transitions occur under discrete conditions for the entire interaction sequence.



- Represent interaction sequence as one-hot vector through semantic ID.
- Conduct discrete diffusion on interaction sequence.

# Discrete diffusion

---

**Algorithm 1** Training of DDSR.

**Input:** historical interaction sequence  $v_{1:n-1} = c_{1:n-1;1:m}$ ; target item  $v_n = c_{n;1:m}$ ; transition matrix  $\mathbf{Q}_t$ ; Approximator  $f_\theta(\cdot)$ .

**Output:** well-trained Approximator  $f_\theta(\cdot)$ .

While not converged do:

1: Sample Diffusion Time:  $t \sim [0, 1, \dots, T]$ ;

2: Calculate  $t$ -step transition probability:  $\overline{\mathbf{Q}}_t = \mathbf{Q}_1 \mathbf{Q}_2 \cdots \mathbf{Q}_t$ ;  $[\mathbf{Q}]_{ij} = \begin{cases} (1 - \beta_t)/(|\mathcal{V}| - 1) & \text{if } i \neq j \\ \beta_t & \text{if } i = j \end{cases}$ .

3: Convert  $c_{n;1:m}$  to one-hot encoding  $\mathbf{x}_{n;1:m}^0$ ;

4: Obtain the discrete state  $x_{n;1:m}^t$  after  $t$  steps by Equation 2, thereby obtaining the 'fuzzy set'  
 $c_{1:n-1;1:m}^t$ ;

5: Modeling  $c_{2:n;1:m}$  based on 'fuzzy sets' through Equation 5;  $\hat{c}_{2:n;1:m} = f_\theta(c_{1:n-1;1:m}^t, t)$ .

6: Take gradient descent step on  $\nabla L_{CE}(\hat{c}_{2:n;1:m}, c_{2:n;1:m})$ .

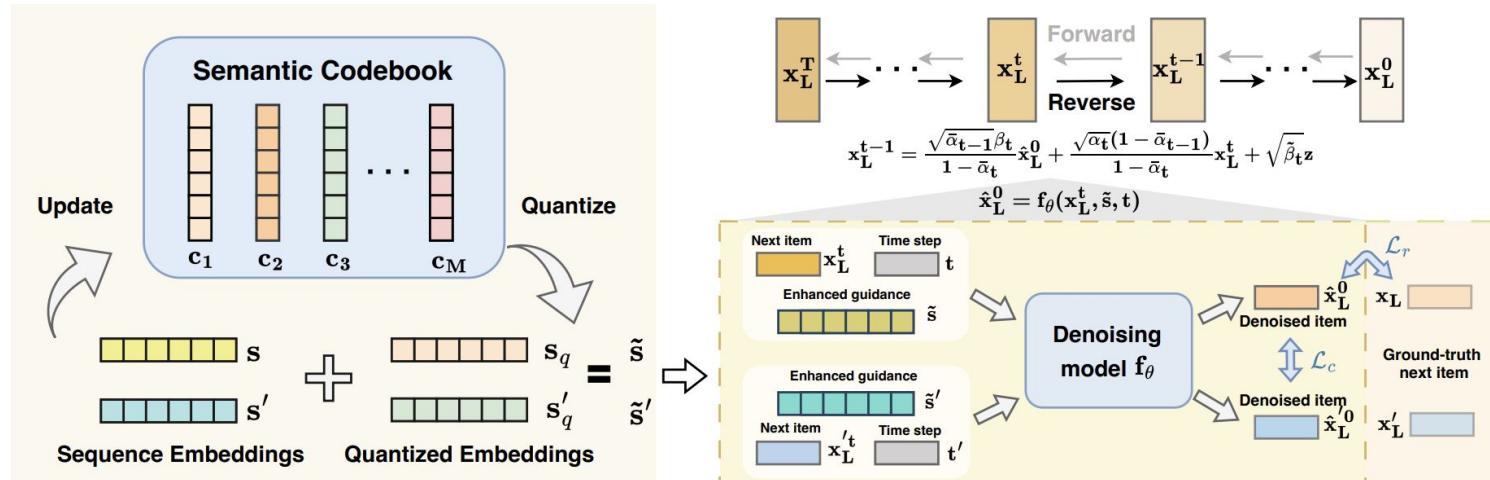
---

Semantic IDs

Forward process

# Discrete diffusion

- Quantization embedding with continuous diffusion.



# Diffusion for recommendation

Use diffusion to enhance traditional recommender

- More robust representation
- Data augmentation

Diffusion as recommender

- Diffuse on the user interaction vector
- Diffuse on item representation
- Discrete diffusion

Diffusion for **personalized content** generation

- Personalized try-on, image,....

# Personalized content generation



Personalized try-on

*A photo of  $\hat{V}$  woman  
shaking hands with Joe Biden*



*A photo of  $\hat{V}$  woman  
piloting a fight jet*



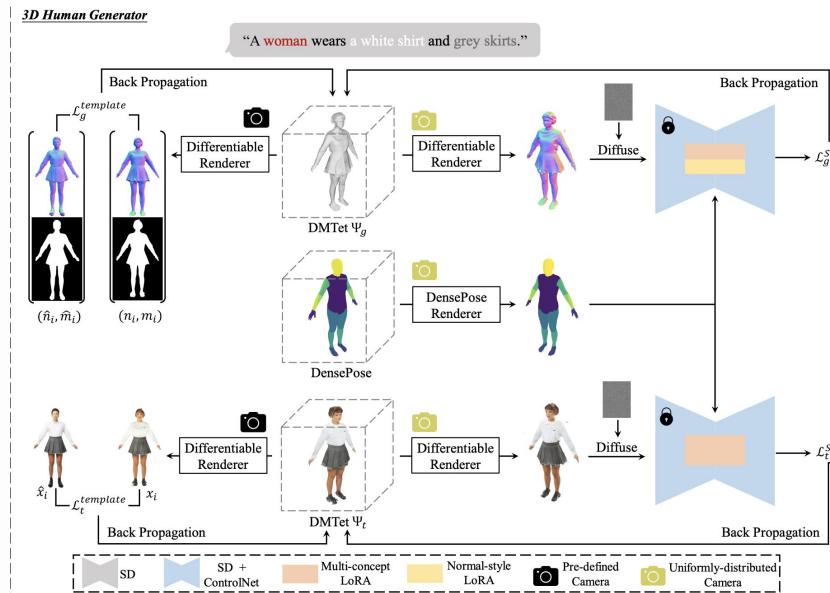
*A photo of mysterious  
 $\hat{V}$  woman witcher at  
night*



Personalized image

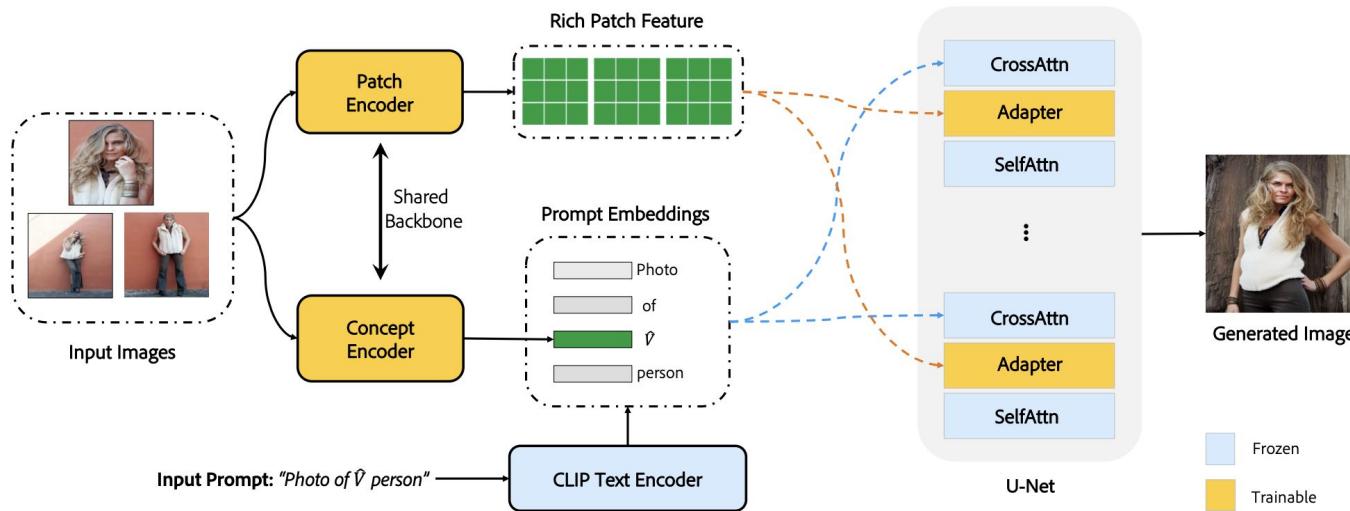
# Personalized Try-on

Generate realistic 3D try-on given person images, clothes images, and a text prompt.



# Personalized Image

Generate personalized image given person images and the desired concept.



05

# Summary

Open Challenges and Beyond

# Summary

## Scaling Law:

- **larger model + larger dataset → better performance**

## Most large models are generative

- (LLMs, Text2Video Models)

# Summary

## Scaling Law:

- **larger model + larger dataset → better performance**

## Most large models are generative

- (LLMs, Text2Video Models)



**Large generative rec models?**

# Summary



## How to get a large generative rec model?

- Pre-trained model (e.g., LLMs) → Adaptation;
- From scratch;

Paradigms

Adapt Pre-trained Models

Train from Scratch

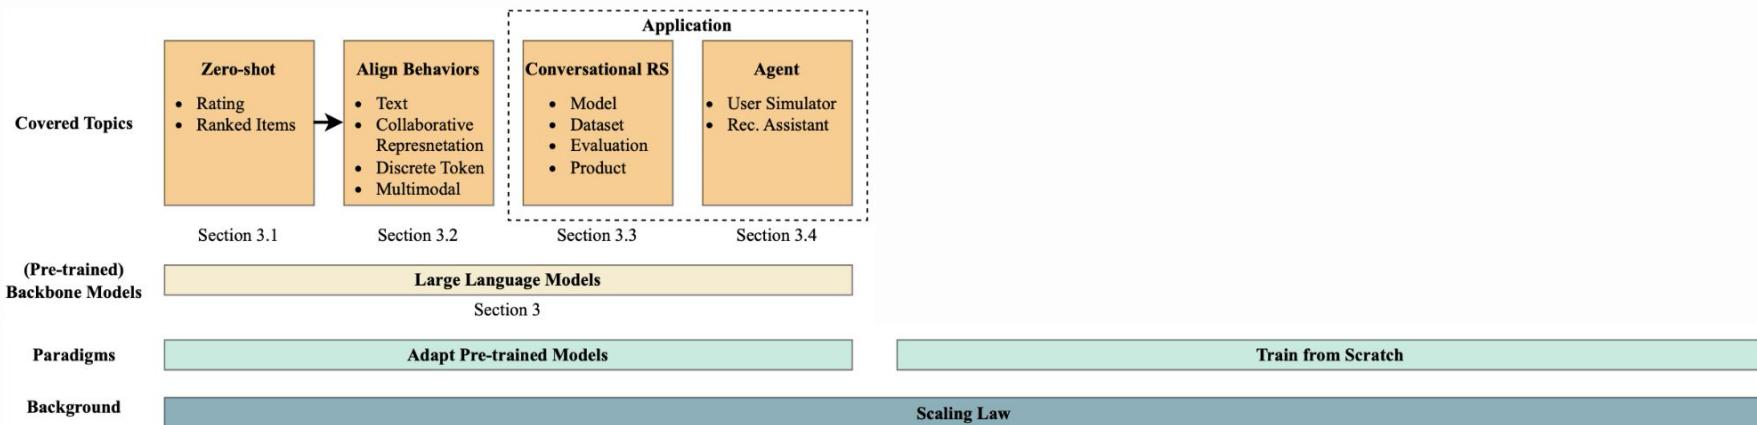
Background

Scaling Law

# Summary

## Adaptation

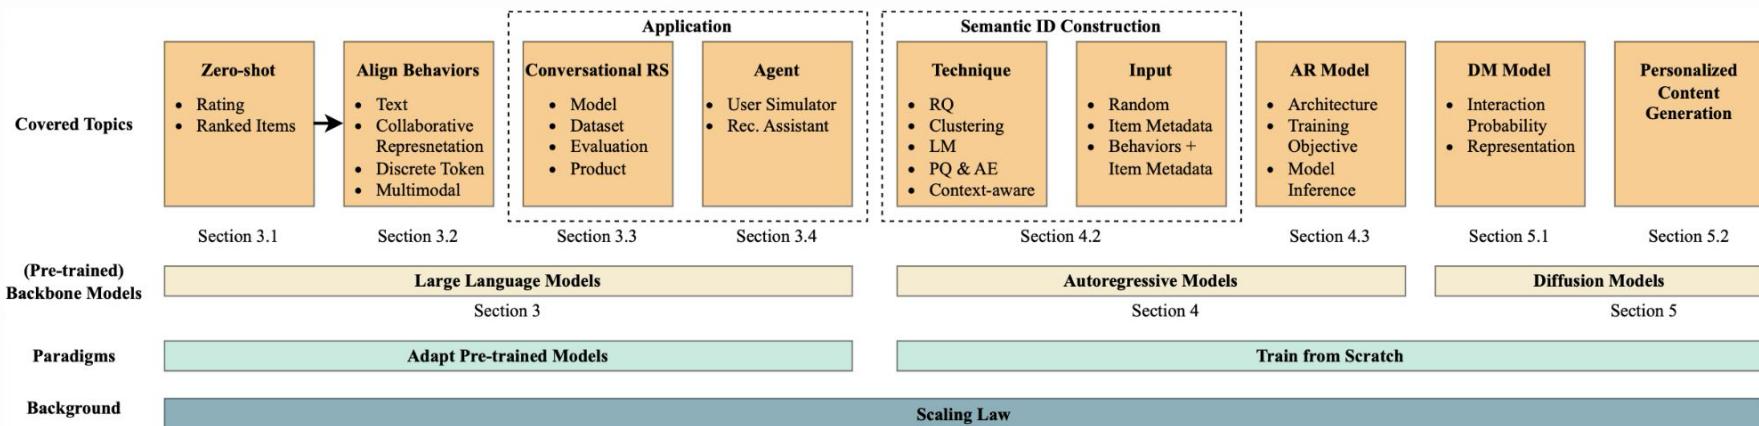
Mainly LLM-based recommendations



# Summary

## From Scratch

- Autoregressive models (e.g., semantic ID-based);
- Diffusion models;



# Summary

## Open Challenges

| Open Challenges               | Inference Efficiency                                                                                                                                                                                                                                      | Model Updating                                                                                                                                                                                                                                             | Item Tokenization                                                                                                                                                                                                                                                                                          | Emergent Ability             | Test-time Scaling & Reasoning                                                                                                   | Unified Retrieval and Ranking                                                                                  |                                        |
|-------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------|---------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|----------------------------------------|
|                               | Section 6.1                                                                                                                                                                                                                                               | Section 6.2                                                                                                                                                                                                                                                | Section 6.3                                                                                                                                                                                                                                                                                                | Section 6.4                  | Section 6.5                                                                                                                     | Section 6.6                                                                                                    |                                        |
| Covered Topics                | <b>Zero-shot</b> <ul style="list-style-type: none"><li>Rating</li><li>Ranked Items</li></ul> → <b>Align Behaviors</b> <ul style="list-style-type: none"><li>Text</li><li>Collaborative Representation</li><li>Discrete Token</li><li>Multimodal</li></ul> | <b>Application</b><br><b>Conversational RS</b> <ul style="list-style-type: none"><li>Model</li><li>Dataset</li><li>Evaluation</li><li>Product</li></ul> <b>Agent</b> <ul style="list-style-type: none"><li>User Simulator</li><li>Rec. Assistant</li></ul> | <b>Semantic ID Construction</b><br><b>Technique</b> <ul style="list-style-type: none"><li>RQ</li><li>Clustering</li><li>LM</li><li>PQ &amp; AE</li><li>Context-aware</li></ul> <b>Input</b> <ul style="list-style-type: none"><li>Random</li><li>Item Metadata</li><li>Behaviors + Item Metadata</li></ul> |                              | <b>AR Model</b> <ul style="list-style-type: none"><li>Architecture</li><li>Training Objective</li><li>Model Inference</li></ul> | <b>DM Model</b> <ul style="list-style-type: none"><li>Interaction Probability</li><li>Representation</li></ul> | <b>Personalized Content Generation</b> |
|                               | Section 3.1                                                                                                                                                                                                                                               | Section 3.2                                                                                                                                                                                                                                                | Section 3.3                                                                                                                                                                                                                                                                                                | Section 4.2                  | Section 4.3                                                                                                                     | Section 5.1                                                                                                    | Section 5.2                            |
| (Pre-trained) Backbone Models | <b>Large Language Models</b>                                                                                                                                                                                                                              |                                                                                                                                                                                                                                                            |                                                                                                                                                                                                                                                                                                            | <b>Autoregressive Models</b> |                                                                                                                                 | <b>Diffusion Models</b>                                                                                        |                                        |
|                               | Section 3                                                                                                                                                                                                                                                 |                                                                                                                                                                                                                                                            |                                                                                                                                                                                                                                                                                                            | Section 4                    |                                                                                                                                 | Section 5                                                                                                      |                                        |
| Paradigms                     | <b>Adapt Pre-trained Models</b>                                                                                                                                                                                                                           |                                                                                                                                                                                                                                                            |                                                                                                                                                                                                                                                                                                            | <b>Train from Scratch</b>    |                                                                                                                                 |                                                                                                                |                                        |
| Background                    | <b>Scaling Law</b>                                                                                                                                                                                                                                        |                                                                                                                                                                                                                                                            |                                                                                                                                                                                                                                                                                                            |                              |                                                                                                                                 |                                                                                                                |                                        |

# Open Challenges

## Part 1: What becomes harder?

Comparing to traditional RecSys, what challenges may large generative models face?

Open Challenges

Inference Efficiency

Section 6.1

Model Updating

Section 6.2

Item Tokenization

Section 6.3

Emergent Ability

Section 6.4

Test-time Scaling & Reasoning

Section 6.5

Unified Retrieval and Ranking

Section 6.6

# Open Challenges

## Part 1: What becomes harder?

Comparing to traditional RecSys, what challenges may large generative models face?

## Part 2: What becomes possible?

What new opportunities may large generative models unlock for recommender systems?

### Open Challenges

#### Inference Efficiency

Section 6.1

#### Model Updating

Section 6.2

#### Item Tokenization

Section 6.3

#### Emergent Ability

Section 6.4

#### Test-time Scaling & Reasoning

Section 6.5

#### Unified Retrieval and Ranking

Section 6.6

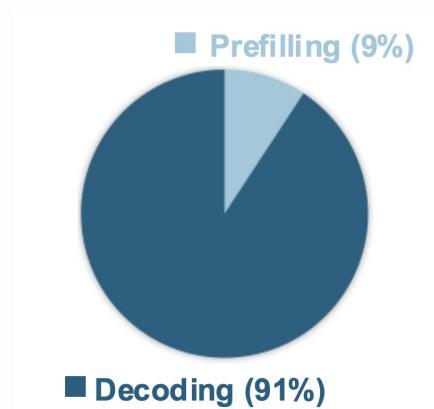
# **Part 1: What Becomes Harder?**

Comparing to traditional RecSys, what challenges may large generative models face?

# Inference Efficiency

Retrieval Models: **K Nearest Neighbor Search**

Generative Models (e.g., AR models): **Beam Search**



# Inference Efficiency

How to accelerate LLMs? **Speculative Decoding**

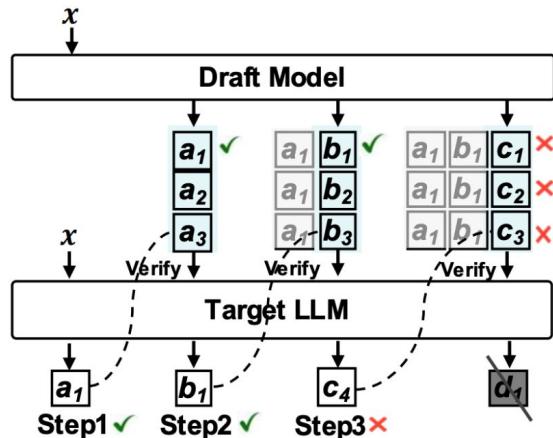
- Use a “cheap” model to generate candidates
- “Expensive” model can **accept** or **reject** (and perform inference if necessary)

```
[START] japan ' s benchmark bond n
[START] japan ' s benchmark nikkei 22 75
[START] japan ' s benchmark nikkei 225 index rose 22 76
[START] japan ' s benchmark nikkei 225 index rose 226 : 69 points
[START] japan ' s benchmark nikkei 225 index rose 226 : 69 points , or 0 1
[START] japan ' s benchmark nikkei 225 index rose 226 . 69 points , or 1 . 5 percent , to 10 , 9859
[START] japan ' s benchmark nikkei 225 index rose 226 . 69 points , or 1 . 5 percent , to 10 , 989 , 79 in
[START] japan ' s benchmark nikkei 225 index rose 226 . 69 points , or 1 . 5 percent , to 10 , 989 . 79 in tokyo late
[START] japan ' s benchmark nikkei 225 index rose 226 . 69 points , or 1 . 5 percent , to 10 , 989 . 79 in late morning trading . [END]
```

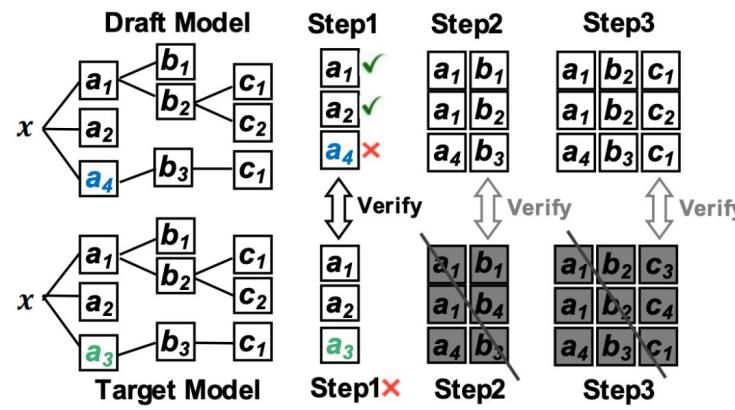
# Inference Efficiency

Speculative decoding for generative rec? 

## N-to-K verification



(b) N-to-1 Verification of Traditional SD (N=3)



(c) N-to-K Verification of SD with Beam Search (N=K=3)

# Inference Efficiency

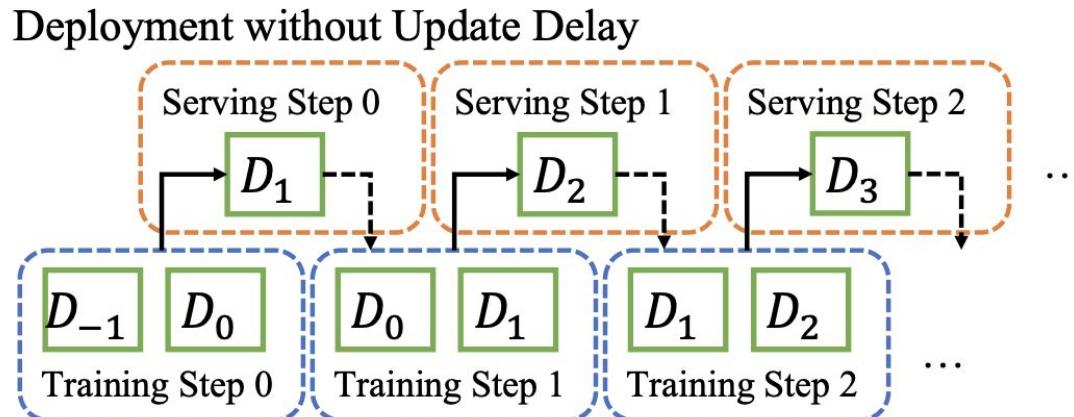
In addition to single-model acceleration methods,  
what about “**serving throughout**”?

Example:  LLM offers solutions for high-throughput  
and memory-efficient inference and serving

What's unique for generative rec?

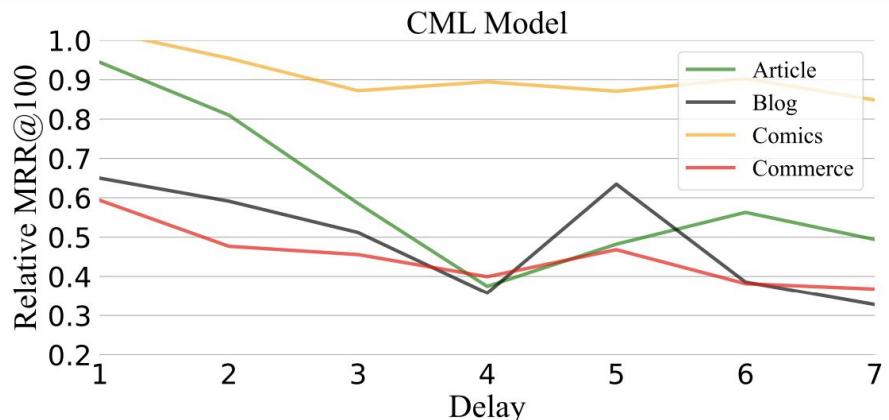
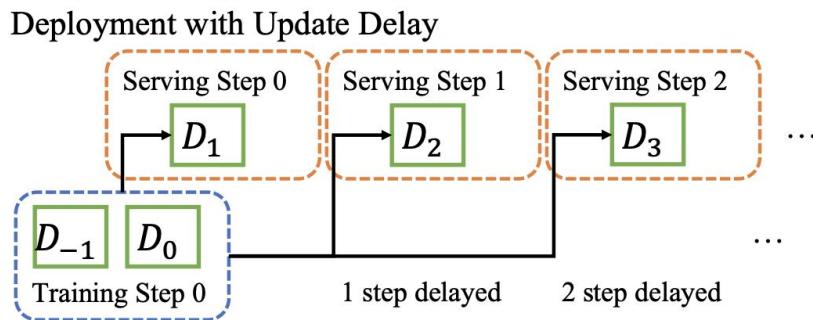
# Timely Model Update

Recommendation models favor **timely updates**



# Timely Model Update

Delayed updates lead to performance degradation



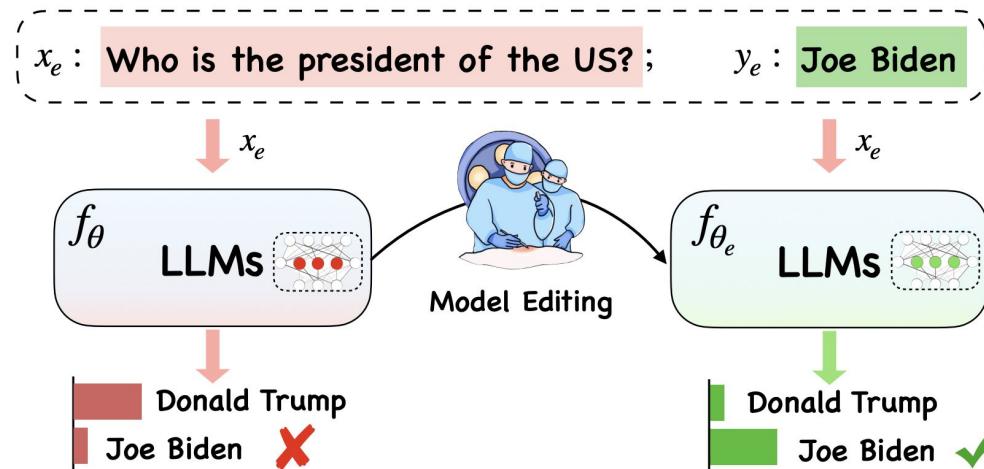
# Timely Model Update

How to update large generative rec models timely?

(Frequently retraining large generative models may be resource consuming)

# Timely Model Update

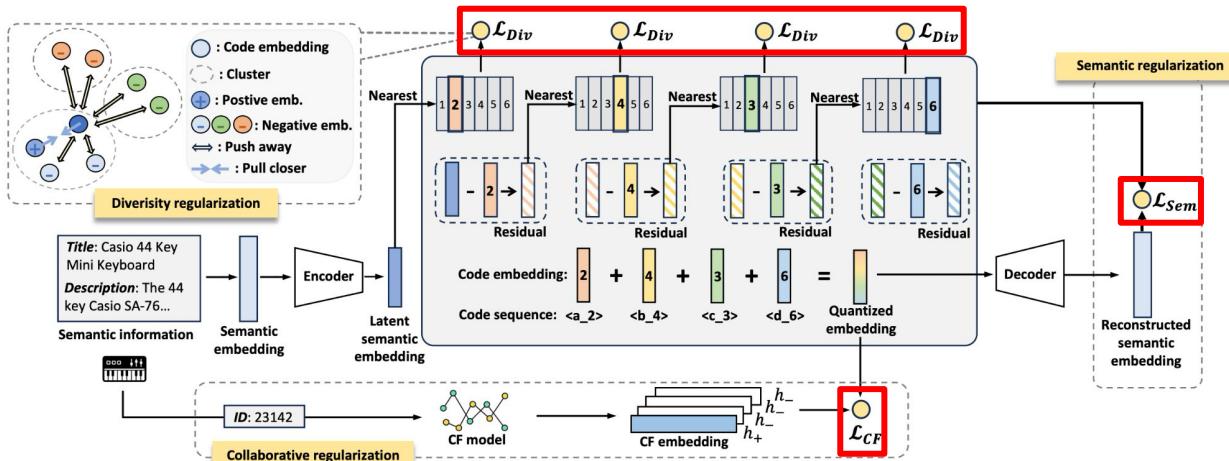
How to update large generative rec models timely?



Knowledge editing?

# Item Tokenization

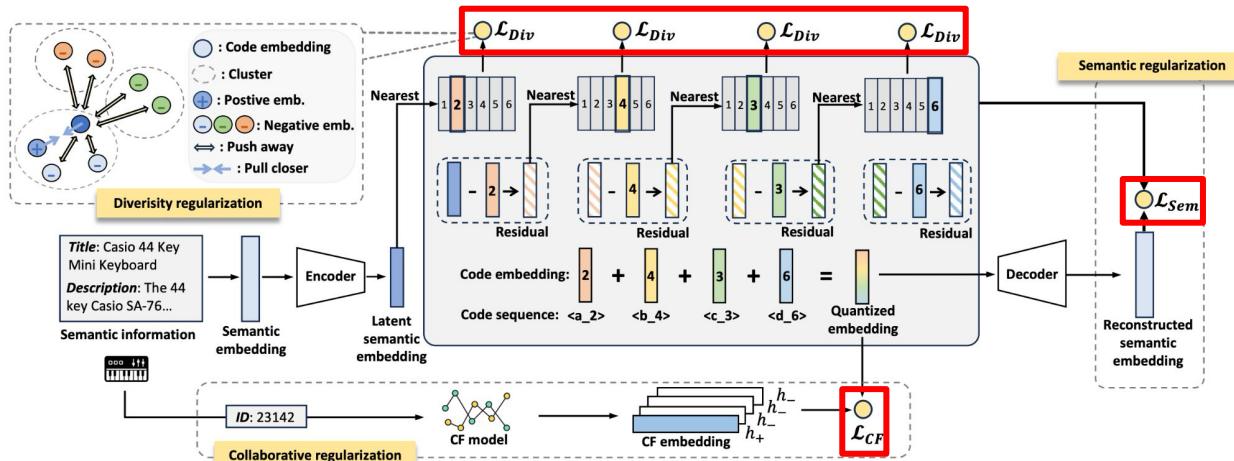
Multiple objectives for optimizing item tokenization ...



# Item Tokenization

Multiple objectives for optimizing item tokenization ...

But **none** of them is **directly related to rec performance**



# Item Tokenization

**reconstruction loss ≠ downstream performance**

How to connect tokenization objective with  
recommendation performance?

*Zipf's distribution? Entropy? Linguistic metrics?*

# Item Tokenization

## Language Tokenization

2014~2015:

Word / Char

---

*Context-independent ⇒ Context-aware*

---

# Item Tokenization

## Language Tokenization

2014~2015:

Word / Char

2016~present:

BPE / WordPiece

---

*Context-independent ⇒ Context-aware*

---

# Item Tokenization

## Language Tokenization

2014~2015:

Word / Char

2016~present:

BPE / WordPiece

---

*Context-independent ⇒ Context-aware*

---

## SemID Construction

2023~2024:

RQ / PQ / Clustering /  
LM-based Generator

# Item Tokenization

## Language Tokenization

2014~2015:

Word / Char

2016~present:

BPE / WordPiece

---

*Context-independent ⇒ Context-aware*

---

## SemID Construction

2023~2024:

RQ / PQ / Clustering /  
LM-based Generator

2025:

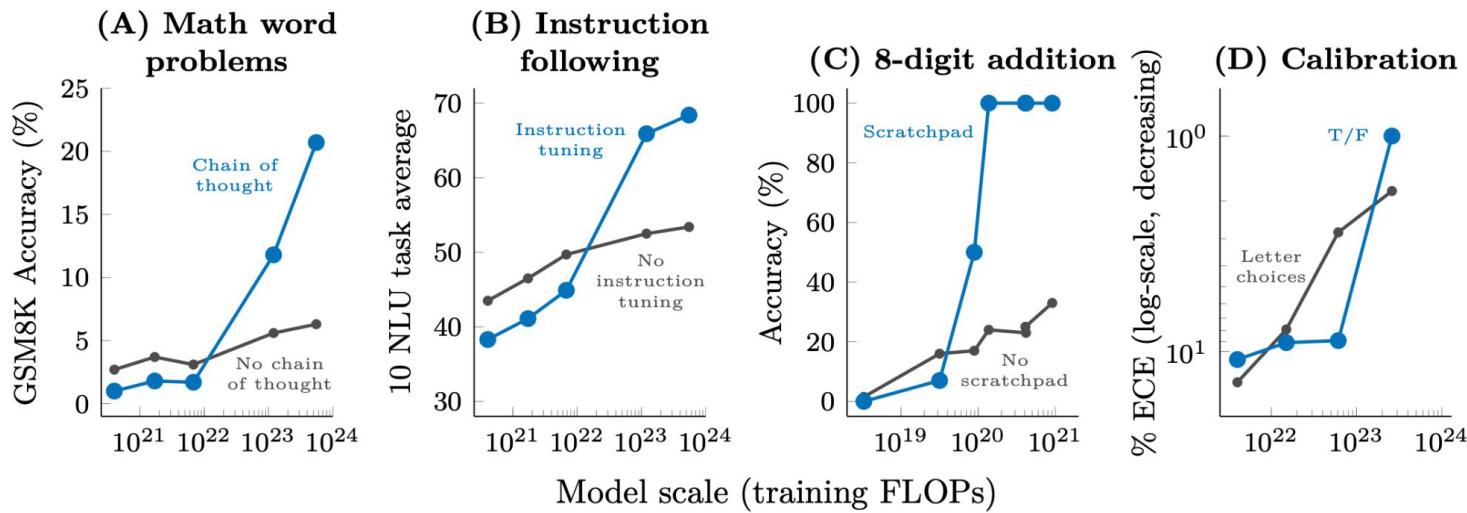
ActionPiece / ?

## **Part 2: What Becomes Possible?**

What new opportunities may large generative models unlock for recommender systems?

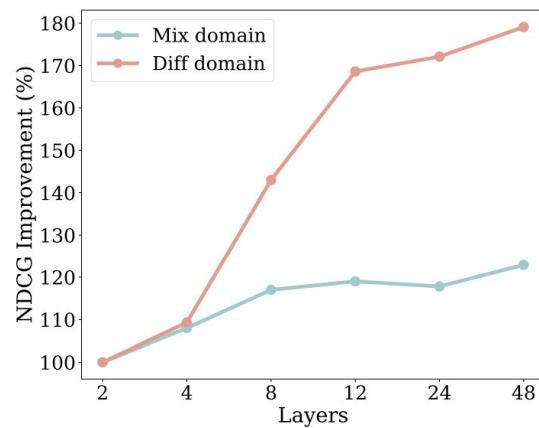
# Emergent Ability

Abilities not present in smaller models but is present in larger models

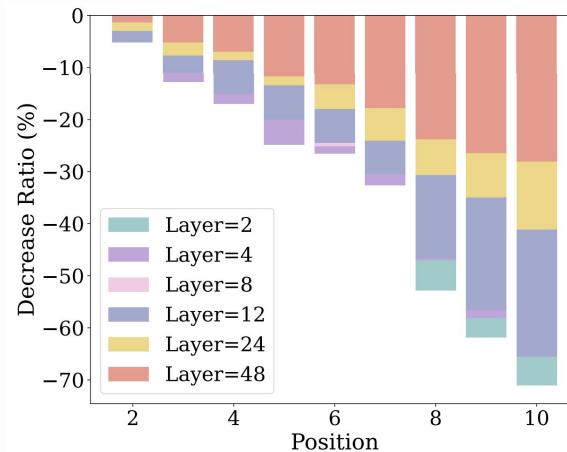


# Emergent Ability

Do we have **emergent abilities** in large generative recommendation models?



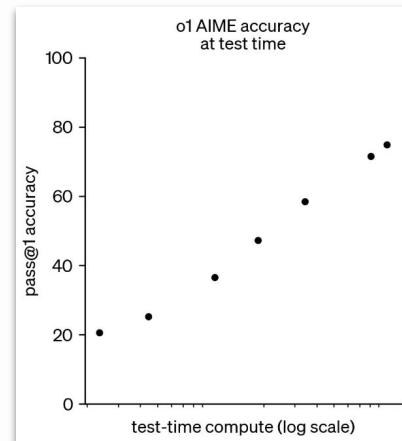
Cross-domain



Trajectory Prediction

# Test-time Scaling

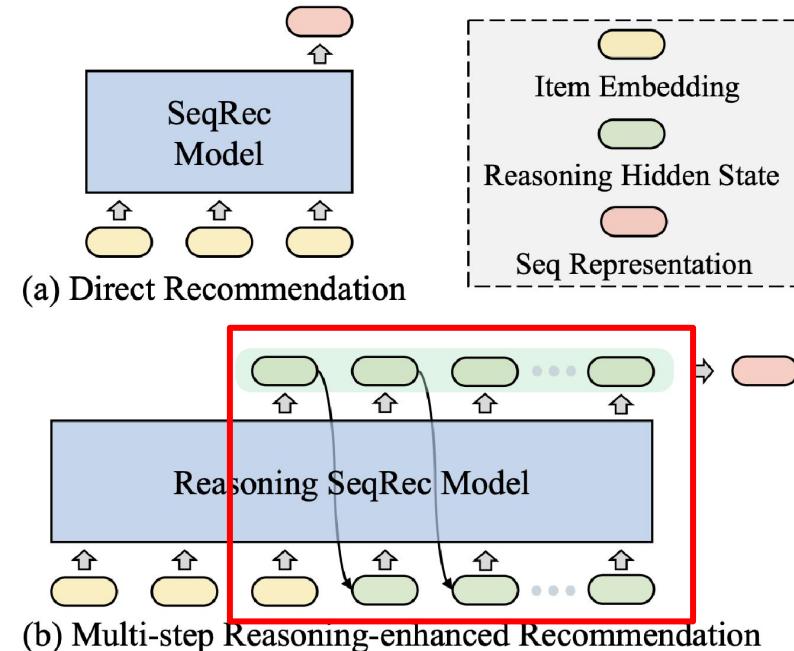
There have been explorations on **model / data scaling** of recommendation models



**Test-time scaling** is still under exploration

# Test-time Scaling

Reasoning over **latent hidden states** to scale up test-time computation.



# Large Diffusion Model

End-to-end personalized content generation with diffusion model.

- Existing work largely uses **pretrained diffusion model** to generate personalized content.
- An **end-to-end training** could bring mere personalization.

# Agentic

## Next: personalization & agentic system

**Customize ChatGPT**

Introduce yourself to get better, more personalized responses ⓘ

**What should ChatGPT call you?**

**What do you do?**

**What traits should ChatGPT have? ⓘ**

I love Iron Man. Recommend some scientific movies for me.

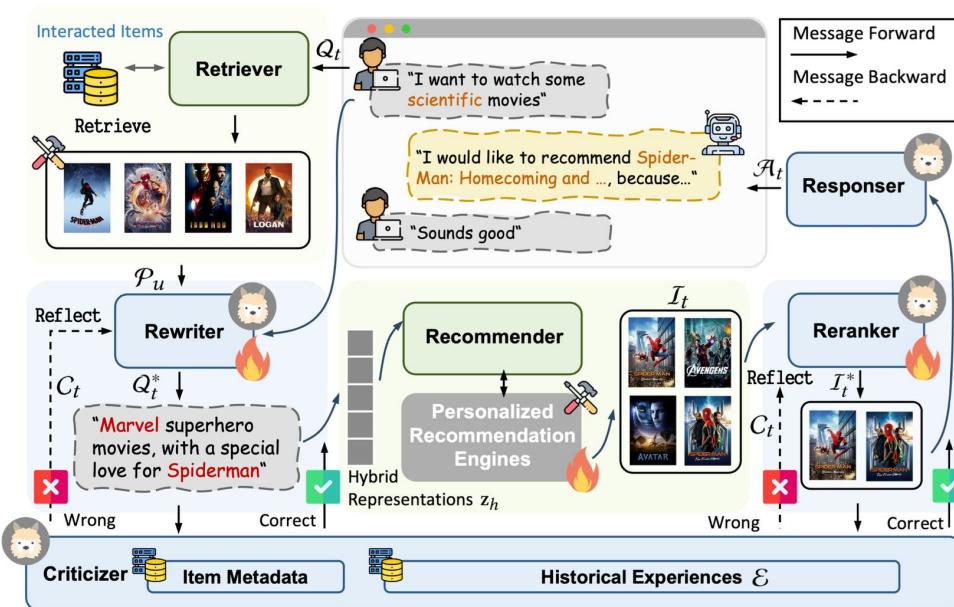
Q. Completed  
44s · 22 sources

- Thinking
- Exploring scientific movies
- Evaluating movie themes

- All these movies have strong scientific elements, either based on real events or accurate concepts. The Imitation Game shows early computing, while Interstellar uses real physics theories.
- I'm checking if they fit the user's taste for adventure and innovation, like Iron Man. The Martian's survival story and 2001's space exploration seem to match well.
- Hidden Figures tells a true story of NASA mathematicians, which is inspiring and scientific. I think these choices will engage the user with their blend of science and drama.
- Let me ensure they're not just fiction but grounded in science. The Martian is known for realism, and 2001 incorporates real science from Arthur C. Clarke.
- The Martian (2015) follows an astronaut using science to survive on Mars. It's problem-solving and engineering, much like Tony Stark's character, and focuses on botany and space exploration.
- Hidden Figures (2016) celebrates African-American women at NASA, focusing on math and aerospace engineering. Its inspiring story aligns with Iron Man's themes of intelligence and overcoming challenges.

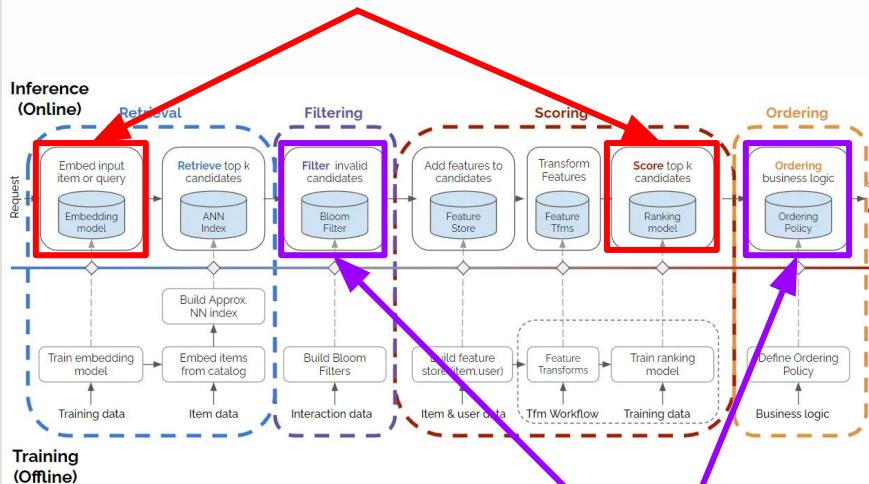
22 web pages

# Agentic



Large agentic system  
for precise user  
modeling and better  
recommendation

# Unified Retrieval & Ranking models



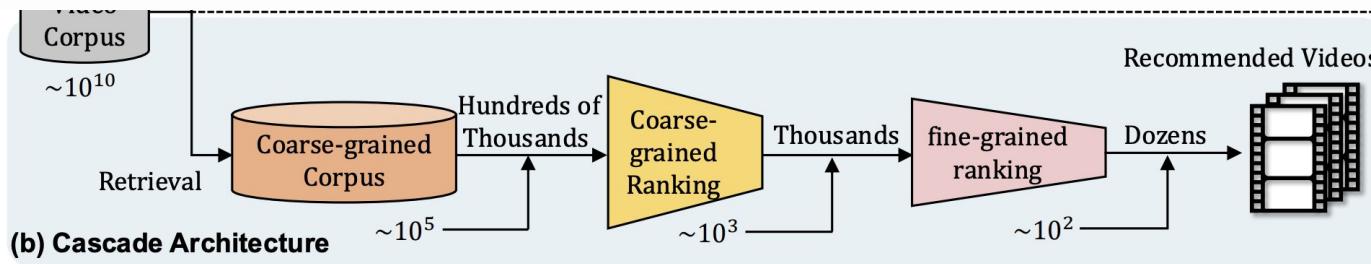
## Complicated Architecture

- Difficult to be optimized in an **end-to-end** way
- Latency** between / within different modules

rules, strategies, heuristics

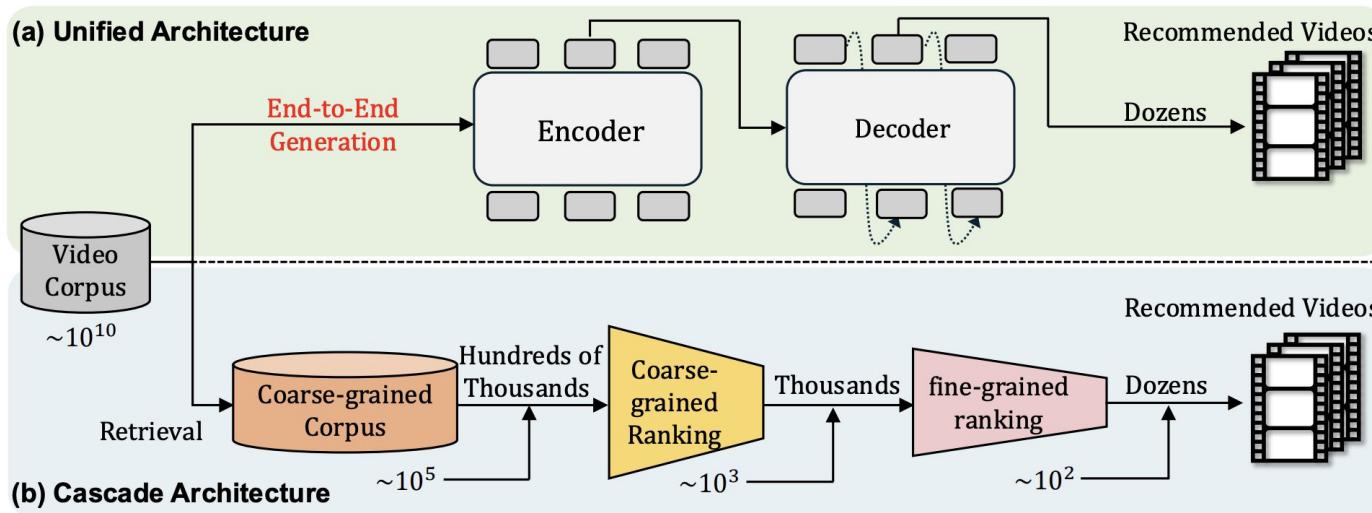
# Unified Retrieval & Ranking

Is it possible to replace traditional cascade architecture



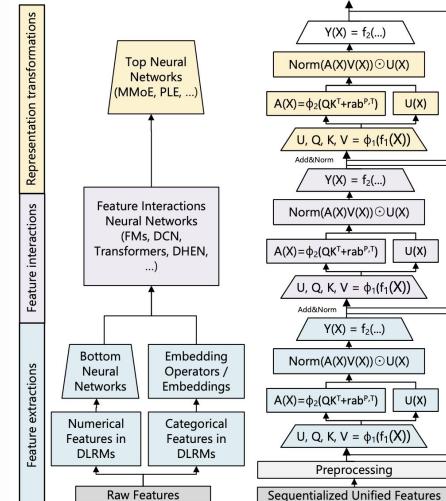
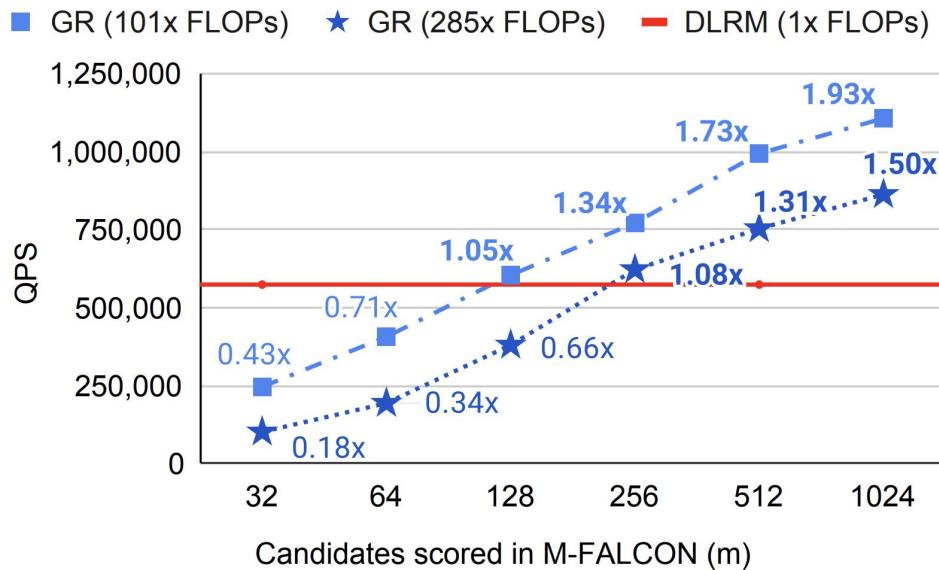
# Unified Retrieval & Ranking

Is it possible to replace traditional cascade architecture with a **unified generative model**?



# Unified Retrieval & Ranking

Better throughout when ranking more candidates



# Q & A

Thank you for coming!



Please refer to

[large-genrec.github.io](https://large-genrec.github.io)

for *slides, paper list, .....*

- [Papers](#)
  - [Surveys](#)
  - [LLM-based Generative Recommendation](#)
    - [LLM as Sequential Recommender](#)
      - [Early Efforts: Zero-shot Recommendation with LLMs](#)
      - [Aligning LLMs for Recommendation](#)
      - [Training Objectives & Inference](#)
    - [LLM as Conversational Recommender & Recommendation Assistant](#)
    - [LLM as User Simulator](#)
  - [Semantic ID-based Generative Recommendation](#)
    - [Semantic ID Construction](#)
      - [Quantization](#)
      - [Hierarchical Clustering](#)
      - [Contextual Action Tokenization](#)
      - [Behavior-aware Tokenization](#)
      - [Language Model-based Generator](#)
  - [Architecture](#)
    - [Dense & Generative Retrieval](#)
    - [Unified Retrieval and Ranking](#)
    - [Aligning with LLMs](#)
  - [Diffusion Model-based Generative Recommendation](#)
    - [Diffusion-enhanced Recommendation](#)
    - [Diffusion as Recommender](#)
    - [Personalized Content Generation with Diffusion](#)
  - [Resources](#)