

03

Semantic ID

-based Generative Recommendation

How to Index an Item in RecSys?

The screenshot shows a web browser window with the URL amazon.com/dp/B097B2YWFX. The page displays the product listing for "The Legend of Zelda: Tears of the Kingdom - Nintendo Switch (US Version)".

Product Information:

- Image:** A large cover art for the game, showing Link standing on a cliff edge.
- Brand:** Nintendo
- Platform:** Nintendo Switch | Rated: Rating Pending
- Rating:** 4.9 ★★★★★ (22,782 ratings)
- Offer:** Amazon's Choice
- Sales:** 3K+ bought in past month

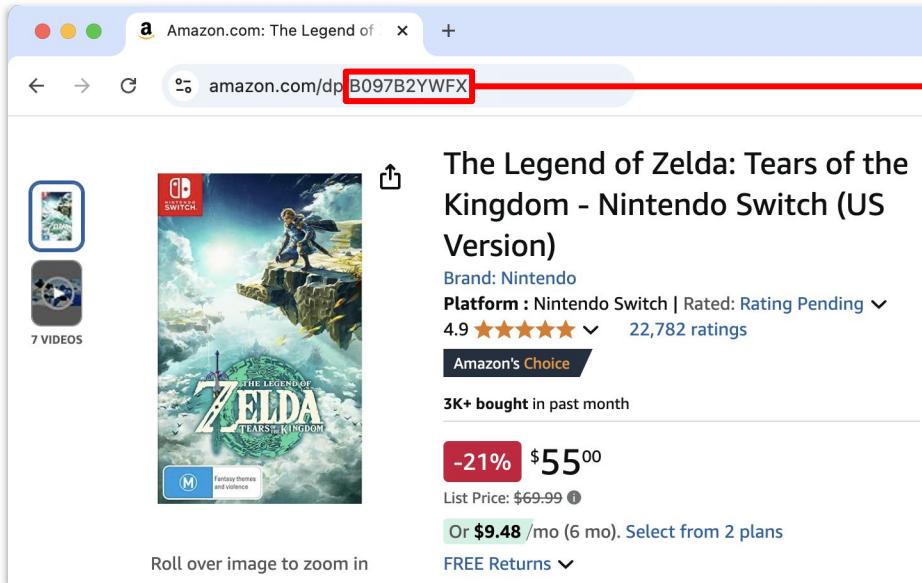
Pricing:

- Price:** \$55⁰⁰ (21% off from \$69.99)
- Plan:** Or \$9.48 /mo (6 mo). Select from 2 plans
- Returns:** FREE Returns

Other Options:

- 7 VIDEOS
- Roll over image to zoom in

How to Index an Item in RecSys?



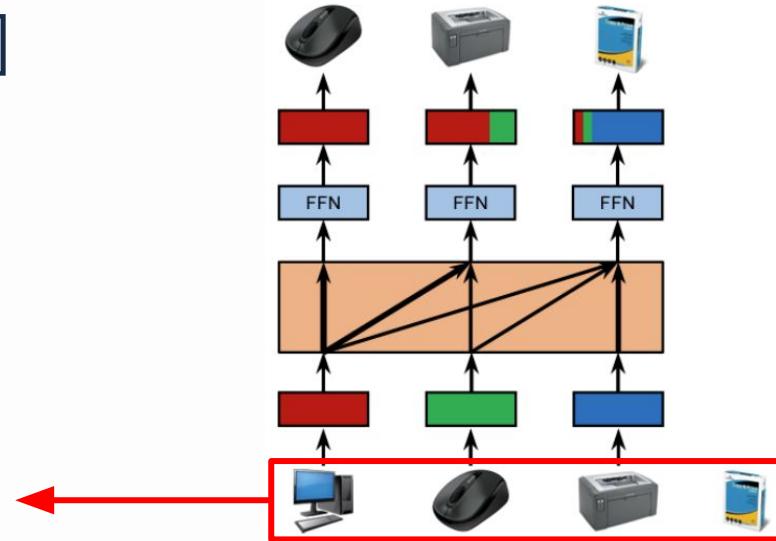
Item ID:

B097B2YWFX

How to Index an Item in RecSys?

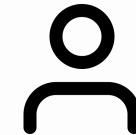
Example: SASRec [*ICDM'18*]

Each item is indexed by
a unique **item ID**



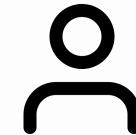
How to Index an Item in LLMs

I wanna some popular Nintendo games



How to Index an Item in LLMs

I wanna some popular Nintendo games



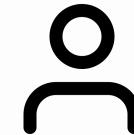
LLMs

How about or ?



How to Index an Item in LLMs

I wanna some popular Nintendo games



LLMs

How about or ?



How to index items in LLMs? **Item ID?**



How to Index an Item in LLMs

How many tokens in LLMs?

 Meta

Llama 3

~128,000

 OpenAI

GPT-4o

~200,000

 Google DeepMind

Gemma 2

~256,000

How to Index an Item in LLMs

How many tokens in LLMs?

 Meta	Llama 3	~128,000
 OpenAI	GPT-4o	~200,000
 Google DeepMind	Gemma 2	~256,000

How many item IDs?

Amazon-Reviews-2023 ~48,200,000

How to Index an Item in LLMs

How many tokens/item IDs in LLMs/RecSys?

**Difficult to align
these vocabularies
given so many tokens**

~128,000
~200,000
~256,000

~48,200,000

How to Index an Item in LLMs

Is there a way to
index a large volume of items
using a **compact vocabulary?**

Semantic IDs

(also called: SemID or SID)

A few tokens that jointly index one item.

t3, t321, t643, t1011



Semantic IDs

(also called: SemID or SID)

A few tokens that jointly index one item.

t3, **t321**, t643, t1011



{t257, t258, ..., t320, **t321**, t322, ..., t511, t512}

Each token from a **vocabulary shared by all items**

Semantic IDs

(also called: SemID or SID)

A few tokens that jointly index one item.

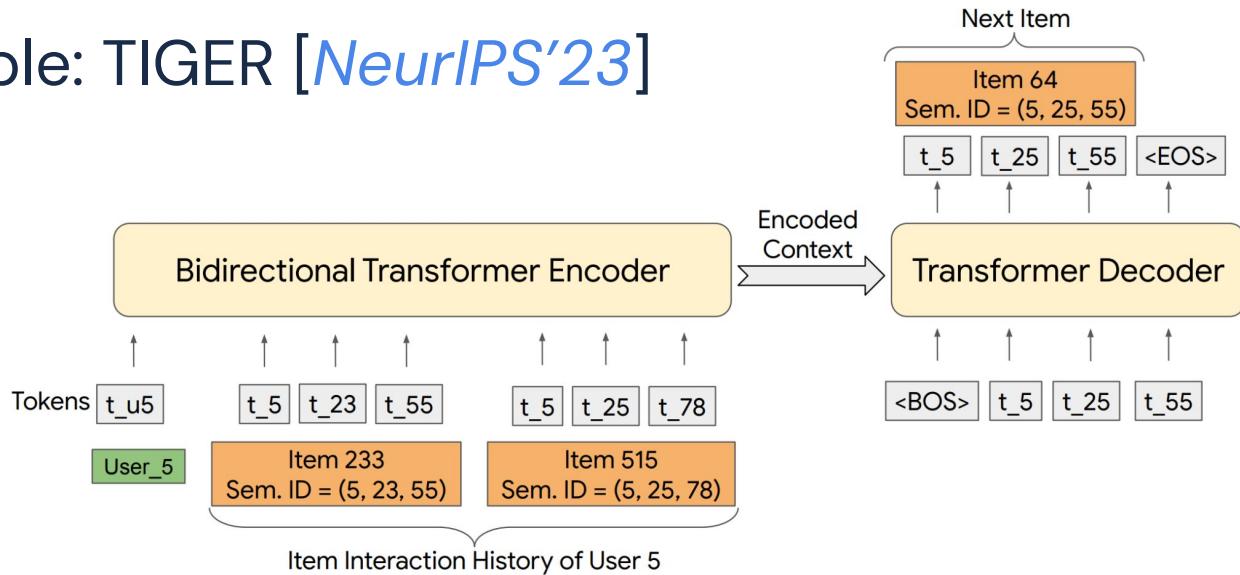
t3, t321, t643, t1011

Can index maximally $256^4 \approx 4.3 \times 10^9$ items with 1024 tokens

(4 tokens per item, each from a vocabulary of 256)

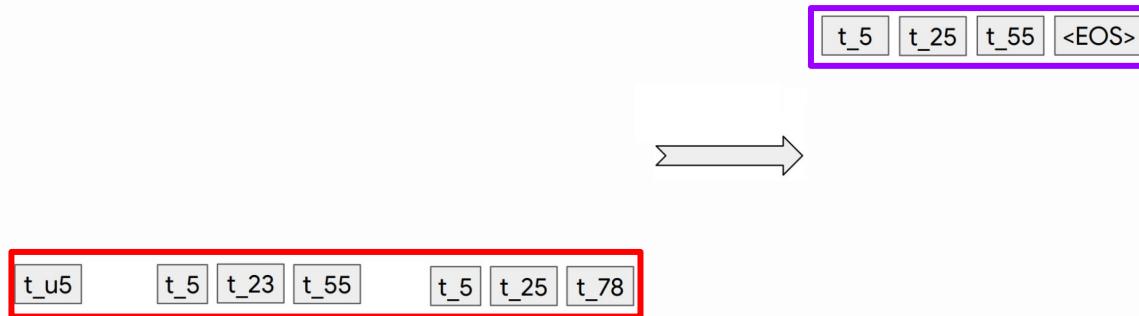
Generative Models based on Semantic IDs

Example: TIGER [NeurIPS'23]



Generative Models based on Semantic IDs

Example: TIGER [*NeurIPS'23*]



Recommendation as a **seq**-to-**seq** generation problem

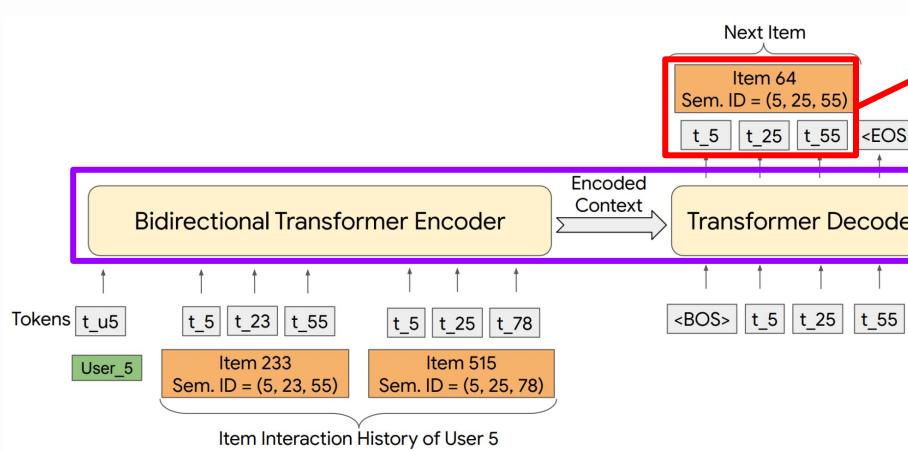
Generative Models based on Semantic IDs

Recommendation as a **seq**-to-**seq** generation problem

Input: user interacted items $\{c_{11}, c_{12}, c_{13}, c_{14}, c_{21}, c_{22}, \dots\}$

Output: next item $\{c_{t1}, c_{t2}, c_{t3}, c_{t4}\}$

SemID-based Generative Recommendation



Part 1:

How to construct **SIDs**

Part 2:

How to build SID-based
Generative Rec **Models**

Part 1: Semantic ID Construction

Semantic ID Construction

Input: all data associated with the item
(description, title, interactions, features, ...)



Output: mapping between **items** \leftrightarrow **Semantic IDs**

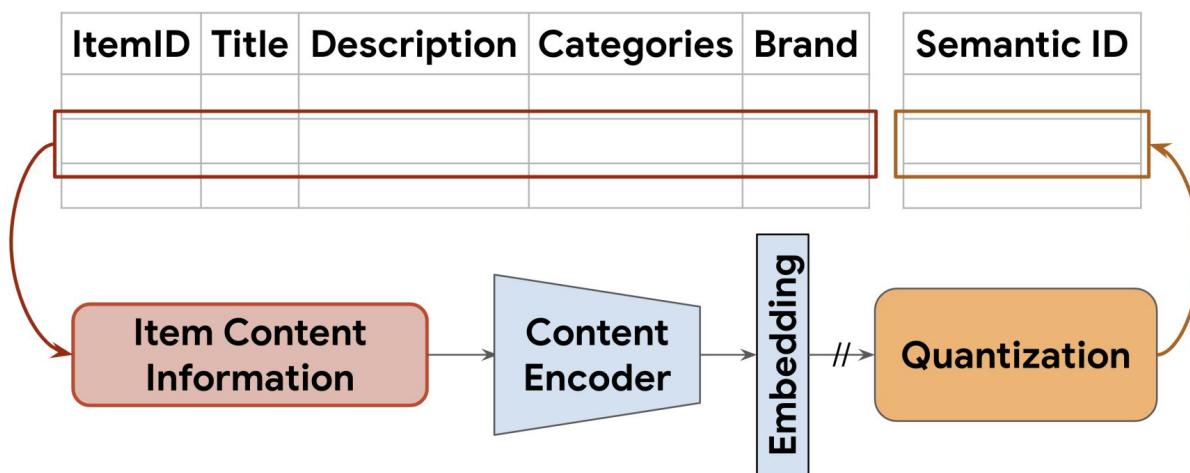
B097B2YWFX \leftrightarrow {t3, t321, t643, t1011} 169

Part 1: Semantic ID Construction

(i) **First example:** TIGER and RQ-VAE-based SemIDs;

SemID Construction – First Example: TIGER

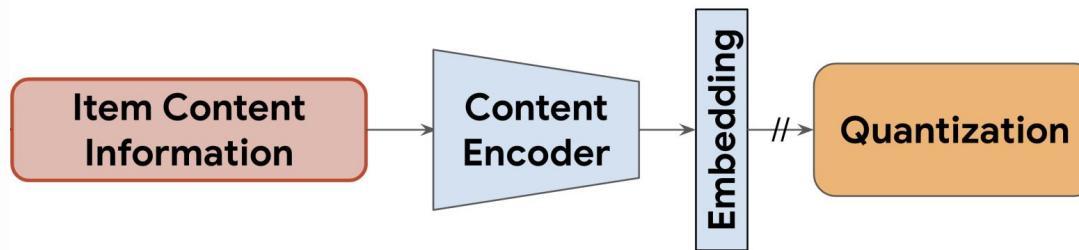
Input: concatenated text features



Output: mapping between **items** \leftrightarrow **Semantic IDs**

SemID Construction – First Example: TIGER

Text ➤ **Vector** ➤ **IDs**



SemID Construction – First Example: TIGER

1. Item Content Information (**Text**)

ItemID	Title	Description
B097B2YWFX.	The Legend of Zelda: Tears of the Kingdom - Nintendo Switch (US Version).	An epic adventure across the land ... threaten the kingdom?

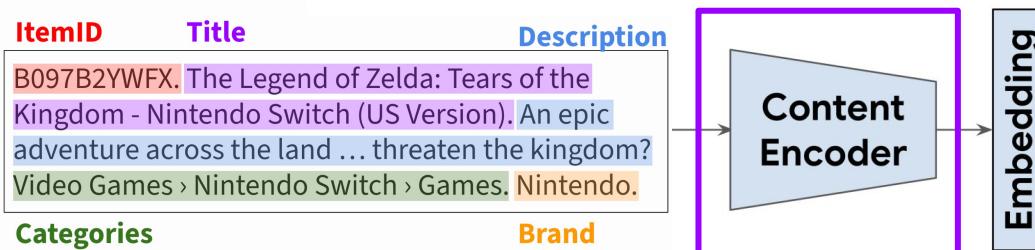
Categories

Brand

SemID Construction – First Example: TIGER

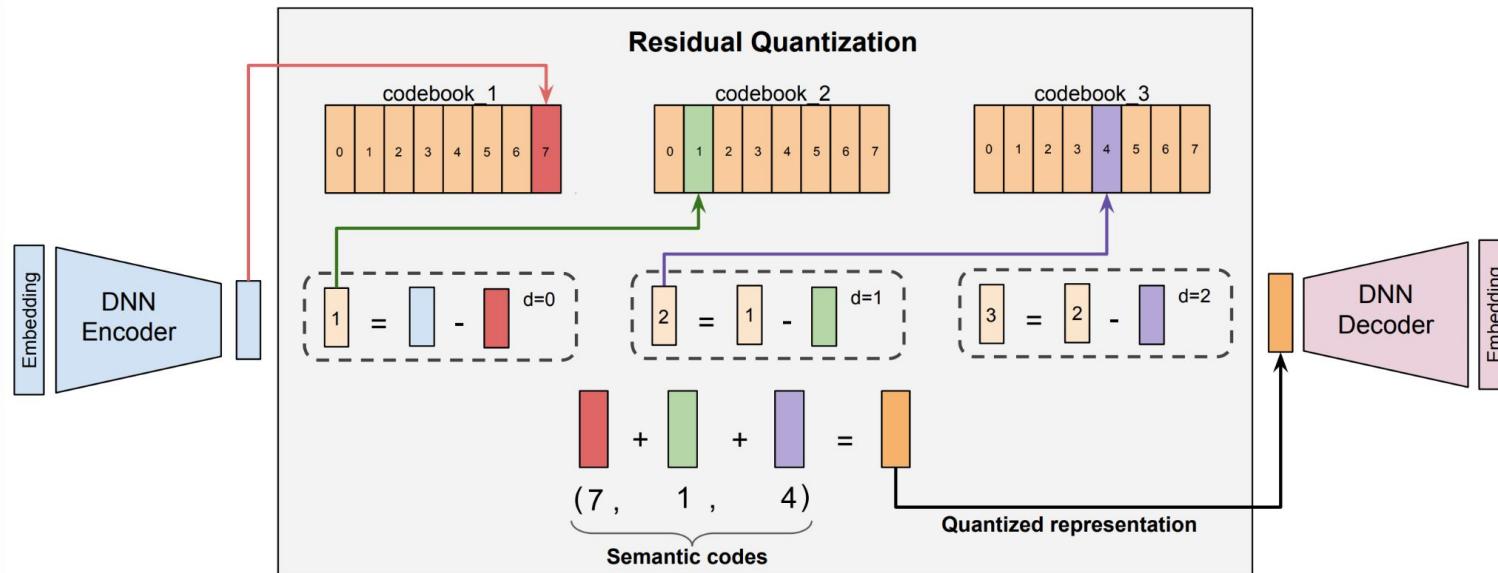
2. Content Encoder + Embedding (**Text** ➤ **Vector**)

Pre-trained (fixed) sentence embedding model
(SentenceT5)



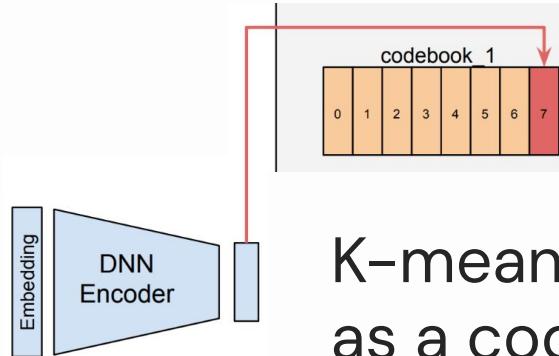
SemID Construction – First Example: TIGER

3. RQ-VAE Quantization (Vector \rightarrow IDs)



SemID Construction – First Example: TIGER

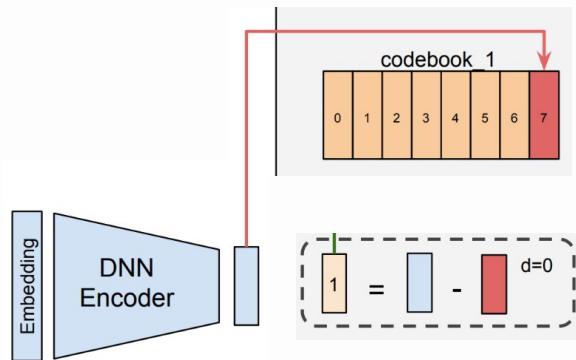
3. RQ-VAE Quantization (Vector \rightarrow IDs)



K-means: cluster center ID
as a code in the codebook

SemID Construction – First Example: TIGER

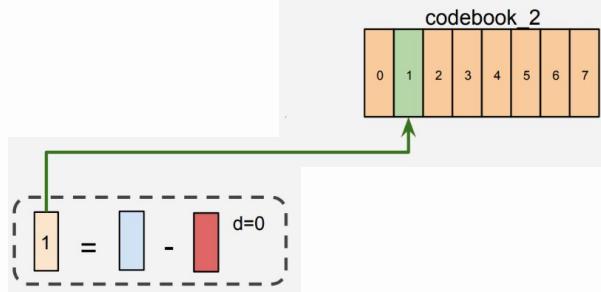
3. RQ-VAE Quantization (Vector \rightarrow IDs)



Residual of “input vector” and
“clustering center vector”

SemID Construction – First Example: TIGER

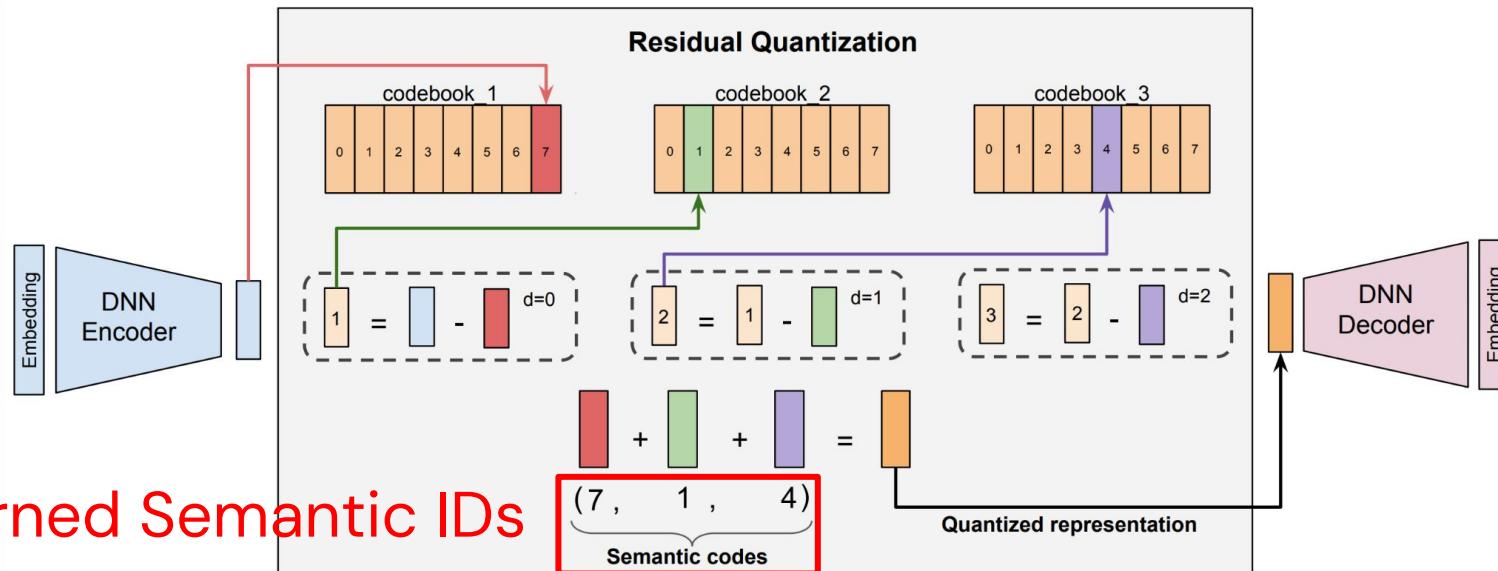
3. RQ-VAE Quantization (Vector \rightarrow IDs)



Residual as next level's input

SemID Construction – First Example: TIGER

3. RQ-VAE Quantization (Vector \rightarrow IDs)

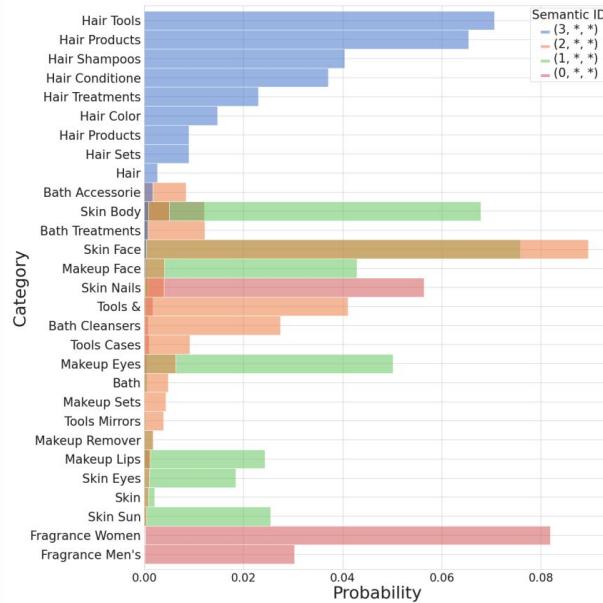


Learned Semantic IDs

SemID Construction – First Example: TIGER

Properties of RQ-VAE-based SemIDs

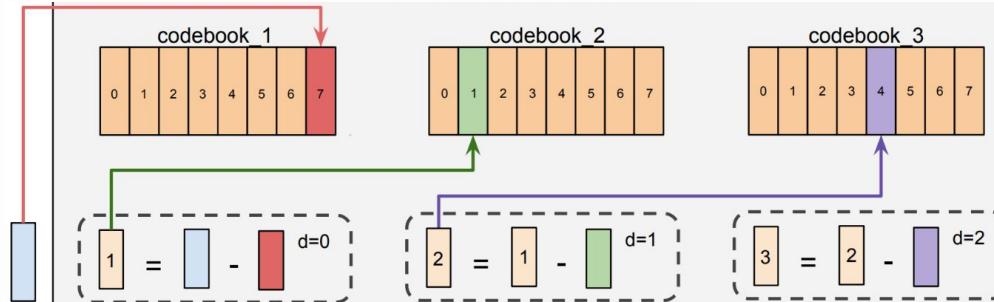
1. Semantic;



SemID Construction – First Example: TIGER

Properties of RQ-VAE-based SemIDs

1. Semantic;
2. Ordered / sequential dependent;



SemID Construction – First Example: TIGER

Collisions



(12, 24, 52)



(12, 24, 52)

SemID Construction – First Example: TIGER

Collisions



(12, 24, 52, 0)



(12, 24, 52, 1)

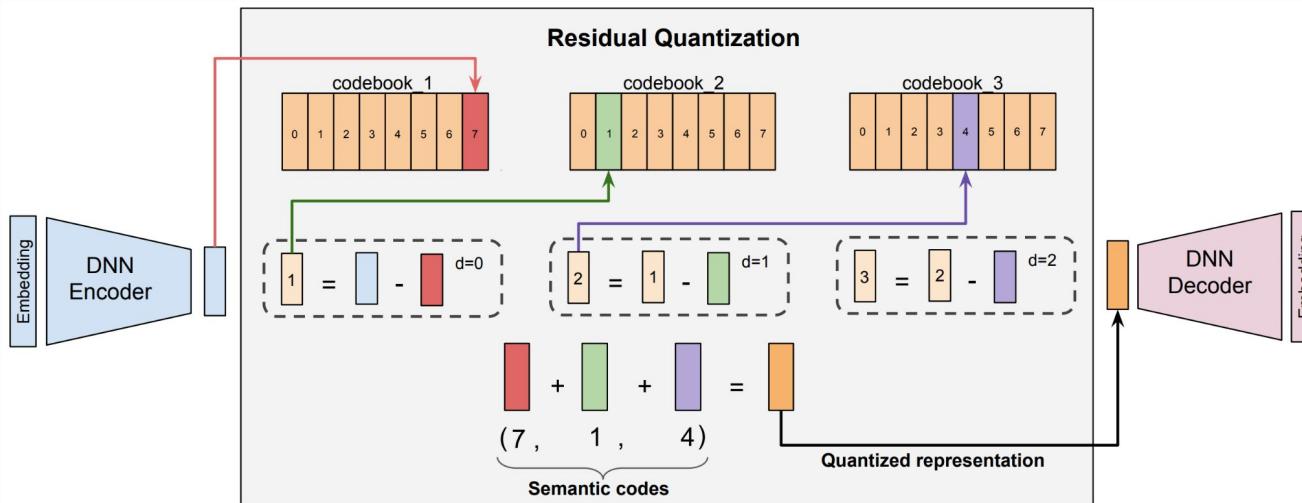
One extra token
to avoid conflicts

Part 1: Semantic ID Construction

- (i) First example: TIGER and RQ-VAE-based SemIDs;
- (ii) **Techniques** to construct SemIDs;

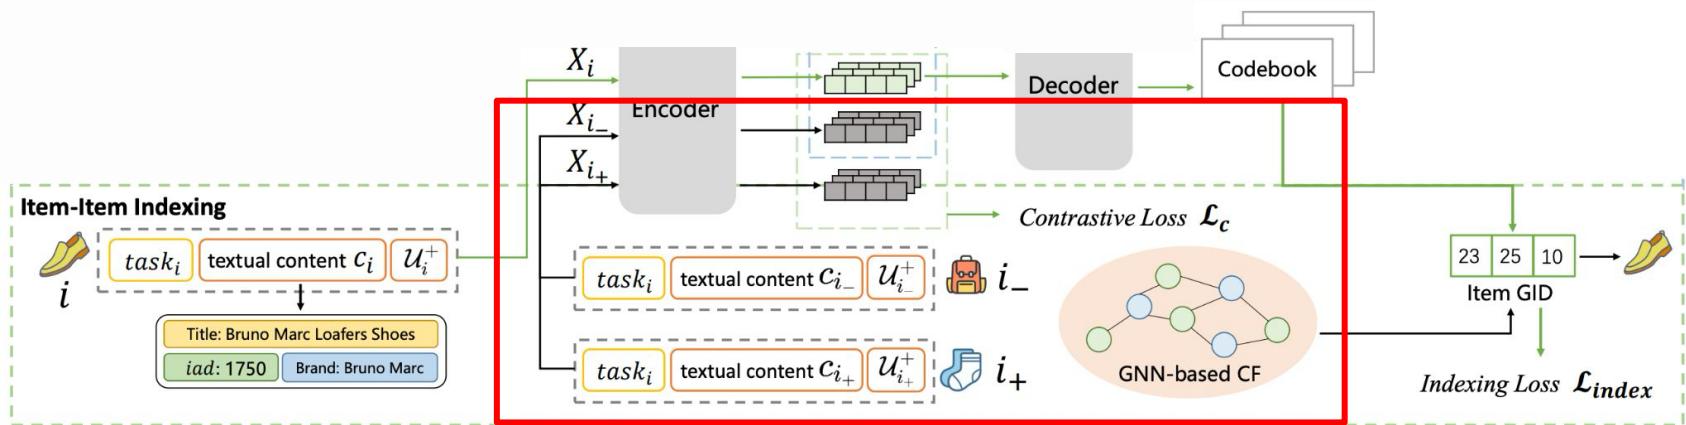
Techniques to Construct SemIDs

Residual Quantization



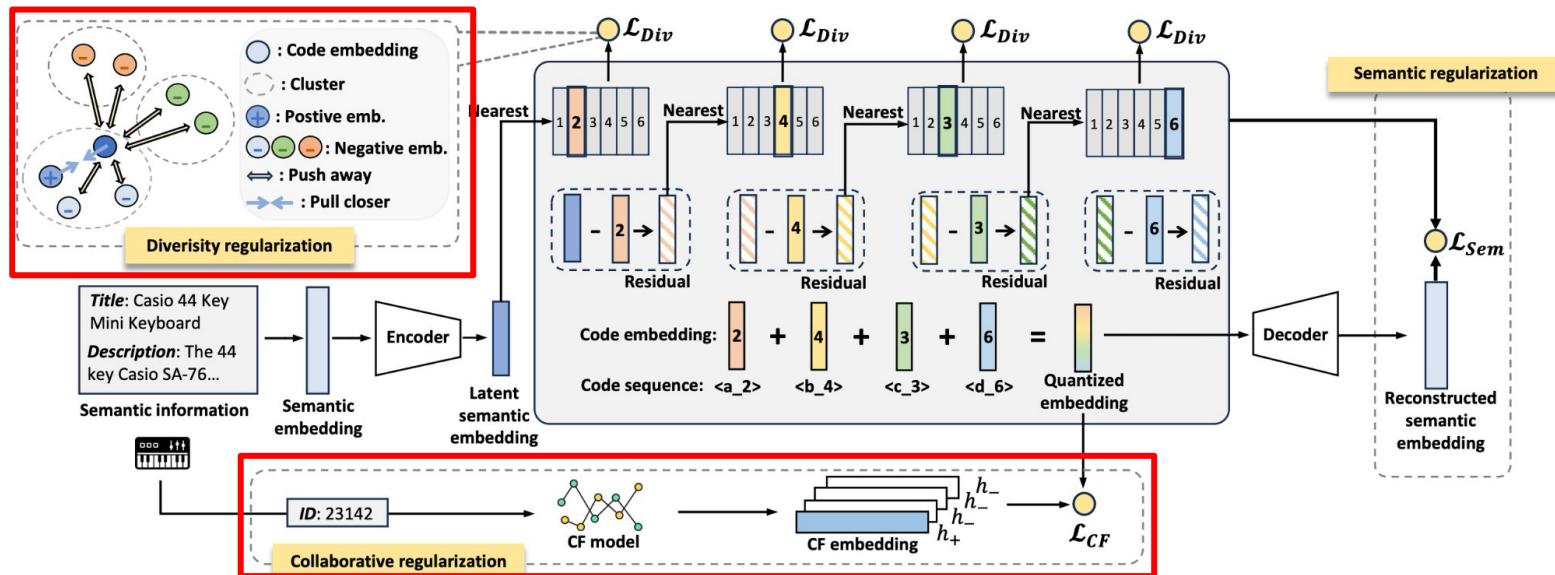
Techniques to Construct SemIDs

Residual Quantization + Item-level Regularization



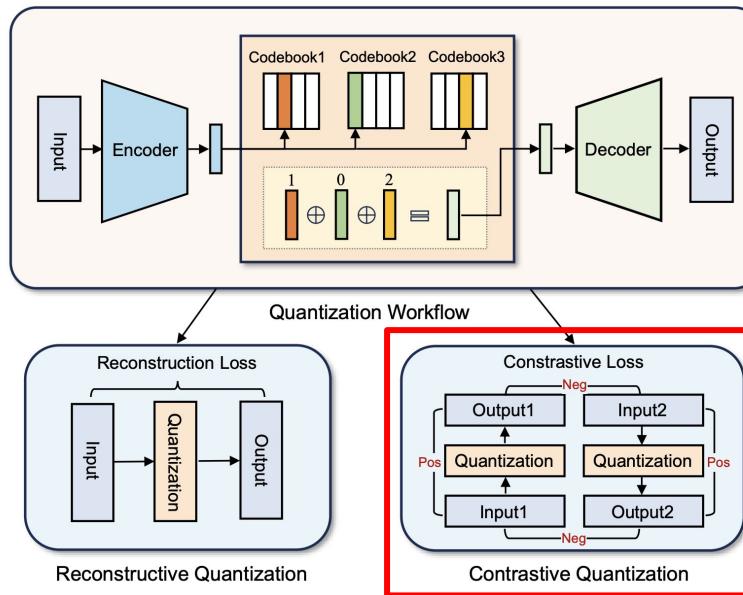
Techniques to Construct SemIDs

Residual Quantization + Item-level Regularization



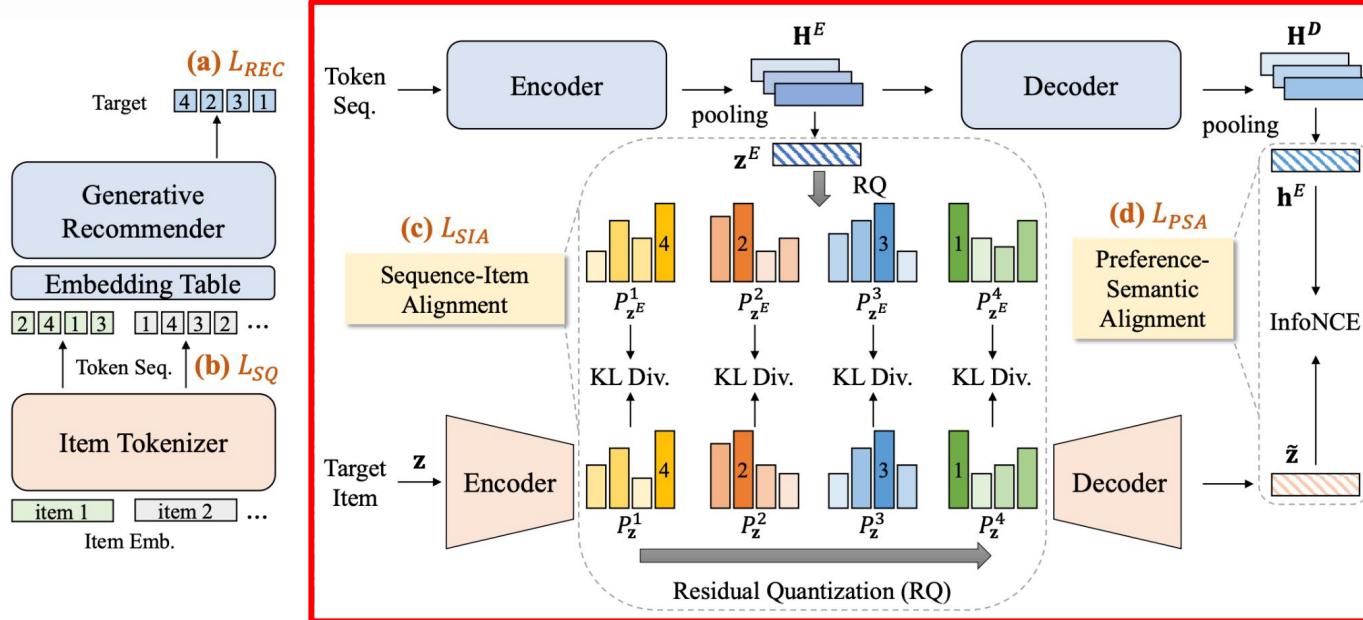
Techniques to Construct SemIDs

Residual Quantization + Item-level Regularization



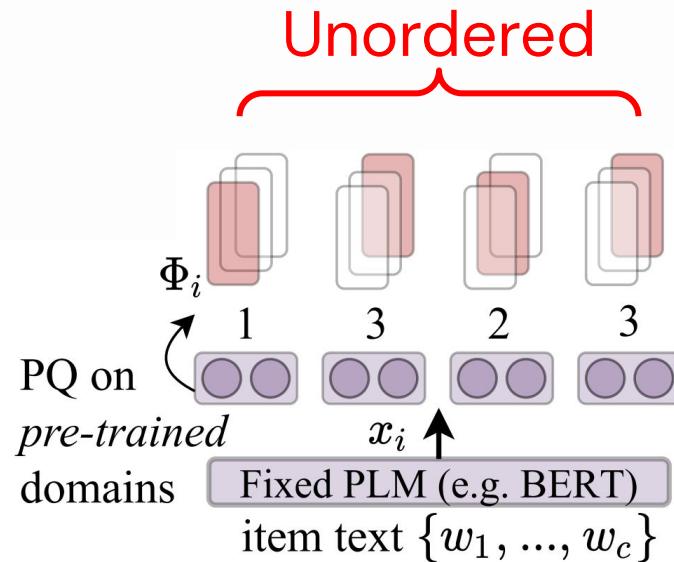
Techniques to Construct SemIDs

Residual Quantization + Recommendation Loss



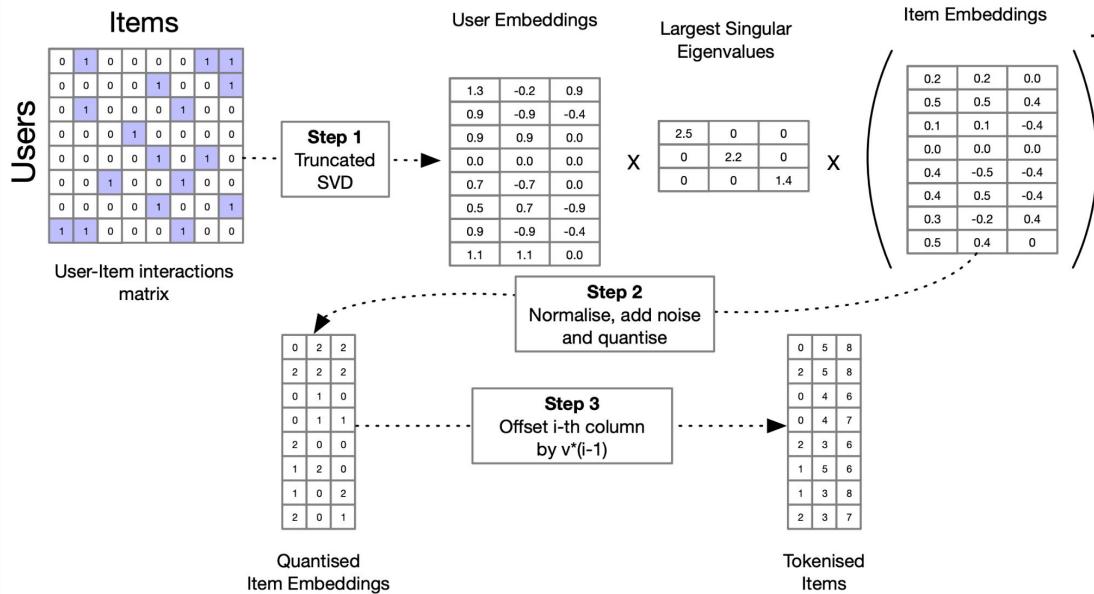
Techniques to Construct SemIDs

Product Quantization



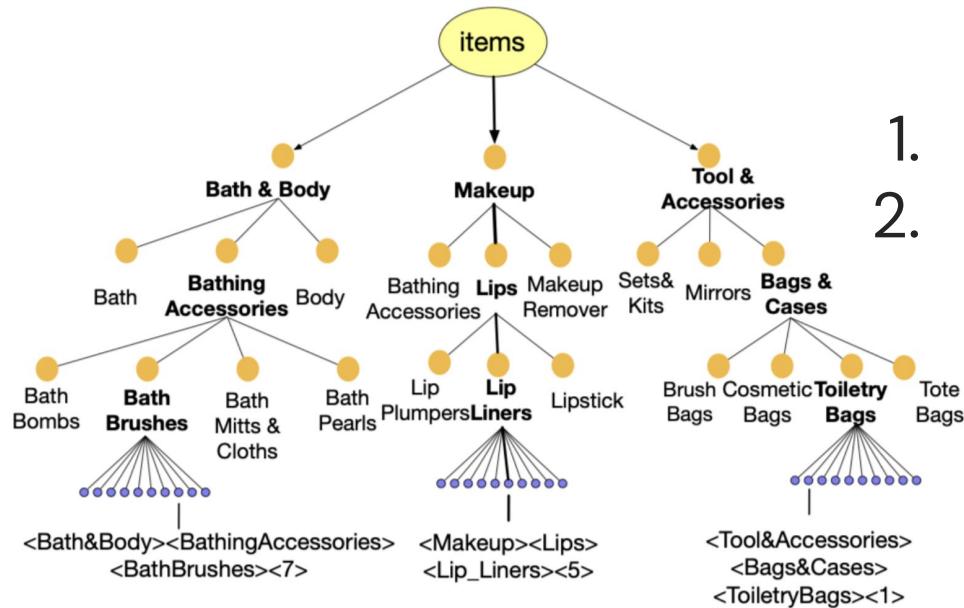
Techniques to Construct SemIDs

Product Quantization



Techniques to Construct SemIDs

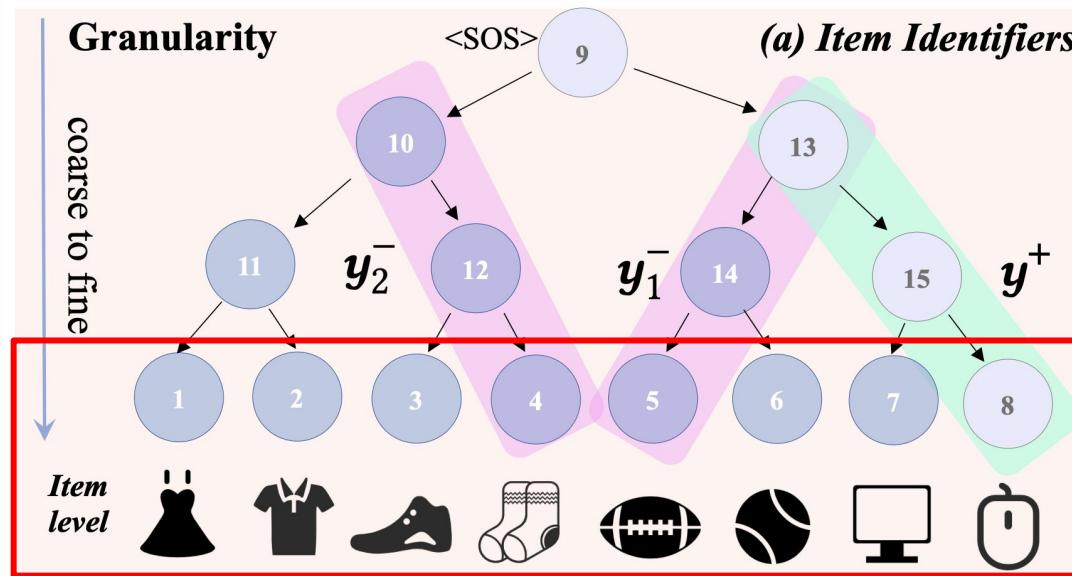
Hierarchical Clustering (**Heuristics**-based)



1. Ordered;
2. **Variable-length** SemIDs;

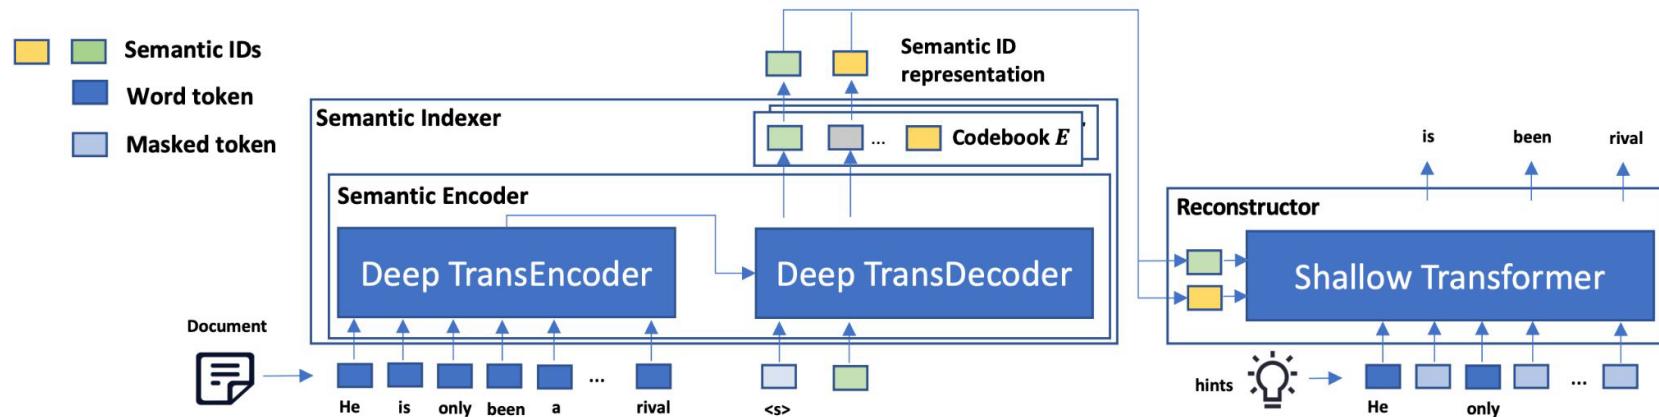
Techniques to Construct SemIDs

Hierarchical Clustering (**Latent**-based)



Techniques to Construct SemIDs

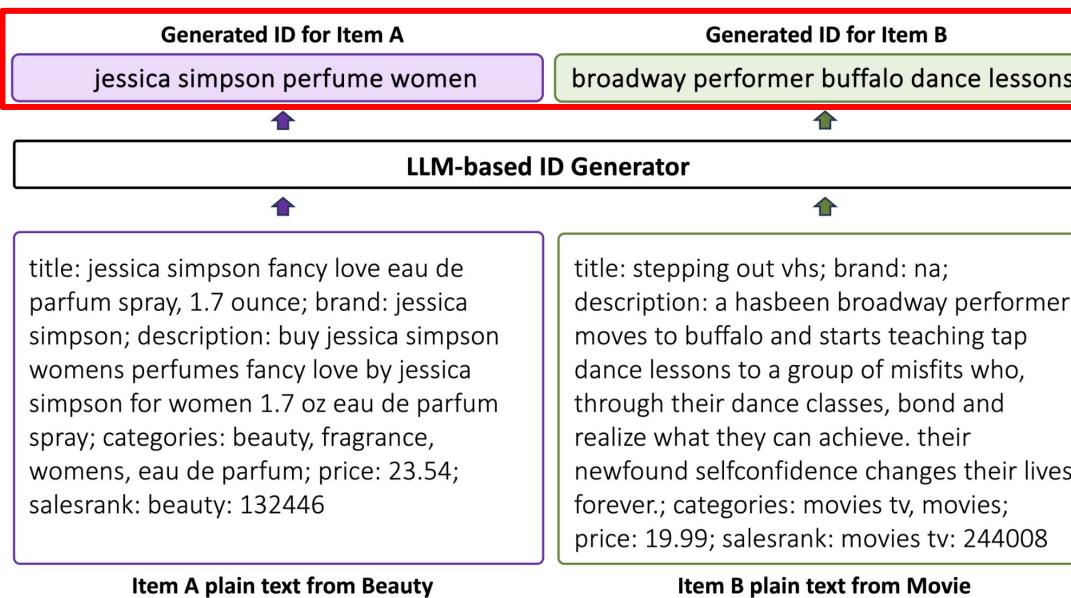
Language Model-based ID Generator



Natural language as inputs;
SemIDs as outputs

Techniques to Construct SemIDs

Language Model-based ID Generator



Words as SemIDs
(like tagging)

Techniques to Construct SemIDs

Context-independent

Action Tokenization	Example	Contextual	Unordered
Product Quantization	VQ-Rec (Hou et al., 2023)	✗	✓
Hierarchical Clustering	P5-CID (Hua et al., 2023)	✗	✗
Residual Quantization	TIGER (Rajput et al., 2023)	✗	✗
Text Tokenization	LMIndexer (Jin et al., 2024)	✗	✗
Raw Features	HSTU (Zhai et al., 2024)	✗	✗
SentencePiece	SPM-SID (Singh et al., 2024)	✗	✗

Same item ⇒ **fixed semIDs** in all sequences

Techniques to Construct SemIDs

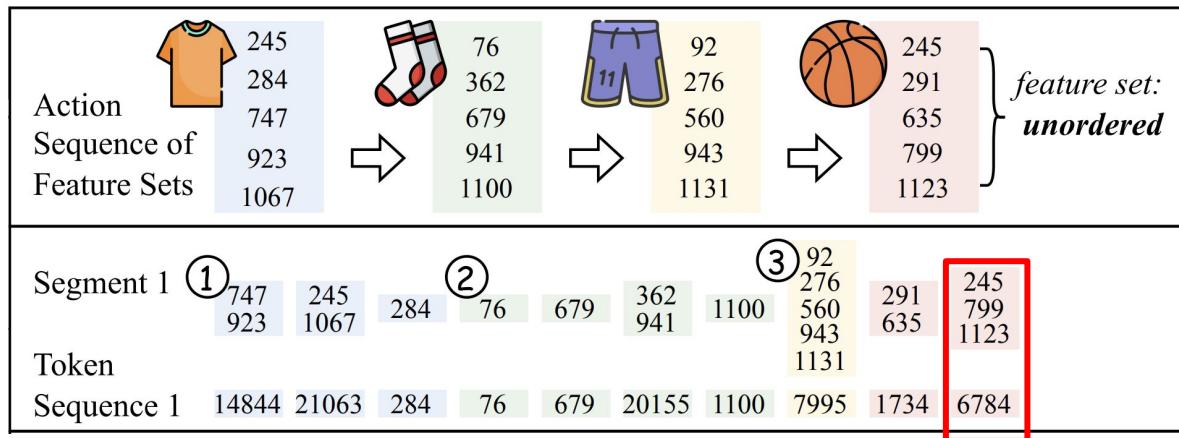
Context-independent \Rightarrow Context-aware

Action	245 284 747 923 1067	Sequence of Feature Sets	76 362 679 941 1100	feature set: <i>unordered</i>
Segment 1	① 747 923	245 1067	284 76 679 362 1100	② 276 560 943 1131 ③ 291 635 799 1123
Token	14844 21063 284	76 679 20155 1100	7995 1734 6784	Sequence 1
Segment 2	747 284 245 1067	923 679 941 76 362	1100 560 943 1131 276 245 799 1123	④ ④
Token	747 284 21063	8316 941 76 362 19895	24364 276 6784 1734	Sequence 2

Same item \Rightarrow
different semIDs
based on context

Techniques to Construct SemIDs

Context-independent \Rightarrow Context-aware



(ActionPiece: “WordPiece” tokenization for generative rec)

Core Idea:
Merge frequently co-occurring features as new tokens

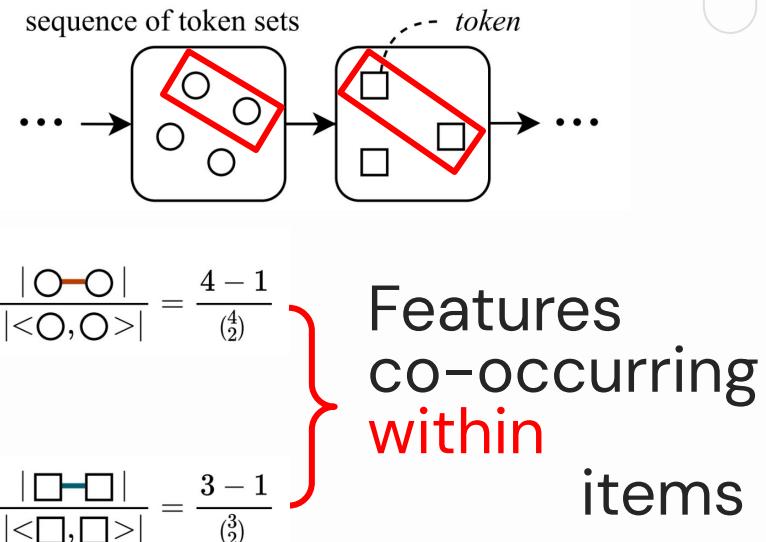
Techniques to Construct SemIDs

Context-independent \Rightarrow Context-aware

Algorithm 1 ActionPiece Vocabulary Construction

input Sequence corpus \mathcal{S}' , initial tokens \mathcal{V}_0 , target size Q
output Merge rules \mathcal{R} , constructed vocabulary \mathcal{V}

```
1: Initialize vocabulary  $\mathcal{V} \leftarrow \mathcal{V}_0$  # each initial token corresponds  
   to one unique item feature  
2:  $\mathcal{R} \leftarrow \emptyset$   
3: while  $|\mathcal{V}| < Q$  do  
4:   # Count: accumulate weighted token co-occurrences  
5:    $\text{count}(\cdot, \cdot) \leftarrow \text{Count}(\mathcal{S}', \mathcal{V})$  # Algorithm 2  
6:   # Update: Merge a frequent token pair into a new token  
7:   Select  $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$   
8:    $\mathcal{S}' \leftarrow \text{Update}(\mathcal{S}', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$  # Algorithm 3  
9:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$  # new merge rule  
10:   $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$  # add new token to the vocabulary  
11: end while  
return  $\mathcal{R}, \mathcal{V}$ 
```



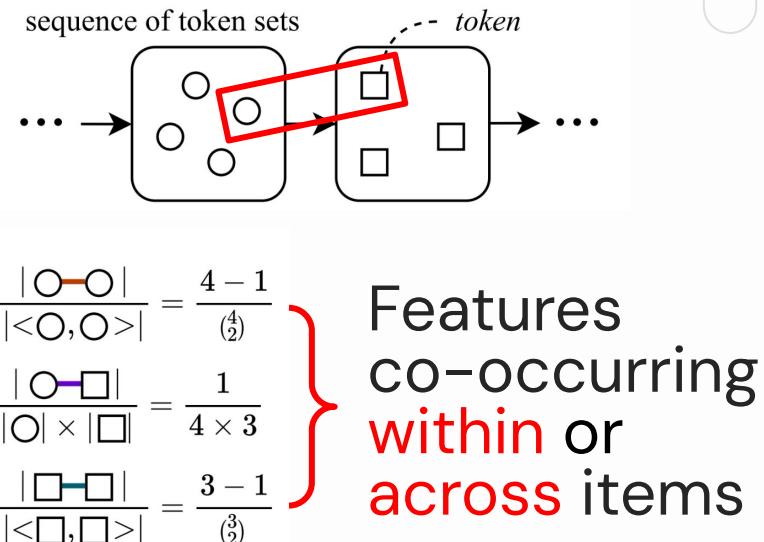
Techniques to Construct SemIDs

Context-independent \Rightarrow Context-aware

Algorithm 1 ActionPiece Vocabulary Construction

input Sequence corpus \mathcal{S}' , initial tokens \mathcal{V}_0 , target size Q
output Merge rules \mathcal{R} , constructed vocabulary \mathcal{V}

```
1: Initialize vocabulary  $\mathcal{V} \leftarrow \mathcal{V}_0$  # each initial token corresponds  
   to one unique item feature  
2:  $\mathcal{R} \leftarrow \emptyset$   
3: while  $|\mathcal{V}| < Q$  do  
4:   # Count: accumulate weighted token co-occurrences  
5:    $\text{count}(\cdot, \cdot) \leftarrow \text{Count}(\mathcal{S}', \mathcal{V})$  # Algorithm 2  
6:   # Update: Merge a frequent token pair into a new token  
7:   Select  $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$   
8:    $\mathcal{S}' \leftarrow \text{Update}(\mathcal{S}', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$  # Algorithm 3  
9:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$  # new merge rule  
10:   $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$  # add new token to the vocabulary  
11: end while  
return  $\mathcal{R}, \mathcal{V}$ 
```



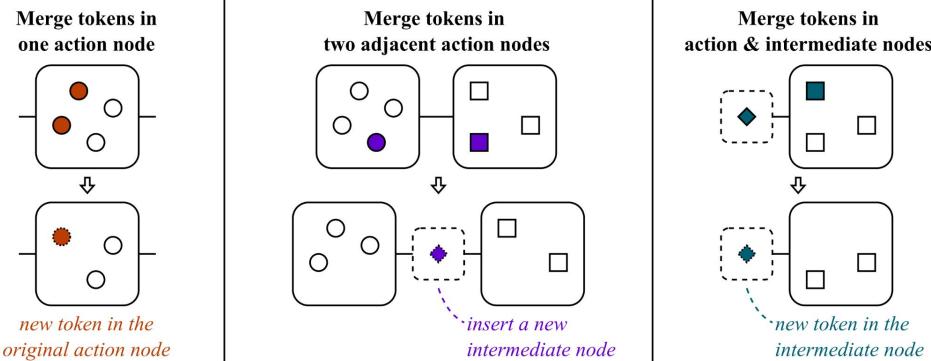
Techniques to Construct SemIDs

Context-independent \Rightarrow Context-aware

Algorithm 1 ActionPiece Vocabulary Construction

input Sequence corpus \mathcal{S}' , initial tokens \mathcal{V}_0 , target size Q
output Merge rules \mathcal{R} , constructed vocabulary \mathcal{V}

```
1: Initialize vocabulary  $\mathcal{V} \leftarrow \mathcal{V}_0$  # each initial token corresponds  
   to one unique item feature  
2:  $\mathcal{R} \leftarrow \emptyset$   
3: while  $|\mathcal{V}| < Q$  do  
4:   # Count: accumulate weighted token co-occurrences  
5:   count( $\cdot, \cdot$ )  $\leftarrow$  Count( $\mathcal{S}', \mathcal{V}$ ) # Algorithm 2  
6:   # Update: Merge a frequent token pair into a new token  
7:   Select  $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$   
8:    $\mathcal{S}' \leftarrow \text{Update}(\mathcal{S}', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$  # Algorithm 3  
9:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$  # new merge rule  
10:   $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$  # add new token to the vocabulary  
11: end while  
return  $\mathcal{R}, \mathcal{V}$ 
```



Summary of Techniques to Construct SemIDs

Context-independent

- Residual Quantization (+ regularization)
- Product Quantization
- Hierarchical Clustering
- LM-based ID Generator

Summary of Techniques to Construct SemIDs

Context-independent

- Residual Quantization (+ regularization)
- Product Quantization
- Hierarchical Clustering
- LM-based ID Generator

Context-aware

Part 1: Semantic ID Construction

- (i) First example: TIGER and RQ-VAE-based SemIDs;
- (ii) Techniques to construct SemIDs;
- (iii) **Inputs** for SemID construction;

Inputs for SemID Construction

Input: all data associated with the item

The screenshot shows a product listing for 'The Legend of Zelda: Tears of the Kingdom' on the Nintendo Switch. On the left, there's a thumbnail image of the game cover, which features Link standing on a cliff edge over a vast ocean. To the left of the thumbnail are two small video thumbnail icons and the text '7 VIDEOS'. Below the thumbnail is a 'M' rating icon and the text 'Fantasy violence'.

The Legend of Zelda: Tears of the Kingdom - Nintendo Switch (US Version)

Brand: Nintendo

Platform: Nintendo Switch | **Rated:** Rating Pending

4.9 ★★★★★ 22,782 ratings

Amazon's Choice

3K+ bought in past month

-21% \$55⁰⁰

List Price: \$69.99 ⓘ

Or **\$9.48** /mo (6 mo). Select from 2 plans

FREE Returns

Roll over image to zoom in

Inputs for SemID Construction

Input: all data associated with the item

What exactly does “all data” mean? 🤔

Inputs for SemID Construction

Text or Multimodal Features



ItemID	Title	Description
B097B2YWFX	The Legend of Zelda: Tears of the Kingdom - Nintendo Switch (US Version)	An epic adventure across the land ... threaten the kingdom?
Video Games	Nintendo Switch	Games
	Nintendo	

Text

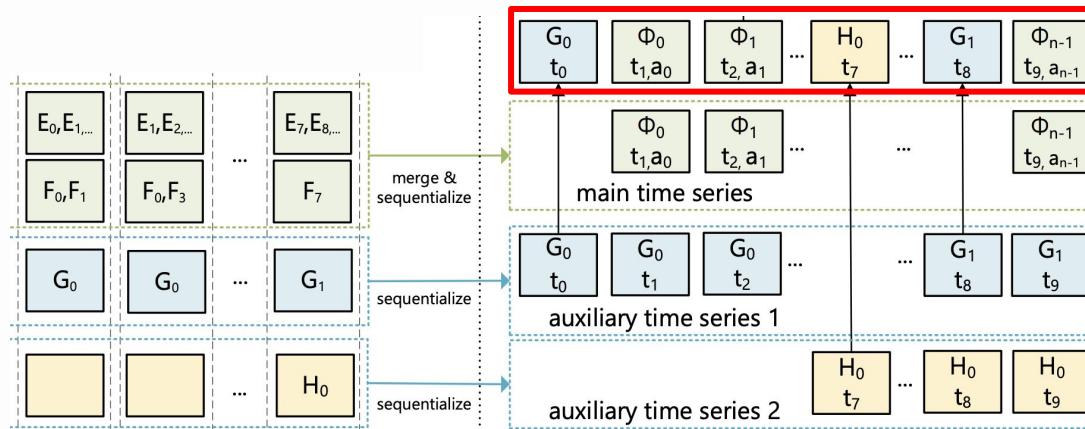


Multimodal

Inputs for SemID Construction

Categorical Features

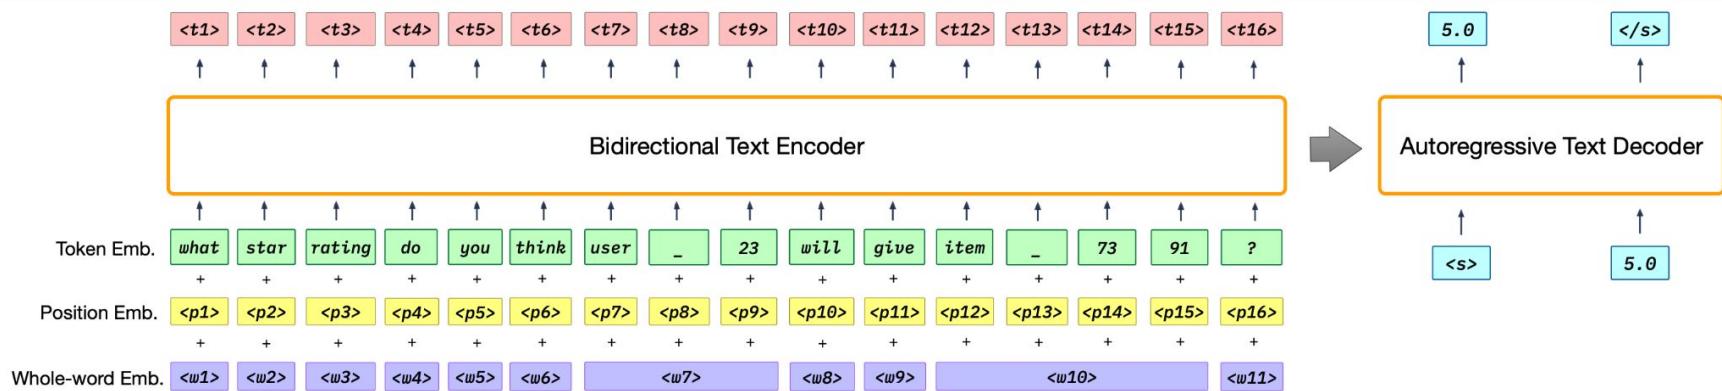
Categorical Features \rightarrow IDs
Merge & Sequentialize



Inputs for SemID Construction

No Features

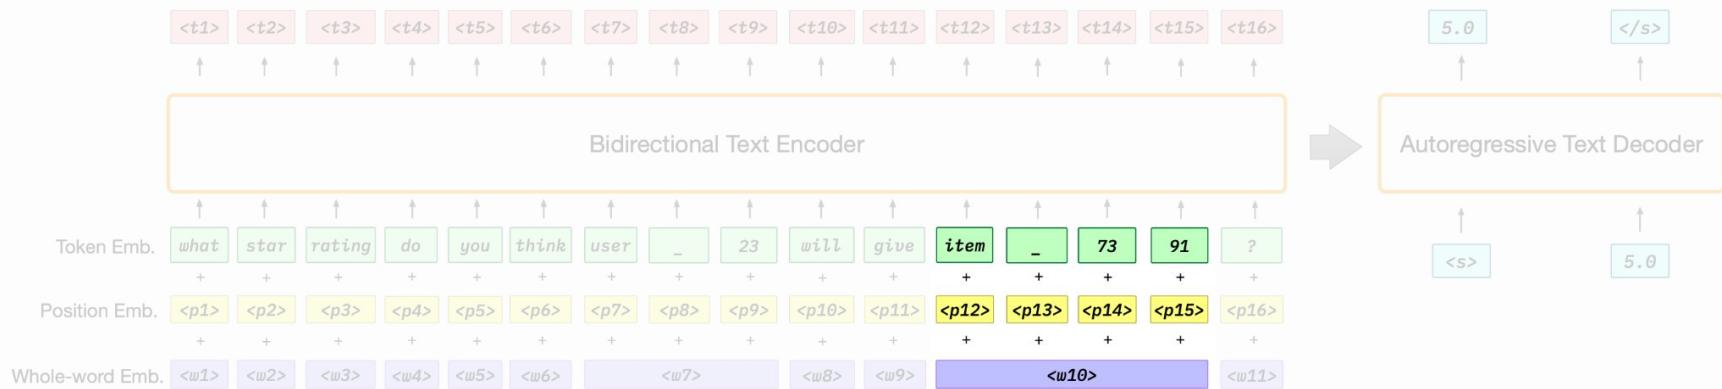
Item ID > IDs
Text Tokenizer



Inputs for SemID Construction

No Features

Item ID > IDs
Text Tokenizer

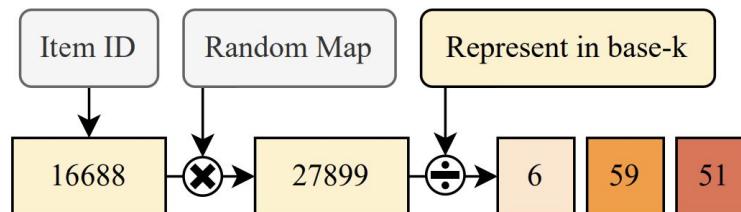


Inputs for SemID Construction

No Features

Random IDs

Balanced Chunked ID



Inputs for SemID Construction

Input: all data associated with the item

(1) Item Metadata

Text / Multimodal / Categorical / No Features

Inputs for SemID Construction

Input: all data associated with the item

(1) Item Metadata

Text / Multimodal / Categorical / No Features

(2) Item Metadata + Behaviors

Inputs for SemID Construction

Input: all data associated with the item

(1) Item Metadata

Text / Multimodal / Categorical / No Features

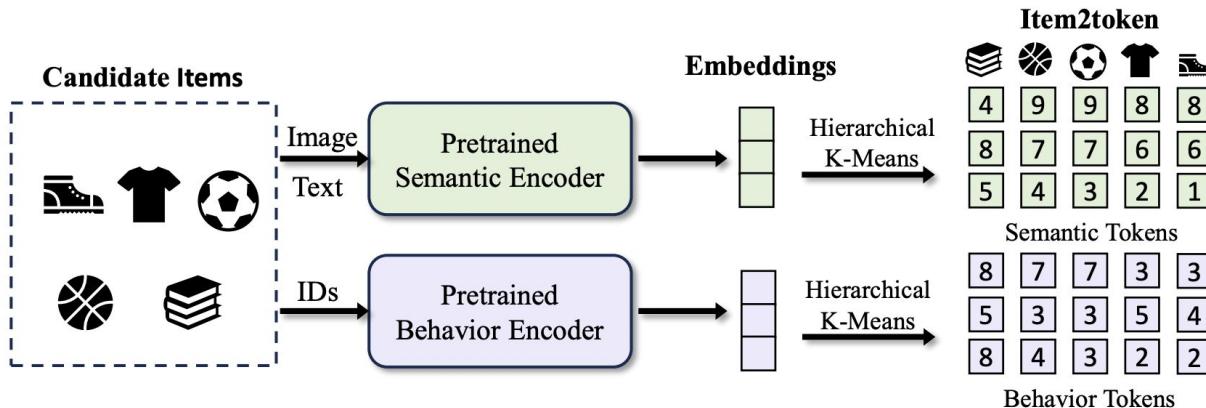
(2) Item Metadata + Behaviors

But how?

Inputs for SemID Construction

Item Metadata + Behaviors

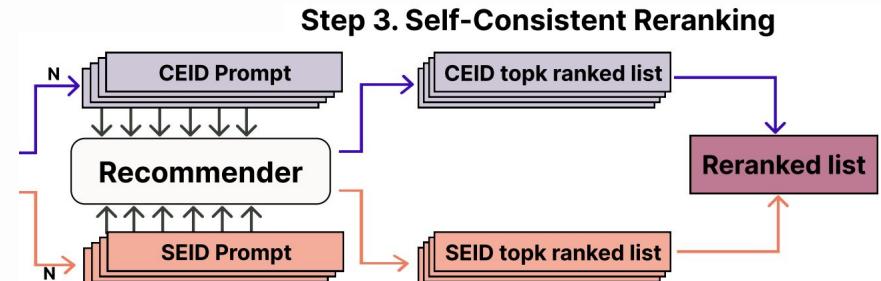
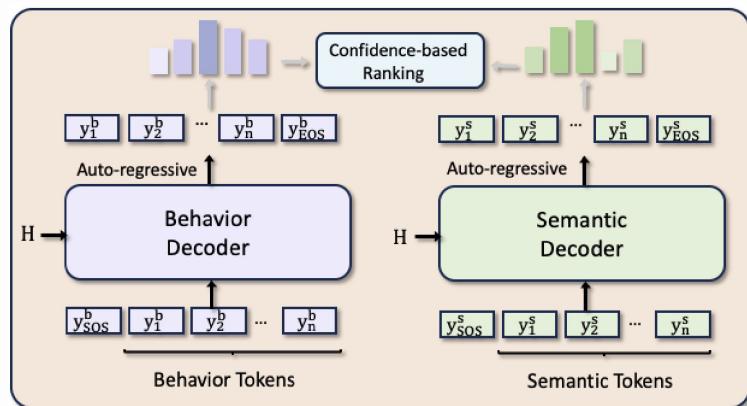
Fused Semantic IDs



Inputs for SemID Construction

Item Metadata + Behaviors

Fused Semantic IDs + Two-stream Generation

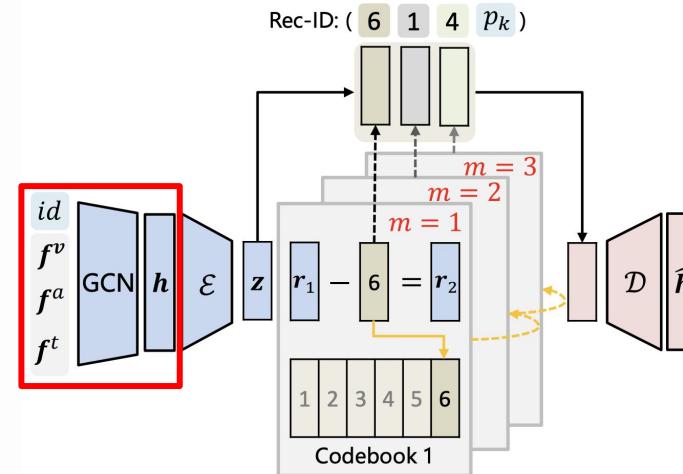


Inputs for SemID Construction

Item Metadata + Behaviors

Fused Representations

User–Item Graph +
Semantic Features



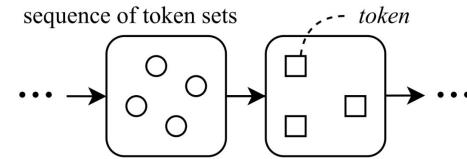
Inputs for SemID Construction

Item Metadata + Behaviors

Train Tokenizer on Behavior Sequence Corpus

Algorithm 1 ActionPiece Vocabulary Construction

```
input Sequence corpus  $S'$ , initial tokens  $\mathcal{V}_0$ , target size  $Q$ 
output Merge rules  $\mathcal{R}$ , constructed vocabulary  $\mathcal{V}$ 
1: Initialize vocabulary  $\mathcal{V} \leftarrow \mathcal{V}_0$  # each initial token corresponds
   to one unique item feature
2:  $\mathcal{R} \leftarrow \emptyset$ 
3: while  $|\mathcal{V}| < Q$  do
4:   # Count: accumulate weighted token co-occurrences
5:    $\text{count}(\cdot, \cdot) \leftarrow \text{Count}(S', \mathcal{V})$  # Algorithm 2
6:   # Update: Merge a frequent token pair into a new token
7:   Select  $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$ 
8:    $S' \leftarrow \text{Update}(S', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$  # Algorithm 3
9:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$  # new merge rule
10:   $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$  # add new token to the vocabulary
11: end while
return  $\mathcal{R}, \mathcal{V}$ 
```



$$P(O, O) = \frac{|O-O|}{|O, O|} = \frac{4-1}{\binom{4}{2}}$$
$$P(O, \square) = \frac{|O-\square|}{|O| \times |\square|} = \frac{1}{4 \times 3}$$
$$P(\square, \square) = \frac{|\square-\square|}{|\square, \square|} = \frac{3-1}{\binom{3}{2}}$$

Features
co-occurring
within or
across items

Inputs for SemID Construction

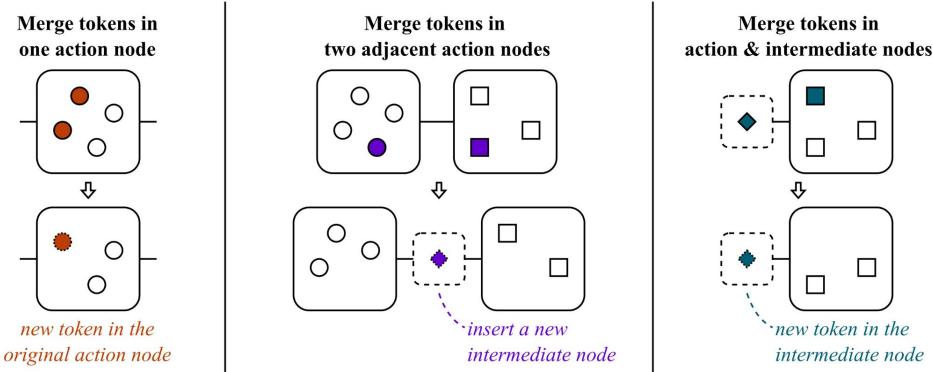
Item Metadata + Behaviors

Train Tokenizer on Behavior Sequence Corpus

Algorithm 1 ActionPiece Vocabulary Construction

input Sequence corpus \mathcal{S}' , initial tokens \mathcal{V}_0 , target size Q
output Merge rules \mathcal{R} , constructed vocabulary \mathcal{V}

- 1: Initialize vocabulary $\mathcal{V} \leftarrow \mathcal{V}_0$ # each initial token corresponds to one unique item feature
- 2: $\mathcal{R} \leftarrow \emptyset$
- 3: **while** $|\mathcal{V}| < Q$ **do**
- 4: **# Count:** accumulate weighted token co-occurrences
- 5: $\text{count}(\cdot, \cdot) \leftarrow \text{Count}(\mathcal{S}', \mathcal{V})$ # Algorithm 2
- 6: **# Update:** Merge a frequent token pair into a new token
- 7: Select $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$
- 8: $\mathcal{S}' \leftarrow \text{Update}(\mathcal{S}', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$ # Algorithm 3
- 9: $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$ # new merge rule
- 10: $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$ # add new token to the vocabulary
- 11: **end while**
- return** \mathcal{R}, \mathcal{V}

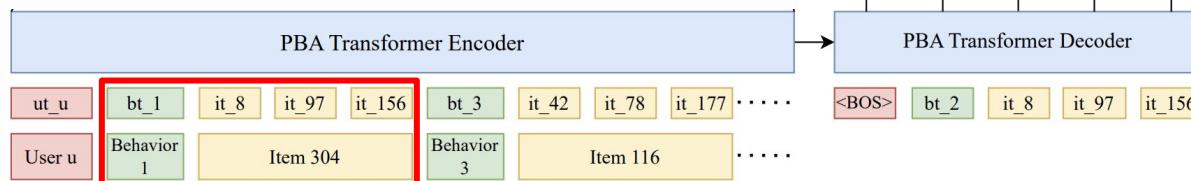


Inputs for SemID Construction

Item Metadata + Behaviors

Multi-Behavior Recommendation

Semantic IDs fused
with **behavior types**



Unified Multi-Task Framework

Target Behavior Item Prediction

bt_1 it_? it_? it_?

Behavior-Specific Item Prediction

bt_* it_? it_? it_?

Behavior-Item Prediction

bt_? it_? it_? it_?

Behavior Prediction

bt_? [] [] []

Next Behavior Next Item

bt_2 it_8 it_97 it_156 <EOS>

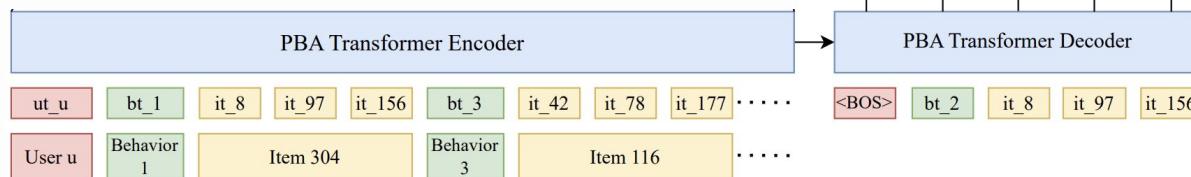
Inputs for SemID Construction

Item Metadata + Behaviors

Multi-Behavior Recommendation

Next Token Prediction as
natural **multi-task learning**

(prompted by behavior type)



Inputs for SemID Construction

Input: all data associated with the item

(1) Item Metadata

Text / Multimodal / Categorical / No Features

(2) Item Metadata + Behaviors

Fused semantic IDs & Representations

Tokenizer trained on behavior sequences

Part 1 Summary – SemID Construction

(1) First Example: TIGER

(2) Construction Techniques

(3) Inputs

Part 1 Summary – SemID Construction

(1) First Example: TIGER

(2) Construction Techniques

Context-independent (PQ, RQ, Clustering,
LM-based generator) → Context-aware

(3) Inputs

Part 1 Summary – SemID Construction

(1) First Example: TIGER

(2) Construction Techniques

Context-independent (PQ, RQ, Clustering,
LM-based generator) → Context-aware

(3) Inputs

Item Metadata (Text, Multimodal, Features)

+ Behaviors (Fused SemIDs / Representations)

Part 2: SemID-based Generative Recommendation Model Architecture

SemID-based Recommender Architecture

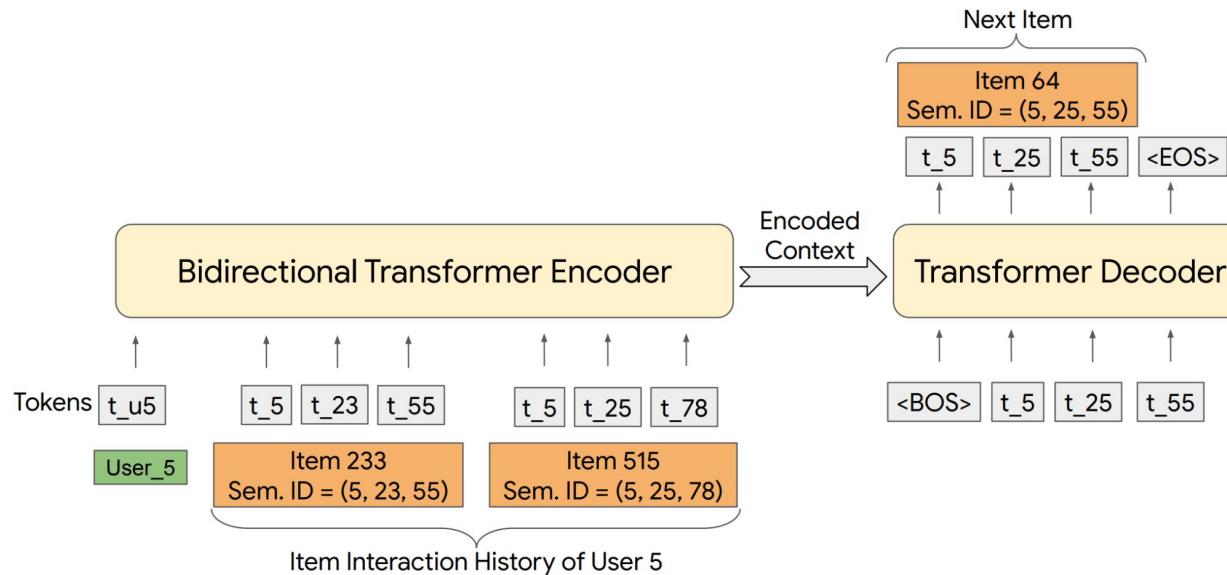
Recommendation as a **seq**-to-**seq** generation problem

Input: user interacted items $\{c_{11}, c_{12}, c_{13}, c_{14}, c_{21}, c_{22}, \dots\}$

Output: next item $\{c_{t1}, c_{t2}, c_{t3}, c_{t4}\}$

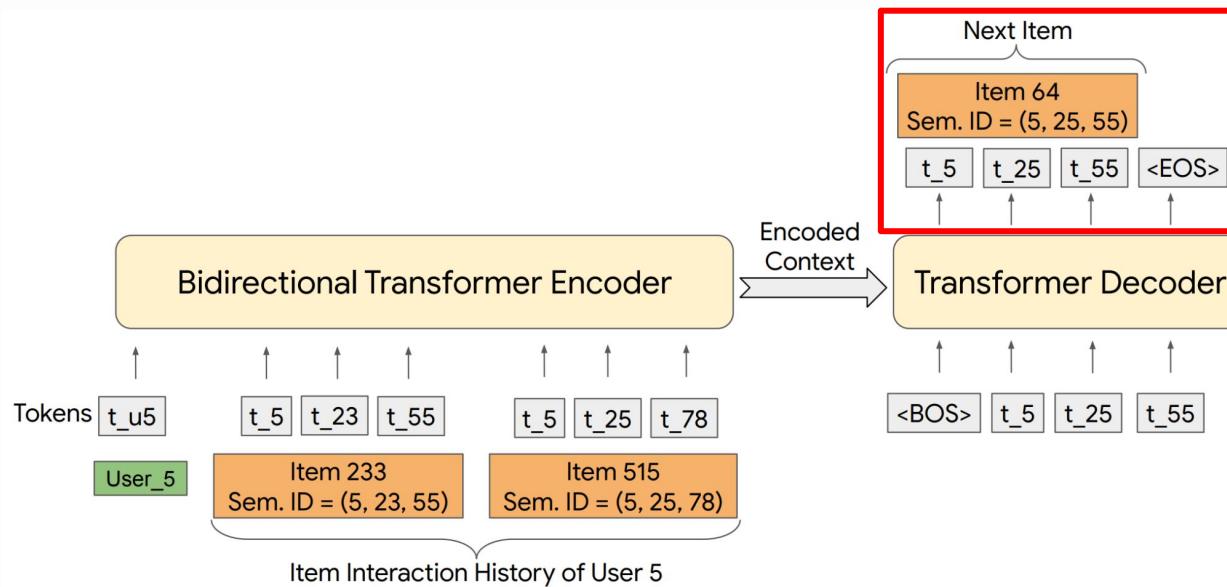
SemID-based Recommender Architecture

Architecture: Decoder-Only / Encoder-Decoder



SemID-based Recommender Architecture

Objective: Next-Token Prediction



SemID-based Recommender Architecture

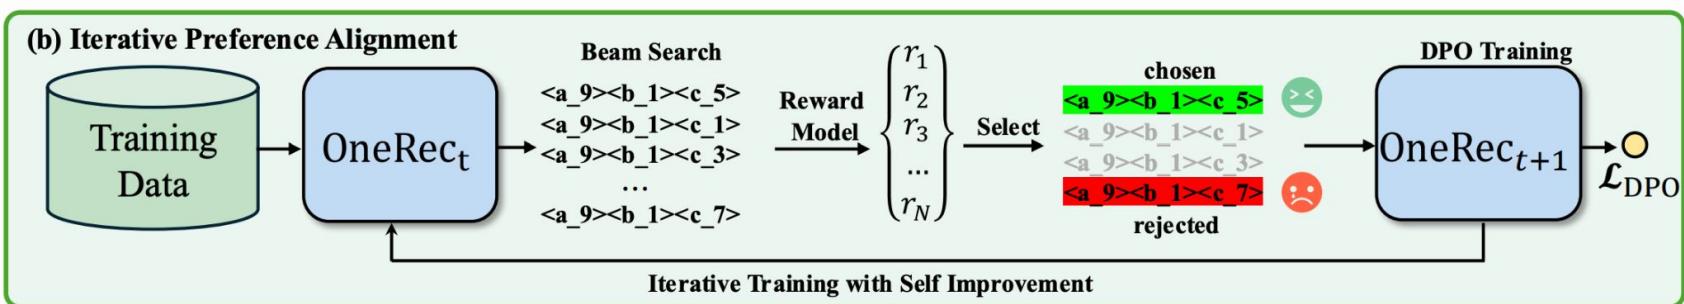
Objective: Next-Token Prediction

Could we add **negative samples** like BPR loss?

SemID-based Recommender Architecture

Objective: Preference Alignment Objective

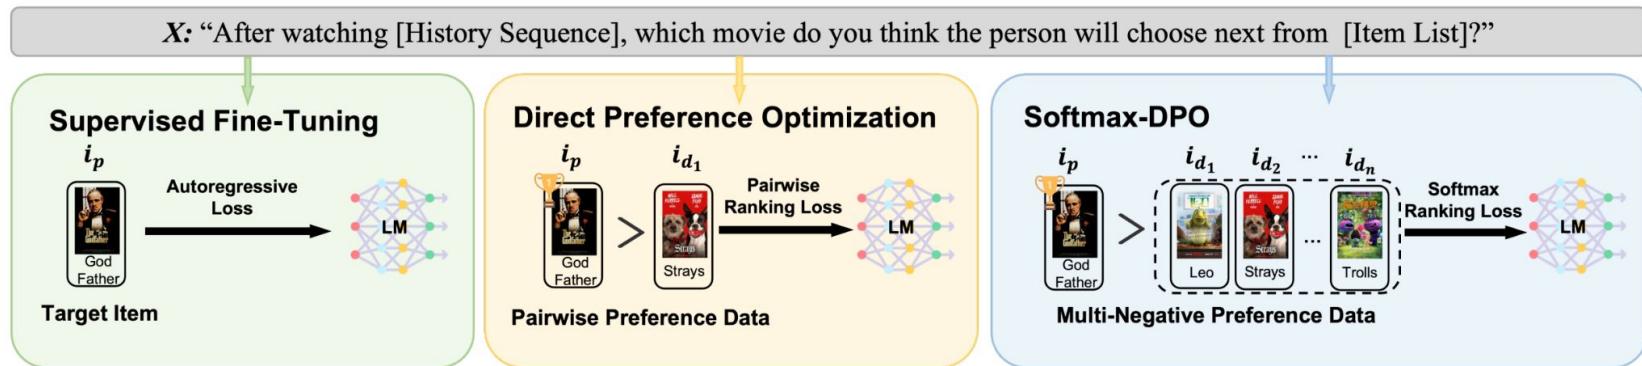
One negative sample per instance



SemID-based Recommender Architecture

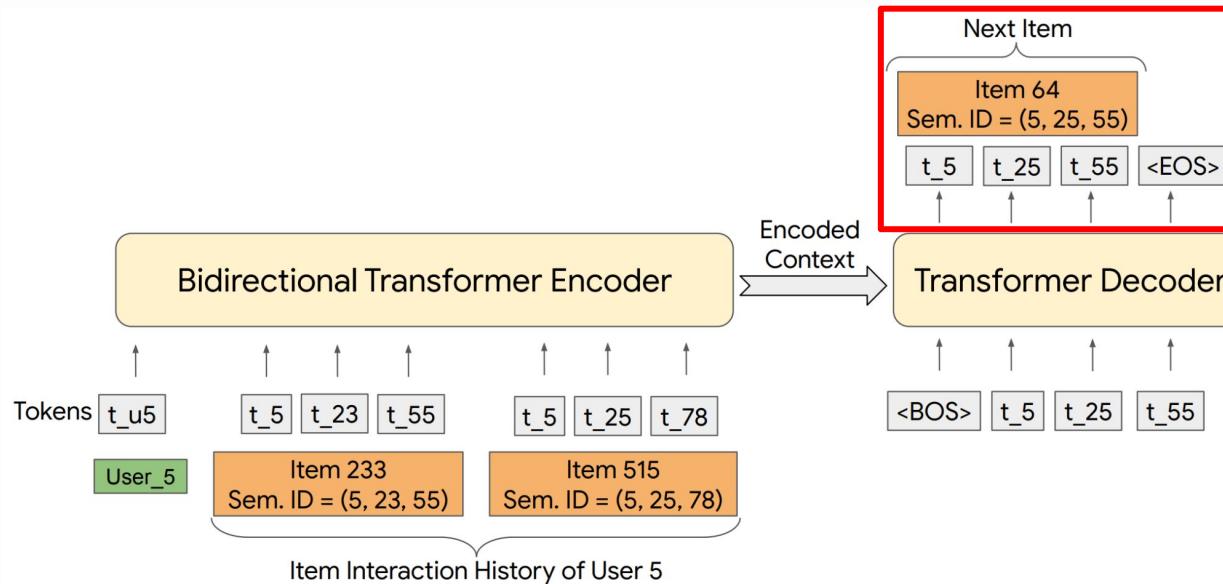
Objective: Preference Alignment Objective

Multiple negative samples per instance



SemID-based Recommender Architecture

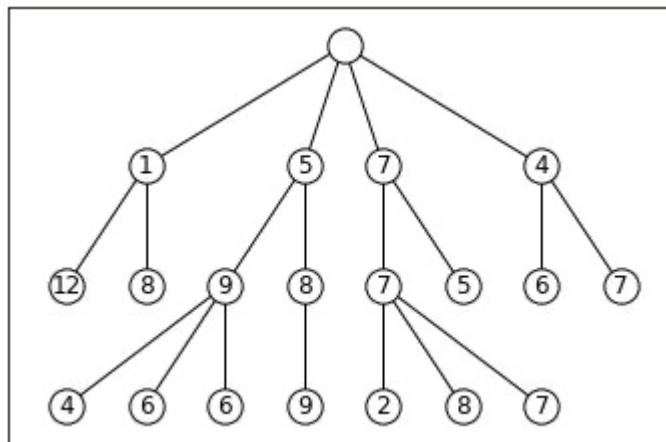
Inference: How to get a **ranking list**?



SemID-based Recommender Architecture

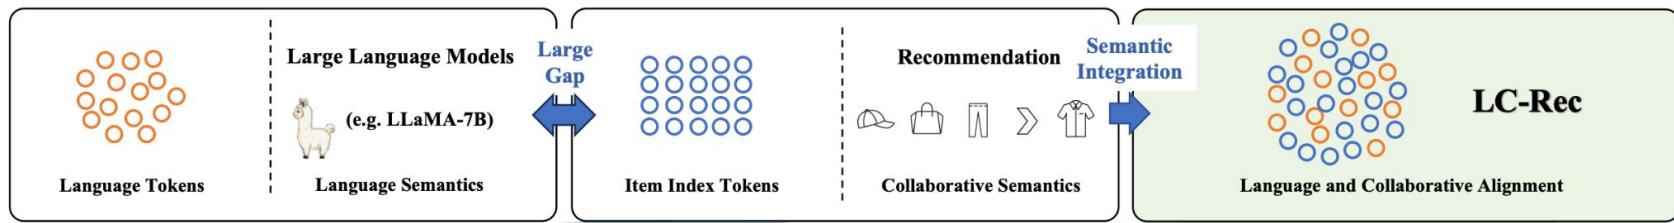
Inference: How to get a **ranking list**?

(Constrained) Beam Search



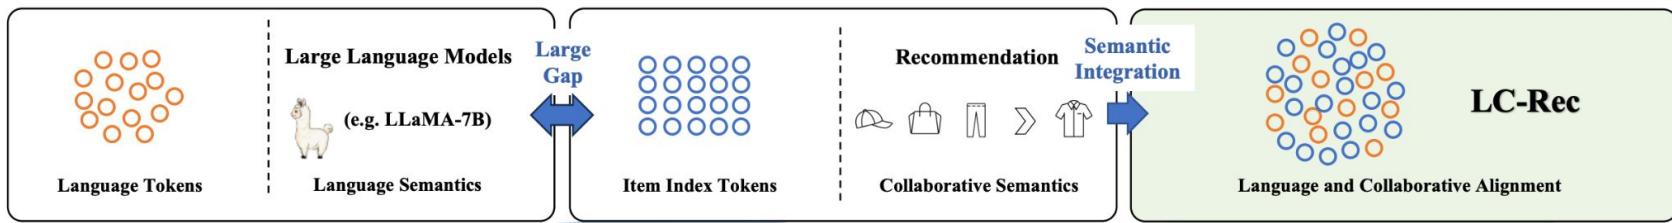
SemID-based Recommender Architecture

Align with LLMS – LC-Rec



SemID-based Recommender Architecture

Align with LLMS – LC-Rec

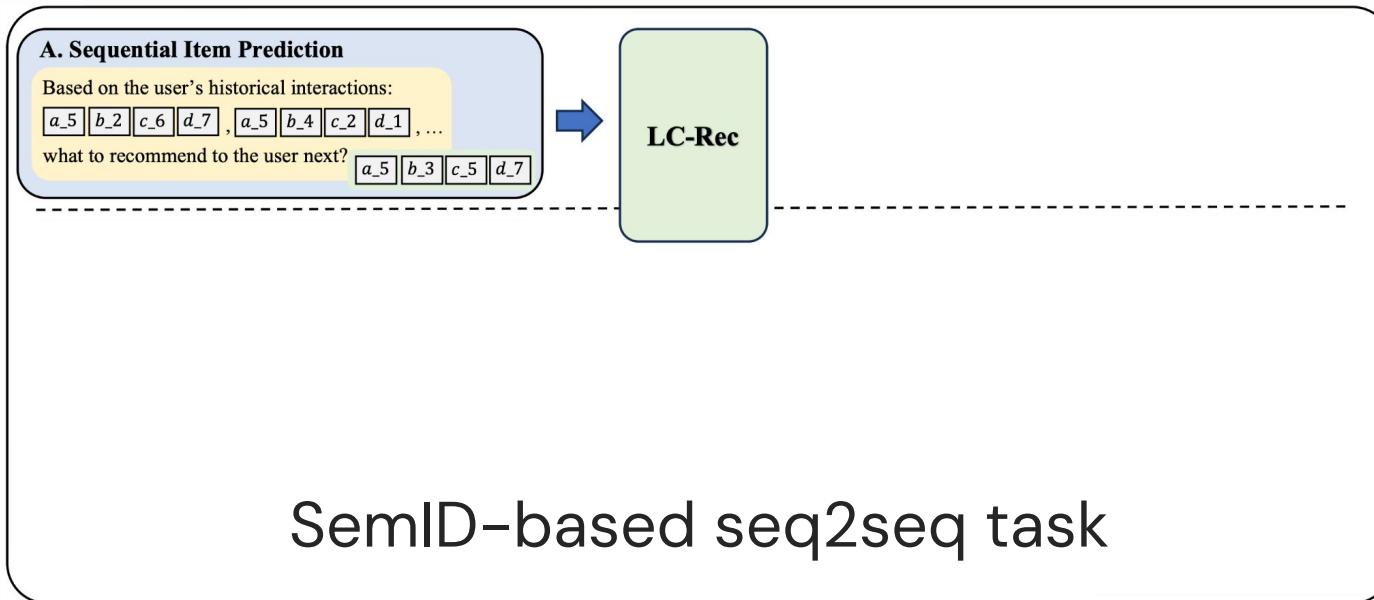


Core Idea:

Construct **instructions** containing
both **Semantic IDs** and **language tokens**

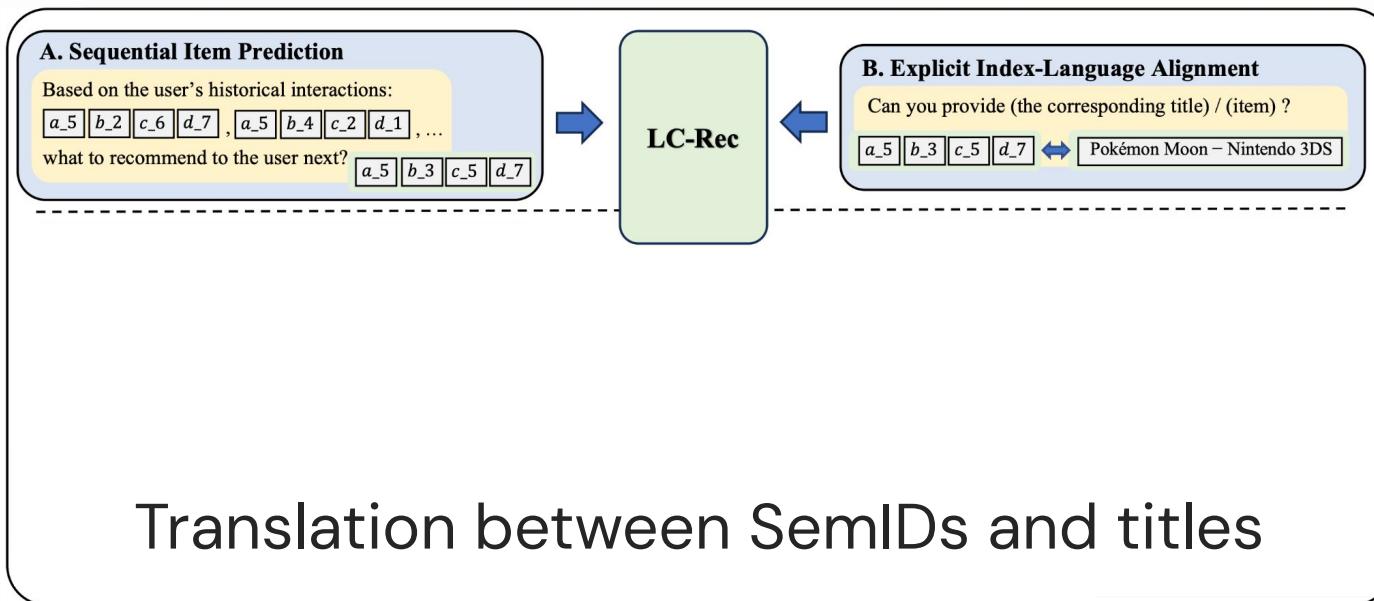
SemID-based Recommender Architecture

Align with LLMS – LC-Rec



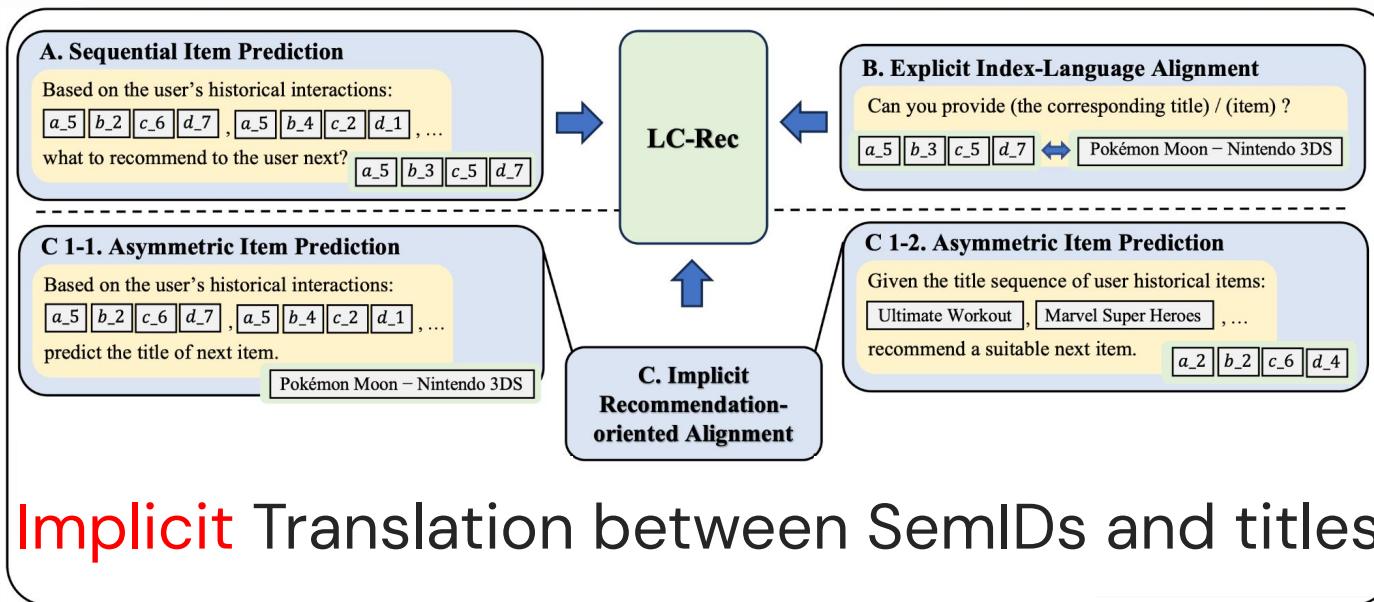
SemID-based Recommender Architecture

Align with LLMS – LC-Rec



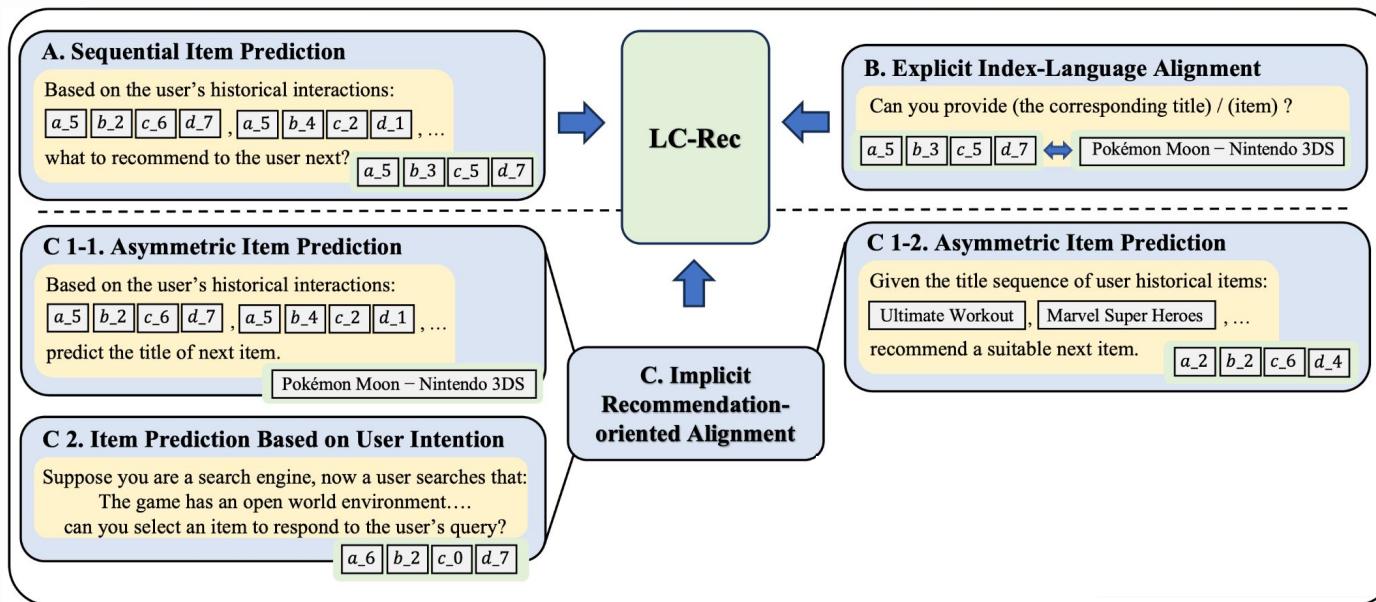
SemID-based Recommender Architecture

Align with LLMS – LC-Rec



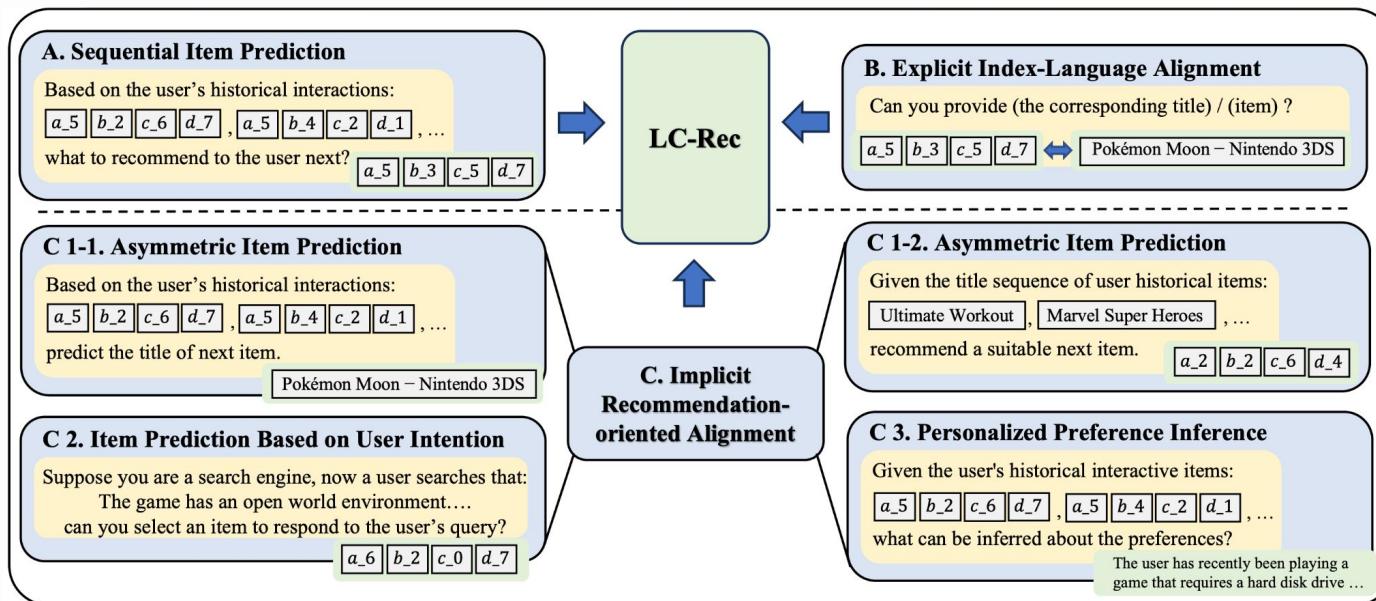
SemID-based Recommender Architecture

Align with LLMS – LC-Rec



SemID-based Recommender Architecture

Align with LLMS – LC-Rec



Part 2 Summary – Architecture

(1) Train from Scratch

(2) Align with LLMs

Part 2 Summary – Architecture

(1) Train from Scratch

Objective (NTP, DPO, S-DPO)

Inference (Beam Search)

(2) Align with LLMs

Part 2 Summary – Architecture

(1) Train from Scratch

Objective (NTP, DPO, S-DPO)

Inference (Beam Search)

(2) Align with LLMs

LC-Rec: Instructions containing both semIDs and language tokens